



Trading Volatility, Correlation, Term Structure and Skew



Colin Bennett

TRADING VOLATILITY

**"A master piece to learn in a nutshell all the essentials about volatility with a practical and lively approach.
A must read!"**

Carole Bernard, Equity Derivatives Specialist at Bloomberg

"This book could be seen as the 'volatility bible'!"

Markus-Alexander Flesch, Head of Sales & Marketing at Eurex

"I highly recommend this book both for those new to the equity derivatives business, and for more advanced readers. The balance between theory and practice is struck At-The-Money"

Paul Stephens, Head of Institutional Marketing at CBOE

"One of the best resources out there for the volatility community"

Paul Britton, CEO and Founder of Capstone Investment Advisors

"Colin has managed to convey often complex derivative and volatility concepts with an admirable simplicity, a welcome change from the all-too-dense tomes one usually finds on the subject"

Edmund Shing PhD, former Proprietary Trader at BNP Paribas

"In a crowded space, Colin has supplied a useful and concise guide"

Gary Delany, Director Europe at the Options Industry Council



ABOUT THE AUTHOR

Colin Bennett is a Managing Director and Head of Quantitative and Derivative Strategy at Banco Santander. Previously he was Head of Delta 1 Research at Barclays Capital, and Head of Convertible and Derivative Research at Dresdner Kleinwort.

Colin started his career in Convertible Bond Research at Merrill Lynch, after studying Mathematics and Electrical Engineering at Cambridge University. In the 1993 National Mathematics Contest Colin came 16th in the UK. He has also worked in Equity Derivative Sales, and as a Desk Analyst for the equity derivative trading desk.

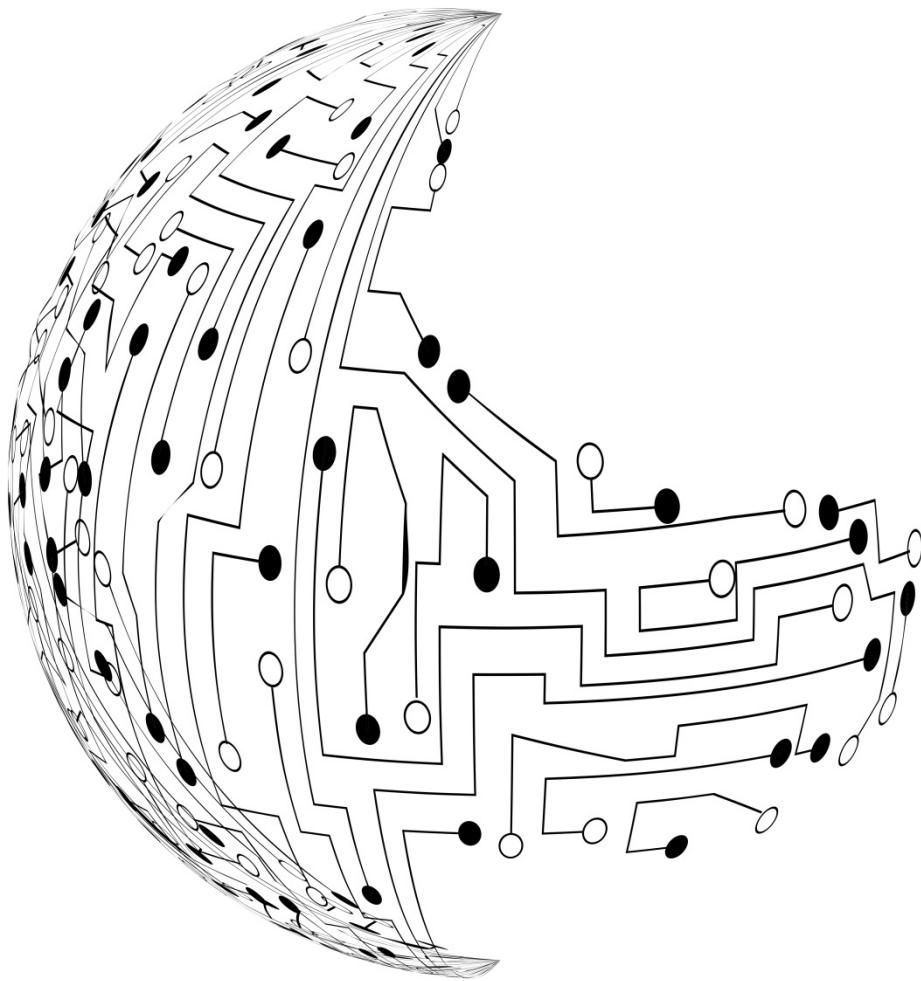
Colin is a regular speaker at CBOE, Eurex, Marcus Evans, Futures and Option World, Risk Magazine and Bloomberg conferences.

www.trading-volatility.com

TRADING VOLATILITY

Trading Volatility, Correlation,
Term Structure and Skew

Colin Bennett



Copyright © 2014 Colin Bennett

All rights reserved. No part of this publication may be reproduced, stored in a retrieval system, or transmitted in any form or by any means – electronic, mechanical, photocopying, recording or otherwise without written permission of the publisher or author.

Trading Volatility: Trading Volatility, Correlation, Term Structure and Skew

ISBN-13: 978-1499206074

Cover Design by Gareth Allen

ACKNOWLEDGEMENTS

I would like to thank James Peattie for persuading me to work at Merrill Lynch, and starting my career within research which I still enjoy to this very day.

Iain Clamp aka “The Guru” deserves special recognition, for explaining the intricacies of volatility trading.

I will always be grateful to Tom Dauterman and Irene Ferrero for their many months of effort proofing this publication. I would also like to thank Mariano Colmenar, without whose support this book would never have been published.

And finally, and most importantly, I would like to thank my wife Claire for her endless patience and understanding with me while I researched and wrote this book over the past 5 years.

NOTE ON CONTENTS

While there are many different aspects to volatility trading, not all of them are suitable for all investors. In order to allow easy navigation, the sections are combined into seven chapters that are likely to appeal to different parts of the equity derivatives client base. The earlier chapters are most suited to equity investors, while later chapters are aimed at hedge funds and proprietary trading desks. The Appendix contains reference material and mathematical detail that has been removed from earlier chapters to enhance readability.

CONTENTS

PREFACE	i
CHAPTER 1	
OPTIONS.....	1
1.1: Option basics	2
1.2: Option trading in practice	4
1.3: Maintenance of option positions	11
1.4: Call overwriting.....	15
1.5: Protection strategies using options.....	23
1.6: Option structures trading.....	31
CHAPTER 2	
VOLATILITY TRADING	35
2.1: Volatility trading using options	36
2.2: Variance is the key, not volatility	49
2.3: Volatility, variance and gamma swaps.....	52
2.4: Options on variance.....	68
CHAPTER 3	
WHY EVERYTHING YOU THINK YOU KNOW ABOUT VOLATILITY IS WRONG.....	73
3.1: Implied volatility should be above realized volatility.....	74
3.2: Long volatility is a poor hedge	77
3.3: Variable annuity hedging lifts long-term vol.....	81
3.4: Structured products vicious circle.....	83
3.5: Stretching black-scholes assumptions	89
CHAPTER 4	
FORWARD STARTING PRODUCTS AND VOLATILITY INDICES	99
4.1: Forward starting products.....	100
4.2: Volatility indices.....	109
4.3: Futures on volatility indices	115
4.4: Volatility future ETN/ETF	122
4.5: Options on volatility futures	129

CHAPTER 5	
LIGHT EXOTICS.....	133
5.1: Barrier options	134
5.2: Worst-of/best-of options.....	140
5.3: Outperformance options.....	143
5.4: Look-back options	146
5.5: Contingent premium options	148
5.6: Composite and quanto options	149
CHAPTER 6	
RELATIVE VALUE AND CORRELATION TRADING.....	153
6.1: Relative value trading.....	154
6.2: Relative value volatility trading.....	158
6.3: Correlation trading	161
6.4: Trading earnings announcements/jumps.....	178
CHAPTER 7	
SKEW AND TERM STRUCTURE TRADING.....	183
7.1: Skew and term structure are linked	184
7.2: Square root of time rule can compare term structures and skews.....	193
7.3: Term structure trading.....	198
7.4: How to measure skew and smile.....	201
7.5: Skew trading	210
APPENDIX	
A.1: Local volatility	230
A.2: Measuring historical volatility	232
A.3: Proof implied jump formula.....	245
A.4: Proof var swaps can be hedged by log contract ($=1/k^2$)	248
A.5: Proof variance swap notional = vega/ 2σ	250
A.6: Modelling volatility surfaces.....	251
A.7: Black-scholes formula.....	255
A.8: Greeks and their meaning	257
A.9: Advanced (practical or shadow) greeks	262
A.10: Shorting stock by borrowing shares	266
A.11: Sortino ratio.....	269
A.12: Capital structure arbitrage	270
INDEX	286

PREFACE

One of the main reasons I decided to write this book, was due to the lack of other publications that deal with the practical issues of using derivatives. This publication aims to fill the void between books providing an introduction to derivatives, and advanced books whose target audience are members of quantitative modelling community.

In order to appeal to the widest audience, this publication tries to assume the least amount of prior knowledge. The content quickly moves onto more advanced subjects in order to concentrate on more practical and advanced topics.

CHAPTER 1

OPTIONS

This chapter is focused on real life uses of options for directional investors, for example using options to replace a long position in the underlying, to enhance the yield of a position through call overwriting, or to provide protection from declines. In addition to explaining these strategies, a methodology to choose an appropriate strike and expiry is shown. Answers to the most common questions are given, such as when an investor should convert an option before maturity, and the difference between delta and the probability that an option expires in the money.

1.1: OPTION BASICS

This section introduces options and the history of options trading. The definition of call and put options, and how they can be used to gain long or short equity exposure, is explained. Key definitions and terminology are given, including strike, expiry, intrinsic value, time value, ATM, OTM and ITM.

HISTORY OF VOLATILITY TRADING

While standardised exchange traded options only started trading in 1973 when the CBOE (Chicago Board Options Exchange) opened, options were first traded in London from 1690. Pricing was made easier by the Black-Scholes-Merton formula (usually shortened to Black-Scholes), which was invented in 1970 by Fischer Black, Myron Scholes and Robert Merton.

Option trading exploded in the 1990s

The derivatives explosion in the 1990s was partly due to the increasing popularity of hedge funds, which led to volatility becoming an asset class in its own right. New volatility products such as volatility swaps and variance swaps were created, and a decade later futures on volatility indices gave investors listed instruments to trade volatility. In this chapter we shall concentrate on option trading.

CALL OPTIONS GIVE RIGHT TO BUY, PUTS RIGHT TO SELL

A European call is a contract that gives the investor the right (but not the obligation) to buy a security at a certain strike price on a certain expiry date (American options can be exercised before expiry). A put is identical except it is the right to sell the security.

Call option gives long exposure, put options give short exposure

A call option profits when markets rise (as exercising the call means the investor can buy the underlying security cheaper than it is trading, and then sell it at a profit). A put option profits when markets fall (as you can buy the underlying security for less, exercise the put and sell the security for a profit). Options therefore allow investors to put on long (profit when prices rise) or short (profit when prices fall) strategies.

SELLING OPTIONS GIVES OPPOSITE EXPOSURE

As a call option gives long exposure to the underlying security, selling a call option results in short exposure to the underlying security. Similarly while a put option is a bearish (profits from decline in the underlying) strategy, selling a put option is a bullish strategy (profits from a rise in the underlying). While the direction of the underlying is the primary driver of profits and losses from buying or selling options, the volatility of the underlying is also a driver.

OPTIONS TRADING GIVES VOLATILITY EXPOSURE

If the volatility of an underlying is zero, then the price will not move and an option's payout is equal to the intrinsic value. Intrinsic value is the greater of zero and the 'spot – strike price' for a call and is the greater of zero and 'strike price – spot' for a put. Assuming that stock prices can move, the value of a call and put will be greater than intrinsic due to the time value (price of option = intrinsic value + time value). If an option strike is equal to spot (or is the nearest listed strike to spot) it is called at-the-money (ATM).

As volatility increases so does the price of call and put options

If volatility is zero, an ATM option has a price of zero (as intrinsic is zero). However, if we assume a stock is €50 and has a 50% chance of falling to €40 and 50% chance of rising to €60, it has a volatility above zero. In this example, an ATM call option with strike €50 has a 50% chance of making €10 (if the price rises to €60 the call can be exercised to buy the stock at €50, which can be sold for €10 profit). The fair value of the ATM option is therefore €5 ($50\% \times €10$); hence, as volatility rises the value of a call rises (a similar argument can be used for puts).

Options have greatest time value when strike is similar to spot (i.e. ATM)

An ATM option has the greatest time value (the amount the option price is above the intrinsic value). This can be seen in the same example by looking at an out-the-money (OTM) call option of strike €60 (an OTM option has strike far away from spot and zero intrinsic value). This OTM €60 call option would be worth zero, as the stock in this example cannot rise above €60.

ITM options trade less than OTM options as they are more expensive

An in-the-money (ITM) option is one which has a strike far away from spot and positive intrinsic value. Due to the positive intrinsic value ITM options are relatively expensive, hence tend to trade less than their cheaper OTM counterparts.

1.2: OPTION TRADING IN PRACTICE

Using options to invest has many advantages over investing in cash equity. Options provide leverage and an ability to take a view on volatility as well as equity direction. However, investing in options is more complicated than investing in equity, as a strike and expiry need to be chosen. This section explains hidden risks, (e.g. dividends) and other practical aspects of option trading such as how to choose the strike, and the difference between delta and the probability an option ends up ITM.

CHOOSING EXPIRY IS THE MOST DIFFICULT DECISION

The biggest difference between using options and cash equities (or delta 1 products) to gain equity exposure is the fact a suitable expiry has to be chosen. Determining if an equity is cheap or expensive is often easier than determining the driver and timing of the likely increase / decrease.

Choosing a far dated expiry gives most opportunity for the expected correction, however far dated options are very expensive. Conversely if a cheaper near dated expiry is chosen, there is little time for the anticipated movement to occur. Usually key dates such as quarterly reporting or elections help determine a suitable expiry.

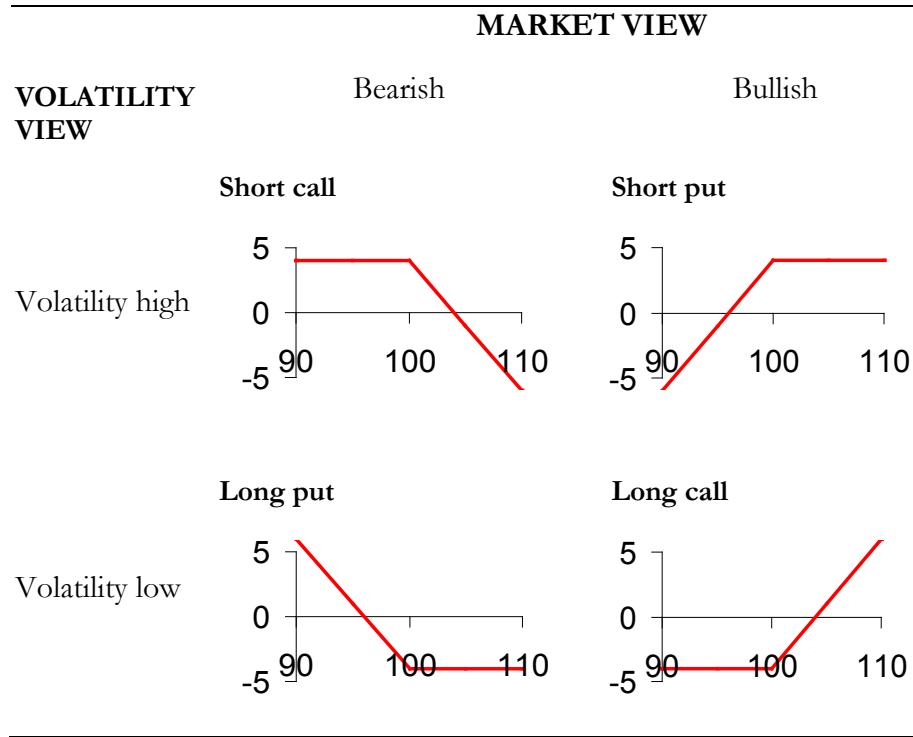
Expiry choice enforces investor discipline

Having to choose an expiry can be seen as a disadvantage of option trading, but some investors see it as an advantage as it enforces investor discipline. The process of choosing an expiry focuses attention on the likely dates a stock will converge with a forecast target price.

If the stock performs as expected the process of expiration forces the profits on the option position to be taken, and ensures a position is not held longer than it should be. Additionally if the anticipated return has not occurred by the expected date, the position expires worthless and forces the investor to make a decision if another position should be initiated. Using options to gain equity exposure therefore prevents “inertia” in a portfolio.

Both an equity and volatility view is needed to trade options

Option trading allows a view on equity and volatility markets to be taken. If implied volatility is seen to be expensive then a short volatility strategy is best (short put for a bullish strategy, short call for a bearish strategy). However if implied volatility is seen to be cheap then a long volatility strategy is best (long call for a bullish strategy, long put for a bearish strategy). The appropriate strategy for a one leg option trade is shown in Figure 1 below. For multiple leg strategies see the section [1.6 Option Structures Trading](#).

Figure 1. Option Strategy for Different Market and Volatility Views

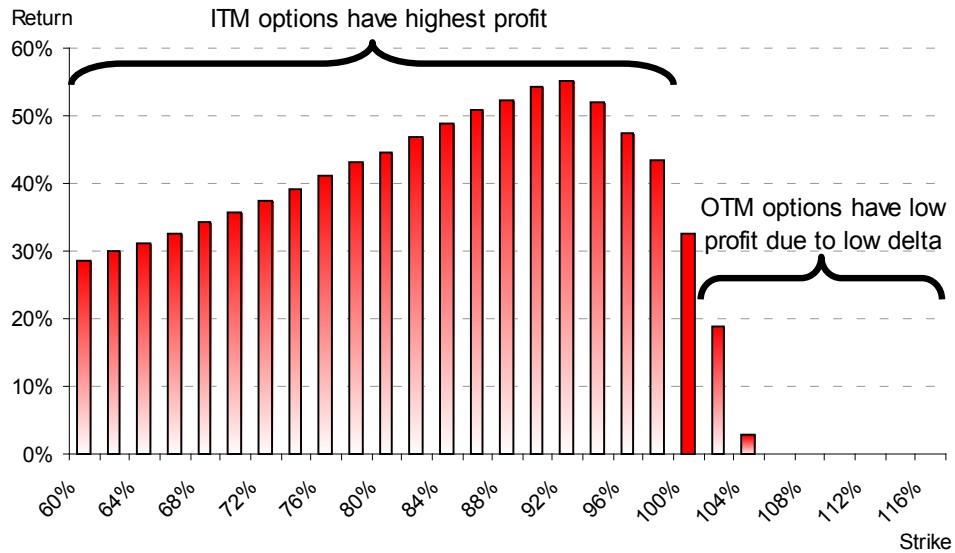
Long vol strategies should have expiry just after key date

Typically if a key date is likely to be volatile then a long volatility strategy (long call or long put) should have an expiry just after this date. Conversely a short volatility strategy (short call or short put) should have an expiry just before the key date. For investors who wish to trade the implied jump of a key date, details of how to trade this implied jump is dealt with in the section [6.4: Trading Earnings Announcements/Jumps](#).

CHOOSING STRIKE OF STRATEGY IS NOT TRIVIAL

While choosing the strike of a strategy is not as difficult as choosing the expiry, it is not trivial. Investors could choose ATM to benefit from greatest liquidity. Alternatively, they could look at the highest expected return (option payout less the premium paid, as a percentage of the premium paid).

Figure 2. Profit of 12 Month Options if Markets Rise 10% by Expiry



ITM options have highest return for “normal” market moves

While choosing a cheap OTM option might be thought of as giving the highest return, the Figure below shows that, in fact, the highest returns come from in-the-money (ITM) options (ITM options have a strike far away from spot and have intrinsic value). This is because an ITM option has a high delta (sensitivity to equity price); hence, if an investor is relatively confident of a specific return, an ITM option has the highest return for relatively “normal” market moves (as trading an ITM option is similar to trading a forward).

Forwards are better than options for pure directional plays

A forward is a contract that obliges the investor to buy a security on a certain expiry date at a certain strike price. A forward has a delta of 100%. An ITM call option has many similarities with being long a forward, as it has a relatively small time value (compared to ATM) and a delta close to 100%. While the intrinsic value does make the option more expensive, this intrinsic value is returned at expiry. However, for an ATM option, the time value purchased is deducted from the returns. For pure directional plays, forwards (or futures, their listed equivalent) are more profitable than options. The advantage of options is in offering convexity: if markets move against the investor the only loss is the premium paid, whereas a forward has a virtually unlimited loss.

OTM options have highest return for “abnormal” moves

Only if the expected return is relatively high (or abnormal) do ATM or OTM options have the highest return. This is because for exceptional returns their low cost and high leverage more than compensates for their lower delta.

LIQUIDITY CAN BE A FACTOR IN CHOOSING STRIKE

If an underlying is relatively illiquid, or if the size of the trade is large, an investor should take into account the liquidity of the maturity and strike of the option. Typically, OTM options are more liquid than ITM options as ITM options tie up a lot of capital. This means that for strikes less than spot, puts are more liquid than calls and vice versa.

Low strike puts are usually more liquid than high strike calls

We note that as low-strike puts have a higher implied than high-strike calls, their value is greater and, hence, traders are more willing to use them. Low strike put options are therefore usually more liquid than high-strike call options. In addition, demand for protection lifts liquidity for low strikes compared with high strikes.

Single stock liquidity is limited for maturities up to two years

For single stock options, liquidity starts to fade after one year and options rarely trade over two years. For indices, longer maturities are liquid, partly due to the demand for long-dated hedges and their use in structured products. While structured products can have a maturity of five to ten years, investors typically lose interest after a few years and sell the product back. The hedging of a structured product, therefore, tends to be focused on more liquid maturities of around three years.

Hedge funds and structured product flow can overlap

Hedge funds tend to focus around the one-year maturity, with two to three years being the longest maturity they will consider. The two-to-three year maturity is where there is greatest overlap between hedge funds and structured desks.

DELTA MEASURES DIVIDEND RISK AND EQUITY RISK

The delta of the option is the amount of equity market exposure an option has. As a stock price falls by the dividend amount on its ex-date, delta is equal to the exposure to dividends that go ex before expiry. The dividend risk is equal to the negative of the delta. For example, if you have a call of positive delta, if (expected or actual) dividends rise, the call is worth less (as the stock falls by the dividend amount).

If a dividend is substantial, it could be in an investor's interest to exercise early. For more details, see the section [1.3 Maintenance of Option Positions](#).

DELTA IS NOT THE PROBABILITY OPTION EXPIRES ITM

A digital call option is an option that pays 100% if spot expires above the strike price (a digital put pays 100% if spot is below the strike price). The probability of such an option expiring ITM is equal to its delta, as the payoff only depends on it being ITM or not (the size of the payment does not change with how much ITM spot is). For a vanilla option this is not the case; hence, there is a difference between the delta and the probability of being ITM. This difference is typically small unless the maturity of the option is very long.

Delta takes into account the amount an option can be ITM

While a call can have an infinite payoff, a put's maximum value is the strike (as spot cannot go below zero). The delta hedge for the option has to take this into account, so a call delta must be greater than the probability of being ITM. Similarly, the absolute value (as put deltas are negative) of the put delta must be less than the probability of expiring ITM. A more mathematical explanation (for European options) is given below:

$$\text{Call delta} > \text{Probability call ends up ITM}$$

$$\text{Abs (Put delta)} < \text{Probability put ends up ITM}$$

Mathematical proof option delta is different from probability of being ITM at expiry

$$\text{Call delta} = N(d_1) \quad \text{Put delta} = N(d_1) - 1$$

$$\text{Call probability ITM} = N(d_2) \quad \text{Put probability ITM} = 1 - N(d_2)$$

where:

Definition of d_1 is the standard Black-Scholes formula for d_1 . For more details, see the section [A.7 Black-Scholes Formula](#).

$$d_2 = d_1 - \sigma \sqrt{T}$$

$$\sigma = \text{implied volatility}$$

$$T = \text{time to expiry}$$

$$N(z) = \text{cumulative normal distribution}$$

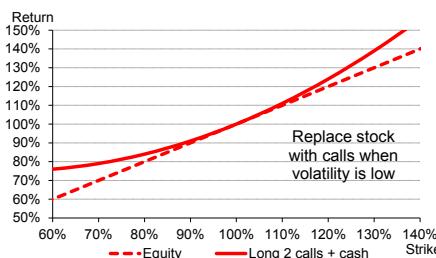
As d_2 is less than d_1 (see above) and $N(z)$ is a monotonically increasing function, this means that $N(d_2)$ is less than $N(d_1)$. Hence, the probability of a call being in the money = $N(d_2)$ is less than the delta = $N(d_1)$. As the delta of a put = delta of call - 1, and the sum of call and put being ITM = 1, the above results for a put must be true as well.

The difference between delta and probability being ITM at expiry is greatest for long-dated options with high volatility (as the difference between d_1 and d_2 is greatest for them).

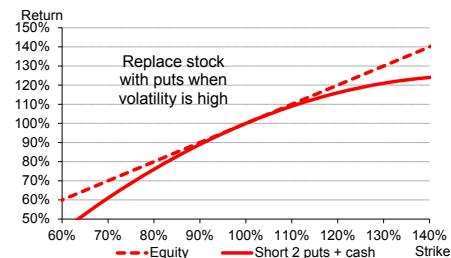
STOCK REPLACING WITH LONG CALL OR SHORT PUT

As a stock has a delta of 100%, the identical exposure to the equity market can be obtained by purchasing calls (or selling puts) whose total delta is 100%. For example, one stock could be replaced by two 50% delta calls, or by going short two -50% delta puts. Such a strategy can benefit from buying (or selling) expensive implied volatility. There can also be benefits from a tax perspective and, potentially, from any embedded borrow cost in the price of options (price of positive delta option strategies is improved by borrow cost). As the proceeds from selling the stock are typically greater than the cost of the calls (or margin requirement of the short put), the difference can be invested to earn interest.

Figure 3. Stock Replacing with Calls



Stock Replacing with Puts



Stock replacing via calls benefits from convexity

As a call option is convex, this means that the delta increases as spot increases and vice versa. If a long position in the underlying is sold and replaced with calls of equal delta, then if markets rise the delta increases and the calls make more money than the long position would have. Similarly, if markets fall the delta decreases and the losses are reduced. This can be seen in Figure 3 above as the portfolio of cash (proceeds from sale of the underlying) and call options is always above the long underlying profile. The downside of using calls is that the position will give a worse profile than the original long position if the underlying does not move much (as call options will fall each day by the theta if spot remains unchanged). Using call options is best when implied volatility is cheap and the investor expects the stock to move by more than currently implied.

Put underwriting benefits from selling expensive implied

Typically the implied volatility of options trades slightly above the expected realised volatility of the underlying over the life of the option (due to a mismatch between supply and demand). Stock replacement via put selling therefore benefits from selling (on average) expensive volatility. Selling a naked put is known as put underwriting, as the investor has effectively underwritten the stock (in the same way investment banks underwrite a rights issue).

Put underwriting pays investors for work that otherwise might be wasted

The strike of put underwriting should be chosen at the highest level at which the investor would wish to purchase the stock, which allows an investor to earn a premium from taking this view (whereas normally the work done to establish an attractive entry point would be wasted if the stock did not fall to that level).

Asset allocators use put underwriting to rebalance portfolios

This strategy has been used significantly recently by asset allocators who are underweight equities and are waiting for a better entry point to re-enter the equity market (earning the premium provides a buffer should equities rally). If an investor does not wish to own the stock and only wants to earn the premium, then an OTM strike should be chosen at a support level that is likely to remain firm.

Put underwriting benefits from selling skew

Put underwriting gives a similar profile to a long stock, short call profile, otherwise known as call overwriting. One difference between call overwriting and put underwriting is that if OTM options are used, then put underwriting benefits from selling skew (which is normally overpriced). For more details on the benefits of selling volatility, see the section [1.4 Call Overwriting](#).

STOCK REPLACEMENT ALTERS DIVIDEND EXPOSURE

It is important to note that the dividend exposure is not the same, as only the owner of a stock receives dividends. While the option owner does not benefit directly, the expected dividend will be used to price the option fairly (hence investors only suffer/benefit if dividends are different from expectations).

1.3: MAINTENANCE OF OPTION POSITIONS

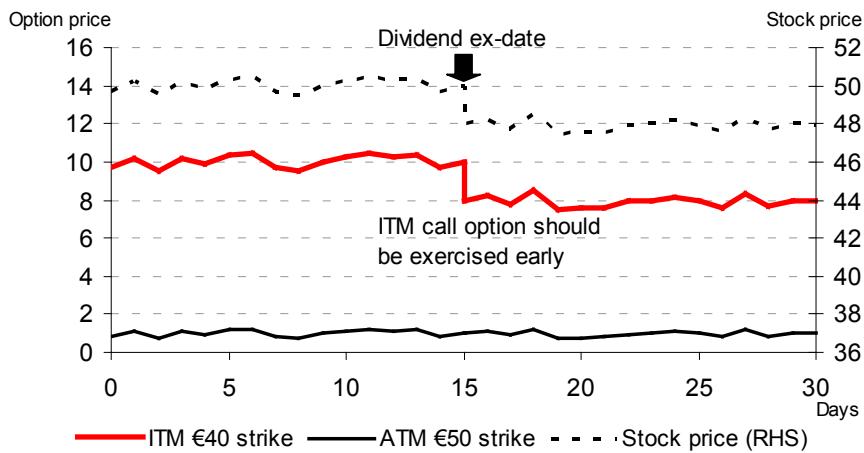
During the life of an American option many events can occur in which it might be preferable to own the underlying shares (rather than the option) and exercise early. In addition to dividends, an investor might want the voting rights, or alternatively might want to sell the option to purchase another option (rolling the option). We investigate these life cycle events and explain when it is in an investor's interest to exercise, or roll, an option before expiry.

CONVERTING OPTIONS EARLY IS RARE

Options on indices are usually European, which means they can only be exercised at maturity. The inclusion of automatic exercise, and the fact it is impossible to exercise before maturity, means European options require only minimal maintenance. Single stock options, however, are typically American (apart from emerging market underlyings). While American options are rarely exercised early, there are circumstances when it is in an investor's interest to exercise an ITM option early. For both calls and puts the correct decision for early exercise depends on the net benefit of doing so (ie, the difference between earning the interest on the strike and net present value of dividends) versus the time value of the option.

- **Calls should be exercised just before the ex-date of a large unadjusted dividend.** In order to exercise a call, the strike price needs to be paid. The interest on this strike price normally makes it unattractive to exercise early. However, if there is a large unadjusted dividend that goes ex before expiry, it might be in an investor's interest to exercise an ITM option early (see Figure 4 above). In this case, the time value should be less than the dividend NPV (net present value) less total interest r ($=e^{rf_{t \times T}} - 1$) earned on the strike price K . In order to maximise 'dividend NPV – Kr ', it is best to exercise just before an ex-date (as this maximises 'dividend NPV' and minimises the total interest r).
- **Puts should be exercised early (preferably just after ex-date) if interest rates are high.** If interest rates are high, then the interest r from putting the stock back at a high strike price K (less dividend NPV) might be greater than the time value. In this case, a put should be exercised early. In order to maximise ' Kr – dividend NPV', a put should preferably be exercised just after an ex-date.

Figure 4. Price of ITM and ATM Call Option with Stock Price over Div Ex-Date



Calls should be exercised early if there is a large dividend

The payout profile of a long call is similar to the payout of a long stock + long put of the same strike . As only ITM options should be exercised and as the strike of an ITM call means the put of the same strike is OTM, we shall use this relationship to calculate when an option should be exercised early.

An American call should only be exercised if it is in an investor's interest to exercise the option and buy a European put of the same strike (a European put of same strike will have the same time value as a European call if intrinsic value is assumed to be the forward).

- **Choice A:** Do not exercise. In this case there is no benefit or cost.
- **Choice B:** Borrow strike K at interest $r (=e^{fr \times T}-1)$ in order to exercise the American call. The called stock will earn the dividend NPV and the position has to be hedged with the purchase of a European put (of cost equal to the time value of a European call).

An investor will only exercise early if choice B > choice A.

$$\begin{aligned} \Rightarrow -Kr + \text{dividend NPV} - \text{time value} &> 0 \\ \Rightarrow \text{dividend NPV} - Kr &> \text{time value for American call to be exercised} \end{aligned}$$

Puts should only be exercised if interest earned (less dividends) exceeds time value

For puts, it is simplest to assume an investor is long stock and long an American put. This payout is similar to a long call of the same strike. An American put should only be exercised against the long stock in the same portfolio if it is in an investor's interest to exercise the option and buy a European call of the same strike.

- **Choice A:** Do not exercise. In this case the portfolio of long stock and long put benefits from the dividend NPV.
- **Choice B:** Exercise put against long stock, receiving strike K, which can earn interest r ($=e^{fr \times T} - 1$). The position has to be hedged with the purchase of a European call (of cost equal to the time value of a European put).

An investor will only exercise early if choice B > choice A

- ⇒ $Kr - \text{time value} > \text{dividend NPV}$
- ⇒ $Kr - \text{dividend NPV} > \text{time value for American put to be exercised}$

Selling ITM options that should be exercised early can be profitable

There have been occasions when traders deliberately sell ITM options that should be exercised early, hoping that some investors will forget. Even if the original counterparty is aware of this fact, exchanges randomly assign the counterparty to exercised options. As it is unlikely that 100% of investors will realise in time, such a strategy can be profitable.

ITM OPTIONS USUALLY EXERCISED AUTOMATICALLY

In order to prevent situations where an investor might suffer a loss if they do not give notice to exercise an ITM option in time, most exchanges have some form of automatic exercise. If an investor (for whatever reason) does not want the option to be automatically exercised, he must give instructions to that effect. The hurdle for automatic exercise is usually above ATM in order to account for any trading fees that might be incurred in selling the underlying post exercise.

Eurex automatic exercise has a higher hurdle than CBOE

For the CBOE, options are automatically exercised if they are US\$0.01 or more ITM (reduced in June 2008 from US\$0.05 or more), which is in line with Euronext-Liffe rules of a €0.01 or GBP0.01 minimum ITM hurdle. Eurex has a higher automatic hurdle, as a contract price has to be ITM by 99.99 or more (eg, for a euro-denominated stock with a contract size of 100 shares this means it needs to be at €0.9999 or more). Eurex does allow an investor to specify an automatic exercise level lower than the automatic hurdle, or a percentage of exercise price up to 9.99%.

CORPORATE ACTIONS CAN ADJUST STRIKE

While options do not adjust for ordinary dividends¹, they do adjust for special dividends. Different exchanges have different definitions of what is a special dividend, but typically it is considered special if it is declared as a special dividend, or is larger than a certain threshold (eg, 10% of the stock price). In addition, options are adjusted in the event of a corporate action, for example, a stock split or rights issue.

Equities and indices can treat bonus share issues differently

Options on equities and indices can treat bonus share issues differently. A stock dividend in lieu of an ordinary dividend is considered an ordinary dividend for options on an equity (hence is not adjusted) but is normally adjusted by the index provider.

Adjustment negates impact of dividend or corporate action

For both special dividends and corporate actions, the adjustment negates the impact of the event (principal of unchanged contract values), so the theoretical price of the options should be able to ignore the event. As the strike post adjustment will be non-standard, typically exchanges create a new set of options with the normal strikes. While older options can still trade, liquidity generally passes to the new standard strike options (particularly for longer maturities which do not have much open interest).

M&A AND SPINOFFS CAN CAUSE PROBLEMS

If a company spins off a subsidiary and gives shareholders shares in the new company, the underlying for the option turns into a basket of the original equity and the spun-off company. New options going into the original company are usually created, and the liquidity of the options into the basket is likely to fade. For a company that is taken over, the existing options in that company will convert into whatever shareholders were offered. If the acquisition was for stock, then the options convert into shares, but if the offer is partly in cash, then options can lose a lot of value as the volatility of cash is zero.

OPTIONS OFTEN ROLLED BEFORE EXPIRY

The time value of an option decays quicker for short-dated options than for far-dated options. To reduce the effect of time decay, investors often roll before expiry. For example, an investor could buy a one-year option and roll it after six months to a new one-year option.

¹ Some option markets adjust for all dividends.

1.4: CALL OVERWRITING

For a directional investor who owns a stock (or index), call overwriting by selling an OTM call is one of the most popular methods of yield enhancement. Historically, call overwriting (otherwise known as buy-write, as the stock is bought but a call is written against it) has been a profitable strategy due to implied volatility usually being overpriced. However, call overwriting does underperform in volatile, strongly rising equity markets. Overwriting with the shortest maturity is best, and the strike should be slightly OTM for optimum returns.

IMPLIED VOLATILITY IS USUALLY OVERPRICED

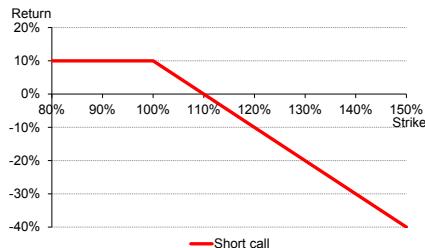
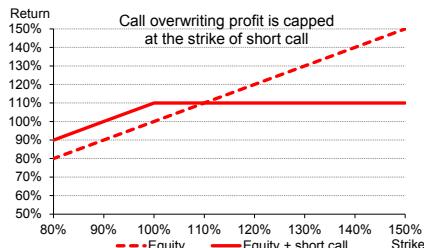
The implied volatility of options is on average 1-2pts above the volatility realised over the life of the option. This ‘implied volatility premium’ is usually greater for indices than for single stocks. As we can see no reason why these imbalances will fade, we expect call overwriting to continue to outperform on average. The key imbalances are:

- **Option buying for protection.** In the same way that no one buys car insurance because they think it is a good investment, investors are happy buying expensive protection to protect against downside risks.
- **Unwillingness to sell low premium options causes market makers to raise their prices** (selling low premium options, like selling lottery tickets, has to be done on a large scale to be attractive).
- **High gamma of near-dated options has a gap risk premium** (risk of stock jumping, either intraday or between closing and opening prices).
- **Index implieds lifted by structured products.** Structured products are often based on an index, and can offer downside protection. This lifts index implied relative to single stock implied. Also protection is usually bought on an index to protect against macros risks. It is rare to protect a single stock position (if an investor is worried about downside in a stock they usually do not buy it to begin with).

BUY-WRITE BENEFITS FROM SELLING EXPENSIVE VOL

Short-dated implied volatility has historically been overpriced² due to the above supply and demand imbalances. In order to profit from this characteristic, a long investor can sell a call against a long position in the underlying of the option. Should the underlying perform well and the call be exercised, the underlying can be used to satisfy the exercise of the call. As equities should be assumed to have, on average, a positive return, it is best to overwrite with a slightly OTM option to reduce the probability of the option sold expiring ITM.

² We note that implied volatility is not necessarily as overpriced as would first appear. For more detail, see the section *3.1 Implied Volatility Should Be Above Realized Volatility*.

Figure 5. Short Call**Call Overwriting (or Buy Write)**

Call overwriting is a useful way to gain yield in flat markets

If markets are range trading, or are approaching a technical resistance level, then selling a call at the top of the range (or resistance level) is a useful way of gaining yield. Such a strategy can be a useful tactical way of earning income on a core strategic portfolio, or potentially could be used as part of an exit strategy for a given target price.

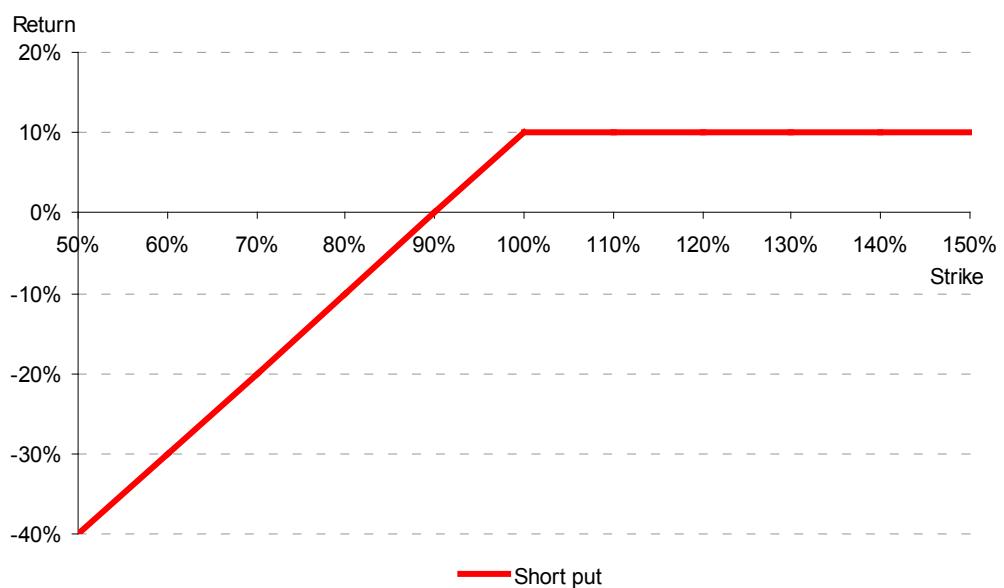
Selling at target price enforces disciplined investing

If a stock reaches the desired target price, there is the temptation to continue to own the strong performer. Over time a portfolio can run the risk of being a collection of stocks that had previously been undervalued, but are now at fair value. To prevent this inertia diluting the performance of a fund, some fund managers prefer to call overwrite at their target price to enforce disciplined investing, (as the stock will be called away when it reaches the target). As there are typically more Buy recommendations than Sell recommendations, call overwriting can ensure a better balance between the purchase and (called away) sale of stocks.

PUT UNDERWRITING HAS SIMILAR PROFILE

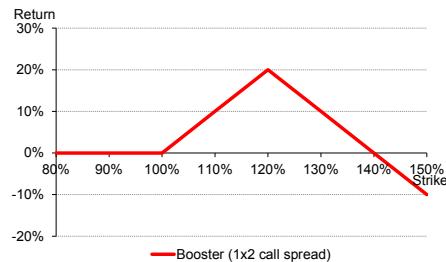
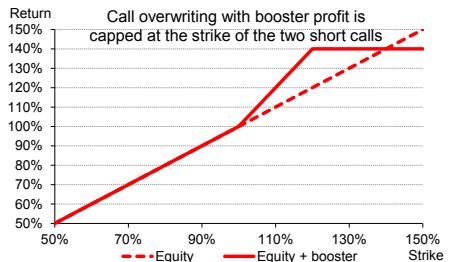
Figure 5 above shows the profiles of a short call and of a long equity with an overwritten call. The resulting profile of call overwriting is similar to that of a short put (Figure 6 below); hence, call overwriting could be considered similar to stock replacement with a short put (or put underwriting). Both call overwriting and put underwriting attempt to profit from the fact that implied volatility, on average, tends to be overpriced. While selling a naked put is seen as risky, due to the near infinite losses should stock prices fall, selling a call against a long equity position is seen as less risky (as the equity can be delivered against the exercise of the call).

Figure 6. Put Underwriting



1×2 call spreads are useful when a bounce-back is expected

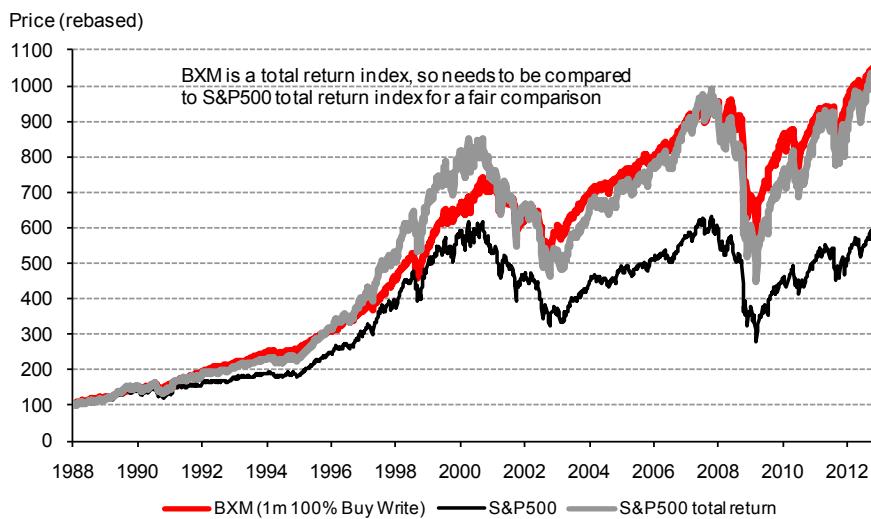
If a near zero cost 1×2 call spread (long 1×ATM call, short 2×OTM calls) is overlaid on a long stock position, the resulting position offers the investor twice the return for equity increases up to the short upper strike. For very high returns the payout is capped, in a similar way as for call overwriting. Such positioning is useful when there has been a sharp drop in the markets and a limited bounce back to earlier levels is anticipated. The level of the bounce back should be in line with or below the short upper strike. Typically, short maturities are best (less than three months) as the profile of a 1×2 call spread is similar to a short call for longer maturities.

Figure 7. Booster (1x2 Call Spread)**Call Overwriting with Booster**

CALL OVERWRITING IS BEST DONE ON AN INDEX

Many investors call overwrite on single stocks. However, single-stock implied volatility trades more in line with realised volatility than index implieds. The reason why index implieds are more overpriced than single-stock implieds is due to the demand from hedgers and structured product sellers. Call overwriting at the index level also reduces trading costs (due to the narrower bid-offer spread).

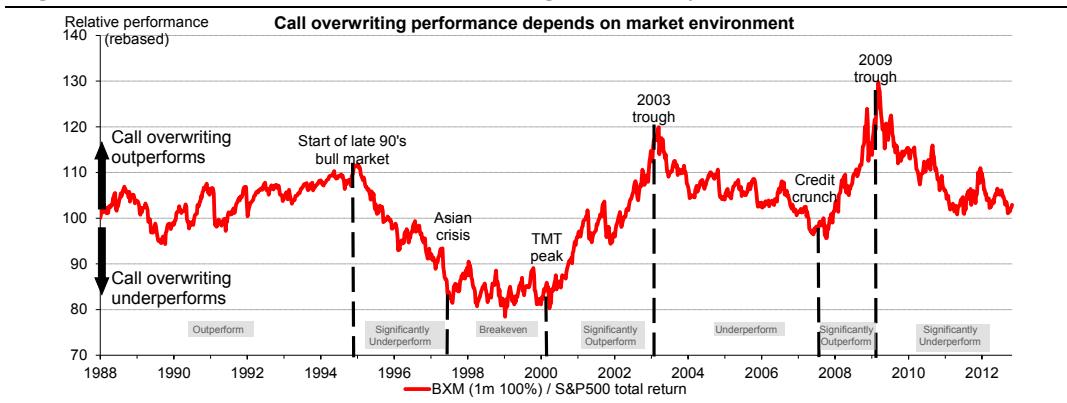
The CBOE has created a one-month call overwriting index on the S&P500 (BXM index), which is the longest call overwriting time series available. It is important to note that the BXM is a total return index; hence, it needs to be compared to the S&P500 total return index (SPXT Bloomberg code) not the S&P500 price return (SPX Bloomberg code). As can be seen in Figure 8 below, comparing the BXM index to the S&P500 price return index artificially flatters the performance of call overwriting.

Figure 8. S&P500 and S&P500 1M ATM Call Overwriting Index (BXM)

Call overwriting performance varies according to conditions

On average, ATM index call overwriting has been a profitable strategy. However, there have been periods of time when it is has been unprofitable. The best way to examine the returns under different market conditions is to divide the BXM index by the total return S&P500 index (as the BXM is a total return index).

Figure 9. S&P500 1M ATM Call Overwriting Divided by S&P500 Total Return



Overwriting underperforms in bull markets with low volatility

Since the BXM index was created, there have been seven distinct periods (see Figure 9 above), each with different equity and volatility market conditions. Of the seven periods, the two in which returns for call overwriting are negative are the bull markets of the mid-1990s and middle of the last decade. These were markets with very low volatility, causing the short call option sold to earn insufficient premium to compensate for the option being ITM.

It is important to note that call overwriting can outperform in slowly rising markets, as the premium earned is in excess of the amount the option ends up ITM. This was the case for the BXM between 1986 and the mid-1990s. It is difficult to identify these periods in advance as there is a very low correlation between BXM outperformance and the earlier historical volatility.

LOWER EXPOSURE TO EQUITY RISK PREMIUM

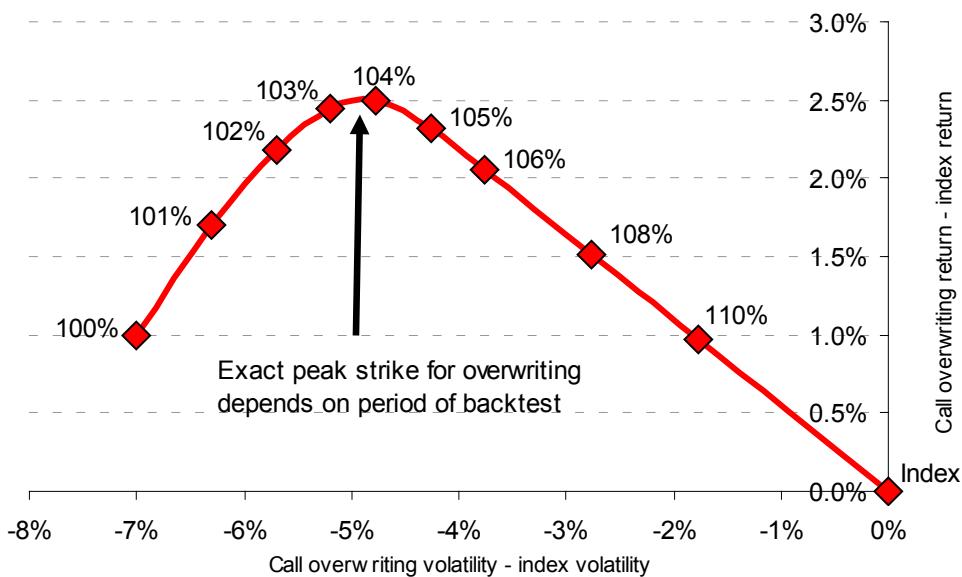
We note that while profits should be earned from selling an expensive call, the delta (or equity sensitivity) of the long underlying short call portfolio is significantly less than 100% (even if the premium from the short call is reinvested into the strategy). Assuming that equities are expected to earn more than the risk free rate (ie, have a positive equity risk premium), this lower delta can mean more money is lost by having a less equity-sensitive portfolio than is gained by selling expensive volatility. On average, call overwriting appears

to be a successful strategy, and its success has meant that it is one of the most popular uses of trading options.

OVERWRITING WITH NEAR-DATED OPTIONS IS BEST

Near-dated options have the highest theta, so an investor earns the greatest carry from call overwriting with short-dated options. It is possible to overwrite with 12 one-month options in a year, as opposed to four three-month options or one 12-month option. While overwriting with the shortest maturity possible has the highest returns on average, the strategy does have potentially higher risk. If a market rises one month, then retreats back to its original value by the end of the quarter, a one-month call overwriting strategy will have suffered a loss on the first call sold but a three-month overwriting strategy will not have had a call expire ITM. However, overwriting with far-dated expiries is more likely to eliminate the equity risk premium the investor is trying to earn (as any outperformance above a certain level will be called away).

Figure 10. Call Overwriting SX5E with One-Month Calls of Different Strikes



BEST RETURNS WITH SLIGHTLY OTM OPTIONS

While overwriting with near-dated expiries is clearly superior to overwriting with far-dated expiries, the optimal choice of strike to overwrite with depends on the market environment. As equities are expected, on average, to post a positive return, overwriting should be done with slightly OTM options. However, if a period of time where equities had a negative return is chosen for a back-test, then a strike below 100% could show the highest return. Looking

at a period of time where the SX5E had a positive return shows that for one-month options a strike between 103%-104% is best (see Figure 10 above).

Typically call overwriting with c25% delta call options is best

For three-month options, the optimal strike is a higher 107%-108%, but the outperformance is approximately half as good as for one-month options. These optimal strikes for overwriting could be seen to be arguably high, as recently there have been instances of severe declines (TMT bubble bursting, Lehman bankruptcy), which were followed by significant price rises afterwards. For single-stock call overwriting, these strikes could seem to be low, as single stocks are more volatile. For this reason, many investors use the current level of volatility to determine the strike or choose a fixed delta option (eg, 25%).

OVERWRITING REDUCES VOLATILITY

While selling an option could be considered risky, the volatility of returns from overwriting a long equity position is reduced by overwriting. This is because the payout profile is capped for equity prices above the strike. An alternative way of looking at this is that the delta of the portfolio is reduced from 100% (solely invested in equity) to 100% less the delta of the call (c50% depending on strike). The reduced delta suppresses the volatility of the portfolio.

Risk reduction less impressive if Sortino ratios are used

We note that the low call overwriting volatility is due to the lack of volatility to the upside, as call overwriting has the same downside risk as a long position. For this reason, using the Sortino ratio (for more details, see the section *A11 Sortino Ratio* in the *Appendix*) is likely to be a fairer measure of call overwriting risk than standard deviation, as standard deviation is not a good measure of risk for skewed distributions. Sortino ratios show that the call overwriting downside risk is identical to a long position; hence, call overwriting should primarily be done to enhance returns and is not a viable strategy for risk reduction.

Optimal strike is similar for single stocks and indices

While this analysis is focused on the SX5E, the analysis can be used to guide single-stock call overwriting (although the strike could be adjusted higher by the single-stock implied divided by SX5E implied).

ENHANCED CALL OVERWRITING IS DIFFICULT

Enhanced call overwriting is the term given when call overwriting is only done opportunistically or the parameters (strike or expiry) are varied according to market conditions. On the index level, the returns from call overwriting are so high that enhanced call overwriting is difficult, as the opportunity cost from not always overwriting is too high. For single stocks, the returns for call overwriting are less impressive; hence, enhanced call overwriting could be more successful. An example of single-stock enhanced call overwriting is to only overwrite when an implied is high compared to peers in the same sector. We note that even with enhanced single-stock call overwriting, the wider bid-offer cost and smaller implied volatility premium to realised means returns can be lower than call overwriting at the index level.

Enhanced call overwriting returns likely to be arbitraged away

Should a systematic way to enhance call overwriting be viable, this method could be applied to volatility trading without needing an existing long position in the underlying. Given the presence of statistical arbitrage funds and high frequency traders, we believe it is unlikely that a simple automated enhanced call overwriting strategy on equity or volatility markets is likely to outperform vanilla call overwriting on an ongoing basis.

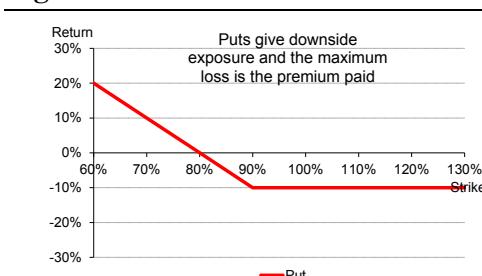
1.5: PROTECTION STRATEGIES USING OPTIONS

For both economic and regulatory reasons, one of the most popular uses of options is to provide protection against a long position in the underlying. The cost of buying protection through a put is lowest in calm, low-volatility markets, but in more turbulent markets the cost can be too high. In order to reduce the cost of buying protection in volatile markets (which is often when protection is in most demand), many investors sell an OTM put and/or an OTM call to lower the cost of the long put protection bought.

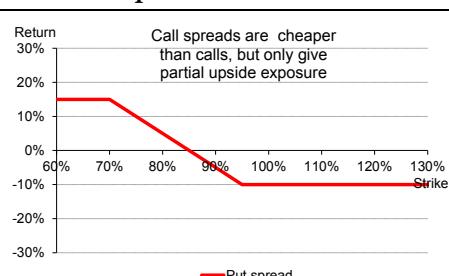
CHEAPEN PROTECTION BY SELLING OTM PUT & CALL

Buying a put against a long position gives complete and total protection for underlying moves below the strike (as the investor can simply put the long position back for the strike price following severe declines). The disadvantage of a put is the relatively high cost, as an investor is typically unwilling to pay more than 1%-2% for protection (as the cost of protection usually has to be made up through alpha to avoid underperforming if markets do not decline). The cost of the long put protection can be cheapened by selling an OTM put (turning the long put into a long put spread), by selling an OTM call (turning put protection into a collar), or both (resulting in a put spread vs call, or put spread collar). The strikes of the OTM puts and calls sold can be chosen to be in line with technical supports or resistance levels.

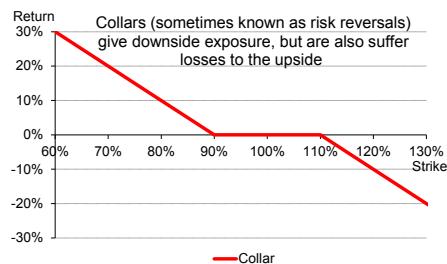
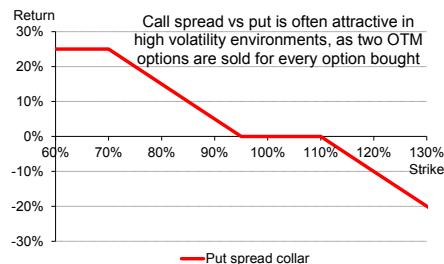
Figure 11. Put



Put spread



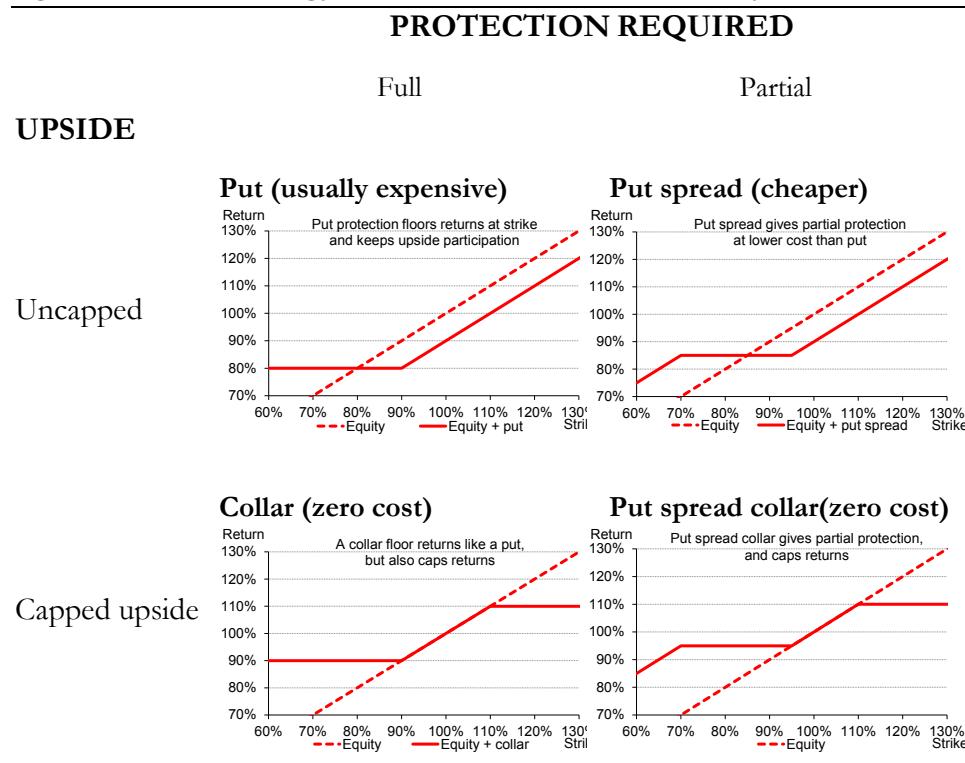
- **Puts give complete protection without capping performance.** As puts give such good protection, their cost is usually prohibitive unless the strike is low. For this reason, put protection is normally bought for strikes around 90%. Given that this protection will not kick in until there is a decline of 10% or more, puts offer the most cost-effective protection only during a severe crash (or if very short-term protection is required).
- **Put spreads only give partial protection but are cost effective.** While puts give complete protection, often only partial protection is necessary, in which case selling an OTM put against the long put (a put spread) can be an attractive protection strategy. The value of the put sold can be used to either cheapen the protection or lift the strike of the long put.
- **Collars can be zero cost as they give up some upside.** While investors appreciate the need for protection, the cost needs to be funded through reduced performance (or less alpha) or by giving up some upside. Selling an OTM call to fund a put (a collar) results in a cap on performance. However, if the strike of the call is set at a reasonable level, the capped return could still be attractive. The strike of the OTM call is often chosen to give the collar a zero cost. Collars can be a visually attractive low (or zero) cost method of protection as returns can be floored at the largest tolerable loss and capped at the target return. A collar is unique among protection strategies in not having significant volatility exposure, as its profile is similar to a short position in the underlying. Collars are, however, exposed to skew.
- **Put spread collars best when volatility is high, as two OTM options are sold.** Selling both an OTM put and OTM call against a long put (a put spread collar) is typically attractive when volatility is high, as this lifts the value of the two OTM options sold more than the long put bought. If equity markets are range bound, a put spread collar can also be an attractive form of protection. Put spread collars are normally structured to be near zero cost (just like a collar).

Figure 12. Collar**Put spread collar**

Portfolio protection is usually done via indices to lower cost

While an equity investor will typically purchase individual stocks, if protection is bought then this is usually done at the index level. This is because the risk the investor wishes to hedge against is the general equity or macroeconomic risk. If a stock is seen as having excessive downside risk, it is usually sold rather than a put bought against it. An additional reason why index protection is more common than single stock protection is the fact that bid-offer spreads for single stocks are wider than for an index.

Figure 13. Option Strategy for Different Market and Volatility Views



Partial protection can give attractive risk reward profile

For six-month maturity options, the cost of a 90% put is typically in line with a 95%-85% put spread (except during periods of high volatility, when the cost of a put is usually more expensive). Put spreads often have an attractive risk-reward profile for protection of the same cost, as the strike of the long put can be higher than the long put of a put spread. Additionally, if an investor is concerned with outperforming peers, then a 10% outperformance given by a 95%-85% put spread should be sufficient to attract investors (there is little incremental competitive advantage in a greater outperformance).

Implied volatility is far more important than skew for put-spread pricing

A rule of thumb is that the value of the OTM put sold should be approximately one-third the value of the long put (if it were significantly less, the cost saving in moving from a put to a put spread would not compensate for giving up complete protection). While selling an OTM put against a near-ATM put does benefit from selling skew (as the implied volatility of the OTM put sold is higher than the volatility of the near ATM long put bought), the effect of skew on put spread pricing is not normally that significant (far more significant is the level of implied volatility).

Collars are more sensitive to skew than implied volatility

Selling a call against a long put suffers from buying skew. The effect of skew is greater for a collar than for a put spread, as skew affects both legs of the structure the same way (whereas the effect of skew on the long and short put of a put spread partly cancels). If skew was flat, the cost of a collar typically reduces by 1% of spot. The level of volatility for near-zero cost collars is not normally significant, as the long volatility of the put cancels the short volatility of the call.

Capping performance should only be used when a long rally is unlikely

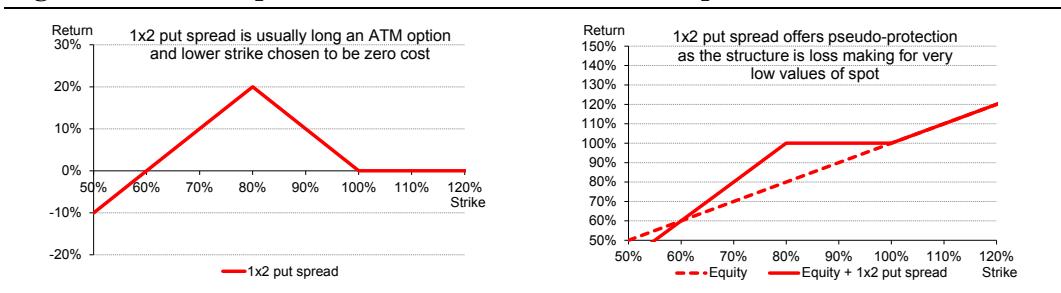
A collar or put spread collar caps the performance of the portfolio at the strike of the OTM call sold. They should only therefore be used when the likelihood of a strong, long-lasting rally (or significant bounce) is perceived to be relatively small.

Bullish investors could sell two puts against long put

If an investor is bullish on the equity market, then a protection strategy that caps performance is unsuitable. Additionally, as the likelihood of substantial declines is seen to be small, the cost of protection via a put or put spread is too high. In this scenario, a zero cost 1×2 put spread could be used as a pseudo-protection strategy. The long put is normally ATM, which means the portfolio is 100% protected against falls up to the lower strike, and gives partial protection below that until the breakeven. A loss is only suffered if the equity market falls below the breakeven.

1×2 put spreads only give pseudo-protection

We do not consider 1×2 put spreads to offer true protection, as during severe declines it will suffer a loss when the underlying portfolio is also heavily loss making. The payout of 1×2 put spreads for maturities of around three months or more is initially similar to a short put, so we consider it to be a bullish strategy. However, for the SX5E a roughly six-month zero-cost 1×2 put spread, whose upper strike is 95%, has historically had a breakeven below 80% and declines of more than 20% in six months are very rare. As 1×2 put spreads do not provide protection when you need it most, they could be seen as a separate long position rather than a protection strategy.

Figure 14. 1x2 Put Spread**1x2 Put Spread Pseudo-Protection**

PROTECTION MUST BE PAID FOR: QUESTION IS HOW?

If an investor seeks protection, the most important decision that has to be made is how to pay for it. The cost of protection can be paid for in one of three ways. Figure 15 below shows when this cost is suffered by the investor, and when the structure starts to provide protection against declines.

Premium. The simplest method of paying for protection is through premium. In this case, a put or put spread should be bought.

Loss of upside. If the likelihood of extremely high returns is small, or if a premium cannot be paid, then giving up upside via collars or put spread collars is the best way to pay for protection.

Potential losses on extreme downside. If an investor is willing to tolerate additional losses during extreme declines, then a 1x2 put spread can offer a zero cost way of buying protection against limited declines in the market.

Figure 15. Protection Strategy Comparison

Equity Performance	Put	Put Spread	Collar	Put Spread Collar	1x2 Put Spread
Bull markets (+10% or more)	Loss of premium	Loss of premium	Loss of upside	Loss of upside	-
Flat markets (±5%)	Loss of premium	Loss of premium	-	-	-
Moderate dip (c-10%)	Loss of premium	Protected	-	Protected	Protected
Correction (c-15%)	Protected	Protected	Protected	Protected	Protected
Bear market (c-20% or worse)	Protected	Partially protected	Protected	Partially protected	Severe loss

BEST STRATEGY DETERMINED BY VOLATILITY LEVEL

The level of volatility can determine the most suitable protection strategy an investor needs to decide how bullish and bearish they are on the equity and volatility markets. If volatility is low, then puts should be affordable enough to buy without offsetting the cost by selling an OTM option. For low to moderate levels of volatility, a put spread is likely to give the best protection that can be easily afforded. As a collar is similar to a short position with limited volatility exposure, it is most appropriate for a bearish investor during average periods of volatility (or if an investor does not have a strong view on volatility). Put spreads collars (or 1×2 put spreads) are most appropriate during high levels of volatility (as two options are sold for every option bought).

MATURITY DRIVEN BY DURATION OF LIKELY DECLINE

The choice of protection strategy is typically driven by an investor's view on equity and volatility markets. Similarly the choice of strikes is usually restricted by the premium an investor can afford. Maturity is potentially the area where there is most choice, and the final decision will be driven by an investor's belief in the severity and duration of any decline. If he wants protection against a sudden crash, a short-dated put is the most appropriate strategy. However, for a long drawn out bear market, a longer maturity is most appropriate.

Figure 16. Types of DAX Declines (of 10% or more) since 1960

Type	Average Decline	Decline Range	Average Duration	Duration Range
Crash	31%	19% to 39%	1 month	0 to 3 months
Correction	14%	10% to 22%	3 months	0 to 1 year
Bear market	44%	23% to 73%	2.5 years	1 to 5 years

Median maturity of protection bought is c4 months

The average choice of protection is c6 months, but this is skewed by a few long-dated hedges. The median maturity is c4 months. Protection can be bought for maturities of one week to over a year. Even if an investor has decided how long he needs protection, he can implement it via one far-dated option or multiple near-dated options. For example, one-year protection could be via a one-year put or via the purchase of a three-month put every three months (four puts over the course of a year). The typical cost of ATM puts for different maturities is given below.

Figure 17. Cost of ATM Put on SX5E

Cost	1 Month	2 Months	3 Months	6 Months	1 Year
Individual premium	2.3%	3.3%	4.0%	5.7%	8.0%
Rolling protection cost per year	27.7%	19.6%	16.0%	11.3%	8.0%

Short-dated puts offer greatest protection but highest cost

If equity markets fall 20% in the first three months of the year and recover to the earlier level by the end of the year, then a rolling three-month put strategy will have a positive payout in the first quarter but a one-year put will be worth nothing at expiry. While rolling near-dated puts will give greater protection than a long-dated put, the cost is higher (see Figure 17 above).

SHORT VOL AGAINST LONG PUT PERFORMS WELL

All protection strategies that combine a long and a short aim to offset the overpriced cost of protection by selling the same overpriced implied volatility for a different maturity and strike. Hence such strategies tend to back-test well as the overall exposure to expensive implied volatility is near zero. As the net cost of such strategies is near zero, while at the same time (usually) decreasing the volatility of the portfolio, their risk adjusted performance can be impressive. Their performance can often be further improved by selling near dated volatility against the long far dated protection.

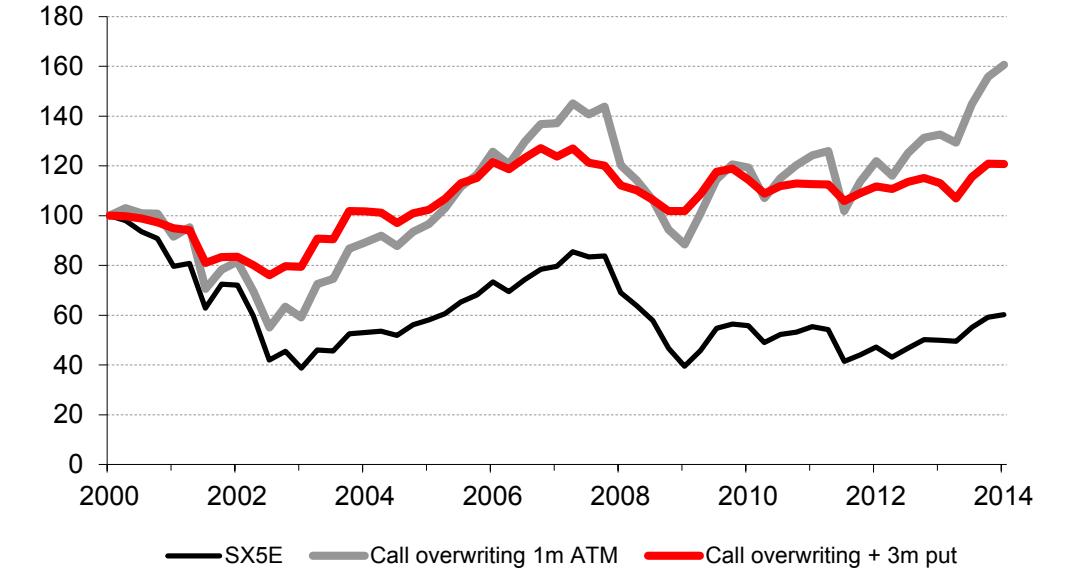
MULTIPLE EXPIRY PROTECTION STRATEGIES

Typically, a protection strategy involving multiple options has the same maturity for all of the options. However, some investors choose a nearer maturity for the options they are short, as more premium can be earned selling a near-dated option multiple times (as near-dated options have higher theta). These strategies are most successful when term structure is inverted, as the volatility for the near-dated option sold is higher. Having a nearer maturity for the long put option and longer maturity for the short options makes less sense, as this increases the cost (assuming the nearer-dated put is rolled at expiry).

Calendar collar effectively overlays call overwriting on a long put position

If the maturity of the short call of a collar is closer than the maturity of the long put, then this is effectively the combination of a long put and call overwriting. For example, the cost of a three-month put can be recovered by selling one-month calls. This strategy outperforms in a downturn and also has a lower volatility (see Figure 18 below).

Figure 18. Performance of 3M Put vs 1M Call Overwriting



Selling a put against calendar collar creates either a calendar put spread collar, or a put vs strangle/straddle

If a put is sold against the position of a calendar collar, then the final position depends on the maturity of the short put. The short put can either have the maturity of the long put (creating a calendar put spread collar), or the maturity of the short call (creating a put vs strangle/straddle).

- **Calendar put spread collar.** If the maturity of the short put is identical to the long (far-dated) put, then the final position is a calendar put spread collar (i.e. far dated put spread funded by sale of short dated calls). The performance of a calendar put spread collar is similar to the calendar collar above.
- **Put vs strangle/straddle.** If the maturity of the short put is the same as the maturity of the short near-dated call, then this position funds the long far-dated put by selling near-dated volatility via a near dated strangle (or straddle if the strikes of the short put and short call are identical).

Short near dated variance swaps is an alternative to selling near dated strangle/straddle

For an investor who is able to trade OTC, a similar strategy involves long put and short near-dated variance swaps.

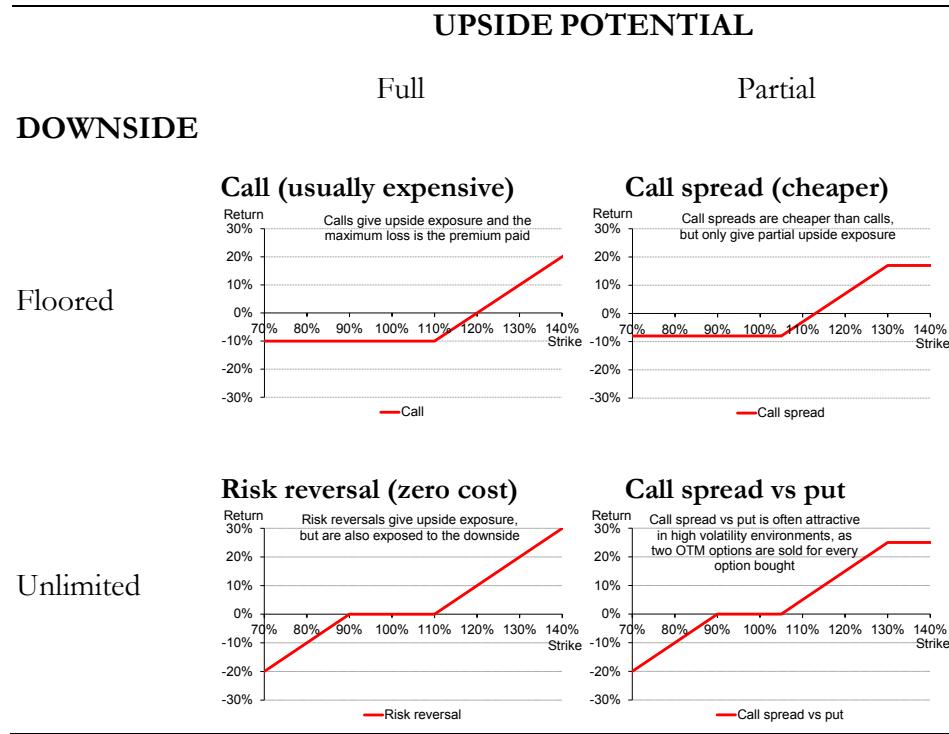
1.6: OPTION STRUCTURES TRADING

While a simple view on both volatility and equity market direction can be implemented via a long or short position in a call or put, a far wider set of payoffs is possible if two or three different options are used. We investigate strategies using option structures (or option combos) that can be used to meet different investor needs.

BULLISH COMBOS ARE REVERSE OF BEARISH

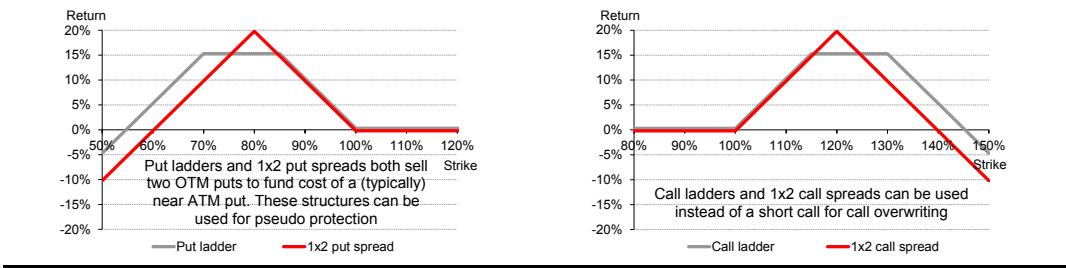
Using option structures to implement a bearish strategy has already been discussed in the section [1.5 Protection Strategies Using Options](#). In the same way a long put protection can be cheapened by selling an OTM put against the put protection (to create a put spread giving only partial protection), a call can be cheapened by selling an OTM call (to create a call spread offering only partial upside). Similarly, the upside exposure of the call (or call spread) can be funded by put underwriting (just as put or put spread protection can be funded by call overwriting). The four option structures for bullish strategies are given below.

- **Calls give complete upside exposure and floored downside.** Calls are the ideal instrument for bullish investors as they offer full upside exposure and the maximum loss is only the premium paid. Unless the call is short dated or is purchased in a period of low volatility, the cost is likely to be high.
- **Call spreads give partial upside but are cheaper.** If an underlying is seen as unlikely to rise significantly, or if a call is too expensive, then selling an OTM call against the long call (to create a call spread) could be the best bullish strategy. The strike of the call sold could be chosen to be in line with a target price or technical resistance level. While the upside is limited to the difference between the two strikes, the cost of the strategy is normally one-third cheaper than the cost of the call.
- **Risk reversals (short put, long call of different strikes) benefit from selling skew.** If a long call position is funded by selling a put (to create a risk reversal), the volatility of the put sold is normally higher than the volatility of the call bought. The higher skew is, the larger this difference and the more attractive this strategy is. Similarly, if interest rates are low, then the lower the forward (which lifts the value of the put and decreases the value of the call) and the more attractive the strategy is. The profile of this risk reversal is similar to being long the underlying.
- **Call spread vs put is most attractive when volatility is high.** A long call can be funded by selling an OTM call and OTM put. This strategy is best when implied volatility is high, as two options are sold.

Figure 19. Upside Participation Strategies

LADDERS HAVE A SIMILAR PROFILE TO 1x2 SPREADS

With a 1x2 call or put spread, two OTM options of the same strike are sold against one (usually near ATM) long option of a different strike. A ladder has exactly the same structure, except the two short OTM options have a different strike. The different strikes of the two OTM options means that the maximum payout of a ladder is lower than for a 1x2, however this maximum payout is earned for a range of values, not a single value at expiry. Additionally the breakeven of a ladder is further away from spot, giving the comfort that markets have to move further for losses to be suffered when compared to a 1x2 call or put spread.

Figure 20. Put Ladders and 1x2**Call Ladders and 1x2**

STRADDLES AND STRANGLES ARE VERY SIMILAR

Using option structures allows a straddle (long call and put of same strike) or strangle (long call and put of different strikes) to be traded. These structures are long (or short) volatility, but do not have any exposure to the direction of the equity market.

STRUCTURES ALLOW WIDER RANGE OF PAYOUTS

Figure 21 below shows the most common structures that can be traded with up to three different options in relation to a view on equity and volatility markets. For simplicity, strangles and ladders are not shown, but they can be considered to be similar to straddles and 1x2 ratio spreads, respectively.

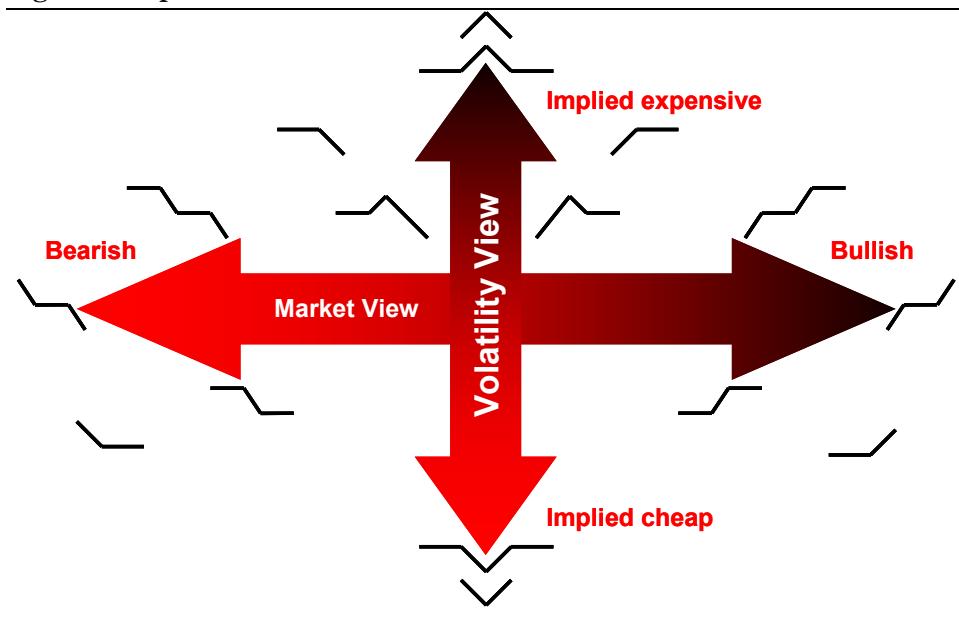
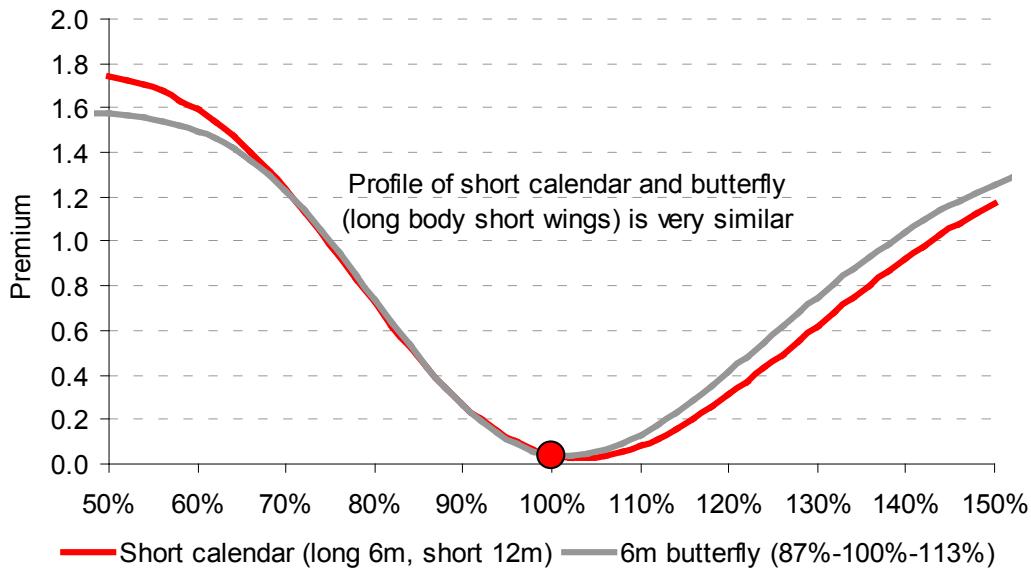
Figure 21. Option Structures

Figure 22. Theoretical Value of Butterfly and Short Calendar



1X1 CALENDAR TRADES ARE SIMILAR TO BUTTERFLY

Butterflies combine a short straddle with a long strangle, which allows a short volatility position to be taken with floored losses. We note the theoretical profile of a short calendar trade is similar to trading a butterfly. If an underlying does not have liquid OTM options, then a calendar can be used as a butterfly substitute (although this approach does involve term structure risk, which a butterfly does not have). A long calendar (short near-dated, long far-dated) is therefore short gamma (as the short near-dated option has more gamma than the far-dated option).

CHAPTER 2

VOLATILITY TRADING

We investigate the benefits and disadvantages of volatility trading via options, volatility swaps, variance swaps and gamma swaps. We also show how these products, correlation swaps, basket options and covariance swaps can give correlation exposure. Recently, options on alternative underlyings have been created, such as options on variance and options on volatility futures. We show how the distribution and skew for options on variance is different from those for equities (options on volatility futures have many similarities to options on variance but are slightly different, and are covered in the section [4.5: Options On Volatility Futures](#)).

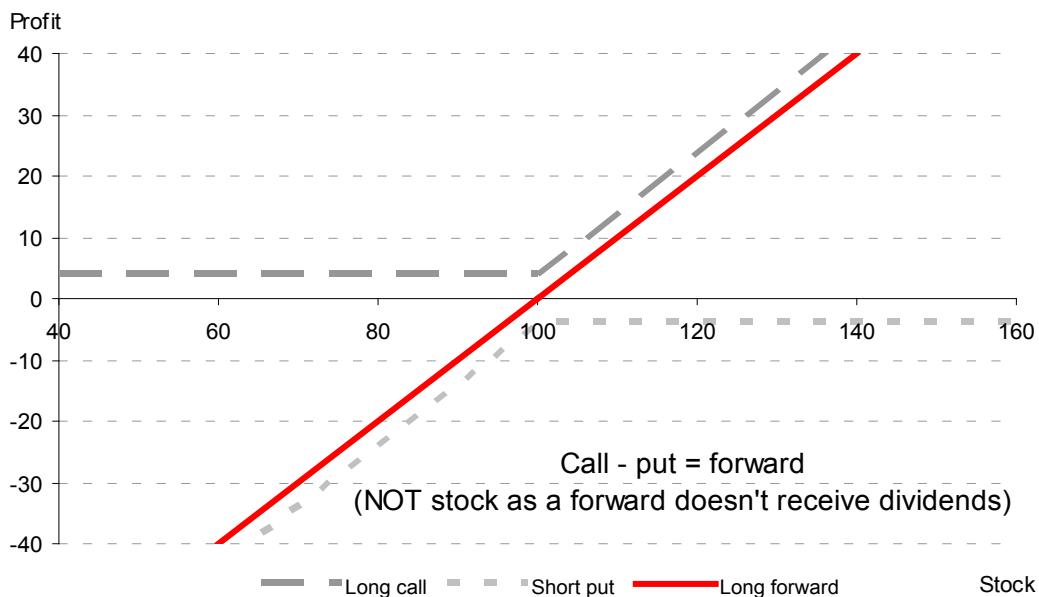
2.1: VOLATILITY TRADING USING OPTIONS

While directional investors typically use options for their equity exposure, volatility investors delta hedge their equity exposure. A delta-hedged option (call or put) is not exposed to equity markets, but only to volatility markets. We demonstrate how volatility investors are exposed to dividend and borrow cost risk and how volatility traders can ‘pin’ a stock approaching expiry. We also show that while the profit from delta hedging is based on percentage move squared (ie, variance or volatility²), it is the absolute difference between realised and implied that determines carry (not the difference between realised² and implied²).

VOLATILITY TRADING VIA CALLS & PUTS IS IDENTICAL

A forward is a contract that obliges the investor to buy (or sell if you have sold the forward) a security on a certain expiry date (but not before) at a certain strike price. A portfolio of a long European call and a short European put of identical expiry and strike is the same as a forward of that expiry and strike, as shown in Figure 23 below. This means that if a call, a put or a straddle is delta hedged with a forward contract (not stock), the end profile is identical.

Figure 23. Put-Call Parity: Call - Put = Long Forward (not long stock)



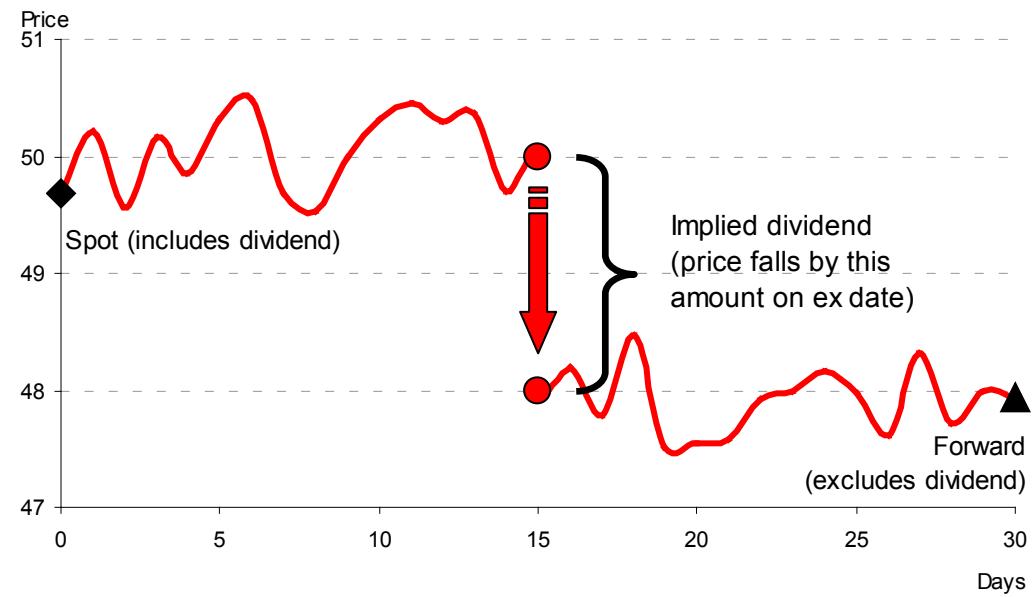
PUT-CALL PARITY IS ONLY TRUE FOR EUROPEAN OPTIONS

We note underlying is only true for European options as European options cannot be exercised before maturity. It is only approximately true for American options, as American options can be exercised before expiry (although in practice they seldom are).

Delta hedging must be done with forward of identical maturity

It is important to note that the delta hedging must be done with a forward of identical maturity to the options. If it is done with a different maturity, or with stock, there will be dividend risk. This is because a forward, like a European call or put, gives the right to a security at maturity but does not give the right to any benefits such as dividends that have an ex date before expiry. A long forward position is therefore equal to long stock and short dividends that go ex before maturity (assuming interest rates and borrow cost are zero or are hedged). This can be seen from the diagram below, as a stock will fall by the value of the dividend (subject to a suitable tax rate) on the ex date. The dividend risk of an option is therefore equal to the delta.

Figure 24. Why Forwards Are Short Dividends



BORROW COST IMPACT ON OPTION PRICING

From a derivative pricing point of view, borrow cost (or repo) can be added to the dividend. This is because it is something that the owner of the shares receives and the owner of a forward does not. While the borrow cost should, in theory, apply to both the bid and offer of calls and puts, in practice an investment bank's stock borrow desk is usually separate from the volatility trading desk (or potentially not all of the long position can be lent out). If the traders on the volatility trading desk do not get an internal transfer of the borrow cost, then only one side of the trade (the side that has positive delta for the volatility trading desk, or negative delta for the client) usually includes the borrow cost. This is shown in Figure 25 below. While the borrow cost is not normally more than 40bp for General Collateral (GC) names, it can be more substantial for emerging market (EM) names. If borrow cost is only included in one leg of pricing, it creates a bid-offer arbitrage channel.

Figure 25. When Borrow Cost Is Usually Included in Calculations

Option	Bid	Ask
Calls	Include borrow	–
Puts	–	Include borrow

Zero delta straddles still need to include borrow cost on one leg

Like dividends, the exposure to borrow cost is equal to the delta. However, a zero delta straddle still has exposure to borrow cost because it should be priced as the sum of two separate trades, one call and one put. As one of the legs of the trade should include borrow, so does a straddle. This is particularly important for EM or other high borrow cost names.

Zero delta straddles have strike above spot

A common misperception is that ATM options have a 50% delta; hence, an ATM straddle has to be zero delta. In fact, a zero delta straddle has to have a strike above spot (an ATM straddle has negative delta). The strike of a zero delta straddle is given below.

$$\text{Strike (\%)} \text{ of zero delta straddle} = e^{(r+\sigma^2/2)T}$$

where r = interest rate, σ = volatility and T = time.

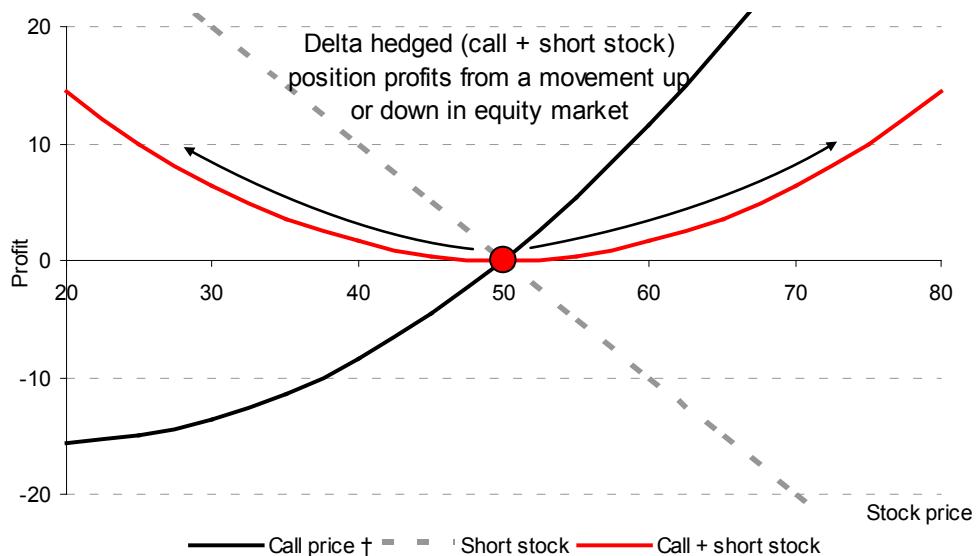
DELTA HEDGING AN OPTION REMOVES EQUITY RISK

If an option is purchased at an implied volatility that is lower than the realised volatility over the life of the option, then the investor, in theory, earns a profit from buying cheap volatility. However, the effect of buying cheap volatility is dwarfed by the profit or loss from the direction of the equity market. For this reason, directional investors are usually more concerned with premium rather than implied volatility. Volatility investors will, however, hedge the equity exposure. This will result in a position whose profitability is solely determined by the volatility (not direction) of the underlying. As delta measures the equity sensitivity of an option, removing equity exposure is called delta hedging (as a portfolio with no equity exposure has delta = 0).

Delta hedging example

As the delta of a portfolio is equal to the sum of the deltas of the securities in the portfolio, a position can be delta hedged by purchasing, or going short, a number of shares (or futures in the case of an index) equal to the delta. For example, if ten call options have been bought with a delta of 40%, then four shares ($10 \times 40\% = 4$) have to be shorted to create a portfolio of zero delta. The shares have to be shorted as a call option has positive delta; hence, the delta hedge has to be negative for the sum of the two positions to have zero delta. If we were long a put (which has negative delta), then we would have to buy stock to ensure the overall delta was zero.

Figure 26. Delta-Hedged Call



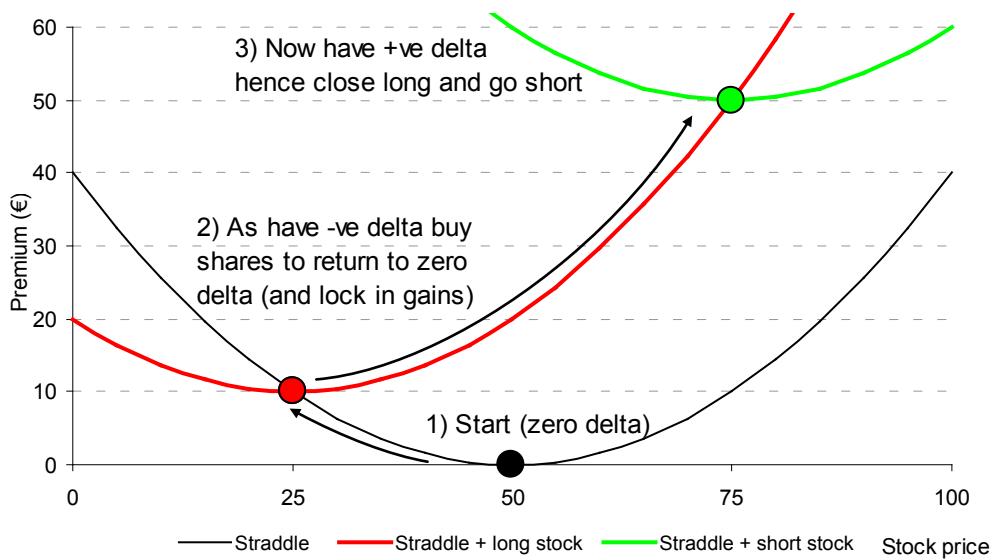
Constant delta hedging is called gamma scalping

The rate delta changes as spot moves is called gamma; hence, gamma is the convexity of the payout. As the delta changes, a volatility investor has to delta hedge in order to ensure there is no equity exposure. Constantly delta hedging in this way is called gamma scalping, as it ensures a long volatility position earns a profit as spot moves.

Gamma scalping (delta re-hedging) locks in profit as underlying moves

We shall assume an investor has purchased a zero delta straddle (or strangle), but the argument will hold for long call or put positions as well. If equity markets fall (from position 1 to position 2 in the chart) the position will become profitable and the delta will decrease from zero to a negative value. In order to lock in the profit, the investor must buy stock (or futures) for the portfolio to return to zero delta. Now that the portfolio is equity market neutral, it will profit from a movement up or down in the equity market. If equity markets then rise, the initial profit will be kept and a further profit earned (movement from position 2 to position 3). At position 3 the stock (or futures) position is sold and a short position initiated to return the position to zero delta.

Figure 27. Locking in Gains through Delta Hedging



Long gamma position can sit on the bid and offer

As shown above, a long gamma (long volatility) position has to buy shares if they fall, and sell them if they rise. Buying low and selling high earns the investor a profit. Additionally, as a gamma scalper can enter bids and offers away from current spot, there is no need to cross the spread (as a long gamma position can be delta hedged by sitting on the bid and offer).

Short gamma position have to cross the bid-offer spread

A short gamma position represents the reverse situation, and requires crossing the spread to delta hedge. While this hidden cost is small, it could be substantial over the long term for underlyings with relatively wide bid-offer spreads.

Best to delta hedge on key dates or on turn of market

If markets have a clear direction (ie, they are trending), it is best to delta hedge less frequently. However, in choppy markets that are range bound it is best to delta hedge very frequently. For more detail on how hedging frequency affects returns and the path dependency of returns, see the section [3.5 Stretching Black-Scholes Assumptions](#). If there is a key announcement (either economic or earnings-related to affect the underlying), it is best to delta hedge just before the announcement to ensure that profit is earned from any jump (up or down) that occurs.

OPTION TRADING RULES OF THUMB

To calculate option premiums and volatility exactly is typically too difficult to do without the aid of a tool. However, there are some useful rules of thumb that can be used to give an estimate. These are a useful sanity check in case an input to a pricing model has been entered incorrectly.

- **Historical volatility roughly equal to $16 \times$ percentage daily move.** Historical volatility can be estimated by multiplying the typical return over a period by the square root of the number of periods in a year (eg, 52 weeks or 12 months in a year). Hence, if a security moves 1% a day, it has an annualised volatility of 16% (as there are c252 trading days and $16 \approx \sqrt{252}$).
- **ATM option premium in percent is roughly $0.4 \times$ volatility \times square root of time.** If one assumes zero interest rates and dividends, then the formula for the premium of an ATM call or put option simplifies to $0.4 \times \sigma \times \sqrt{T}$. Therefore, a one-year ATM option on an underlying with 20% implied is worth c8% ($= 0.4 \times 20\% \times \sqrt{1}$). OTM options can be calculated from this estimate using an estimated 50% delta.
- **Profit from delta hedging is proportional to square of return.** Due to the convexity of an option, if the volatility is doubled the profits from delta hedging are multiplied by a factor of four. For this reason, variance (squared returns) is a better measure of deviation than volatility.

- **Difference between implied and realised determines carry.** While variance is the driver of profits if the implied volatility of an option is constant, the carry is determined by the absolute difference between realised and implied (ie, the same carry is earned by going long a 20% implied option that realises 21% as by going long a 40% implied option that realises 41%).

ANNUALISED VOL = 16 × PERCENTAGE DAILY MOVE

Volatility is defined as the annualised standard deviation of log returns (where return = P_i / P_{i-1}). As returns are normally close to 1 (=100%) the log of returns is very similar to ‘return – 1’ (which is the percentage change of the price). Hence, to calculate the annualised volatility for a given percentage move, all that is needed is to annualise the percentage change in the price. This is done by multiplying the percentage move by the square root of the number of samples in a year (as volatility is the square root of variance). For example, market convention is to assume there are approximately 252 trading days a year. If a stock moves 1% a day, then its volatility is $1\% \times \sqrt{252}$, which is approximately $1\% \times 16 = 16\%$ volatility. Similarly, if a stock moves 2% a day it has 32% volatility.

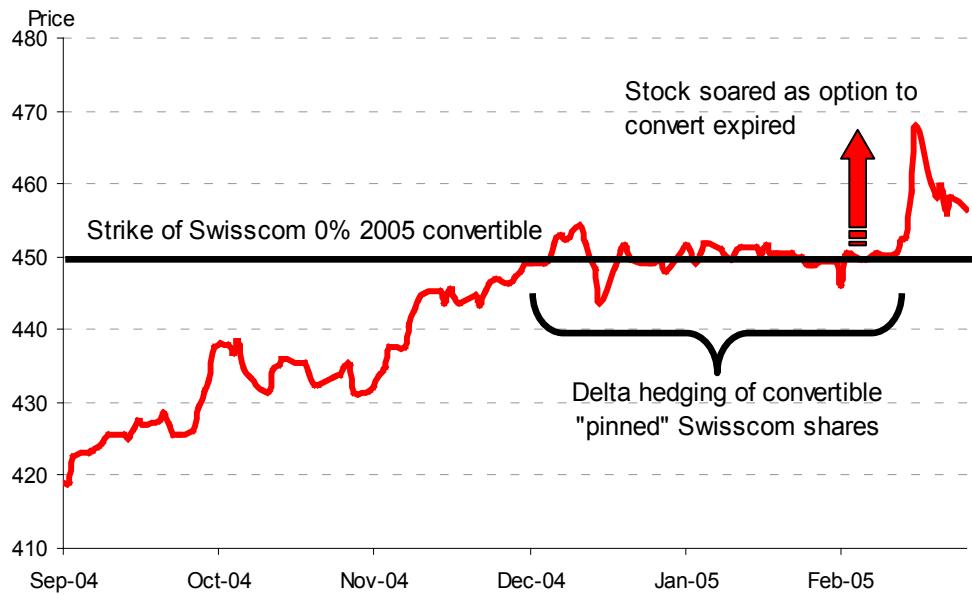
$$\text{Num of trading days in year} = 252 \Rightarrow \text{Multiply daily returns by } \sqrt{252} \approx 16$$

$$\text{Num of weeks in year} = 52 \Rightarrow \text{Multiply weekly returns by } \sqrt{52} \approx 7$$

$$\text{Num of months in year} = 12 \Rightarrow \text{Multiply monthly returns by } \sqrt{12} \approx 3.5$$

HEDGING CAN ‘PIN’ A STOCK APPROACHING EXPIRY

As an investor who is long gamma can delta hedge by sitting on the bid and offer, this trade can pin an underlying to the strike. This is a side effect of selling if the stock rises above the strike, and buying if the stock falls below the strike. The amount of buying and selling has to be significant compared with the traded volume of the underlying, which is why pinning normally occurs for relatively illiquid stocks or where the position is particularly sizeable. Given the high trading volume of indices, it is difficult to pin a major index. Pinning is more likely to occur in relatively calm markets, where there is no strong trend to drive the stock away from its pin.

Figure 28. Pinning of Swisscom Stock Approaching Convertible Expiry

Large size of Swisscom convertible pinned underlying for many months

One of the most visible examples of pinning occurred in late 2004/early 2005, due to a large Swiss government debt issue, (Swisscom 0% 2005) convertible into the relatively illiquid Swisscom shares. As the shares traded close to the strike approaching maturity, the upward trend of the stock was broken. Swisscom was pinned for two to three months until the exchangeable expired.

Pinned stocks snap back to fair value after expiration

After expiration, the stock snapped back to where it would have been if the upward trend had not been paused. A similar event occurred to AXA in the month preceding the Jun05 expiry, when it was pinned close to €20 despite the broader market rising (after expiry AXA rose 4% in four days).

ATM OPTION PREMIUM (%) = $0.4 \times \text{VOLATILITY} \times \sqrt{\text{TIME}}$

$$\text{Call price} = S N(d_1) - K N(d_2) e^{-rT}$$

Assuming zero interest rates and dividends ($r = 0$)

- \Rightarrow ATM call price = $S N(\sigma \times \sqrt{T} / 2) - S N(-\sigma \times \sqrt{T} / 2)$ as $K=S$ (as ATM)
- \Rightarrow ATM call price = $S \times \sigma \times \sqrt{T} / \sqrt{2\pi}$
- \Rightarrow ATM call price = $\sigma \times \sqrt{T} / \sqrt{2\pi}$ in percent
- \Rightarrow ATM call price $\approx 0.4 \times \sigma \times \sqrt{T}$ in percent

where:

Definition of d_1 and d_2 is the standard Black-Scholes formula.

σ = implied volatility

S = spot

K = strike

R = interest rate

T = time to expiry

$N(z)$ = cumulative normal distribution

Example 1

A 1 year ATM option on an underlying with 20% implied is worth c.8% ($=0.4 \times 20\% \times \sqrt{1}$)

Example 2

A 3 month ATM option on an underlying with 20% implied is worth c.4% ($=0.4 \times 20\% \times \sqrt{0.25} = 0.4 \times 20\% \times 0.5$)

OTM options can be calculated by assuming 50% delta

If an index is 3000pts and has a 20% implied then the price of a 1Y ATM option is approximately 240pts ($3000 \times 8\%$ as calculated in Example 1 above). A 3200 call is therefore approximately $240 - 50\% (3200-3000) = 240 - 100 = 140$ pts assuming a 50% delta. Similarly, a 3200 put is approximately 340pts ($240 + 100$).

PROFIT PROPORTIONAL TO PERCENT MOVE SQUARED

Due to the convexity of an option, if the volatility is doubled, the profits from delta hedging are multiplied by a factor of four. For this reason, variance (which looks at squared returns) is a better measure of deviation than volatility. Assuming constant volatility, zero interest rates and dividend, the daily profit and loss (P&L) from delta hedging an option is given below.

$$\text{Daily P\&L} = \text{Delta P\&L} + \text{Gamma P\&L} + \text{Theta P\&L}$$

$$\Rightarrow \text{Daily P\&L} = S\delta + S^2\gamma / 2 + t\theta$$

where S is change in Stock

$$\Rightarrow \text{Daily P\&L} - S\delta = + S^2\gamma / 2 + t\theta = \text{Delta hedged P\&L}$$

$$\Rightarrow \text{Delta hedged P\&L} = S^2\gamma / 2 + \text{cost term}$$

($t\theta$ does not depend on stock price)

where:

δ = delta

γ = gamma

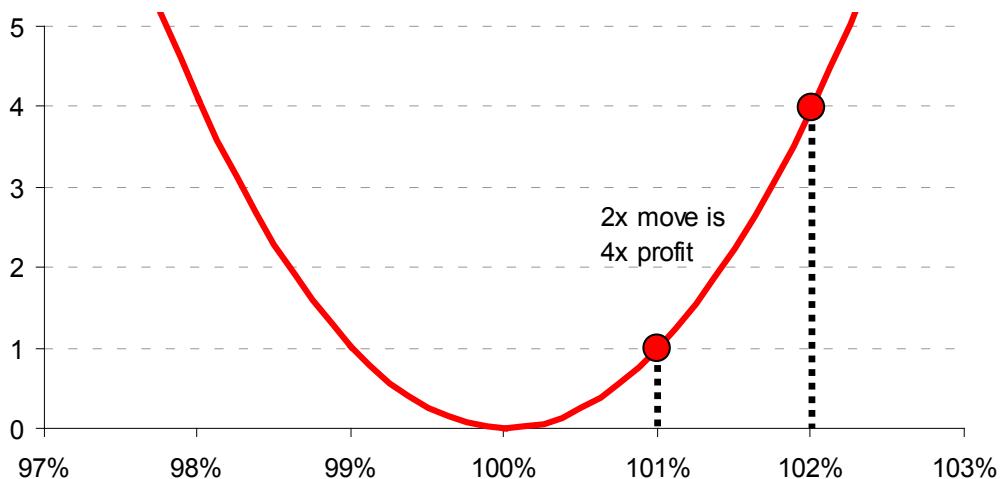
t = time

θ = theta

If the effect of theta is ignored (as it is a cost that does not depend on the size of the stock price movement), the profit of a delta hedged option position is equal to a scaling factor ($\gamma/2$) multiplied by the square of the return. This means that the profit from a 2% move in a stock price is four times ($2^2=4$) the profit from a 1% move in stock price.

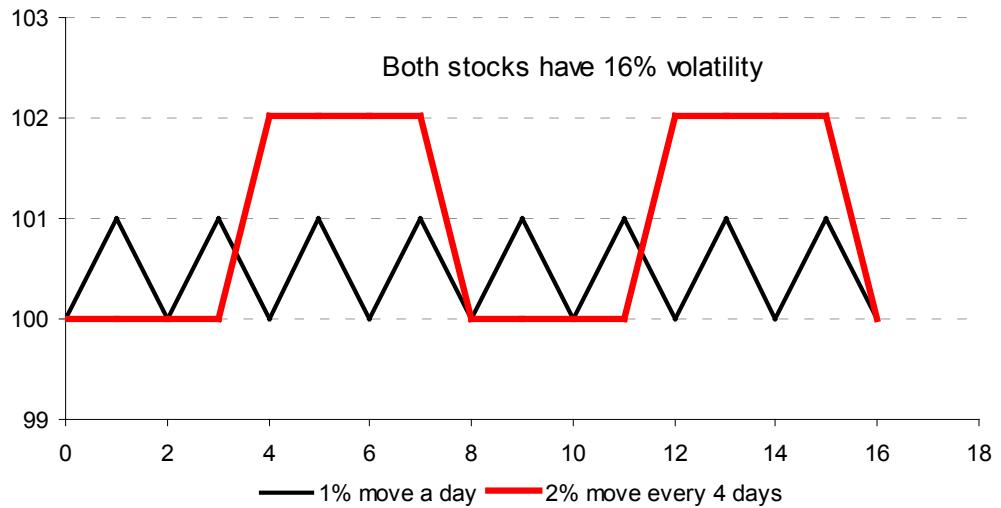
This can also be seen from Figure 29 below, as the additional profit from the move from 1% to 2% is three times the profit from 0% to 1% (for a total profit four times the profit for a 1% move).

Figure 29. Profile of a Delta-Hedged Option



Make same delta hedge profit with 1% a day move as 2% every four days

Let's assume there are two stocks: one of them moves 1% a day and the other 2% every four days (see Figure 30). Both stocks have the same 16% volatility and delta hedging them earns the same profit (as four times as much profit is earned on the days the stock moves 2% as when it moves 1%).

Figure 30. Two Stocks with the Same Volatility

IMPLIED - REALISED = CARRY

Assuming all other inputs are constant, the payout of a delta-hedged option is based on the variance (return squared). However, when examining how much carry is earned, or lost, when delta hedging an option, it is the difference between realised and implied (not realised² - implied²) that matters. This is because the gamma of an option is proportional to $1/\sigma$; hence, if volatility doubles the gamma halves. Thus, the same carry (profit from gamma less cost of theta) is earned by going long a 20% implied option that realises 21% as by going long a 40% implied option that realises 41%. The proof of this is below.

$$\text{Delta hedged P\&L} = \text{Dollar gamma} \times (\text{return}^2 - \sigma^2 dt)$$

where:

$$\sigma = \text{implied volatility}$$

$$\gamma = -N'(d_1) / (S \times \sigma \times \sqrt{T})$$

$$\text{Dollar gamma} = 0.5 \times \gamma \times S^2 \approx \text{constant} / \sigma \text{ for constant spot } S \text{ and time } T$$

$$\Rightarrow \text{Daily P\&L} \approx \text{constant} \times (\text{return}^2 - \sigma^2 dt) / \sigma$$

If we define return to be similar to volatility, then $\text{return} = (\sigma + x)dt$ where x is small

$$\Rightarrow \text{Daily P\&L} \approx \text{constant} \times dt \times ((\sigma + x)^2 - \sigma^2) / \sigma$$

- ⇒ Daily P&L \approx constant \times dt \times $((\sigma^2 + 2\sigma x + x^2) - \sigma^2) / \sigma$
- ⇒ Daily P&L \approx constant \times dt \times $(2x + x^2/\sigma)$
- ⇒ Daily P&L \approx constant \times dt \times $2x$ as x is small

Daily P&L proportional to x , where $x = \text{realised volatility} - \sigma$

Hence, it is the difference between realised and implied volatility that is the key to daily P&L (or carry).

2.2: VARIANCE IS THE KEY, NOT VOLATILITY

Partly due to its use in Black-Scholes, historically, volatility has been used as the measure of deviation for financial assets. However, the correct measure of deviation is variance (or volatility squared). Volatility should be considered to be a derivative of variance. The realisation that variance should be used instead of volatility led volatility indices, such as the VIX, to move away from ATM volatility (VXO index) towards a variance-based calculation.

VAR, NOT VOL, IS CORRECT MEASURE FOR DEVIATION

There are three reasons why variance, not volatility, should be used as the correct measure for volatility. However, despite these reasons, even variance swaps are normally quoted as the square root of variance for an easier comparison with the implied volatility of options (but we note that skew and convexity mean the fair price of variance should always trade above ATM options).

- **Variance takes into account implied volatility at all stock prices.** Variance takes into account the implied volatility of all strikes with the same expiry (while ATM implied will change with spot, even if volatility surface does not change).
- **Deviations need to be squared to avoid cancelling.** Mathematically, if deviations were simply summed then positive and negative deviations would cancel. This is why the sum of squared deviations is taken (variance) to prevent the deviations from cancelling. Taking the square root of this sum (volatility) should be considered a derivative of this pure measure of deviation (variance).
- **Profit from a delta-hedged option depends on the square of the return.** Due to the convexity of an option, if the volatility is doubled, the profits from delta hedging are multiplied by a factor of four. For this reason, variance (which looks at squared returns) is a better measure of deviation than volatility.

(1) VARIANCE TAKES INTO ACCOUNT VOLATILITY AT ALL STOCK PRICES

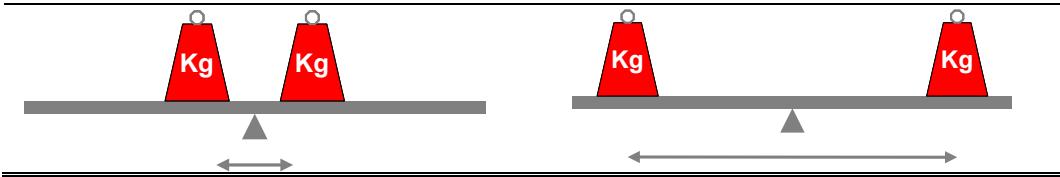
When looking at how rich or cheap options with the same maturity are, rather than looking at the implied volatility for a certain strike (ie, ATM or another suitable strike) it is better to look at the implied variance as it takes into account the implied volatility of all strikes. For example, if an option with a fixed strike that is initially ATM is bought, then as soon as spot moves it is no longer ATM. However, if a variance swap (or log contract³ of options in the absence of a variance swap market) is bought, then its traded level is applicable no matter what the level of spot. The fact a variance swap (or log contract) payout depends only on the realised variance and is not path dependent makes it the ideal measure for deviation.

(2) DEVIATIONS NEED TO BE SQUARED TO AVOID CANCELLING

If a seesaw has two weights on it and the weights are the same distance either side from the pivot, the weights are balanced as the centre of the mass is in line with the pivot (see graph on left hand side below). If the weights are further away from the pivot the centre of the mass (hence the average/expected distance of the weights) is still in line with the pivot (see graph on right hand side below). If the deviation of the two weights from the pivot is summed together, in both cases they would be zero (as one weight's deviation from the pivot is the negative of the other). In order to avoid the deviation cancelling this way, the square of the deviation (or variance) is taken, as the square of a number is always positive.

Figure 31. Low Deviation

High Deviation



(3) PROFIT FROM DELTA HEDGING PROPORTIONAL TO RETURN SQUARED

Assuming constant volatility, zero interest rates and dividend, the daily profit and loss (P&L) from delta hedging an option is given below:

$$\text{Delta-hedged P\&L from option} = S^2 \gamma / 2 + \text{cost term}$$

where:

$$\gamma \equiv \text{gamma}$$

³ For more details, see the section 2.3 *Volatility, Variance and Gamma Swaps*.

This can also be seen from Figure 29 Profile of a Delta-Hedged Option in the previous section (page 46), as the additional profit from the move from 1% to 2% is three times the profit from 0% to 1% (for a total profit four times the profit for a 1% move).

VOL SHOULD BE CONSIDERED A DERIVATIVE OF VAR

The three examples above show why variance is the natural measure for deviation. Volatility, the square root of variance, should be considered a derivative of variance rather than a pure measure of deviation. It is variance, not volatility, that is the second moment of a distribution (the first moment is the forward or expected price). For more details on moments, read the section [7.4 How to Measure Skew and Smile](#).

VIX CALCULATION MOVED FROM ATM VOL TO VAR

Due to the realisation that variance, not volatility, was the correct measure of deviation, on Monday, September 22, 2003, the VIX index moved away from using ATM implied towards a variance-based calculation. Variance-based calculations have also been used by other volatility index providers. The old VIX, renamed VXO, took the implied volatility for strikes above and below spot for both calls and puts. As the first two-month expiries were used, the old index was an average of eight implied volatility measures as $8 = 2 \text{ (strikes)} \times 2 \text{ (put/call)} \times 2 \text{ (expiry)}$. We note that the use of the first two expiries (excluding the front month if it was less than eight calendar days) meant the maturity was on average 1.5 months, not one month as for the new VIX.

VDAX calculation also moved from ATM vol to var

Similarly, the VDAX index, which was based on 45-day ATM-implied volatility, has been superseded by the V1X index, which, like the new VIX, uses a variance swap calculation. All recent volatility indices, such as the vStoxx (V2X), VSMI (V3X), VFTSE, VNKY and VHSI, use a variance swap calculation, although we note the recent VIMEX index uses a similar methodology to the old VIX (potentially due to illiquidity of OTM options on the Mexbol index).

VARIANCE TERM STRUCTURE IS NOT ALWAYS FLAT

While average variance term structure should be flat in theory, in practice supply and demand imbalances can impact variance term structure. The buying of protection at the long end should mean that variance term structure is on average upward sloping, but in turbulent markets it is usually inverted.

2.3: VOLATILITY, VARIANCE AND GAMMA SWAPS

In theory, the profit and loss from delta hedging an option is fixed and is based solely on the difference between the implied volatility of the option when it was purchased and the realised volatility over the life of the option. In practice, with discrete delta hedging and unknown future volatility, this is not the case, leading to the creation of volatility, variance and gamma swaps. These products also remove the need to continuously delta hedge, which can be very labour-intensive and expensive. Until the credit crunch, variance swaps were the most liquid of the three, but now volatility swaps are more popular for single stocks.

VOL, VAR & GAMMA SWAPS GIVE PURE VOL EXPOSURE

As spot moves away from the strike of an option the gamma decreases, and it becomes more difficult to profit via delta hedging. Second-generation volatility products, such as volatility swaps, variance swaps and gamma swaps, were created to give volatility exposure for all levels of spot and also to avoid the overhead and cost of delta hedging. While volatility and variance swaps have been traded since 1993, they became more popular post-1998, when Russia defaulted on its debts and Long-Term Capital Management (LTCM) collapsed.

Variance and gamma swaps usually quoted in volatility terms

Variance and gamma swaps are normally quoted as the square root of variance. This allows easier comparison with the options market.

VOLATILITY SWAP \leq GAMMA SWAP \leq VARIANCE SWAP

The square root of the variance swap is always above volatility swaps (and ATMf implied as volatility swaps \approx ATMf implied). This is due to the fact a variance swap payout is convex (hence, will always be greater than or equal to volatility swap payout of identical vega, which is explained later in the section). Only for the unrealistic case of no vol of vol (ie, future volatility is constant and known) will the price of a volatility swap and variance swap (and gamma swap) be the same⁴. The fair price of a gamma swap is between volatility swaps and variance swaps.

⁴ A variance swap payout is based on cash return assuming zero mean, whereas a delta-hedged option variance payout is based on a forward. Hence, a variance swap fair price will be slightly above a constant and flat volatility surface if the drift is non-zero (as close-to-close cash returns will be lifted by the drift).

- **Volatility swaps.** Volatility swaps were the first product to be traded significantly and became increasingly popular in the late 1990s until interest migrated to variance swaps. Following the collapse of the single-stock variance market in the credit crunch, they are having a renaissance due to demand from dispersion traders. A theoretical drawback of volatility swaps is the fact that they require a volatility of volatility (vol of vol) model for pricing, as options need to be bought and sold during the life of the contract (which leads to higher trading costs). However, in practice, the vol of vol risk is small and volatility swaps trade roughly in line with ATM forward (ATMf) implied volatility.
- **Variance swaps.** The difficulty in hedging volatility swaps drove liquidity towards the variance swap market, particularly during the 2002 equity collapse. As variance swaps can be replicated by delta hedging a static portfolio of options, it is not necessary to buy or sell options during the life of the contract. The problem with this replication is that it assumes options of all strikes can be bought, but in reality very OTM options are either not listed or not liquid. Selling a variance swap and only hedging with the available roughly ATM options leaves the vendor short tail risk. As the payout is on variance, which is volatility squared, the amount can be very significant. For this reason, liquidity on single-stock variance disappeared in the credit crunch.
- **Gamma swaps.** Dispersion traders profit from overpriced index-implied volatility by going long single-stock variance and short index variance. The portfolio of variance swaps is not static; hence, rebalancing trading costs are incurred. Investment banks attempted to create a liquid gamma swap market, as dispersion can be implemented via a static portfolio of gamma swaps (and, hence, it could better hedge the exposure of their books from selling structured products). However, liquidity never really took off due to limited interest from other market participants.

(1) VOLATILITY SWAPS

The payout of a volatility swap is simply the notional, multiplied by the difference between the realised volatility and the fixed swap volatility agreed at the time of trading. As can be seen from the payoff formula below, the profit and loss is completely path independent as it is solely based on the realised volatility. Volatility swaps were previously illiquid, but are now more popular with dispersion traders, given the single stock variance market no longer exists post credit crunch. Unless packaged as a dispersion, volatility swaps rarely trade. As dispersion is short index volatility, long single stock volatility, single stock volatility swaps tend to be bid only (and index volatility swaps offered only).

Volatility swap payoff

$$(\sigma_F - \sigma_S) \times \text{volatility notional}$$

\rightarrow Notional ~~Fee~~

where:

σ_F = future volatility (that occurs over the life of contract)

σ_S = swap rate volatility (fixed at the start of contract)

Volatility notional = Vega = notional amount paid (or received) per volatility point

(2) VARIANCE SWAPS

Variance swaps are identical to volatility swaps except their payout is based on variance (volatility squared) rather than volatility. Variance swaps are long skew (more exposure to downside put options than upside calls) and convexity (more exposure to OTM options than ATM). Typically variance swaps trade in line with the 30 delta put (if skew is downward sloping as normal). One-year variance swaps are the most frequently traded.

Variance swap payoff

$$(\sigma_F^2 - \sigma_S^2) \times \text{variance notional}$$

where:

Variance notional = notional amount paid (or received) per variance point

NB: Variance notional = Vega / $(2 \times \sigma_S)$ where σ_S = current variance swap price

CAPPED VAR SWAPS ARE SHORT OPTION ON VAR

Variance swaps on single stocks and emerging market indices are normally capped at 2.5 times the strike, in order to prevent the payout from rising towards infinity in a crisis or bankruptcy. A cap on a variance swap can be modelled as a vanilla variance swap less an option on variance whose strike is equal to the cap. More details can be found in the section 2.4 Options on Variance.

Capped var should be hedged with OTM calls, not OTM puts

The presence of a cap on a variance swap means that if it is to be hedged by only one option it should be a slightly OTM call, not an OTM (approx delta 30) put. This is to ensure the option is so far OTM when the cap is hit that the hedge disappears. If this is not done, then if a trader is long a capped variance swap he would hedge by going short an OTM put. If

markets fall with high volatility hitting the cap, the trader would be naked short a (now close to ATM) put. Correctly hedging the cap is more important than hedging the skew position.

S&P500 variance market is increasing in liquidity

The payout of volatility swaps and variance swaps of the same vega is similar for small payouts, but for large payouts the difference becomes very significant due to the quadratic (ie, squared) nature of variance. The losses suffered in the credit crunch from the sale of variance swaps, particularly single stock variance (which, like single stock volatility swaps now, was typically bid only), have weighed on their subsequent liquidity. Now variance swaps only trade for indices (usually without cap, but sometimes with). The popularity of VIX futures has raised awareness of variance swaps, which has helped S&P500 variance swaps become more liquid than they were before the credit crunch. SX5E variance liquidity used to be in line with the S&P500, but is now far less liquid.

CORRIDOR VARIANCE SWAPS ARE NOT LIQUID

As volatility and spot are correlated, volatility buyers would typically only want exposure to volatility levels for low values of spot. Conversely, volatility sellers would only want exposure for high values of spot. To satisfy this demand, corridor variance swaps were created. These only have exposure when spot is between spot values A and B. If A is zero, then it is a down variance swap. If B is infinity, it is an up variance swap. There is only a swap payment on those days the spot is in the required range, so if spot is never in the range there is no payment. Because of this, a down variance swap and up variance swap with the same spot barrier is simply a vanilla variance swap. The liquidity of corridor variance swaps was always far lower than for variance swaps and, post credit crunch, they are rarely traded.

Corridor variance swap payoff

$$(\sigma_{F \text{ when in corridor}}^2 - \sigma_S^2) \times \text{variance notional} \times \text{percentage of days spot is within corridor}$$

where:

$$\sigma_{F \text{ when in corridor}} = \text{future volatility (of returns } P_i/P_{i-1} \text{ which occur when } B_L < P_{i-1} \leq B_H)$$

B_L and B_H , are the lower and higher barriers, where B_L could be 0 and B_H could be infinity.

(3) GAMMA SWAPS

The payout of gamma swaps is identical to that of a variance swap, except the daily P&L is weighted by spot (price_n) divided by the initial spot (price_0). If spot range trades after the position is initiated, the payouts of a gamma swap are virtually identical to the payout of a variance swap. Should spot decline, the payout of a gamma swap decreases. Conversely, if spot increases, the payout of a gamma swap increases. This spot-weighting of a variance swap payout has the following attractive features:

- Spot weighting of gamma swap payout makes it unnecessary to have a cap, even for single stocks (if a company goes bankrupt with spot dropping close to zero with very high volatility, multiplying the payout by spot automatically prevents an excessive payout).
- If a dispersion trade uses gamma swaps, the amount of gamma swaps needed does not change over time (hence, the trade is ‘fire and forget’, as the constituents do not have to be rebalanced as they would if variance swaps were used).
- A gamma swap can be replicated by a static portfolio of options (although a different static portfolio to variance swaps), which reduces hedging costs. Hence, no volatility of volatility model is needed (unlike volatility swaps).

Gamma swap market has never had significant liquidity

A number of investment banks attempted to kick start a liquid gamma swap market, partly to satisfy potential demand from dispersion traders and partly to get rid of some of the exposure from selling structured products (if the product has less volatility exposure if prices fall, then a gamma swap better matches the change in the vega profile when spot moves). While the replication of the product is as trivial as for variance swaps, it was difficult to convince other market participants to switch to the new product and liquidity stayed with variance swaps (although after the credit crunch, single-stock variance liquidity moved to the volatility swap market). If the gamma swap market ever gains liquidity, long skew trades could be put on with a long variance-short gamma swap position (as this would be long downside volatility and short upside volatility, as a gamma swap payout decreases/increases with spot).

Gamma swap payoff

$$(\sigma_G^2 - \sigma_S^2) \times \text{variance notional}$$

where:

$$\sigma_G^2 = \text{future spot weighted (ie, multiplied by } \frac{\text{price}_n}{\text{price}_0} \text{)} \text{ variance}$$

$$\sigma_S^2 = \text{swap rate variance (fixed at the start of contract)}$$

PAYOUT OF VOL, VAR AND GAMMA SWAPS

The payout of volatility swaps, variance swaps and gamma swaps is the difference between the fixed and floating leg, multiplied by the notional. The calculation for volatility assumes zero mean return (or zero drift) to make the calculation easier and to allow the variance calculation to be additive.

- **Fixed leg.** The cost (or fixed leg) of going long a volatility, variance or gamma swap is always based on the swap price, σ_S (which is fixed at inception of the contract). The fixed leg is σ_S for volatility swaps, but is σ_S^2 for variance and gamma swaps).
- **Floating leg.** The payout (or floating leg) for volatility and variance swaps is based on the same variable σ_F (see equation below). The only difference is that a volatility swap payout is based on σ_F , whereas for a variance swap it is σ_F^2 . The gamma swap payout is based on a similar variable σ_G^2 , which is σ_F^2 multiplied by $\text{price}_n/\text{price}_0$.

$$\sigma_F = 100 \times \sqrt{\frac{\sum_{i=1}^T [\ln(\text{return}_i)]^2}{T_{\text{exp}}}} \times \text{number business days in year}$$

$$\sigma_G = 100 \times \sqrt{\frac{\sum_{i=1}^T \frac{\text{price}_i}{\text{price}_0} [\ln(\text{return}_i)]^2}{T_{\text{exp}}}} \times \text{number business days in year}$$

$$\text{return}_i = \frac{\text{price}_i}{\text{price}_{i-1}} \text{ for indices}$$

$$\text{return}_i = \frac{\text{price}_i + \text{dividend}_i}{\text{price}_{i-1}} \text{ for single stocks (dividend}_i \text{ is dividend going ex on day n)}$$

where:

number of business days in year = 252 (usual market practice)

T_{exp} = Expected value of N (if no market disruption occurs). A market disruption is when shares accounting for at least 20% of the index market cap have not traded in the last 20 minutes of the trading day.

Basis risk between cash and futures can cause traders problems

We note that the payout of variance swaps is based on the cash close, but traders normally delta hedge using futures. The difference between the cash and futures price is called the basis, and the risk due to a change in basis is called basis risk. Traders have to take this basis risk between the cash close and futures close, which can be significant as liquidity in the futures market tends to be reduced after the cash market closes.

VOL, VAR AND GAMMA SWAPS ARE ACTUALLY FORWARDS

The naming of volatility swaps, variance swaps and gamma swaps is misleading, as they are in fact forwards. This is because their payoff is at maturity, whereas swaps have intermediate payments.

VARIANCE IS ADDITIVE WITH ZERO MEAN ASSUMPTION

Normally, standard deviation or variance looks at the deviation from the mean. The above calculations assume a zero mean, which simplifies the calculation (typically, one would expect the mean daily return to be relatively small). With a zero mean assumption, variance is additive. A mathematical proof of the formula below is given in the section A2 Measuring Historical Volatility in the *Appendix*.

Past variance + future variance = total variance

Lack of dividend adjustment for indices affects pricing

The return calculation for a variance swap on an index does not adjust for any dividend payments that go ex. This means that the dividend modelling method can affect the pricing. Near-dated and, hence, either known or relatively certain dividends should be modelled discretely rather than as a flat yield. The changing exposure of the variance swap to the volatility on the ex date can be as large as 0.5 volatility points for a three-year variance swap (if all other inputs are kept constant, discrete (ie, fixed) dividends lift the value of both calls and puts, as proportional dividends simply reduce the volatility of the underlying by the dividend yield).

Calculation agents have discretion as to when a market disruption occurs

Normally, the investment bank is the calculation agent for any variance swaps traded. As the calculation agent normally has some discretion over when a market disruption event occurs, this can lead to cases where one calculation agent believes a market disruption occurs and another does not. This led to a number of disputes in 2008, as it was not clear if a market or exchange disruption had occurred. Similarly, if a stock is delisted, the estimate of future volatility for settlement prices is unlikely to be identical between firms, which can lead to issues if a client is long and short identical products at different investment banks. These problems are less of an issue if the counterparties are joint calculation agents.

HEDGING OF VAR CAN IMPACT EQUITY & VOL MARKET

Hedging volatility, variance and gamma swaps always involve the trading of a strip of options of all strikes and delta hedging at the close. The impact the hedging of all three products has on equity and volatility markets is similar, but we shall use the term variance swaps, as it has by far the most impact of the three (the same arguments will apply for volatility swaps and gamma swaps).

Short end of volatility surfaces is now pinned to realised

If there is a divergence between short-dated variance swaps and realised volatility, hedge funds will put on variance swap trades to profit from this divergence. This puts pressure on the short-dated end of volatility surfaces to trade close to the current levels of realised volatility. Due to the greater risk of unexpected events, it is riskier to attempt a similar trade at the longer-dated end of volatility surfaces.

Skew levels affected by direction of volatility trading

As variance swaps became a popular way to express a view of the direction of implied volatility, they impacted the levels of skew. This occurred as variance swaps are long skew (explained below) and, if volatility is being sold through variance swaps, this weighs on skew. This occurred between 2003 and 2005, which pushed skew to a multiple-year low. As volatility bottomed, the pressure from variance swap selling abated and skew recovered.

Delta hedge can suppress or exaggerate market moves

As the payout of variance swaps is based on the close-to-close return, they all have an intraday delta (which is equal to zero if spot is equal to the previous day's close). As this intraday delta resets to zero at the end of the day, the hedging of these products requires a delta hedge at the cash close. A rule of thumb is that the direction of hedging flow is in the direction that makes the trade the least profit (ensuring that if a trade is crowded, it makes less money). This flow can be hundreds of millions of US dollars or euros per day, especially when structured products based on selling short-dated variance are popular (as they were in 2006 and 2007, less so since the high volatility of the credit crunch).

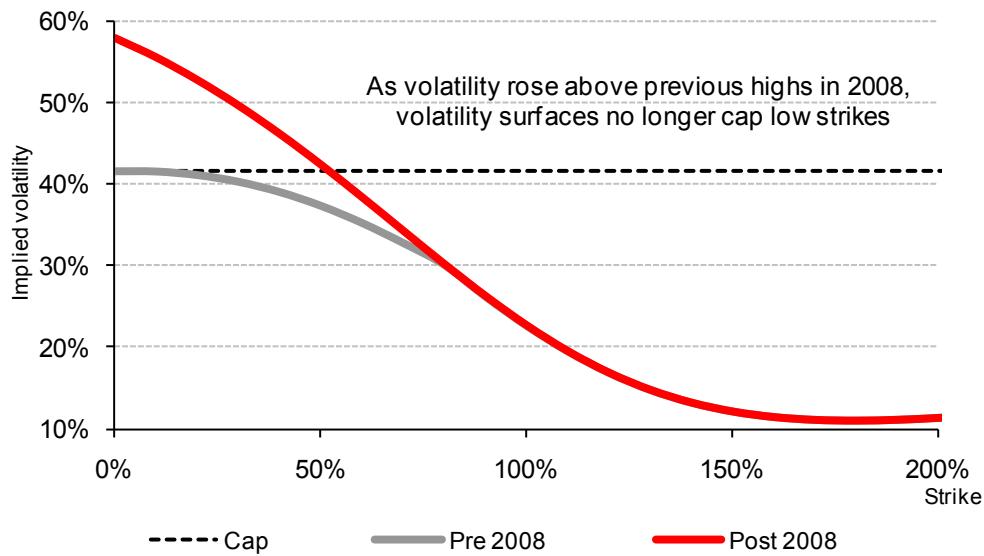
- **Variance buying suppresses equity market moves.** If clients are net buyers of variance swaps, they leave the counterparty trader short. The trader will hedge this short position by buying a portfolio of options and delta hedging them on the close. If spot has risen over the day the position (which was originally delta-neutral) has a positive delta (in the same way as a delta-hedged straddle would have a positive delta if markets rise). The end of day hedge of this position requires selling the underlying (to become delta-flat), which suppresses the rise of spot. Similarly, if markets fall, the delta hedge required is to buy the underlying, again suppressing the market movement.

- **Variance selling exaggerates equity market moves.** Should clients be predominantly selling variance swaps, the hedging of these products exaggerates market moves. The argument is simply the inverse of the argument above. The trader who is long a variance swap (as the client is short) has hedged by selling a portfolio of options. If markets rise, the delta of the position is negative and, as the variance swap delta is reset to zero at the end of the day, the trader has to buy equities at the same time (causing the close to be lifted for underlyings that have increased in value over the day). If markets fall, then the trader has to sell equities at the end of the day (as the delta of a short portfolio of options is positive). Movements are therefore exaggerated, and realised volatility increases if clients have sold variance swaps.

VAR PRICING CHANGED POST THE 2008 SPIKE IN VOL

The turmoil seen in 2008 caused 3-month realised volatility to spike above 70%. This was higher than the mid-60's high reached during the Great Depression. Before the Lehman bankruptcy, volatility traders used to cap implied volatility surfaces at a level similar to the all-time highs of realised volatility. The realisation that there could be an event that occurs in the future that has not occurred in the past, a so called 'black swan', has removed this cap (as it is now understood that volatility can spike above historical highs in a severe crisis).

Figure 32. Volatility Surfaces Pre- and Post-2008



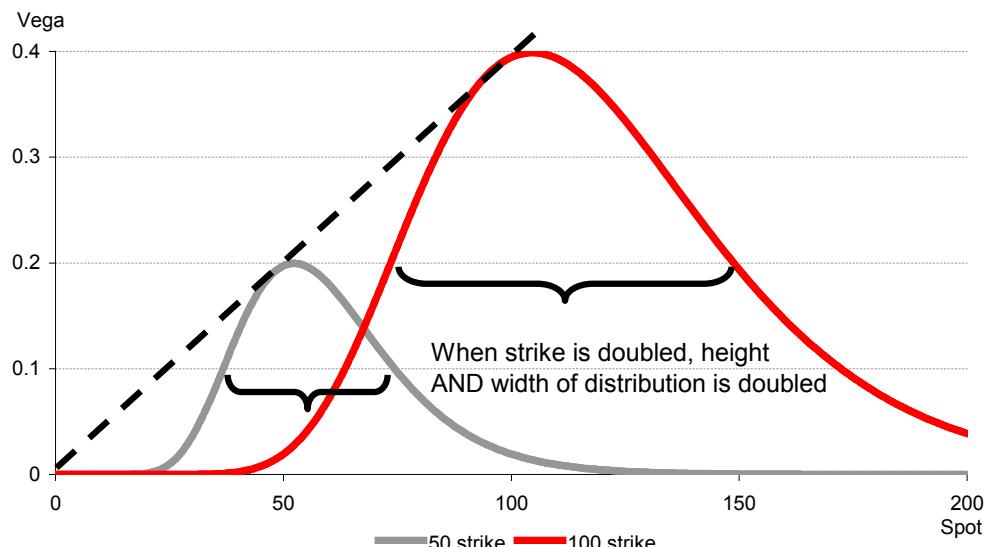
Removal of the implied vol cap has lifted var swap levels

The removal of the cap on implied volatility has caused low strike puts to be priced with a far higher implied volatility⁵. While the effect on premium for vanilla options (where the time value of very low strike puts is small) is small, for variance swaps the effect is very large. As variance swaps are more sensitive to low strike implied volatility (shown below), the removal of the cap lifted levels of variance swaps from c2pts above ATMf to c7pts above (which has since fallen to 3-4pts above).

HEDGING VOL, VAR AND GAMMA SWAPS

As volatility, variance and gamma swaps give volatility exposure for all values of spot, they need to be hedged by a portfolio of options of every strike. An equal-weighted portfolio is not suitable, as the vega profile of an option increases in size and width as strike increases (ie, an option of strike 2K has a peak vega double the peak vega of an option of strike K and is also twice the width). This is shown below.

Figure 33. Vega of Options of Different Strikes



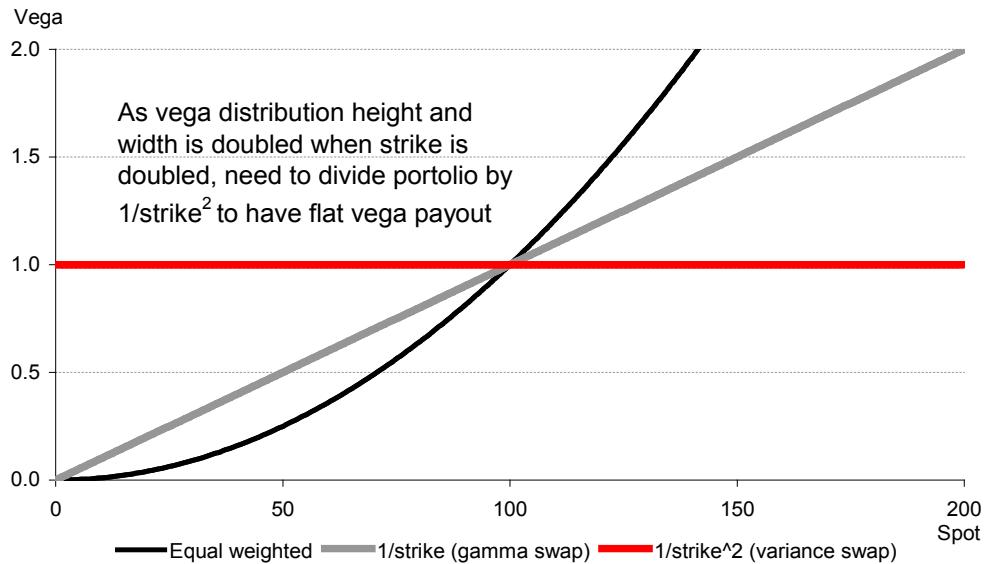
⁵ Note the slope of $\ln(\text{strike})$ cannot become steeper as spot declines without arbitrage occurring

Variance swaps are hedged with portfolio weighted $1/K^2$

Because a variance swap has a flat vega profile, the correct hedge is a portfolio of options weighted $1/K^2$ (where K is the strike of the option, ie, each option is weighted by 1 divided by its own strike squared). The reason why this is the correct weighting is due to the fact the vega profile doubles in height and width if the strike is doubled. The portfolio has to be divided by strike K once, to correct for the increase in height, and again to compensate for the increase in width (for a combined weight of $1/K^2$). A more mathematical proof of why the hedge for a variance swap is a portfolio of options weighted $1/K^2$ (a so-called log contract) is given in the section [A4 Proof Variance Swaps Should Be Hedged by a Log Contract \(\$= 1/K^2\$ \)](#) in the *Appendix*.

As a gamma swap payout is identical to a variance swap multiplied by spot, the weighting is $1/K$ (multiplying by spot cancels one of the K 's on the denominator). The vega profile of a portfolio weighted $1/K$ and $1/K^2$ is shown below, along with an equal-weighted portfolio for comparison. We note that although the vega profile of a variance swap is a flat line, the value is not constant and it moves with volatility (variance swap vega = variance notional $\times 2\sigma$). The vega profile of a volatility swap is of course a flat line (as vega is equal to the volatility notional).

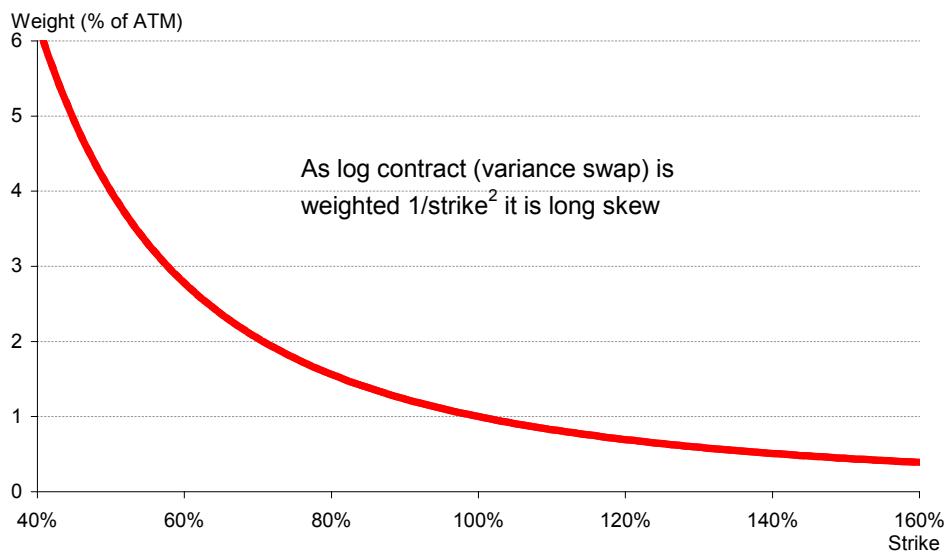
Figure 34. Vega of Portfolio of Options of All Strikes



Var swaps are long skew and volatility surface curvature

The $1/K^2$ weighting means a larger amount of OTM puts are traded than OTM calls (c60% is made up of puts). This causes a log contract (portfolio of options weighted $1/K^2$) to be long skew. This weighting means the wings (very out-of-the-money options) have a greater weighting than the body (near ATM options), which means a log contract is long volatility surface curvature⁶.

Figure 35. Weight of Options in Log Contract (Variance Swap)

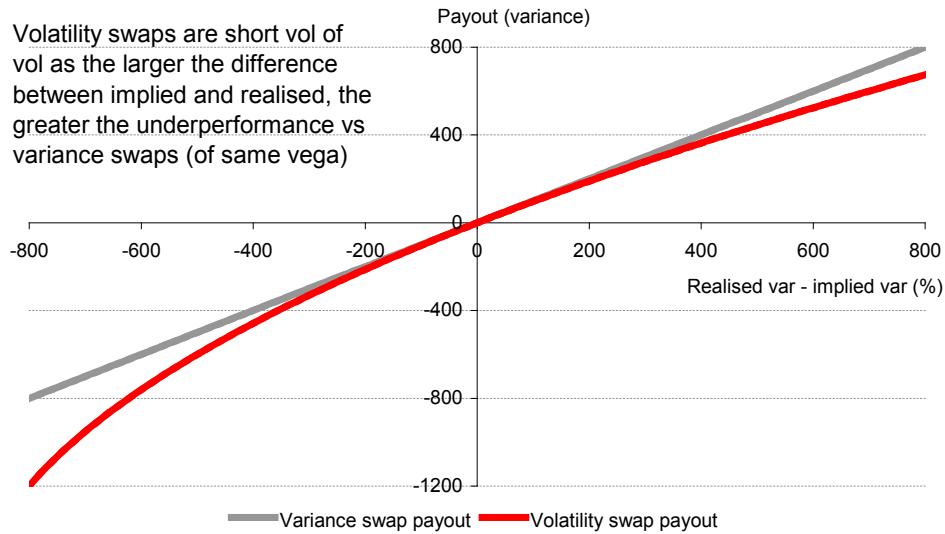


VOL SWAPS CAN BE HEDGED WITH VAR SWAPS

Unlike variance swaps (or gamma swaps), volatility swaps cannot be hedged by a static portfolio of options. Volatility swaps can be hedged with variance swaps as, for small moves, the payout can be similar (see Figure 36 below). The vega of a variance swap is equal to $\text{variance notional} \times 2\sigma$. For example, for $\sigma=25$ the vega is $2 \times 25 = 50$ times the size of the variance swap notional. So, a volatility swap of vega 'V' can be hedged with $V/2\sigma$ variance notional of a variance swap. As a variance swap is normally quoted in vega, the vega / 2σ formula is used to calculate the variance notional of the trade.

⁶ The inclusion of OTM (and hence convex) options mean the log contract is also long volga (or vega convexity), but they are not the same thing. Long OTM (wing) options is long vega convexity, but not volatility surface curvature (unless they are shorting the ATM or body at the same time). The curvature of the volatility surface can be defined as the difference between 90-100 skew and 100-110 skew (ie, the value of 90% + 110% – 2×100% implied volatilities).

Figure 36. Payout of Variance Swap and Volatility Swap



VOLATILITY SWAPS ARE SHORT VOL OF VOL

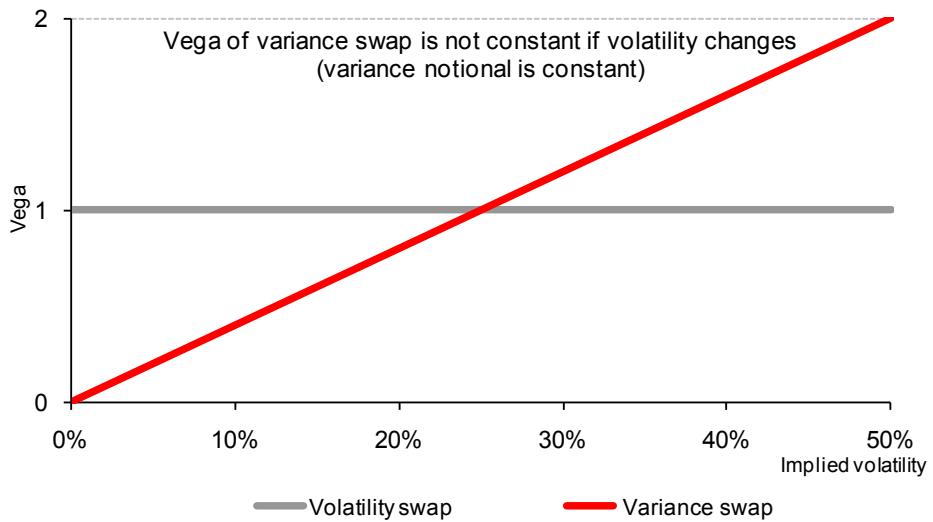
Figure 36 above shows that the payout of a variance swap is always in excess of the payout of a volatility swap of the same vega. This is why the fair level of a variance swap is usually one or two volatility points above volatility swaps. The negative convexity of the payout (compared to a variance swap) shows that volatility swaps are short vol of vol.

A volatility swap being short vol of vol can also be shown by the fact the identical vega of a variance swap has to be weighted $1/(2\sigma)$. If a trader is long a volatility swap and has hedged with a short variance swap position weighted $1/(2\sigma)$, then as volatility decreases more variance swaps have to be sold (as σ decreases, $1/(2\sigma)$ rises). Conversely, as volatility rises, variance swaps have to be bought (to decrease the short). Having to sell when volatility declines and buy when it rises shows that volatility swaps are short vol of vol.

$$\text{Variance notional} = \text{Vega} / (2\sigma)$$

VAR SWAP VEGA NOT CONSTANT IF VOLATILITY CHANGES

We note that although the vega profile of a variance swap against spot is a flat line, this value is not constant and it moves with volatility (variance swap vega = variance notional $\times 2\sigma$). The vega profile of a volatility swap against volatility is, of course, a constant flat line (as vega is equal to the volatility notional). Therefore, variance swaps have constant vega for changes in spot (but not changes in volatility), while volatility swaps have constant vega for changes in volatility and spot.

Figure 37. Vega of Variance Swaps and Volatility Swaps

VAR AND VOL DIFFERENCE CAN BE APPROXIMATED

Given that the difference between variance and volatility swap prices is due to the fact volatility swaps are short vol of vol, it is possible to derive the formula below, which approximates the difference between variance swap and volatility swap prices (as long as the maturity and vol of vol are not both excessive, which tends not to happen as longer maturities have less vol of vol). Using the formula, the price of a volatility swap can be approximated by the price of a variance swap less the convexity adjustment c . Using this formula, the difference between variance and volatility swaps is graphed in Figure 38 below.

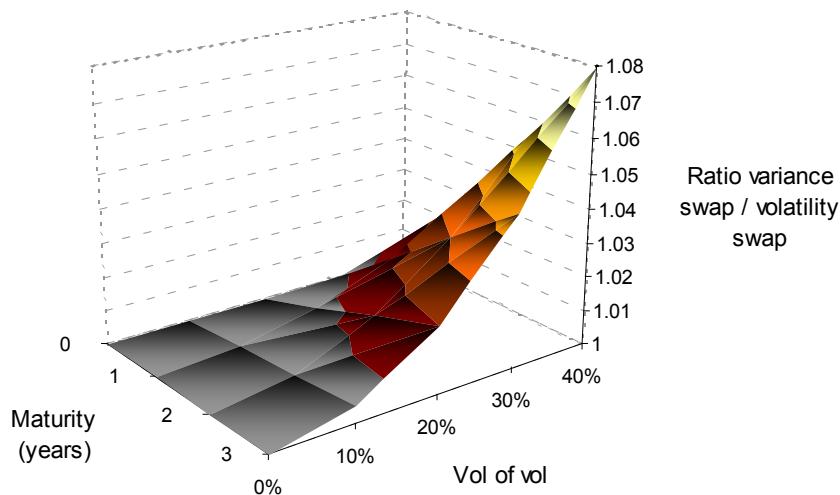
$$c \approx \frac{1}{6} \omega^2 T \sqrt{\text{variance swap price} \times e^{rT}}$$

where:

v = variance swap price

ω = volatility of volatility

Figure 38. Difference between Variance and Volatility Swap Prices



Model risk of vol of vol is small vs tail risk of variance swap

Hedging vol of vol raises trading costs and introduces model risk. Since the credit crunch, however, single-stock variance no longer trades and dispersion is now quoted using volatility swaps instead. Investment banks are happier taking the small model risk of vol of vol rather than being short the tail risk of a variance swap. As can be seen in Figure 39 below, variance swaps trade one or two volatility points above volatility swaps (for the most popular maturities). A simpler rule of thumb is that volatility swaps trade roughly in line with ATM implied volatilities.

Figure 39. Typical Values of Vol of Vol and the Effect on Variance and Volatility Swap Pricing

Maturity	3 Month	6 Month	1 Year	2 Year
Vol of vol	85%	70%	55%	40%
Ratio var/vol	1.030	1.041	1.050	1.053
Difference var - vol (for 30% vol)	0.90	1.23	1.51	1.60

$$\text{Max loss of variance swap} = \text{swap level} \times \text{vega} / 2$$

The notional of a variance swap trade is $\text{vega} / 2\sigma_S$ (σ_S is traded variance swap level) and the payoff is $(\text{realised}^2 - \sigma_S^2) \times \text{Notional}$. The maximum loss of a variance swap is when realised variance is zero, when the loss is $\sigma_S^2 \times \text{Notional} = \sigma_S^2 \times \text{vega} / 2\sigma_S = \sigma_S \times \text{vega} / 2$.

GREEKS OF VOL, VAR AND GAMMA SWAPS

As a volatility swap needs a vol of vol model, the Greeks are dependent on the model used. For variance swaps and gamma swaps, there is no debate as to the Greeks. However, practical considerations can introduce ‘shadow Greeks’. In theory, a variance swap has zero delta, but in practice it has a small ‘shadow delta’ due to the correlation between spot and implied volatility (skew). Similarly, theta is not necessarily as constant as it should be in theory, as movements of the volatility surface can cause it to change.

Variance swap vega decays linearly with time

As variance is additive, the vega decays linearly with time. For example, 100K vega of a one year variance swap at inception will have 75K vega after three months, 50K after six months and 25K after nine months.

Variance swaps offer constant cash gamma, gamma swaps have constant share gamma

Share gamma is the number of shares that need to be bought (or sold) for a given change in spot (typically 1%). It is proportional to the Black-Scholes gamma (second derivative of price with respect to spot) multiplied by spot. Cash gamma (or dollar gamma) is the cash amount that needs to be bought or sold for a given movement in spot; hence, it is proportional to share gamma multiplied by spot (ie, proportional to Black-Scholes gamma multiplied by spot squared). Variance swaps offer a constant cash gamma (constant convexity), whereas gamma swaps offer constant share gamma (hence the name gamma swaps).

$$\gamma \quad = \text{gamma} \quad = \text{number of shares bought (or sold) per } €1 \text{ spot move}$$

$$\gamma \times S \quad = \text{number of shares bought (or sold) per } 100\% \text{ spot } (S \times €1) \text{ move}$$

$$\gamma \times S / 100 \quad = \text{share gamma} = \text{number of shares bought (or sold) per } 1\% \text{ move}$$

$$\gamma \times S^2 / 100 \quad = \text{cash/dollar gamma} = \text{notional cash value traded per } 1\% \text{ move}$$

2.4: OPTIONS ON VARIANCE

As the liquidity of the variance swap market improved in the middle of the last decade, market participants started to trade options on variance. As volatility is more volatile at high levels, the skew is positive (the inverse of the negative skew seen in the equity market). In addition, volatility term structure is inverted, as volatility mean reverts and does not stay elevated for long periods of time.

OPTIONS ON VAR EXPIRY = EXPIRY OF VAR SWAP

An option on variance is a European option (like all exotics) on a variance swap whose expiry is the same expiry as the option. As it is an option on variance, a volatility of volatility model is needed in order to price the option. At inception, the underlying is 100% implied variance, whereas at maturity the underlying is 100% realised variance (and in between it will be a blend of the two). As the daily variance of the underlying is locked in every day, the payoff could be considered to be similar to an Asian (averaging) option.

Options on variance are quoted in volatility points

Like a variance swap, the price of an option on variance is quoted in volatility points. The typical 3-month to 18-month maturity of the option is in line with the length of time it takes 3-month realised volatility to mean revert after a crisis. The poor liquidity of options on variance, and the fact the underlying tends towards a cash basket over time, means a trade is usually held until expiry.

Option on variance swap payoff

$$\text{Max}(\sigma_F^2 - \sigma_K^2, 0) \times \text{Variance notional}$$

where:

σ_F = future volatility (that occurs over the life of contract)

σ_K = strike volatility (fixed at the start of contract)

Variance notional = notional amount paid (or received) per variance point

NB: Variance notional = Vega / (2 σ_S) where σ_S = variance swap reference (current fair price of variance swap, not the strike)

PUT CALL PARITY APPLIES TO OPTIONS ON VARIANCE

As variance swaps have a convex volatility payout, so do options on variance. As options on variance are European, put call parity applies. The fact a long call on variance and short put on variance (of the same strike) is equal to a forward on variance (or variance swap) gives the following result for options on variance whose strike is not the current level of variance swaps.

$$\text{Call Premium}_{\text{variance points}} - \text{Put Premium}_{\text{variance points}} = \text{PV}(\text{Current Var Price}^2 - \text{Strike}^2)$$

where:

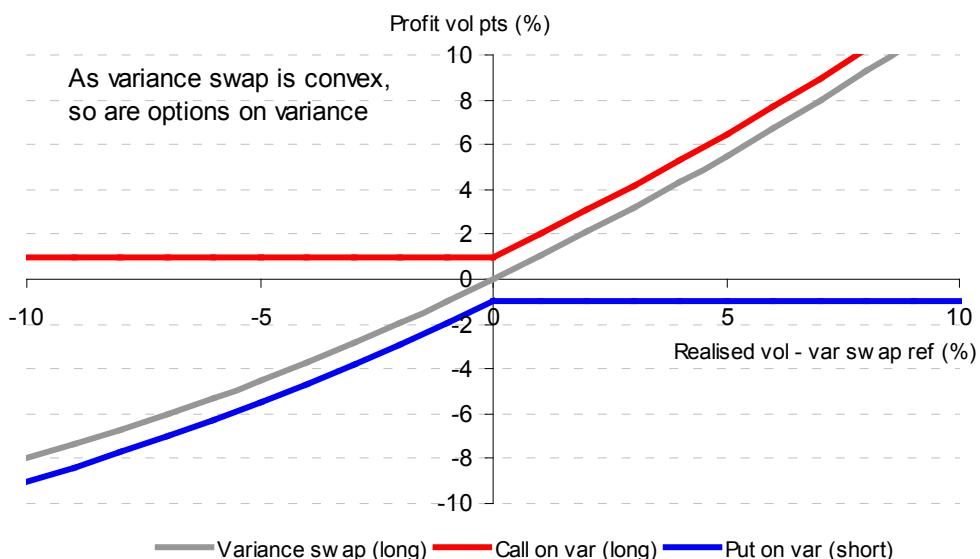
$$\text{Premium}_{\text{variance points}} = 2\sigma_S \times \text{Premium}_{\text{volatility points}} \text{ where } \sigma_S = \text{var swap reference}$$

PREMIUM PAID FOR OPTION = VEGA × PREMIUM IN VOL PTS

The premium paid for the option can either be expressed in terms of vega, or variance notional. Both are shown below:

$$\text{Fixed leg cash flow} = \text{Variance notional} \times \text{Premium}_{\text{variance points}} = \text{Vega} \times \text{Premium}_{\text{volatility points}}$$

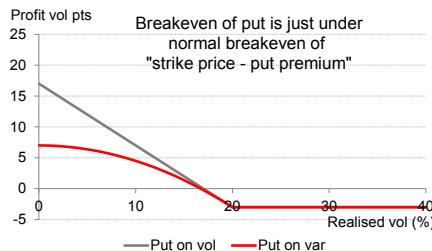
Figure 40. Variance Swap, ATM Call on Variance and ATM Put on Variance



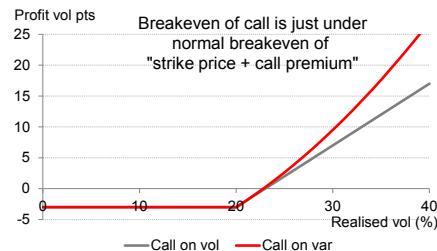
BREAKEVENS ARE NON-TRIVIAL DUE TO CONVEXITY

The convexity of a variance swap means that a put on a variance swap has a lower payout than a put on volatility and a call on variance swap has a higher payout than a call on volatility (see Figure 41 below). Similarly, it also means the maximum payout of a put on variance is significantly less than the strike. This convexity also means the breakevens for option on variance are slightly different from the breakevens for option on volatility (strike – premium for puts, strike + premium for calls).

Figure 41. Put on Variance Swap



Call on Variance Swap



BREAKEVENS ARE SIMILAR BUT NOT IDENTICAL TO OPTIONS ON VOLATILITY

In order to calculate the exact breakevens, the premium paid (premium P in vol points \times Vega) must equal the payout of the variance swap.

Premium paid = payout of variance swap

$$\text{For call option on variance: } P \times \text{Vega} = (\sigma_{\text{Call Breakeven}}^2 - \sigma_K^2) \times \frac{\text{Vega}}{2\sigma_S}$$

$$\sigma_{\text{Call Breakeven}} = \sqrt{\sigma_K^2 + 2\sigma_S P} \leq \sigma_K + P = \text{Call on volatility breakeven}$$

$$\text{Similarly } \sigma_{\text{Put Breakeven}} = \sqrt{\sigma_K^2 - 2\sigma_S P} \leq \sigma_K - P = \text{Put on volatility breakeven}$$

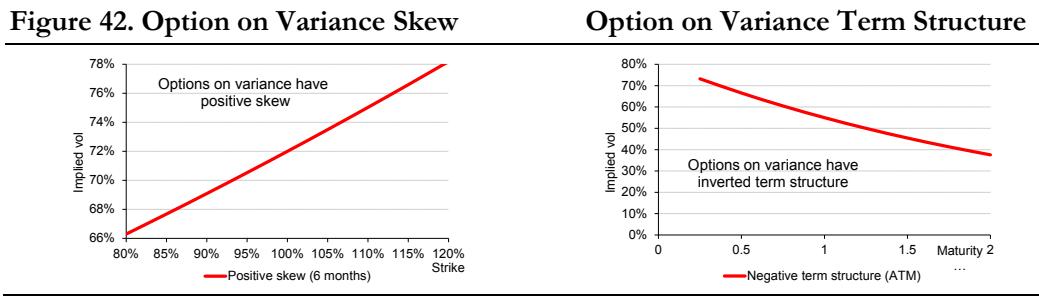
OPTIONS ON VARIANCE HAVE POSITIVE SKEW

Volatility (and hence variance) is relatively stable when it is low, as calm markets tend to have low and stable volatility. Conversely, volatility is more unstable when it is high (as turbulent markets could get worse with higher volatility, or recover with lower volatility). For this reason, options on variance have positive skew, with high strikes having higher implied volatility than low strikes.

Implied variance term structure is inverted, but not as inverted as realised variance

As historical volatility tends to mean revert in an eight-month time horizon (on average), the term structure of options on variance is inverted (while volatility can spike and be high for short periods of time, over the long term it trades in a far narrower range). We note that, as the highest volatility occurs due to unexpected events, the peak of implied volatility (which is based on the market's expected future volatility) is lower than the peak of realised volatility. Hence, the volatility of implied variance is lower than the volatility of realised variance, especially for short maturities.

Figure 42. Option on Variance Skew



CAPPED VAR HAVE EMBEDDED OPTION ON VAR

While options on variance swaps are not particularly liquid, their pricing is key for valuing variance swaps with a cap. Capped variance swaps are standard for single stocks and emerging market indices and can be traded on regular indices as well. When the variance swap market initially became more liquid, some participants did not properly model the cap, as it was seen to have little value. The advent of the credit crunch and resulting rise in volatility made the caps more valuable, and now market participants fail to model them at their peril.

Variance Swap with Cap C = Variance Swap - Option on Variance with Cap C

Option on Variance with Cap C = Variance Swap - Variance Swap with Cap C

While value of cap is small at inception, it can become more valuable as market moves

A capped variance swap can be modelled as a vanilla variance swap less an option on variance, whose strike is the cap. This is true as the value of an option on variance at the cap will be equal to the difference between the capped and uncapped variance swaps. Typically, the cap is at $2.5 \times$ the strike and, hence, is not particularly valuable at inception. However, as the market moves, the cap can become closer to the money and more valuable.

OPTIONS ON VAR STRATEGIES SIMILAR TO OPTIONS

Strategies that are useful for vanilla options have a read-across for options on variance. For example, a long variance position can be protected or overwritten. The increased liquidity of VIX options allows relative value trades to be put on.

Selling straddles on options on variance can also be a popular strategy, as volatility can be seen to have a floor above zero. Hence, strikes can be chosen so that the lower breakeven is in line with the perceived floor to volatility.

Options on variance can also be used to hedge a volatility swap position, as an option on variance can offset the vol of vol risk embedded in a volatility swap.

OPTIONS ON VOL FUTURES ARE SLIGHTLY DIFFERENT

While options on volatility futures have many similarities to options on variance, they are slightly different. For more details see the section [4.5: Options On Volatility Futures](#).

CHAPTER 3

WHY EVERYTHING YOU THINK YOU KNOW ABOUT VOLATILITY IS WRONG

Assuming implied volatility is an unbiased estimate of future realised volatility is an easy mistake to make, however the fair price of implied volatility is above the expected future realised volatility. In addition, the impact of hedging both structured products and variable annuity products can cause imbalances in the volatility market. These distortions can create opportunities for investors willing to take the other side. We examine the opportunities from these imbalances and dispel the myth of using volatility as an equity hedge.

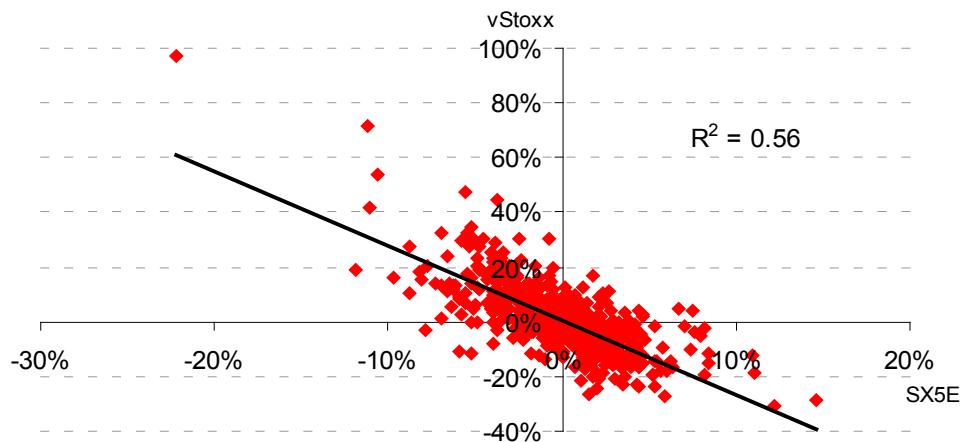
3.1: IMPLIED VOLATILITY SHOULD BE ABOVE REALIZED VOLATILITY

Selling implied volatility is one of the most popular trading strategies in equity derivatives. Empirical analysis shows that implied volatility or variance is, on average, overpriced. However, as volatility is negatively correlated to equity returns, a short volatility (or variance) position is implicitly long equity risk. As equity returns are expected to return an equity risk premium over the risk-free rate (which is used for derivative pricing), this implies short volatility should also be abnormally profitable. Therefore, part of the profits from short volatility strategies can be attributed to the fact equities are expected to deliver returns above the risk-free rate.

SHORT VOL HAS SIMILAR RISK & RETURN TO LONG EQUITY

As implied volatility tends to trade at a higher level than realised volatility, a common perception is that implied volatility is overpriced. While there are supply and demand imbalances that can cause volatility to be overpriced, part of the overpricing is due to the correlation between volatility and equity returns. A short volatility position is positively correlated to the equity market (as volatility typically increases when equities decline). As equities' average return is greater than the risk-free rate, this means that the risk-neutral implied volatility should be expected to be above the true realised volatility. Even taking this into account, volatility appears to be overpriced. We believe that implied volatility is overpriced on average due to the demand for hedges.

Figure 43. Correlation vStoxx (volatility) and SX5E



Far-dated options are most overpriced, due to upward sloping term structure

Volatility selling strategies typically involve selling near-dated volatility (or variance). Examples include call overwriting or selling near-dated variance (until the recent explosion of volatility, this was a popular hedge fund strategy that many structured products copied). As term structure is on average upward sloping, this implies that far-dated implieds are more expensive than near dated implieds. The demand for long-dated protection (eg, from variable annuity providers) offers a fundamental explanation for term structure being upward sloping (see section [3.3 Variable Annuity Hedging Lifts Long-Term Vol](#)). However, as 12× one month options (or variance swaps) can be sold in the same period of time as 1× one-year option (or variance swap), greater profits can be earned from selling the near-dated product despite it being less overpriced. We note the risk is greater if several near-dated options (or variance swap) are sold in any period.

VOLATILITY OVERPRICING IS UNLIKELY TO DISAPPEAR

There are several fundamental reasons why volatility, and variance, is overpriced. Since these reasons are structural, we believe that implied volatility is likely to remain overpriced for the foreseeable future. Given variance exposure to overpriced wings (and low strike puts) and the risk aversion to variance post credit crunch, we view variance as more overpriced than volatility.

- **Demand for put protection.** The demand for hedging products, either from investors, structured products or providers of variable annuity products, needs to be offset by market makers. As market makers are usually net sellers of volatility, they charge margin for taking this risk and for the costs of gamma hedging.
- **Demand for OTM options lifts wings.** Investors typically like buying OTM options as there is an attractive risk-reward profile (similar to buying a lottery ticket). Market makers therefore raise their prices to compensate for the asymmetric risk they face. As the price of variance swaps is based on options of all strikes, this lifts the price of variance.
- **Index implieds lifted from structured product demand.** The demand from structured products typically lifts index implied compared to single-stock implied. This is why implied correlation is higher than it should be.

SELLING VOL SHOULD BE LESS PROFITABLE THAN BEFORE

Hedge funds typically aim to identify mispricings in order to deliver superior returns. However, as both hedge funds and the total hedge fund marketplace grow larger, their opportunities are gradually being eroded. We believe that above-average returns are only possible in the following circumstances:

- **A fund has a unique edge** (eg, through analytics, trading algorithms or proprietary information/analysis).
- **There are relatively few funds in competition**, or it is not possible for a significant number of competitors to participate in an opportunity (either due to funding or legal restrictions, lack of liquid derivatives markets or excessive risk/time horizon of trade).
- **There is a source of imbalance in the markets** (eg, structured product flow or regulatory demand for hedging), causing a mispricing of risk.

All of the above reasons have previously held for volatility selling strategies (eg, call overwriting or selling of one/three-month variance swaps). However, given the abundance of publications on the topic in the past few years and the launch of several structured products that attempt to profit from this opportunity, we believe that volatility selling could be less profitable than before. The fact there remains an imbalance in the market due to the demand for hedging should mean volatility selling is, on average, a profitable strategy. However, we would caution against using a back test based on historical data as a reliable estimate of future profitability.

3.2: LONG VOLATILITY IS A POOR HEDGE

An ideal hedging instrument for a security is an instrument with -100% correlation to that security and zero cost. As the return on variance swaps have a c-70% correlation with equity markets, adding long volatility positions (either through variance swaps or futures on volatility indices such as VIX or vStoxx) to an equity position could be thought of as a useful hedge. However, as volatility is on average overpriced, the cost of this strategy far outweighs any diversification benefit.

VOLATILITY HAS UP TO 70% CORRELATION WITH EQUITY

Equity markets tend to become more volatile when they decline and less volatile when they rise. A fundamental reason for this is the fact that gearing increases as equities decline⁷. As both gearing and volatility are measures of risk, they should be correlated; hence, they are negatively correlated to equity returns. While short term measures of volatility (eg, vStoxx) only have an R² of 50%-60% against the equity market, longer dated variance swaps (purest way to trade volatility) can have up to c70% R².

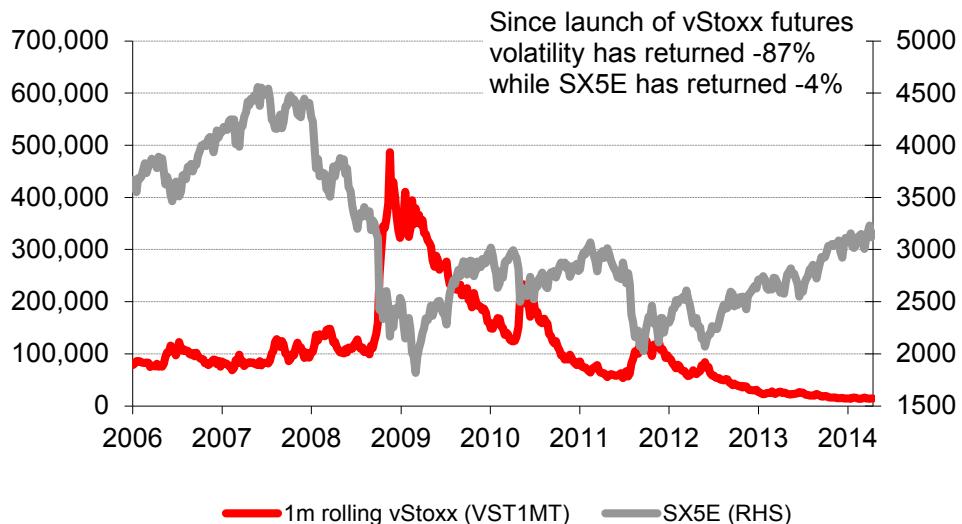
1 YEAR VOLATILITY HAS HIGHEST CORRELATION TO EQUITY

There are two competing factors to the optimum maturity for a volatility hedge. The longer the maturity, the more likely the prolonged period of volatility will be due to a decline in the market. This should give longer maturities higher equity volatility correlation, as the impact of short-term noise is reduced. However, for long maturities (years), there is a significant chance that the equity market will recover from any downturn, reducing equity volatility correlation. The optimum correlation between the SX5E and variance swaps, is for returns between nine months and one year. This is roughly in line with the c8 months it takes realised volatility to mean revert after a crisis.

SHORT-DATED VOLATILITY FUTURES ARE A POOR HEDGE

Recently, there have been several products based on rolling VIX or vStoxx futures whose average maturity is kept constant. As these products have to continually buy far-dated futures and sell near-dated futures (to keep average maturity constant as time passes), returns suffer from upward sloping term structure. Since the launch of vStoxx futures, rolling one-month vStoxx futures have had negative returns (see Figure 51 below). This is despite the SX5E also having suffered a negative return, suggesting that rolling vStoxx futures are a poor hedge. For more details on futures on volatility indices, see the section [4.1 Forward Starting Products](#).

⁷ Assuming no rights issues, share buybacks, debt issuance or repurchase/redemption.

Figure 44. SX5E and One-Month Rolling vStoxx Futures

LONG VOLATILITY HAS NEGATIVE RETURNS ON AVERAGE

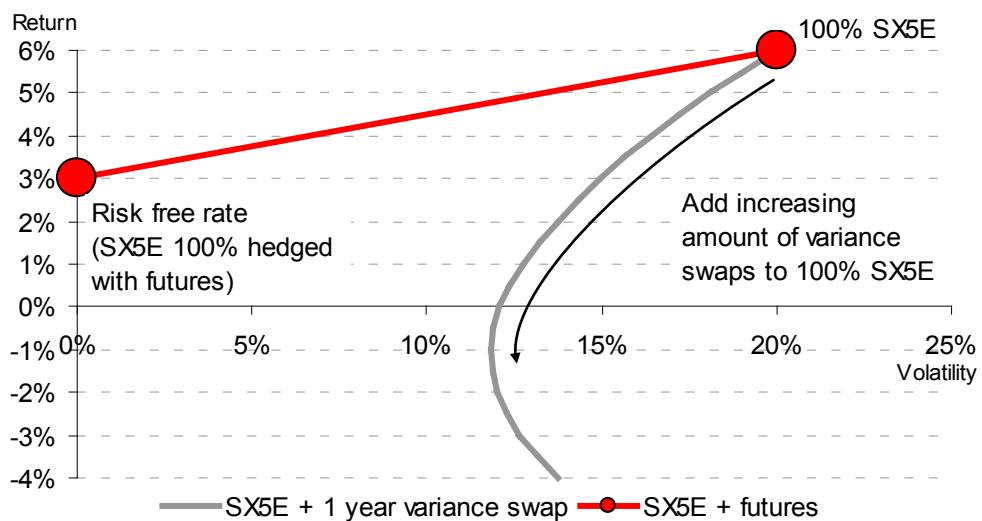
Long volatility strategies, on average, have negative returns. This overpricing can be broken down into two components:

- **Correlation with equity market.** As equity markets are expected to return an equity risk premium over the risk-free rate, strategies that are implicitly long equity risk should similarly outperform (and strategies that are implicitly short equity risk should underperform). As a long volatility strategy is implicitly short equity risk, it should underperform. We note this drawback should affect all hedging instruments, as a hedging instrument by definition has to be short the risk to be hedged.
- **Overpricing of volatility.** Excessive demand for volatility products has historically caused implied volatility to be overpriced. As this demand is not expected to significantly decrease, it is likely that implied volatility will continue to be overpriced (although volatility will probably not be as overpriced as in the past).

VOLATILITY IS A POOR HEDGE COMPARED TO FUTURES

While all hedging instruments can be expected to have a cost (due to being implicitly short equities and assuming a positive equity risk premium), long variance swaps have historically had an additional cost due to the overpricing of volatility. This additional cost makes long variance swaps an unattractive hedge compared to reducing the position (or shorting futures). This is shown in Figure 45 below by adding an additional variance swap position to a 100% investment in equities (we optimistically assume zero margining and other trading costs to the variance swap position).

Figure 45. SX5E hedged with Variance Swaps or Futures



Vol as a hedge suffers from overpricing, and less than 100% correlation

While the risk of the long equity and long variance swap position initially decreases as the long variance position increases in size, the returns of the portfolio are less than the returns for a reduced equity position of the same risk (we assume the proceeds from the equity sale are invested in the risk-free rate, which should give similar returns to hedging via short futures). Unlike hedging with futures, there comes a point at which increasing variance swap exposure does not reduce risk (and, in fact, increases it) due to the less than 100% correlation with the equity market.

Equities need to have positive return over hedging back-testing period

While we acknowledge that there are periods of time in which a long volatility position is a profitable hedge, these tend to occur when equity returns are negative (and short futures are usually a better hedge). We believe that the best back-testing periods for comparing hedging strategies are those in which equities have a return above the risk-free rate (if returns below the risk-free rate are expected, then investors should switch allocation away from equities into risk-free debt). For these back-testing periods, long volatility strategies struggle to demonstrate value as a useful hedging instrument. Hence, we see little reason for investors to hedge with variance swaps rather than futures given the overpricing of volatility, and less than 100% correlation between volatility and equity returns.

HEDGING WITH VARIANCE IS NOT COMPABLE TO PUTS

Due to the lack of convexity of a variance swap hedge, we believe it is best to compare long variance hedges to hedging with futures rather than hedging with puts. Although variance hedges might be cheaper than put hedges, the lack of convexity for long volatility makes this an unfair comparison, in our view.

3.3: VARIABLE ANNUITY HEDGING LIFTS LONG-TERM VOL

Since the 1980s, a significant amount of variable annuity products have been sold, particularly in the USA. The size of this market is now over US\$1trn. From the mid-1990s, these products started to become more complicated and offered guarantees to the purchaser (similar to being long a put). The hedging of these products increases the demand for long-dated downside strikes, which lifts long-dated implied volatility and skew.

VARIABLE ANNUITIES OFTEN EMBED A ‘PUT’ OPTION

With a fixed annuity, the insurance company that sold the product invests the proceeds and guarantees the purchaser a guaranteed fixed return. Variable annuities, however, allow the purchaser to pick which investments they want to put their funds into. The downside to this flexibility is the unprotected exposure to a decline in the market. To make variable annuities more attractive, from the 1990s many were sold with some forms of downside protection (or put). The different types of protection are detailed below in order of the risk to the insurance company.

- **Return of premium.** This product effectively buys an ATM put in addition to investing proceeds. The investor is guaranteed returns will not be negative.
- **Roll-up.** Similar to return of premium; however, the minimum guaranteed return is greater than 0%. The hedging of this product buys a put which is ITM with reference to spot, but OTM compared with the forward.
- **Ratchet (or maximum anniversary value).** These products return the highest value the underlying has ever traded at (on certain dates). The hedging of these products involves payout look-back options, more details of which are in the section [5.4 Look-Back Options](#).
- **Greater of ‘ratchet’ or ‘roll-up’.** This product returns the greater of the ‘roll-up’ or ‘ratchet’ protection.

Hedging of variable annuity products lifts index term structure and skew

The hedging of variable annuity involves the purchase of downside protection for long maturities. Often the products are 20+ years long, but as the maximum maturity with sufficient liquidity available on indices can only be 3-5 years, the position has to be dynamically hedged with the shorter-dated option. This constant bid for long-dated protection lifts index term structure and skew, particularly for the S&P500 but also affects other major indices (potentially due to relative value trading). The demand for protection (from viable annuity providers or other investors), particularly on the downside and for longer maturities, could be considered to be the reason why volatility (of all strikes and maturities), skew (for all maturities) and term structure are usually overpriced.

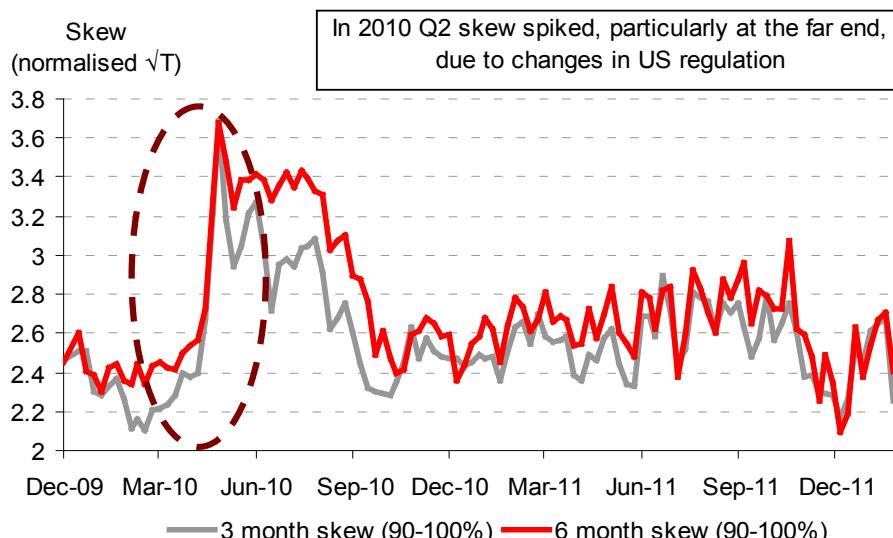
CREDIT CRUNCH HAS HIT VARIABLE ANNUITY PROVIDERS

Until the TMT bubble burst, guarantees embedded in variable annuity products were often seen as unnecessary ‘bells and whistles’. The severe declines between 2000 and 2003 made guarantees in variable annuity products more popular. When modelling dynamic strategies, insurance companies need to estimate what implied volatility will be in the future (eg, if hedging short 20-year options with 5-year options). The implied volatility chosen will be based on a confidence interval, say 95%, to give only a 1-in-20 chance that implieds are higher than the level embedded in the security. As the credit crunch caused realised volatility to reach levels that by some measures were higher than in the Great Depression, implied volatility rose to unprecedented heights. This increase in the cost of hedging has weighed on margins.

NO PROP DESK + MOVE TO EXCHANGE = HEDGE COST RISE

The passing of the Dodd-Frank Act in mid-2010 was designed to improve the transparency of derivatives by moving them onto an exchange. However, this would increase the margin requirements of long-dated options, which were previously traded OTC. This made it more expensive to be the counterparty to variable annuity providers. As the act also included the ‘Volcker Rule’, which prohibits proprietary trading, the number of counterparties shrank (as prop desks with attractive funding levels were a common counterparty for the long-dated protection required by variable annuity hedgers). The combination of the spinoff of prop desks, and movement of OTC options onto an exchange caused skew to rise in mid-2010, particularly at the far-dated end of volatility surfaces.

Figure 46. SX5E Skew Multiplied by the Square Root of Time ($R^2=83\%$)



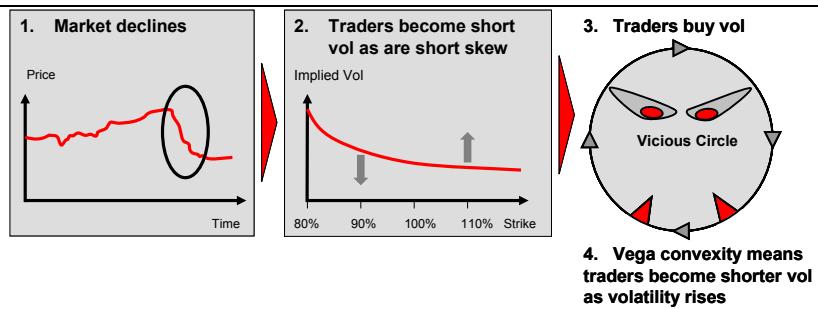
3.4: STRUCTURED PRODUCTS VICIOUS CIRCLE

The sale of structured products leaves investment banks with a short skew position (eg, short an OTM put in order to provide capital-protected products). Whenever there is a large decline in equities, this short skew position causes the investment bank to be short volatility (eg, as the short OTM put becomes more ATM, the vega increases). The covering of this short vega position lifts implied volatility further than would be expected. As investment banks are also short vega convexity, this increase in volatility causes the short vega position to increase in size. This can lead to a ‘structured products vicious circle’ as the covering of short vega causes the size of the short position to increase. Similarly, if equity markets rise and implied volatility falls, investment banks become long implied volatility and have to sell. Structured products can therefore cause implied volatility to undershoot in a recovery, as well as overshoot in a crisis.

IMPLIED VOL OVER AND UNDERSHOOTS BIG VOL MOVES

The sale of structured products causes investment banks to have a short skew and short vega convexity position⁸. Whenever there is a significant decline in equities and a spike in implied volatility, or a recovery in equities and a collapse in implied volatility, the position of structured product sellers can exaggerate the movement in implied volatility. This can cause implied volatility to overshoot (in a crisis) or undershoot (in a recovery post-crisis). There are four parts to the ‘structured products vicious circle’ effect on implied volatilities, which are shown in Figure 47 below.

Figure 47. Four Stages Towards Implied Volatility Overshoot



⁸ There is more detail on the position of structured product sellers at the end of this section.

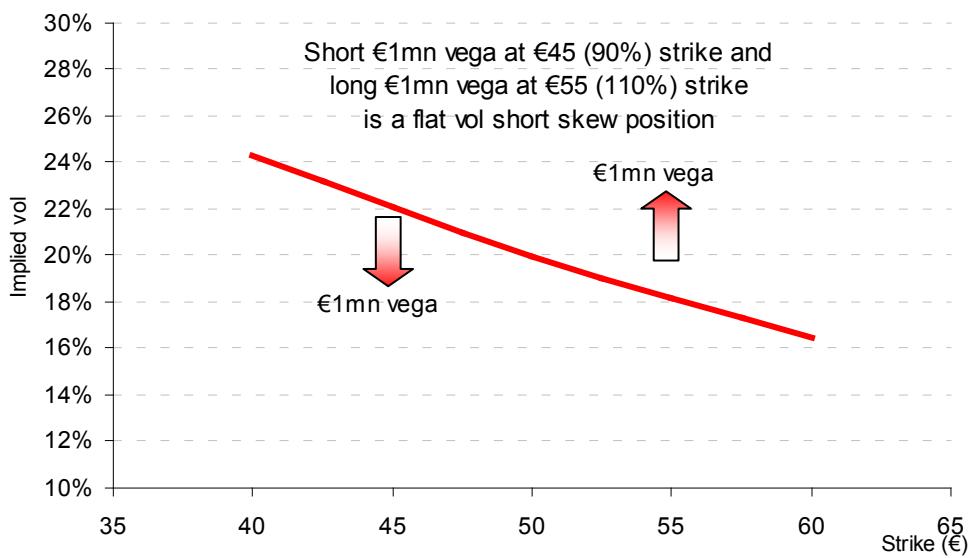
(1) EQUITY MARKET DECLINES

While implied volatility moves – in both directions – are exaggerated, for this example we shall assume that there is a decline in the markets and a rise in implied volatility. If this decline occurs within a short period of time, trading desks have less time to hedge positions, and imbalances in the market become more significant.

(2) DESKS BECOME SHORT IMPLIED VOLATILITY (DUE TO SHORT SKEW)

Investment banks are typically short skew from the sale of structured products. This position causes trading desks to become short implied volatility following declines in the equity market. To demonstrate how this occurs, we shall examine a short skew position through a vega flat risk reversal (short 90% put, long 110% call)⁹.

Figure 48. Short Skew Position Due to 90%-110% Risk Reversal (vega flat)

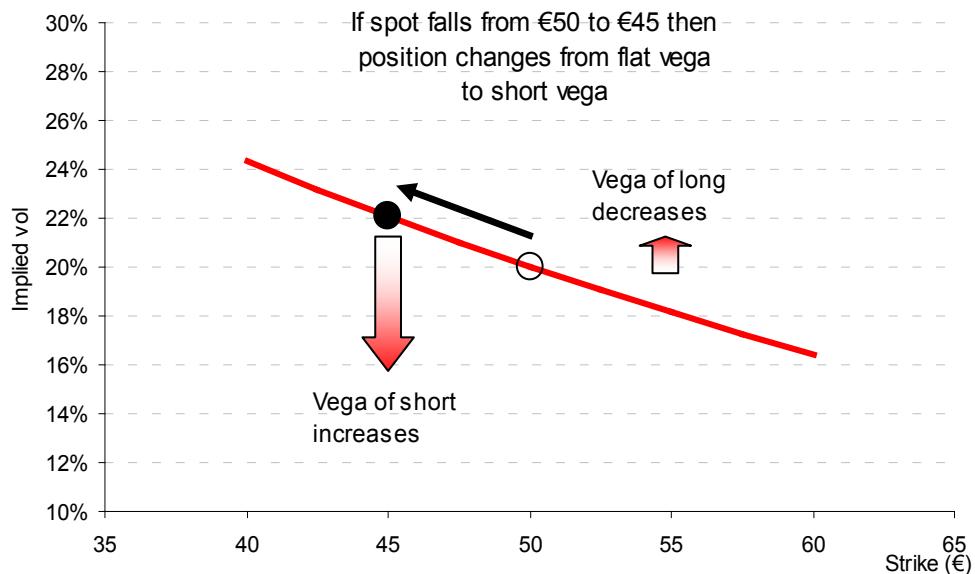


⁹ This simple example is very different from the position of structured product sellers. We note a vega flat risk reversal is not necessarily 1-1, as the vega of the put is likely to be lower than the vega of the call.

Short skew + equity markets decline = short vega (ie, short implied vol)

If there is a 10% decline in equity markets, the 90% put becomes ATM and increases in vega. As the risk reversal is short the 90% put, the position becomes short vega (or short implied volatility). In addition, the 110% call option becomes more OTM and further decreases the vega of the position (increasing the value of the short implied volatility position).

Figure 49. Change in Vega of 90%-110% Risk Reversal If Markets Decline 10%



Even if skew was flat, markets declines cause short skew position to become short vega

The above example demonstrates that it is the fact options become more or less ATM that causes the change in vega. It is not the fact downside put options have a higher implied than upside call options. If skew was flat (or even if puts traded at a lower implied than calls), the above argument would still hold. We therefore need a measure of the rate of change of vega for a given change in spot, and this measure is called vanna.

$$\text{Vanna} = \frac{d\text{Vega}}{d\text{Spot}}$$

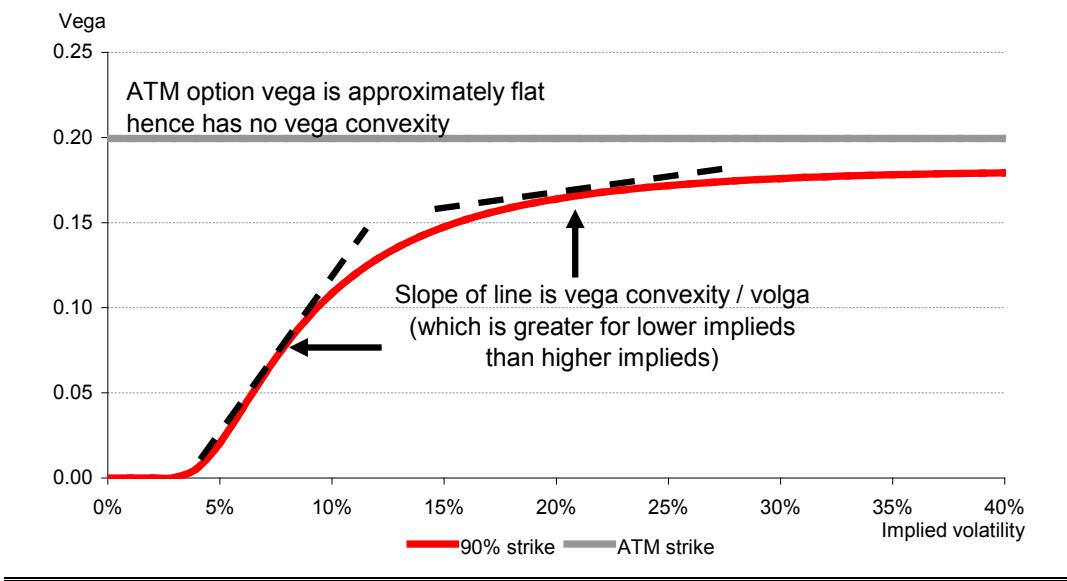
Vanna measures size of skew position, skew measures value of skew position

Vanna can be thought of as the size of the skew position (in a similar way that vega is the size of a volatility position), while skew (eg, 90%-100% skew) measures the value of skew (in a similar way that implied volatility measures the value of a volatility position). For more details on different Greeks, including vanna, see the section [A8 Greeks and Their Meaning](#) in the *Appendix*.

(3) SHORT COVERING OF SHORT VEGA POSITION LIFTS IMPLIED VOL

As the size of trading desks' short vega position increases during equity market declines, this position is likely to be covered. As all trading desks have similar positions, this buying pressure causes an increase in implied volatility. This flow is in addition to any buying pressure due to an increase in realised volatility and hence can cause an overshoot in implied volatility.

Figure 50. Vega of ATM and OTM Options Against Implied (Vega Convexity)



(4) SHORT VEGA POSITION INCREASES DUE TO VEGA CONVEXITY

Options have their peak vega when they are (approximately) ATM. As implied volatility increases, the vega of OTM options increases and converges with the vega of the peak ATM option. Therefore, as implied volatility increases, the vega of OTM options increases (see Figure 50 above). The rate of change of vega given a change in volatility is called volga (VOL-Gamma) or vomma, and is known as vega convexity.

$$\text{Volga} = d\text{Vega}/d\text{Vol}$$

Vega convexity causes short volatility position to increase

As the vega of options rises as volatility increases, this increases the size of the short volatility position that needs to be hedged. As trading desks' volatility short position has now increased, they have to buy volatility to cover the increased short position, which leads to further gains in implied volatility. This starts a vicious circle of increasing volatility, which we call the 'structured products vicious circle'.

VEGA CONVEXITY HIGHEST FOR LOW-TO-MEDIUM IMPLIEDS

As Figure 56 above shows, the slope of vega against volatility is steepest (ie, vega convexity is highest) for low-to-medium implied volatilities. This effect of vega convexity is therefore more important in volatility regimes of c20% or less; hence, the effect of structured products can have a similar effect when markets rise and volatilities decline. In this case, trading desks become long vega, due to skew, and have to sell volatility. Vega convexity causes this long position to increase as volatility declines, causing further volatility sellings. This is typically seen when a market recovers after a volatile decline (eg, in 2003 and 2009, following the end of the tech bubble and credit crunch, respectively).

IMPACT GREATEST FOR FAR-DATED IMPLIEDS

While this position has the greatest impact at the far end of volatility surfaces, a rise in far-dated term volatility and skew tends to be mirrored to a lesser extent for nearer-dated expiries. If there is a disconnect between near- and far-dated implied volatilities, this can cause a significant change in term structure.

STRUCTURED PRODUCT GUARANTEE IS LONG AN OTM PUT

The capital guarantee of many structured products leaves the seller of the product effectively short an OTM put. A short OTM put is short skew and short vega convexity (or volga). This is a simplification, as structured products tend to buy visually cheap options (ie, OTM options) and sell visually expensive options (ie, ATM options), leaving the seller with a long ATM and short OTM volatility position. As OTM options have more volga (or vega convexity) than ATM options (see the section *A8 Greeks and Their Meaning* in the *Appendix*) the seller is short volga. The embedded option in structured products is floored, which causes the seller to be short skew (as the position is similar to being short an OTM put).

3.5: STRETCHING BLACK-SCHOLES ASSUMPTIONS

The Black-Scholes model assumes an investor knows the future volatility of a stock, in addition to being able to continuously delta hedge. In order to discover what the likely profit (or loss) will be in reality, we stretch these assumptions. If the future volatility is unknown, the amount of profit (or loss) will vary depending on the path, but buying cheap volatility will always show some profit. However, if the position is delta-hedged discretely, the purchase of cheap volatility may reveal a loss. The variance of discretely delta-hedged profits can be halved by hedging four times as frequently. We also show why traders should hedge with a delta calculated from expected – not implied – volatility, especially when long volatility.

MODEL ASSUMES KNOWN VOL AND CONTINUOUS HEDGING

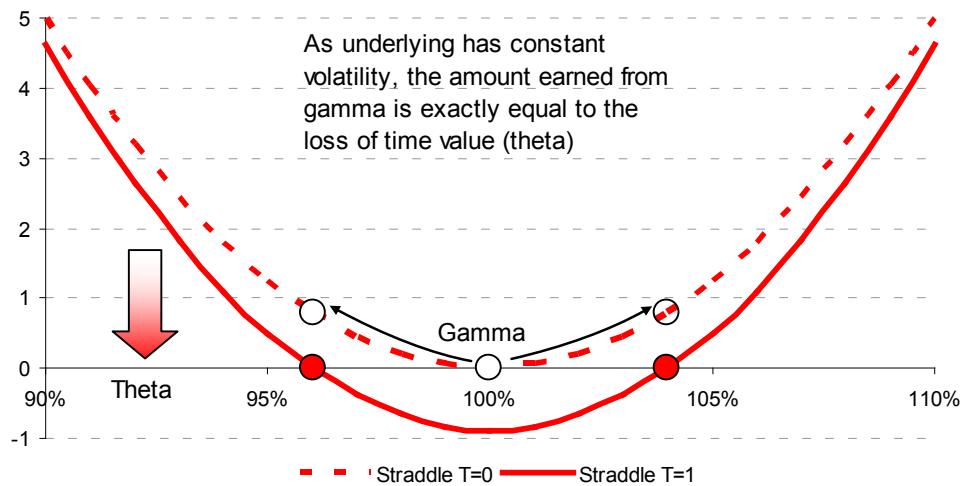
While there are a number of assumptions behind Black-Scholes, the two which are the least realistic are: (1) a continuous and known future realised volatility; and (2) an ability to delta hedge continuously. There are, therefore, four different scenarios to investigate. We assume that options are European.

- **Continuous delta hedging with known volatility.** In this scenario, the profit (or loss) from volatility trading is fixed. If the known volatility is constant, then the assumptions are identical to Black-Scholes. Interestingly, the results are the same if volatility is allowed not to be constant (while still being known).
- **Continuous delta hedging with unknown volatility.** If volatility is unknown, then typically traders hedge with the delta calculated using implied volatility. However, as implied volatility is not a perfect predictor of future realised volatility, this causes some variation in the profit (or loss) of the position. However, with these assumptions, if realised volatility is above the implied volatility price paid, it is impossible to suffer a loss.
- **Discrete delta hedging with known volatility.** As markets are not open 24/7, continuous delta hedging is arguably an unreasonable assumption. The path dependency of discrete delta hedging adds a certain amount of variation in profits (or losses), which can cause the purchase of cheap volatility (implied less than realised) to suffer a loss. The variance of the payout is inversely proportional to the frequency of the delta hedging. For example, the payout from hedging four times a day has a variance that is a quarter of the variance that results if the position is hedged only once a day.
- **Discrete delta hedging with unknown volatility.** The most realistic assumption is to hedge discretely with unknown volatility. In this case, the payout of volatility trading is equal to the sum of the variance due to hedging with unknown volatility plus the variance due to discretely delta hedging.

(1) CONTINUOUS DELTA HEDGING WITH KNOWN VOLATILITY

In a Black-Scholes world, the volatility of a stock is constant and known. While a trader is also able to continuously delta hedge, Figure 51 below will assume we hedge discretely but in an infinitesimally small amount of time. In each unit of time, the stock can either go up or down. As the position is initially delta-neutral (ie, delta is zero), the gamma (or convexity) of the position gives it a profit for both downward and upward movements. While this effect is always profitable, the position does lose time value (due to theta). If an option is priced using the actual fixed constant volatility of the stock, the two effects cancel each other and the position does not earn an abnormal profit or loss as the return is equal to the risk-free rate. There is a very strong relationship between gamma and theta (theta pays for gamma)¹⁰.

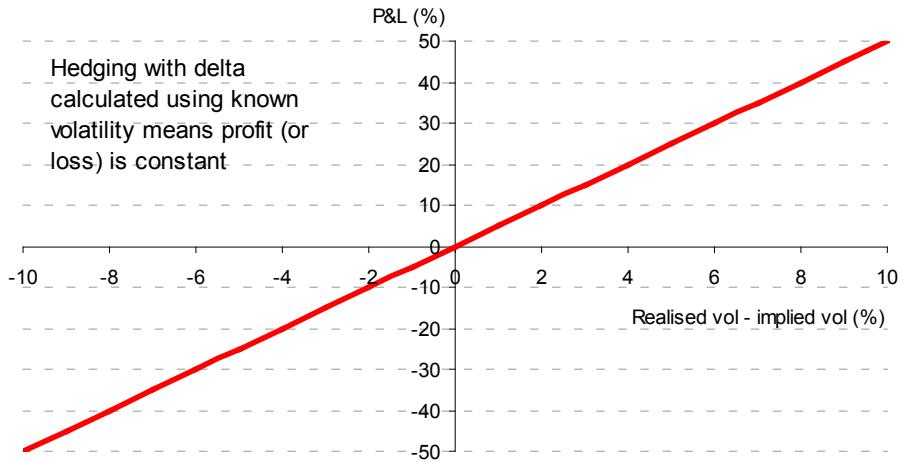
Figure 51. Constant and Known Realised Volatility to Calculate Delta



Profit from delta hedging is equal to the difference between price and theoretical price

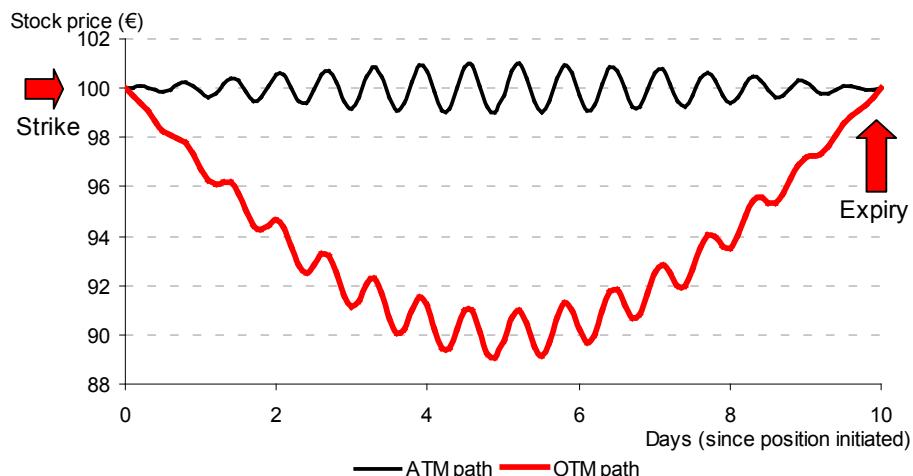
The theoretical price of an option, using the known volatility, can be extracted by delta hedging. Should an option be bought at an implied volatility less than realised volatility, the difference between the theoretical price and the actual price will equal the profit of the trade. Figure 52 below shows the profit vs the difference in implied and realised volatility. As there is no path dependency, the profit (or loss) of the trade is fixed and cannot vary.

¹⁰ They are not perfectly correlated, due to the interest paid on borrowing the shares (which varies with spot).

Figure 52. Profit (or Loss) from Continuously Delta-Hedging Known Volatility

As theta and gamma are correlated, profits are not path dependent

If a position is continuously delta hedged with the correct delta (calculated from the known future volatility over the life of the option), then the payout is not path dependent. Figure 53 below shows two paths with equal volatility and the same start and end point. Even though one path is always ATM while the other has most volatility OTM, delta hedging gives the same profit for both. This is due to the fact that, while the ATM option earns more due to delta hedging, the total theta cost is also higher (and exactly cancels the delta hedging profit).

Figure 53. Two Security Paths with Identical Volatility, Start and End Points

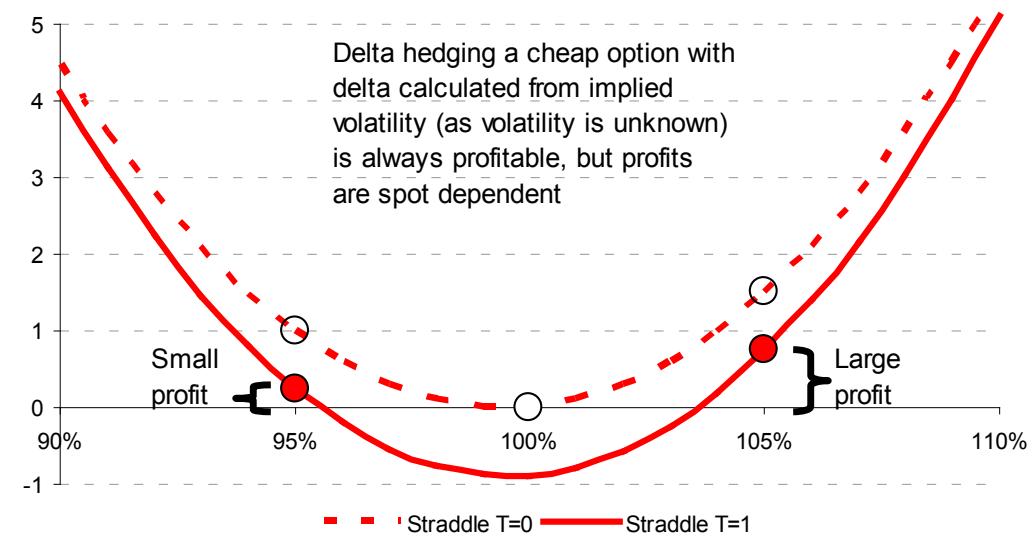
Profits are path independent, even if vol is not constant (but still known)

While Black-Scholes assumes a constant known volatility, there are similar results for non-constant known volatility. This result is due to the fact that a European option payout depends only on the stock price at expiry. Therefore, the volatility over the life of the option is the only input to pricing. The timing of this volatility is irrelevant.

(2) CONTINUOUS DELTA HEDGING WITH UNKNOWN VOLATILITY

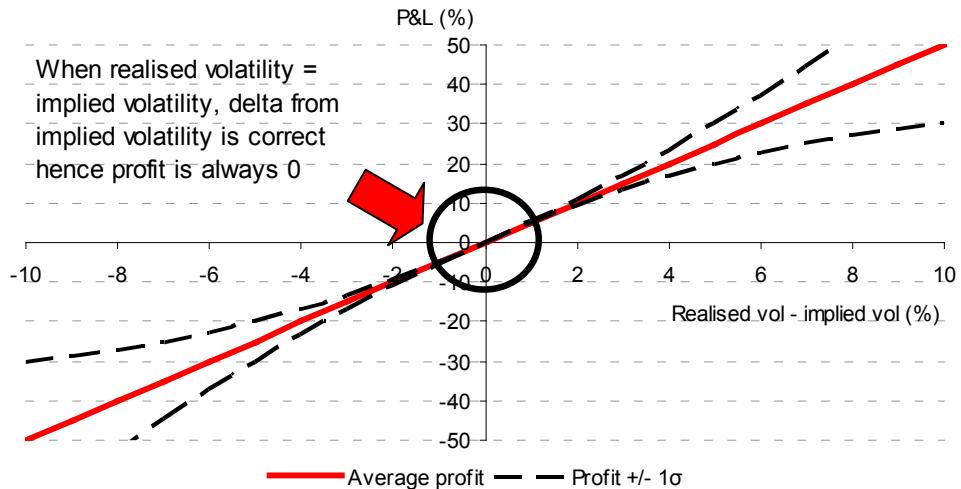
As it is impossible to know in advance what the future volatility of a security will be, the implied volatility is often used to calculate deltas. Delta hedging using this estimate causes the position to have equity market risk and, hence, it becomes path dependent (although the average or expected profit remains unchanged). Figure 54 below shows that the profits from delta hedging are no longer independent of the direction in which the underlying moves. The fact that there is a difference between the correct delta (calculated using the remaining volatility to be realised over the life of the option) and the delta calculated using the implied volatility means returns are dependent on the direction of equity markets.

Figure 54. Profit from Cheap Options Is Not Constant if Volatility Is Not Known



If implied volatility = realised volatility, profits are path independent

If the implied volatility is equal to the realised volatility, then the estimated delta calculated from the implied will be equal to the actual delta (calculated from the realised). In this case, profits from hedging will exactly match the theta cost for all paths, so it is path independent.

Figure 55. Profit (or Loss) from Continuously Delta Hedging Unknown Volatility

With continuous hedging buying a cheap option is always profitable

If there is a difference between the actual delta and estimated delta, there is market risk but not enough to make a cheap option unprofitable (or an expensive option profitable). This is because in each infinitesimally small amount of time a cheap option will always reveal a profit from delta hedging (net of theta), although the magnitude of this profit is uncertain. The greater the difference between implied and realised, the greater the market risk and the larger the potential variation in profit.

(3) DISCRETE DELTA HEDGING WITH KNOWN VOLATILITY

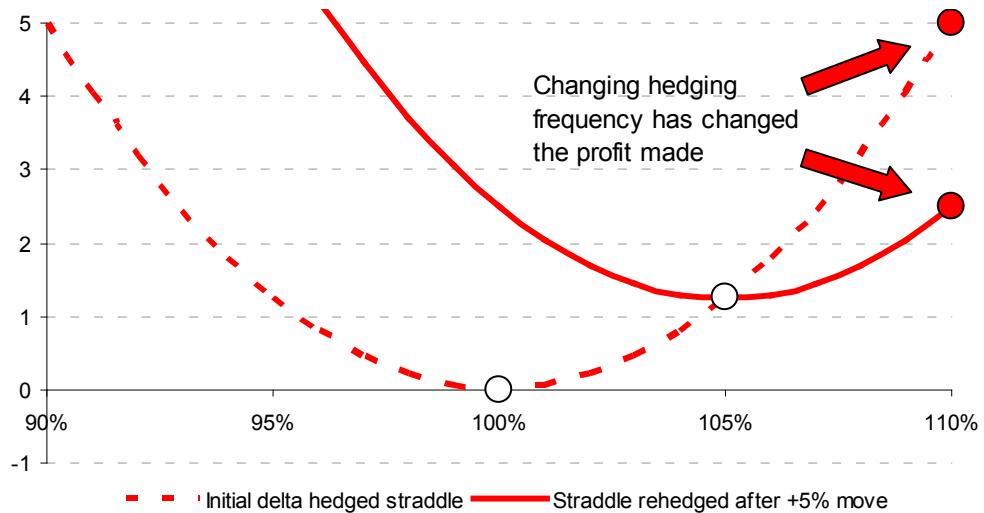
While assuming continuous delta hedging is mathematically convenient, it is impossible in practice. Issues such as the cost of trading and minimum trading size (even if this is one share) make continuous trading impossible, as do fundamental reasons, such as trading hours (if you cannot trade 24 hours then it is impossible to trade overnight and prices can jump between the close of one day and start of another) and weekends.

Discrete hedging errors can be reduced by increasing the frequency of hedging

The more frequent the discrete hedging, the less variation in the returns. If 24-hour trading were possible, then with an infinite frequency of hedging with known volatility the returns converge to the same case as continuous hedging with known volatility (ie, Black-Scholes).

In order to show how the frequency of hedging can affect the payout of delta hedging, we shall examine hedging for every 5% and 10% move in spot.

Figure 56. Profit from Discrete Delta Hedging with Different Frequencies



Hedging every 5% move in spot

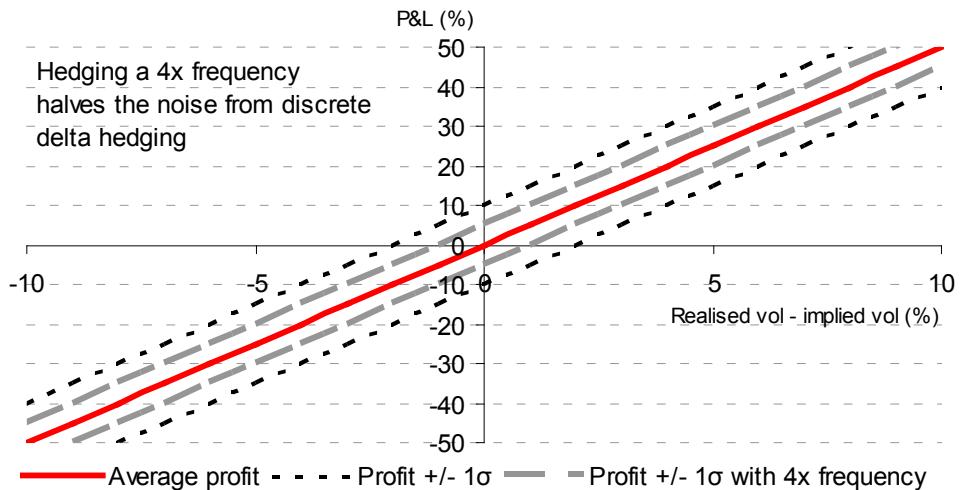
If an investor delta hedges every 5% move in spot, then an identical profit is earned if the underlying rises 10% as if the underlying rises 5% and then returns to its starting point. This shows that the hedging frequency should ideally be frequent enough to capture the major turning points of an underlying.

Hedging every 10% move in spot

If the investor is hedged for every 10% move in the underlying, then no profit will be earned if the underlying rises 5% and then returns to its starting point. However, if the underlying rises 10%, a far larger profit will be earned than if the position was hedged every 5%. This shows that in trending markets it is more profitable to let positions run than to re-hedge them frequently.

Hedging error is independent of average profitability of trade

As the volatility of the underlying is known, there is no error due to the calculation of delta. As the only variation introduced is essentially ‘noise’, the size of this noise, or variation, is independent from the average profitability (or difference between realised vol and implied vol) of the trade.

Figure 57. Profit (or Loss) from Discrete Delta Hedging Known Volatility

With discrete hedging, cheap options can lose money

With continuous delta hedging (with known or unknown volatility) it is impossible to lose money on a cheap option (an option whose implied volatility is less than the realised volatility over its life). However, as the error from discrete hedging is independent from the profitability of the trade, it is possible to lose money on a cheap option (and make money on an expensive option).

Hedging error is halved if frequency of hedging increased by factor of four

The size of the hedging error can be reduced by increasing the frequency of hedging. An approximation (shown below) is that if the frequency of hedging is increased by a factor of four, the hedging error term halves. This rule of thumb breaks down for very high-frequency hedging, as no frequency of hedging can eliminate the noise from non-24x7 trading (it will always have noise, due to the movement in share prices from one day's close to the next day's open).

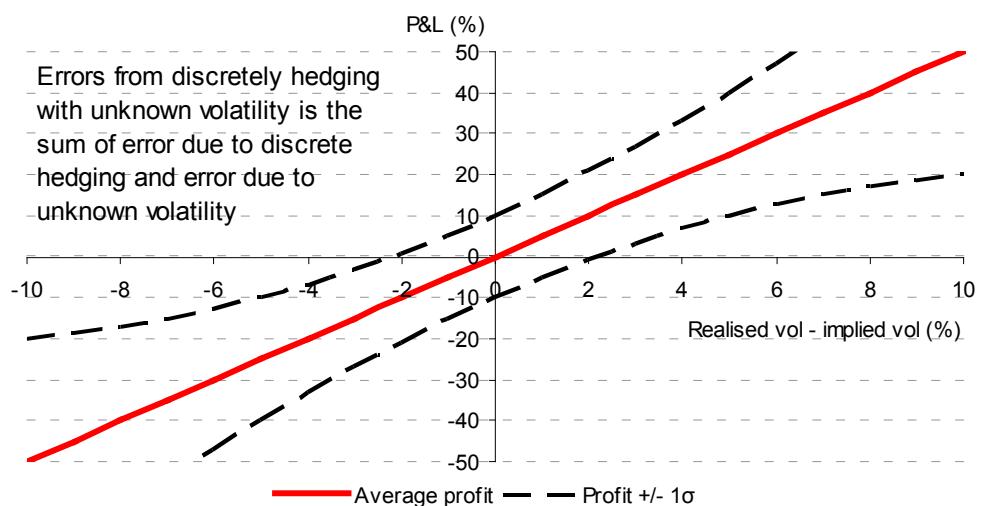
$$\sigma_{P\&L} \approx \sigma \times vega \times \sqrt{\frac{\pi}{4N}}$$

where N is the number of times position is hedged in a year

(4) DISCRETE DELTA HEDGING WITH UNKNOWN VOLATILITY

The most realistic assumption for profitability comes from the combination of discrete delta hedging and unknown volatility. Trading hours and trading costs are likely to limit the frequency at which a trader can delta hedge. Equally, the volatility of a stock is unknown, so implied volatility is likely to be used to calculate the delta. The variation in the profit (or loss) is caused by the variation due to discrete hedging and the inaccuracy of the delta (as volatility is unknown). Figure 58 below shows this combined variation in profit (or loss) and, as for discrete hedging with known volatility, it is possible for a delta hedged cheap option to reveal a loss.

Figure 58. Profit (or Loss) from Discrete Delta Hedging with Unknown Volatility



CAN YOU LOSE MONEY BUYING AN OPTION 20 PTS CHEAP?

As an example, let us assume a Dec08 SX5E ATM straddle was purchased in April 2008 when the implied volatility was a low 22%. This position was delta hedged at the close using the delta calculated from the implied volatility of the option (using implied volatility as an estimate of future volatility is standard market practice for calculating Greeks). In theory, this strategy would have been very profitable as the future realised volatility between April 2008 and Dec08 expiry was 42% (20 points higher, and almost double, the implied volatility paid for the straddle). As markets did not anticipate the fact that the future realised volatility would (by some measures) rise to a greater level than in the great depression, the implied volatility of the option was significantly different from the actual future realised volatility. Using implied volatility as an incorrect future volatility assumption to calculate the delta led to a significant loss to what should have been a highly profitable strategy.

Why delta hedging an option 20pts cheap can still cause a loss

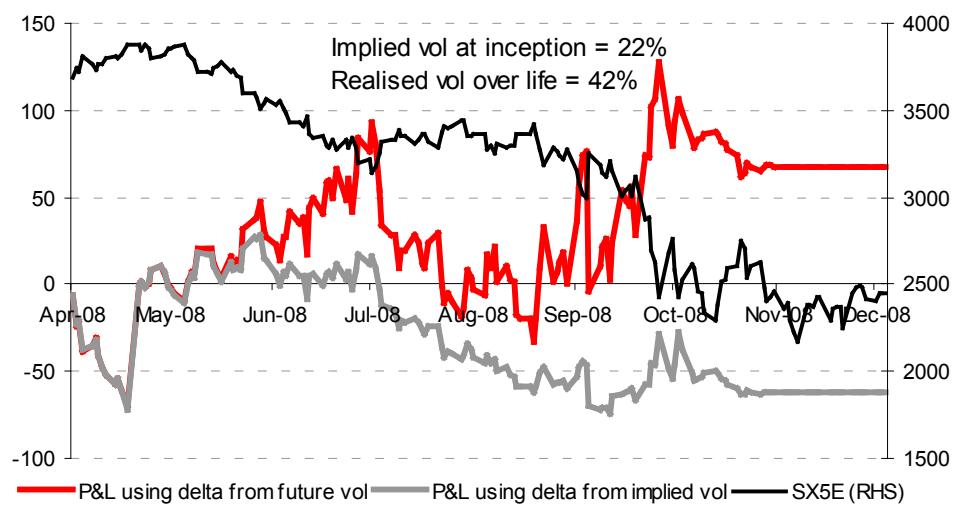
Delta hedging a long Dec08 ATM straddle at 22% implied when the future realised volatility is 42% should have been a very profitable strategy. However, most of the volatility came after the Lehman bankruptcy, which occurred towards the end of the option's life. As equity markets had declined since April 2008, the strike of the straddle would be significantly above spot when Lehman went bankrupt. If implied volatility was used to calculate the delta, then the time value would be assumed to be near zero. The delta of the straddle would therefore be $\approx -100\%$ (the call would be OTM with a delta ≈ 0 , while the put would be ITM with a delta $\approx 100\%$).

To be delta-hedged, the investor would then buy 100% of the underlying per straddle. If the delta was calculated using the actual volatility (which was much higher), then the time value would be higher and the delta greater than -100% (eg, -80%). As the delta-hedged investor would have bought 20% less of the underlying per straddle, when the market fell after Lehman collapsed this position outperformed hedging with implied volatility.

Investors should use expected vol, not implied vol, to calculate Greeks

These results can be seen in Figure 59 below, which gives a clear example of why traders should hedge with the delta calculated from expected volatility rather than implied volatility. Because of the extreme volatility at the end of 2008, the two deltas differed at times by 24% (60% vs 84%).

Figure 59. Discrete Delta Hedging a SX5E Dec08 Straddle



Hedging with delta using implied volatility is bad for long vol strategies

Typically, when volatility rises there is often a decline in the markets. The strikes of the option are therefore likely to be above spot when actual volatility is above implied. This reduces the profits of the delta-hedged position as the position is actually long delta when it appears to be delta flat. Alternatively, the fact that the position hedged with the realised volatility over the life of the option is profitable can be thought of as due to the fact it is properly gamma hedged, as it has more time value than is being priced into the market. Hence, if a trader buys an option when the implied looks 5pts too cheap, then the delta should be calculated from a volatility 5pts above current implied volatility. Using the proper volatility assumption to calculate the delta means the profit from delta hedging an option is approximately the difference between the theoretical value of the option at inception (ie, using actual future realised volatility in pricing) and the price of the option (ie, using implied volatility in pricing).

While this analysis has concentrated on delta, a similar logic applies for the calculation of all Greeks.

CHAPTER 4

FORWARD STARTING PRODUCTS AND VOLATILITY INDICES

Forward starting options are a popular method of trading forward volatility and term structure as there is no exposure to near-term volatility and, hence, zero theta (until the start of the forward starting option). Recently, trading forward volatility via volatility futures such as VIX and vStoxx futures has become increasingly popular. However, as is the case with forward starting options, there are modelling issues.

4.1: FORWARD STARTING PRODUCTS

Forward starting options are a popular method of trading forward volatility and term structure as there is no exposure to near-term volatility and, hence, zero theta (until the start of the forward starting option). As the exposure is to forward volatility rather than volatility, more sophisticated models need to be used to price them than ordinary options. Forward starting options will usually have wider bid-offer spreads than vanilla options, as their pricing and hedging is more complex. Recently, trading forward volatility via VIX and vStoxx futures has become increasingly popular, however there are modelling issues. Forward starting variance swaps are easier to price as the price is determined by two variance swaps.

ZERO THETA IS ADVANTAGE OF FWD STARTING PRODUCTS

The main attraction of forward starting products is that they provide investors with long-term volatility (or vega) exposure, without having exposure to short-term volatility (or gamma)¹¹. As there is zero gamma until the forward (fwd) starting product starts, the product does not have to pay any theta. Forward starting products are most appropriate for investors who believe that there is going to be volatility in the future (eg, during a key economic announcement or a reporting date) but that realised volatility is likely to be low in the near term (eg, over Christmas or the summer lull).

Forward starting products are low cost, but also lower payout

We note that while forward starting products have a lower theta cost than vanilla options, if there is a rise in volatility surfaces before the forward starting period is over, they are likely to benefit less than vanilla options (this is because the front end of volatility surfaces tends to move the most, and this is the area to which forward start has no sensitivity). Forward starting products can therefore be seen as a low-cost, lower-payout method of trading volatility.

TERM STRUCTURE PENALISES FWD STARTING PRODUCTS

While forward starting products have zero mathematical theta, they do suffer from the fact that volatility and variance term structure is usually expensive and upward sloping. The average implied volatility of a forward starting product is likely to be higher than a vanilla product, which will cause the long forward starting position to suffer carry as the volatility is re-marked lower¹² during the forward starting period.

SKEW CAUSES NEGATIVE SHADOW DELTA

¹¹ We shall assume that the investor wishes to be long a forward starting product.

¹² If a 3-month forward starting option is compared to a 3-month vanilla option, then during the forward starting period the forward starting implied volatility should, on average, decline.

The presence of skew causes a correlation between volatility and spot. This correlation produces a negative shadow delta for all forward starting products (forward starting options have a theoretical delta of zero). The rationale is similar to the argument that variance swaps have negative shadow delta due to skew.

FIXED DIVIDENDS ALSO CAUSES SHADOW DELTA

If a dividend is fixed, then the dividend yield tends to zero as spot tends to infinity, which causes a shadow delta (which is positive for calls and negative for puts).

Proportional dividends reduce volatility of underlying

Options, variance swaps and futures on volatility indices gain in value if dividends are fixed, as proportional dividends simply reduce the volatility of an underlying.

THERE ARE 3 MAIN METHODS TO TRADE FORWARD VOL

Historically, forward volatility could only be traded via forward starting options, which had to be dynamically hedged and, hence, had high costs and wide bid-offer spreads. When variance swaps became liquid, this allowed the creation of forward starting variance swaps (as a forward starting variance can be perfectly hedged by a long and short position in two vanilla variance swaps of different maturity, which is explained later). The client base for trading forward volatility has recently been expanded by the listing of forwards on volatility indices (such as the VIX or vStoxx). The definition of the three main forward starting products is given below:

- (1) **Forward starting options.** A forward starting option is an option whose strike will be determined at the end of the forward starting period. The strike will be quoted as a percentage of spot. For example, a one-year ATM option three-month forward start, bought in September 2012, will turn into a one-year ATM option in December 2012 (ie, expiry will be December 2013 and the strike will be the value of spot in December 2012). Forward starting options are quoted OTC. For flow client requests, the maturity of the forward starting period is typically three months and with an option maturity no longer than a year. The sale of structured products creates significant demand for forward starting products, but of much longer maturity (2-3 years, the length of the structured product). Investment banks will estimate the size of the product they can sell and buy a forward starting option for that size. While the structured product itself does not incorporate a forward start, as the price for the product needs to be fixed for a period of 1-2 months (the marketing period), the product needs to be hedged with a forward start before marketing can begin.
- (2) **Forward starting variance swaps.** The easiest forward starting product to trade is a variance swap, as it can be hedged with two static variance swap positions (one long, one short). Like plain variance swaps, these products are traded OTC and their maturities

can be up to a similar length (although investors typically ask for quotes up to three years).

- (3) **Futures on volatility index.** A forward on a volatility index works in the same way as a forward on an equity index: they both are listed and both settle against the value of the index on the expiry date. While forwards on volatility indices such as the VIX and vStoxx have been quoted for some time, their liquidity has only recently improved to such an extent that they are now a viable method for trading. This improvement has been driven by increasing structured issuance and by options on volatility indices (delta hedging of these options has to be carried out in the forward market). Current listed maturities for the VIX and vStoxx exist for expiries under a year.

HIGHER HEDGING RISKS INCREASE COST

While forward starting options do not need to be delta hedged before the forward starting period ends, they have to be vega hedged with vanilla straddles (or very OTM strangles if they are liquidity enough, as they also have zero delta and gamma). A long straddle has to be purchased on the expiry date of the option, while a short straddle has to be sold on the strike fixing date. As spot moves the strikes will need to be rolled, which increases costs (which are likely to be passed on to clients) and risks (unknown future volatility and skew) to the trader.

Pricing of futures on vol indices tends to be slanted against long investors

Similarly, the hedging of futures on volatility indices is not trivial, as (like volatility swaps) they require a volatility of volatility model. While the market for futures on volatility indices has become more liquid, as the flow is predominantly on the buy side, forwards on volatility indices have historically been overpriced. They are a viable instrument for investors who want to short volatility, or who require a listed product.

Forward starting var swaps have fewer imbalances than other products

The price – and the hedging – of a forward starting variance swap is based on two vanilla variance swaps (as it can be constructed from two vanilla variance swaps). The worst-case scenario for pricing is therefore twice the spread of a vanilla variance swap. In practice, the spread of a forward starting variance swap is usually slightly wider than the width of the widest bid-offer of the variance swap legs (ie, slightly wider than the bid-offer of the furthest maturity).

(1) FORWARD STARTING OPTIONS

A forward starting option can be priced using Black-Scholes in a similar way to a vanilla option. The only difference is that the forward volatility (rather than volatility) is needed as an input¹³. The three different methods of calculating the forward volatility, and examples of how the volatility input changes, are detailed below:

- **Sticky delta (or moneyness) and relative time.** This method assumes volatility surfaces never change in relative dimensions (sticky delta and relative time). This is not a realistic assumption unless the ATM term structure is approximately flat.
- **Additive variance rule.** Using the additive variance rule takes into account the term structure of a volatility surface. This method has the disadvantage that the forward skew is assumed to be constant in absolute (fixed) time, which is not usually the case. As skew is normally larger for shorter-dated maturities, it should increase approaching expiry.
- **Constant smile rule.** The constant smile rule combines the two methods above by using the additive variance rule for ATM options (hence, it takes into account varying volatility over time) and applying a sticky delta skew for a relative maturity. It can be seen as ‘bumping’ the current volatility surface by the change in ATM forward volatility calculated using the additive variance rule.

STICKY DELTA & RELATIVE TIME USES CURRENT VOL SURFACE

If the relative dimensions of a volatility surface are assumed to never change, then the volatility input for a forward starting option can be priced with the current volatility surface. For example, a three-month 110% strike option forward starting after a period of time T can be priced using the implied volatility of a current three-month 110% strike option (the forward starting time T is irrelevant to the volatility input¹⁴). As term structure is normally positive, this method tends to underprice forward starting options. An example of a current relative volatility surface, which can be used for pricing forward starting options under this method, is shown below:

¹³ Forwards of the other inputs, for example interest rates, are generally trivial to compute.

¹⁴ Hence, the price of the three-month 110% option forward start will only be significantly different from the price of the vanilla three-month 110% option if there is a significant difference in interest rates or dividends.

Figure 60. Relative Dimensions / Current Implied Volatility Surface

Strike	End	Start	Now	Now	Now	Now
		1 Year	2 Years	3 Years	4 Years	
80%		24.0%	23.4%	23.2%	23.0%	
90%		22.0%	22.0%	22.0%	22.0%	
100%		20.0%	20.6%	20.8%	21.0%	
110%		18.0%	19.2%	19.7%	20.0%	
120%		16.0%	17.8%	18.5%	19.0%	

ADDITIVE VARIANCE RULE CALCULATES FWD VOL

As variance time weighted is additive, and as variance is the square of volatility, the forward volatility can be calculated mathematically. Using these relationships to calculate forward volatilities is called the additive variance rule and is shown below.

$$\sigma_2^2 T_2 = \sigma_1^2 T_1 + \sigma_{12}^2 (T_2 - T_1) \text{ as variance time weighted is additive}$$

$$\sigma_{12} = \sqrt{\frac{\sigma_2^2 T_2 - \sigma_1^2 T_1}{T_2 - T_1}} = \text{forward volatility } T_1 \text{ to } T_2$$

where σ_i is the implied volatility of an option of maturity T_i

The above relationship can be used to calculate forward volatilities for the entire volatility surface. This calculation does assume that skew in absolute (fixed) time is fixed. An example, using the previous volatility surface, is shown below.

Figure 61. One Year Additive Variance Rule Forward Volatility Surface

Strike	End	Start	Now	1 Year	2 Years	3 Years
		1 Year	2 Years	3 Years	4 Years	
80%		24.0%	22.8%	22.6%	22.5%	
90%		22.0%	22.0%	22.0%	22.0%	
100% AVR 1 Year Fwd Volatility		20.0%	21.2%	21.4%	21.5%	
110%		18.0%	20.3%	20.7%	20.9%	
120%		16.0%	19.4%	20.0%	20.3%	

ATM ADDITIVE VAR + STICKY DELTA = CONST SMILE RULE

Using a relative time rule has the advantage of pricing forward skew in a reasonable manner, but it does not price the change in term structure correctly. While pricing using the additive variance rule gives improved pricing for ATM options, for OTM options the skew used is likely to be too low (as the method uses forward skew, which tends to decay by square root of time). The constant smile rule combines the best features of the previous two approaches, with ATM options priced using the additive variance rule and the skew priced using sticky delta.

Figure 62. Current Volatility Surface

Strike	1 Year	2 Years	3 Years	4 Years	1 Year Skew
80%	24.0%	23.4%	23.2%	23.0%	4.0%
90%	22.0%	22.0%	22.0%	22.0%	2.0%
100%	20.0%	20.6%	20.8%	21.0%	0.0%
110%	18.0%	19.2%	19.7%	20.0%	-2.0%
120%	16.0%	17.8%	18.5%	19.0%	-4.0%

One Year Additive Variance Rule (AVR) Forward Volatility

	Start	Now	1 Year	2 Years	3 Years
	End	1 Year	2 Years	3 Years	4 Years
100% AVR 1 Year Fwd Volatility		20.0%	21.2%	21.4%	21.5%



One Year Constant Smile Rule Forward Volatility Surface

Strike	Start	Now	1 Year	2 Years	3 Years	1 Year skew
	End	1 Year	2 Years	3 Years	4 Years	
80%		24.0%	25.2%	25.4%	25.5%	4.0%
90%		22.0%	23.2%	23.4%	23.5%	2.0%
100% AVR 1 year Fwd Volatility		20.0%	21.2%	21.4%	21.5%	0.0%
110%		18.0%	19.2%	19.4%	19.5%	-2.0%
120%		16.0%	17.2%	17.4%	17.5%	-4.0%

Constant smile rule bumps sticky delta relative time volatility surface

The above diagrams show how the constant smile rule has the same ATM forward volatilities as the additive variance rule. The static delta (relative time) skew is then added to these ATM options to create the entire surface. An alternative way of thinking of the surface is that it takes the current volatility surface, and shifts (or bumps) each maturity by the exact amount required to get ATM options to be in line with the additive variance rule. The

impact of having a relative time skew on a fixed ATM volatility can be measured by volatility slide theta (see the section *A9.Advanced (Practical or Shadow) Greeks* in the *Appendix*).

CONSTANT SMILE RULE IS THE BEST MODEL OF THE THREE

Pricing with static delta and relative time usually underprices forward volatility (as volatility term structure is normally upward sloping, and long-dated forward volatility is sold at the lower levels of near-dated implied volatility). While additive variance correctly prices forward volatility, this rule does mean future skew will tend towards zero (as skew tends to decay as maturity increases and the additive variance rule assumes absolute – fixed – time for skew). While this rule has been used in the past, the mispricing of long-dated skew for products such as cliques has led traders to move away from this model. The constant smile rule would appear to be the most appropriate.

(2) FORWARD STARTING VAR SWAPS

In the section *A2.Measuring Historical Volatility* in the *Appendix* we show that variance is additive (variance to time T_2 = variance to time T_1 + forward variance T_1 to T_2). This allows the payout of a forward starting variance swap between T_1 and T_2 to be replicated via a long variance swap to T_2 , and short variance swap to T_1 . We define N_1 and N_2 to be the notional of the variance swaps to T_1 and T_2 , respectively. It is important to note that N_1 and N_2 are the notional of the variance swap, not the vega ($N = \text{vega} \div 2\sigma$). As the variance swap payout of the two variance swaps must cancel up to T_1 , the following relationship is true (we are looking at the floating leg of the variance swaps, and ignore constants that cancel such as the annualisation factor):

$$\text{Payout long var to } T_2 = \text{Payout long var to } T_1 + \text{Payout long var } T_1 \text{ to } T_2$$

$$\Leftrightarrow N_2 \frac{\sum_{i=1}^{T_2} [\ln(\text{return}_i)]^2}{T_2} = N_2 \frac{\sum_{i=1}^{T_1} [\ln(\text{return}_i)]^2}{T_2} + N_2 \frac{\sum_{i=T_1+1}^{T_2} [\ln(\text{return}_i)]^2}{T_2}$$

$$\Leftrightarrow \text{Payout short variance} = N_2 \frac{\sum_{i=1}^{T_1} [\ln(\text{return}_i)]^2}{T_2} (= \text{Payout long variance to } T_1)$$

$$\Leftrightarrow N_1 \frac{\sum_{i=1}^{T_1} [\ln(\text{return}_i)]^2}{T_1} = N_2 \frac{\sum_{i=1}^{T_1} [\ln(\text{return}_i)]^2}{T_2}$$

$$\Rightarrow \frac{N_1}{T_1} = \frac{N_2}{T_2}$$

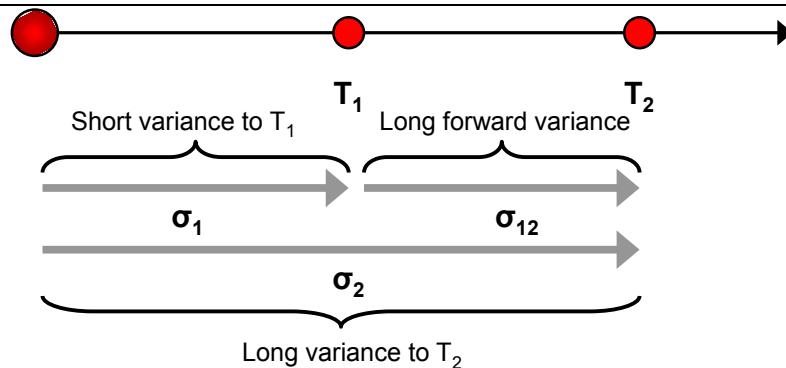
$\Rightarrow N_1 = \frac{T_1}{T_2} N_2$ = (Notional of near dated variance as factor of far dated variance notional)

$\Rightarrow 0 < N_1 < N_2$ (as $T_1 > T_2$)

Notional of the near-dated variance is smaller than notional of far-dated

The above proof shows that in order to construct a forward starting variance swap from two vanilla variance swaps, the near-dated variance should have a notional of T_1/T_2 (which is less than 1) of the notional of the far-dated variance. Intuitively, this makes sense as the near-dated variance swap to T_1 only needs to cancel the overlapping period of the longer-dated variance swap to T_2 . The notional of the near-dated variance swap to T_1 therefore has to be scaled down, depending on its relative maturity to T_2 . For example, if T_1 is 0, then there is no need to short any near-dated variance and N_1 is similarly zero. In addition, if $T_1 = T_2$, then the two legs must cancel, which occurs as $N_1 = N_2$. The notional N_{12} must be equal to the difference of the notionals of the two vanilla variance swaps that hedge it (ie, $N_{12} = N_2 - N_1$) by considering the floating legs and having constant realised volatility ($N_2 \sigma_2^2 = N_1 \sigma_1^2 + N_{12} \sigma_{12}^2$, hence $N_2 = N_1 + N_{12}$ if volatility σ^2 is constant).

Figure 63. Constructing Forward Variance from Vanilla Variance Swaps



CALCULATING FORWARD VARIANCE

The additive variance rule allows the level of forward variance to be calculated.

$$\sigma_2^2 T_2 = \sigma_1^2 T_1 + \sigma_{12}^2 (T_2 - T_1)$$

$$\Rightarrow \sigma_{12}^2 = \frac{\sigma_2^2 T_2 - \sigma_1^2 T_1}{T_2 - T_1} = \text{forward volatility from } T_1 \text{ to } T_2$$

Forward starting var swaps have fewer imbalances than other fwd products

The price – and the hedging – of a forward starting variance swap is based on two vanilla variance swaps (as it can be constructed from two vanilla variance swaps). The worst-case scenario for pricing is therefore twice the spread of a vanilla variance swap. In practice, the spread of a forward starting variance swap is usually slightly wider than the width of the widest bid-offer of the variance swap legs (ie, slightly wider than the bid-offer of the furthest maturity).

(3) FUTURE ON VOLATILITY INDEX

Futures on volatility indices have become one of the most popular forward starting products. For more details on both volatility indices, and futures on those indices, please see the following two sections [4.2: Volatility Indices](#) and [4.3: Futures On Volatility Indices](#).

4.2: VOLATILITY INDICES

While volatility indices were historically based on ATM implied, most providers have swapped to a variance swap-based calculation. The price of a volatility index will, however, typically be 0.2-0.7pts below the price of a variance swap of the same maturity as the calculation of the volatility index typically chops the tails to remove illiquid prices. Each volatility index provider has to use a different method of chopping the tails in order to avoid infringing the copyright of other providers.

THERE ARE 2 WAYS OF CALCULATING A VOLATILITY INDEX

Historically, volatility indices (old VIX and VDAX) were based on ATM implied volatility. This level is virtually identical to the fair price of a volatility swap (as volatility swaps \approx ATMf implied). This methodology has the advantage that it uses the most liquid strikes, and it is still used by some providers in less liquid markets for this reason. Due to the realisation that variance, not volatility, was the correct measure of deviation, on September 22, 2003, the VIX index moved away from using ATM implied towards a variance-based calculation (and also moved from using the S&P100 to the S&P500). While the calculation is variance-based, the index is quoted as the square root of variance for an easier comparison with the implied volatility of options (but we note that skew and convexity mean the fair price of variance swaps and volatility indices should always trade above ATM options).

Volatility indices based on ATM implied usually average 8 different options

The old VIX, renamed VXO, took the implied volatility for the S&P100 strikes above and below spot for both calls and puts. As the first two-month expiries were used, the old index was calculated using eight implied volatility measures as $8 = 2 \text{ (strikes)} \times 2 \text{ (put/call)} \times 2 \text{ (expiry)}$. Similarly, the VDAX index, which was based on DAX 45-day ATM-implied volatility, has been superseded by the V1X index, which, like the new VIX, uses a variance-based calculation.

MOST INDICES NOW USE VARIANCE-BASED CALCULATIONS

Variance-based calculations have also been used by other volatility index providers. All recent volatility indices, such as the vStoxx (V2X), VSMI (V3X), VFTSE, VNKY and VHSI, use a variance swap calculation, although we note the recent VIMEX index uses a similar methodology to the old VIX (due to illiquidity of OTM options on the Mexican index). While the formula for a variance is a mathematical formula and hence not subject to copyright, if this formula is modified to exclude tails (eg, requiring a non-zero bid and/or offer price, excluding strikes too far away from spot, etc), then this calculation becomes proprietary and is subject to copyright. This is the reason why different volatility index providers have chosen different calculation methods.

DIFFERENCES BETWEEN VOL INDICES AND VAR SWAPS

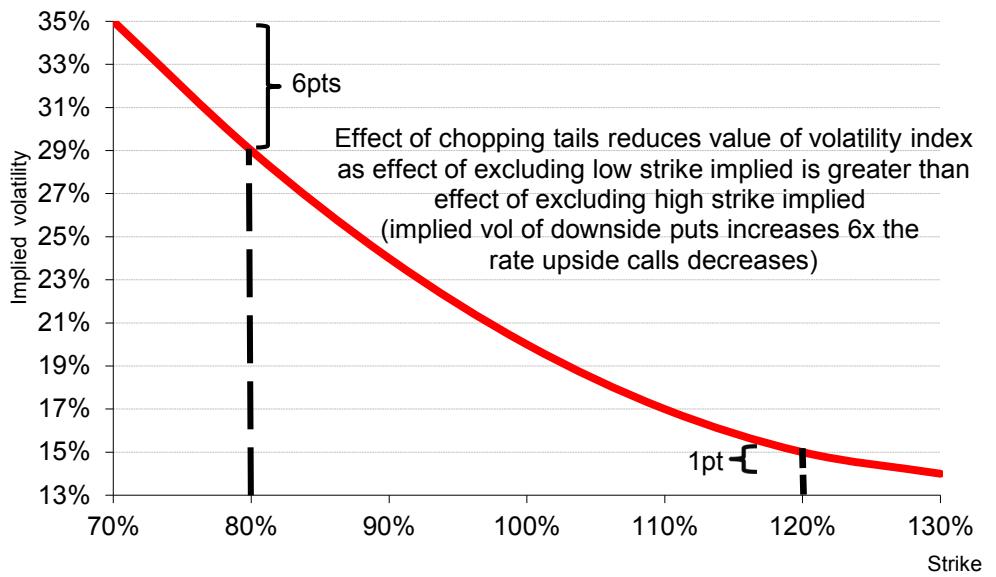
While the calculation of a volatility index might be based on a variance swap calculation, the price of a volatility index will typically be lower than that of a variance swap. The magnitude of the difference depends on the calculation itself, the number of strikes with available prices and the difference in width between strikes.

- **Excluding very high and very low strikes.** To increase the stability of the calculation, volatility indices exclude the implied volatility of options with very high or very low strikes. Given the importance of low strike implied volatility to variance swap pricing, chopping the wings of low strike implieds has a greater impact than removing high strike implieds, hence the level of a volatility index is below the variance swap price (typically between 0.2 and 0.7 volatility points).
- **Discrete sampling by using only listed strikes.** When pricing a variance swap, the value of a parameterised volatility surface is used. This surface is completely continuous and, hence, is not subject to errors due to using discrete data. As a volatility index has to rely on data from listed strikes, this introduces a small error which causes the level of the implied volatility index to be slightly below the variance swap price.
- **Noise due to rolling expiries.** If a volatility index does not interpolate between expiries then the implied volatility will ‘jump’ when the maturity rolls from one expiry to another. This difference can be c2 volatility points. Some indices only interpolate over a few days and take an exact maturity the rest of the time, which smoothes this effect (but does not fully remove it). Indices calculated by the CBOE move from interpolation to extrapolation which will cause similar noise, but has a much smaller effect than rolling. The average value from a volatility index that uses rolling is below the value of a variance swap as term structure is normally positive.
- **Linear interpolation between expiries (should be square root of time).** Linearly interpolating between expiries assumes a flat volatility term structure. In reality, a volatility surface follows a ‘square root of time’ rule, which means that the slope of term structure is steeper for near-dated maturities than for far dated ones. As a volatility surface is normally upward sloping this means a volatility index is on average below the level of a variance swap (up to c0.8 volatility points).

Excluding very high and very low strikes lowers the value of a vol index

ATM options are the most liquid, as they have the most time value. For very OTM options, not only is liquidity typically poor but a small change in price can have a large effect on the implied volatility. To improve reliability of calculation, the very high and very low strikes are excluded. This is either done via a fixed rule (ie, only use strikes between 80% and 120%) or by insisting on a bid price above zero. Requiring the existence of both bid and offer prices implicitly chops the wings as well. Excluding the tails excludes high implied volatility low strike options; which causes the level of the volatility index to be below the fair price of variance swaps. The difference depends on the size of the tail that is chopped. If only strikes between 80% and 120% are used this can cause a discount of c0.7 volatility points between the 1-month volatility index and 1-month variance swaps. If all strikes with a liquid price are used then typically prices can be available for strikes between c60% and c120%, as downside puts are more liquid than upside calls (due to increased demand from hedging and due to the higher premium value given higher implied volatility). Using strikes between c60% and c120% has a small discount to variance swaps of c0.2 volatility points. The VIX requires a non-zero bid and an offer, and stops when two consecutive options have no price.

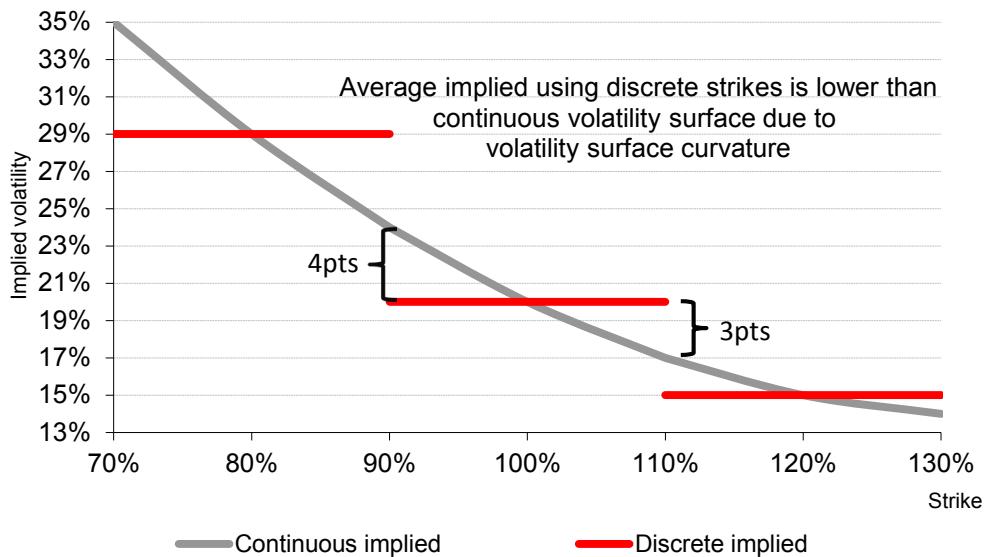
Figure 64. Chopping wings of vol surface is most important for low strikes



Discrete sampling by using only listed strikes lowers value of vol index

Even if prices were available for all strikes, a volatility index would give a slightly lower quote than variance swaps due to discretely sampling the implied volatility. The effect of discretely sampling a volatility surface can be modelled as a continuous volatility surface whose implied volatility is flat near the listed strikes (and jumps in between the listed strikes). Due to volatility surface curvature, this effect causes the value of a volatility index to be lower than a variance swap (as can be seen in Figure 65 below). The effect of discretely sampling depends on the number of strikes available (ie, the price difference between strikes), but is very small compared to the effect of chopping the tails (but both are caused by volatility surface curvature).

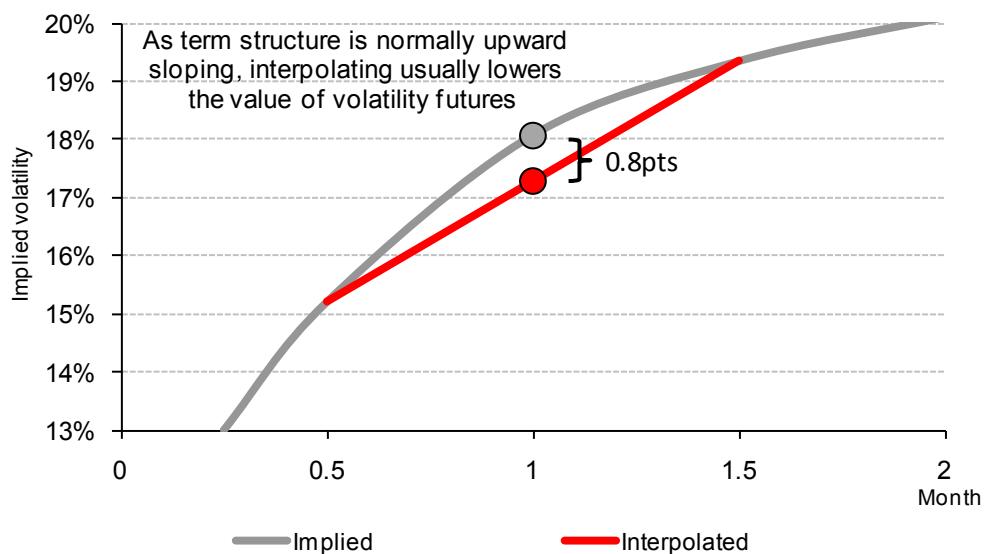
Figure 65. Discretely Sampling a Volatility Surface



Noise due to rolling depends on the calculation method

There are volatility indices that instead of linearly interpolating between expiries roll from one maturity to the other. For similar reasons as to why linearly interpolating a volatility index usually gives a lower value than a variance swap (as there is a greater difference between 0.5-month and 1-month implied than between 1-month and 1.5-month implied, i.e. the slope of term structure flattens as maturity increases), the average value for a volatility index that rolls is similarly too low. There will, however, be greater volatility for the index, due to the jump when the maturity rolls from one expiry to another. The difference between the implied volatility of the front two expiries can be c2 volatility points. Some indices smooth this effect by interpolating for a few days, while having a fixed un-interpolated value the rest of the time. Even a smoothed calculation will have a higher volatility over rolling than a fully interpolated based calculation. The CBOE ignores the front-month expiry in the final week before expiry and extrapolates from the second and third expiry. While this is a fully interpolated based calculation, jumping from using interpolation between the first and second expiry to extrapolation between the second and third can add some noise (but less noise than a roll-based calculation). We note that, as 30-day volatility futures expire exactly 30 days before the vanilla expiry, no interpolation by maturity is necessary for settlement.

Figure 66. Discretely Sampling a Volatility Surface



Interpolation between expiries usually lowers the value of a volatility index

The slope of near-dated implieds is typically steeper than far-dated implieds, as volatility surfaces often move in a ‘square root of time’ manner (near-dated implieds fall more than far-dated implieds when volatility declines). Given the steeper slope of near-dated implieds, linearly interpolating underestimates the implied volatility for positive term structure (and similarly overestimates it for negative term structure). Typically, the demand for long-dated hedges and risk aversion causes far-dated implieds to be greater than near-dated implieds, hence term structure is normally positive. The effect of linearly interpolating between maturities for a fixed maturity volatility index therefore causes a volatility index to normally underestimate a variance swap level. This effect will be greatest when the (typically 1-month) maturity of the volatility index is exactly in between listed expiries, and an extreme example is given in the graph above where the difference is c0.8 volatility points. We note that should the maturity of a volatility index lie close to a listed maturity the error due to interpolating between expiries will be close to zero.

4.3: FUTURES ON VOLATILITY INDICES

While futures on volatility indices were first launched on the VIX in March 2004, it has only been since the more recent launch of structured products and options on volatility futures that liquidity has improved enough to be a viable method of trading volatility. As a volatility future payout is based on the square root of variance, the payout is linear in volatility not variance. The fair price of a future on a volatility index is between the forward volatility swap, and the square root of the forward variance swap. Volatility futures are, therefore, short vol of vol, just like volatility swaps. It is therefore possible to get the implied vol of vol from the listed price of volatility futures.

PRICE IS BETWEEN FORWARD VAR AND FORWARD VOL

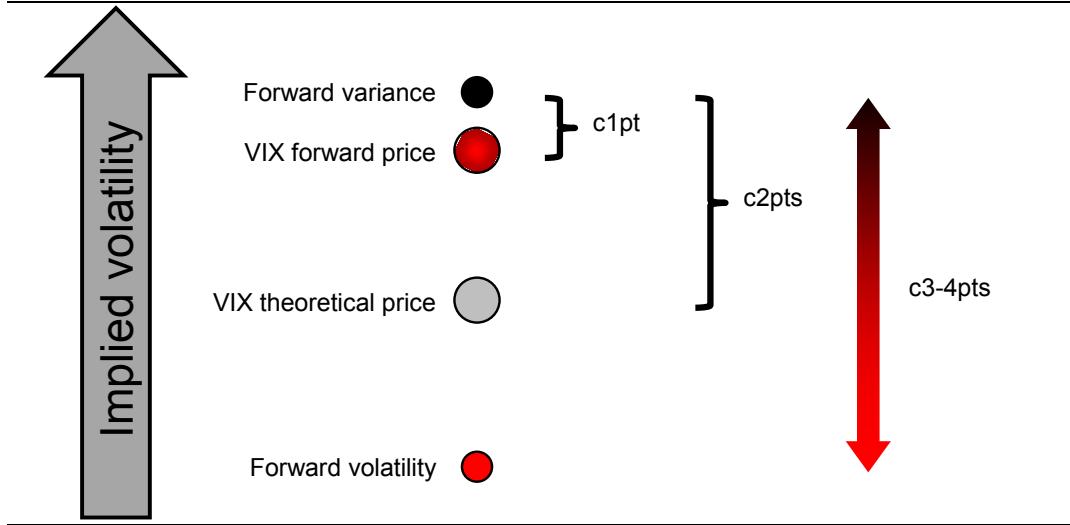
A future on a volatility index functions in exactly the same way as a future on an equity index. However, as volatility future is a forward (hence linear) payout of the square root of variance, the payoff is different from a variance swap (whose payout is on variance itself). The price of a forward on a volatility index lies between the fair value of a forward volatility swap and the square root of the fair value of a forward variance swap.

$$\sigma_{\text{Forward volatility swap}} \leq \text{Future on volatility index} \leq \sigma_{\text{Forward variance swap}}$$

FUTURES ON VOLATILITY INDICES ARE SHORT VOL OF VOL

A variance swap can be hedged by delta hedging a portfolio of options (the portfolio is known as a log contract, where the weight of each option is $1/K^2$ where K is the strike). As the portfolio of options does not change, the only hedging costs are the costs associated with delta hedging. A volatility swap has to be hedged through buying and selling variance swaps (or a log contract of options); hence, it needs to have a volatility of volatility model. As a variance swap is more convex than a volatility swap (variance swap payout is on volatility squared), a volatility swap is short convexity compared to a variance swap. A volatility swap is, therefore, short volatility of volatility (vol of vol) as a variance swap has no vol of vol risk. As the price of a future on a volatility index is linear in volatility, a future on a volatility index is short vol of vol (like volatility swaps).

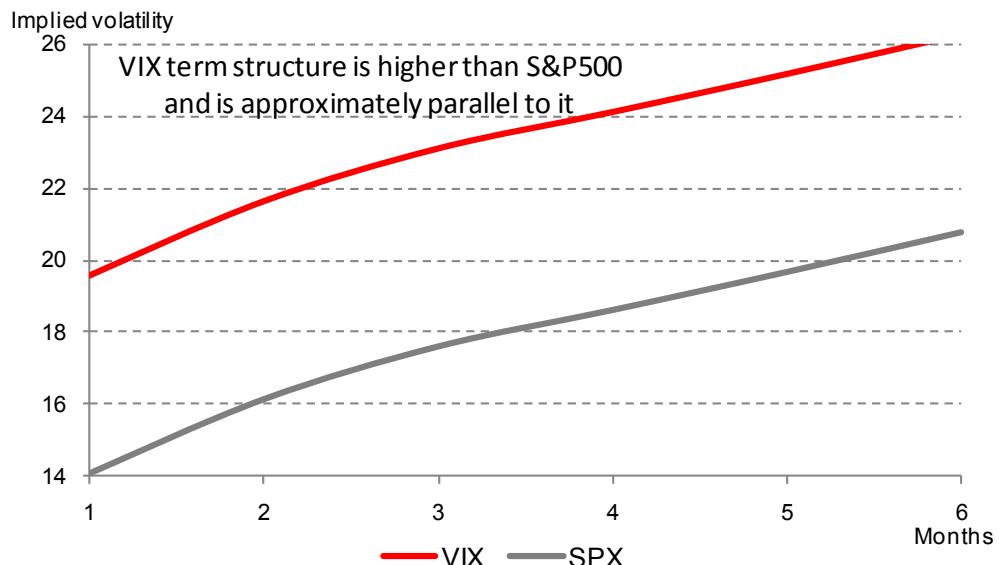
Figure 67. Theoretical (stochastic local vol) and Actual Prices of 6-Month VIX Futures



As vol of vol is underpriced, futures on volatility indices are overpriced

While the price of volatility futures should be well below that of forward variance swaps, retail demand and potential lack of knowledge of the client base means that they have traded at similar levels. Using a stochastic local volatility model, we found that VIX futures¹⁵ should trade roughly half way between a variance future and a volatility future (in fact, slightly closer to forward volatility, as is to be expected for a product linear in volatility). While this means VIX futures should be c2pts below forward variance, they appear to only trade c1pt below. Similarly, VIX futures should be only c1-2pts above ATMf implied volatility, but during 2012 they were c5-6pts above ATMf implied (see Figure 68 below). This overpricing of volatility futures means that volatility of volatility is underpriced in these products. Being short volatility futures and long forward variance is a popular trade to arbitrage this mispricing.

¹⁵ We assume a volatility index calculation matches that of a variance swap, ie, no chopping of tails.

Figure 68. VIX and S&P500 Average Term Structures

VIX and S&P500 term structures are roughly parallel to each other

We note that while the VIX futures term structure lies above S&P500 term structure, they are approximately parallel to each other for the same reason variance term structure is parallel to ATMf term structure. While variance swaps are long skew and skew is lower for far-dated implieds, as OTM options gain more time value as maturity increases, these effects cancel each other out (hence, in the absence of supply and demand imbalances, variance and implied volatility term structure should be roughly parallel to each other). As equities typically have an upward sloping term structure, volatility futures term structure is typically upward sloping as well. Volatility futures, like variance swaps, are also long volatility surface curvature. Volatility futures will have the same seasonality of vol as the underlying security (eg, dips over Christmas and year-end).

VOL OF VOL CAN BE BACKED OUT FROM VOL FUTURE PRICE

A forward on a volatility future is short vol of vol. This means it is possible to back out the implied vol of vol from the price of this volatility future. This implied vol of vol can be used to price options on variance or even options on volatility futures themselves¹⁶.

¹⁶ Assuming the volatility of volatility is log normally distributed.

EUREX WAS THE FIRST EXCHANGE TO LIST VOL FUTURES

While futures on the VIX (launched by the CBOE in March 2004) are the oldest currently traded, the DTB (now Eurex) was the first exchange to list volatility futures, in January 1998. These VOLAX futures were based on 3-month ATM implieds but they ceased trading in December of the same year. More recently, futures based on the Russell 2000 traded from 2007 until their delisting in 2010.

Figure 69. Volatility Indices with Listed Futures

Ticker	Underlying
V2X	SX5E
VIX	SPX
VXD	DJIA
VXN	NDX
VXTH	VIX Tail Hedge (long SPX and long 30 delta VIX calls)

VOL FUTURES EXPIRE ON THIRD OR FOURTH WEDNESDAY

Normal options expiry is on the third Friday of the month (but not always; for example, Nikkei uses the second Friday), to be close to the end of the month and on a day that does not usually have a bank holiday (sometimes, however, Good Friday will fall on the third Friday). For US markets, the expiration is on the Saturday after the third Friday (to give more time for the administration of expiration); however, as the last trade date is the Thursday before (expiration is based on opening prices on the Friday), this is irrelevant for cash-settled derivatives such as volatility futures. We note the extra time value due to expiry being on Saturday rather than Friday could be relevant for physical delivery, as an investor could use the performance of different markets to estimate the weekend movements and hence the likely opening price on Monday. As the only volatility futures currently listed are for volatility indices whose maturity is 30 days, the expiry of volatility futures is either the third or fourth Wednesday of the previous month (so the 30-day VIX calculation for settlement price is based on only one maturity, rather than an interpolation between two maturities). This ensures that the settlement price is not subject to interpolation errors (see previous section) that affect volatility indices on other days.

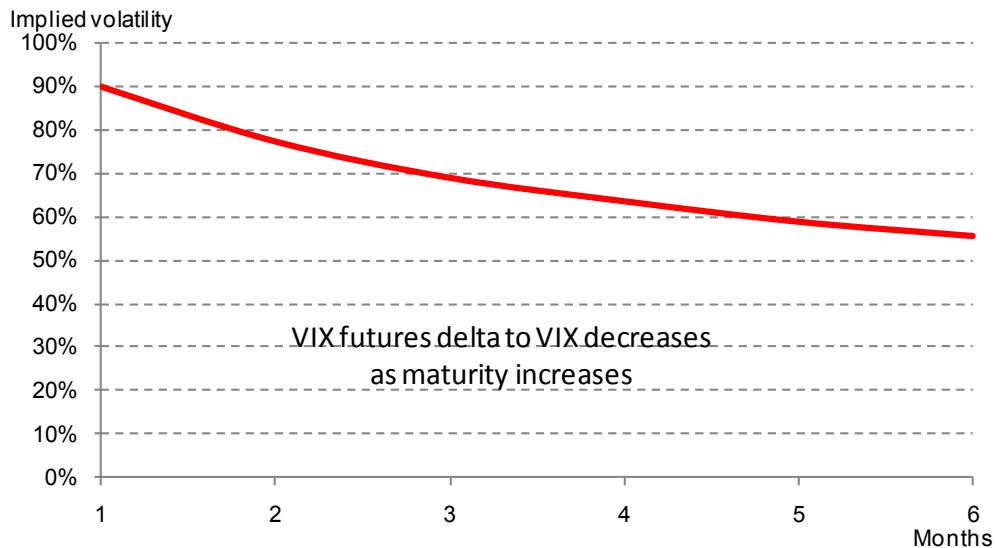
VIX SETTLEMENT PRICE COULD BE MANIPULATED

As the settlement price for the VIX is based on opening trades on S&P500 options, it is far easier to manipulate than the settlement price for the vStoxx, which is based on a 30-minute average ending midday (CET). As shown earlier, the payout of a variance swap is based on a portfolio of options of all strikes weighted $1/K^2$, where K is the strike of the option. This means the calculation of variance swaps is very sensitive to the price of downside puts. Typically, there are offers for downside puts of all strikes at the tick value (ie, the smallest possible non-zero price), as these puts have near zero theoretical value. As the VIX calculation requires a non-zero bid, these offers are usually excluded for strikes below c50%-60%. By entering the minimum bid of US\$0.05 (= tick value), these prices will be included in the calculation and could lift the settlement price by c1pt (as the implied volatility for these low strike puts will be very large). There have been times when the VIX settlement (on the open) has been significantly different to both the close of the day, and the close of the previous day.

MEAN REVERSION MEANS VOL FUTURES HAVE DELTA <1

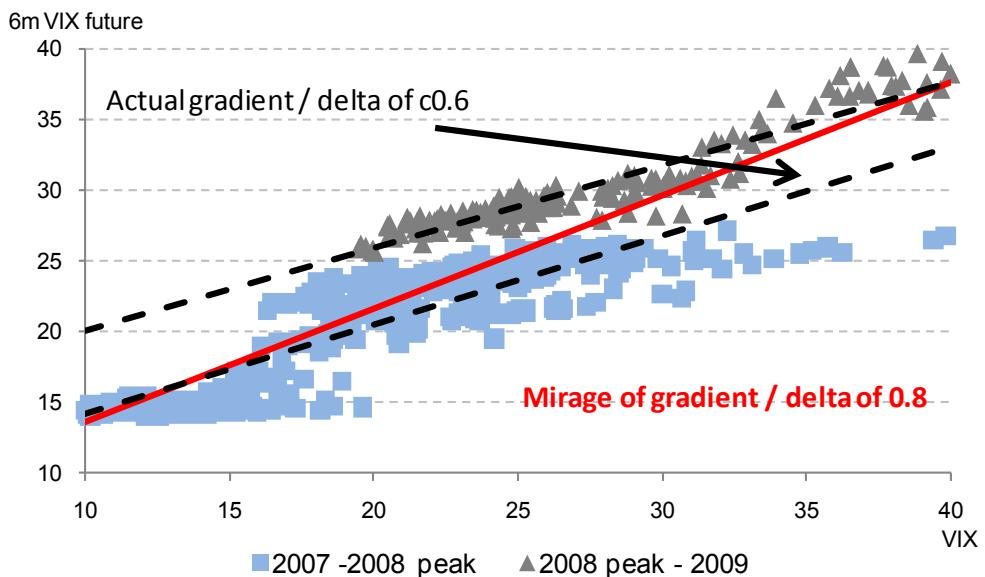
Unlike normal futures, volatility futures are not linear in the underlying index (as the mean reversion of volatility has an effect) and, hence, have deltas significantly lower than 100%. An equity future has near a 100% delta. While the front month VIX future has a high 90% delta (delta vs the VIX), the 6-month VIX future has a lower 55% delta. The lower delta is due to the mean reversion of volatility, as 6-month VIX futures will not trade at 80% even if the VIX trades at 80% (as the VIX only briefly went above 80% post Lehman bankruptcy and swiftly declined, it is highly unlikely to still be at 80% in 6 months' time). The empirical deltas of VIX futures by maturity are shown in Figure 70 below. These values decline in a similar way but less rapidly than they would if volatility solely obeyed a square root of time rule, which is to be expected as volatility surfaces sometimes move in parallel.

Figure 70. VIX Futures Delta to VIX by Maturity



VIX futures delta is both time and volatility dependent

While Figure 70 above gives the average delta of VIX futures over a 9-year period, over a shorter time period the delta can be significantly different. For values of the VIX above 40% the delta of a 6-month VIX future can be close to zero, as the market does not expect the high 40+% volatility environment to last six months as it is likely to be a temporary spike. For values of the VIX below 30%-40%, there is a far higher delta, as moves in volatility within this range are more likely to be part of a volatility regime change that could be longer lasting. When estimating the delta of a volatility future care has to be taken when choosing the period of time used to find the estimate. Figure 71 below shows how the delta for a 6-month VIX future over the 3-year period 2007-09 could be estimated to be 80%, however, in fact, the delta is actually c60% as there was a regime shift assuming higher future volatility post the 2008 peak in volatility post the Lehman bankruptcy. Note also how in the lead-up to the 2008 peak in volatility the 6-month VIX future appeared capped in the high 20s (and had near zero delta for values of VIX above 30%), whereas afterwards no such cap existed (and had high delta for value of VIX above 30%).

Figure 71. 6 Month VIX Futures Delta to VIX

R² of VIX futures declines as maturity increases

The length of time to mean revert, and the level of volatility to which the VIX will mean revert, changes over time. Hence the R² between the VIX and a VIX future decreases as maturity increases. While the front-month VIX future has a high R² of 0.97, the 6-month VIX future has a lower R² of 0.70.

4.4: VOLATILITY FUTURE ETN/ETF

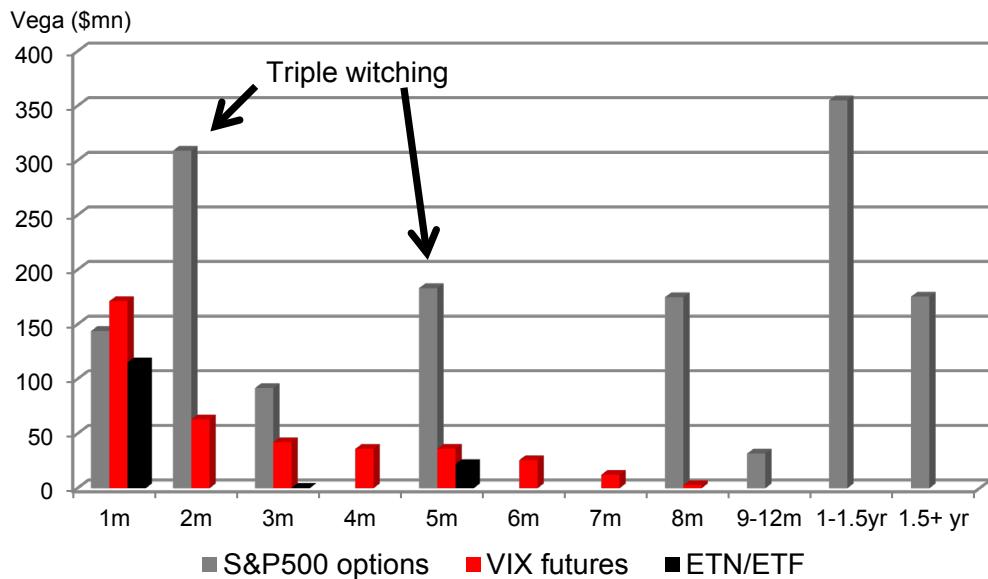
Structured products based on constant maturity volatility futures have become increasingly popular and in the US have at times had a greater size than the underlying volatility futures market. As a constant maturity volatility product needs to sell near-dated expiries and buy far-dated expiries, this flow supports term structure for volatility futures and the underlying options on the index itself. The success of VIX-based products has led to their size being approximately two-thirds of the vega of the relevant VIX futures market (which is a similar size to the net listed S&P500 market) and, hence, appears to be artificially lifting near-dated term structure. The size of vStoxx products is not yet sufficient to significantly impact the market, hence they are a more viable method of trading volatility in our view. We recommend shorting VIX-based structured products to profit from this imbalance, potentially against long vStoxx based products as a hedge. Investors who wish to be long VIX futures should consider the front-month and fourth-month maturities, as their values are likely to be depressed from structured flow.

STRUCTURED PRODUCTS IMPROVED FUTURES LIQUIDITY

As it is impossible to have a product (perpetual or otherwise) whose payout is the volatility index itself, volatility futures were launched to give investors an easy method of trading volatility. Initially, VIX and vStoxx futures had limited liquidity, potentially as they are not perpetual; however, the creation of perpetual structured products has improved the liquidity of volatility futures. Similarly, the introduction of options on these futures has increased the need to delta hedge using these futures, also increasing liquidity. In the US, the size of structured products on VIX futures is so large at times it was bigger than the underlying VIX futures market and appears to have moved the underlying S&P500 market itself.

VIX PRODUCTS ARE 2/3 OF THE SIZE OF FUTURES MARKET

The size in vega of the US market for vanilla S&P500 options, VIX futures and VIX-based ETN/ETF is shown in Figure 72 below. As can be seen, the size of VIX-based ETNs/ETFs is approximately two-thirds of the size of the relevant VIX future.

Figure 72. S&P500 Vega by Maturity for Options, Vol Futures and ETN/ETFs

VIX FUTURES NOW SIMILAR SIZE TO NET OPTIONS MARKET

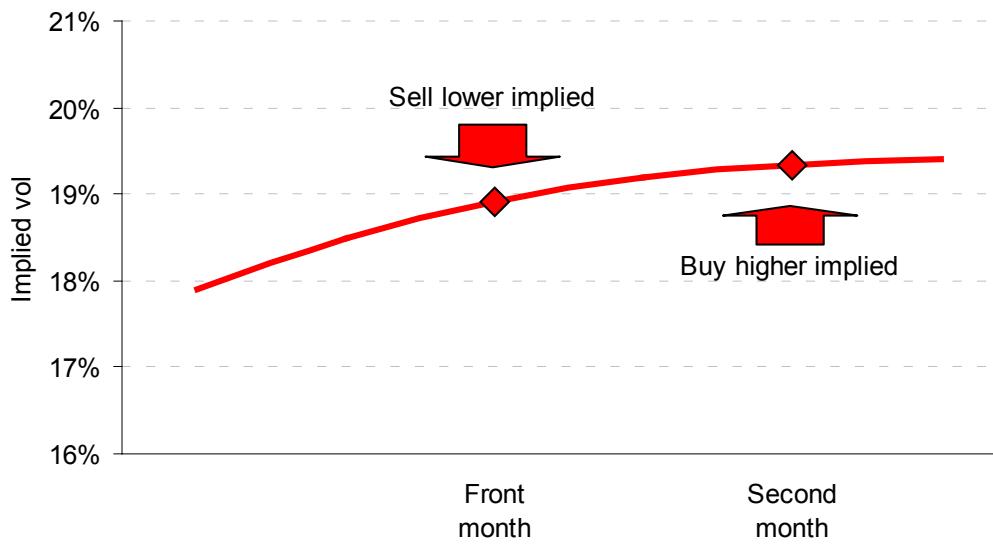
While the size in vega of the S&P500 options market is c4 times bigger than the VIX futures market, this is a theoretical maximum for the market. In practice, if investors trade a spread (eg, call/put spread or ladder) the vega of these structures is far less than the combined vega of the individual legs (vega is the difference between the long and short legs, not the sum of the legs). There is also significant trade in synthetics (long call short put as a substitute for long future) for non-triple witching expiries as S&P500 futures only exist for quarterly expiries. In addition, as one cannot cross futures, when trading on swap (trade volatility structure delta hedged so price is not affected by movements in spot) the delta hedge is done via synthetics. Interest rate trades such as box spreads (long synthetic of one maturity and short synthetic of another maturity) also have no volatility component. A reasonable assumption is that the size of the net vega of the S&P500 listed options market is c25% of the theoretical maximum. Hence, the size of the net listed S&P500 vega is similar to that of the VIX futures market.

VIX futures market comparable to S&P500 vol market if OTC is included

However, we estimate that the OTC market for the S&P500 is 50%-100% of the size of the listed market. This is due to the significant long-term hedging (eg, from variable annuity programs) which cannot be done on exchange (as only maturities up to 2-3 years are listed on the S&P500). Additionally, the size of the variance swap market adds to the size of the

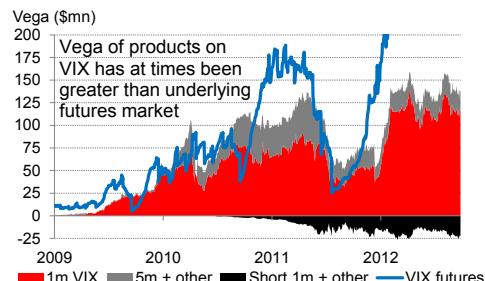
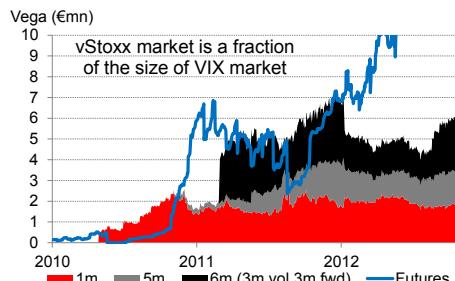
OTC market. Hence, we estimate the vega of VIX futures would be 50%-100% of the total (listed and OTC) size of the S&P500 market.

Figure 73. Volatility Future Term Structure



OPEN-ENDED VOL PRODUCTS STEEPEN TERM STRUCTURE

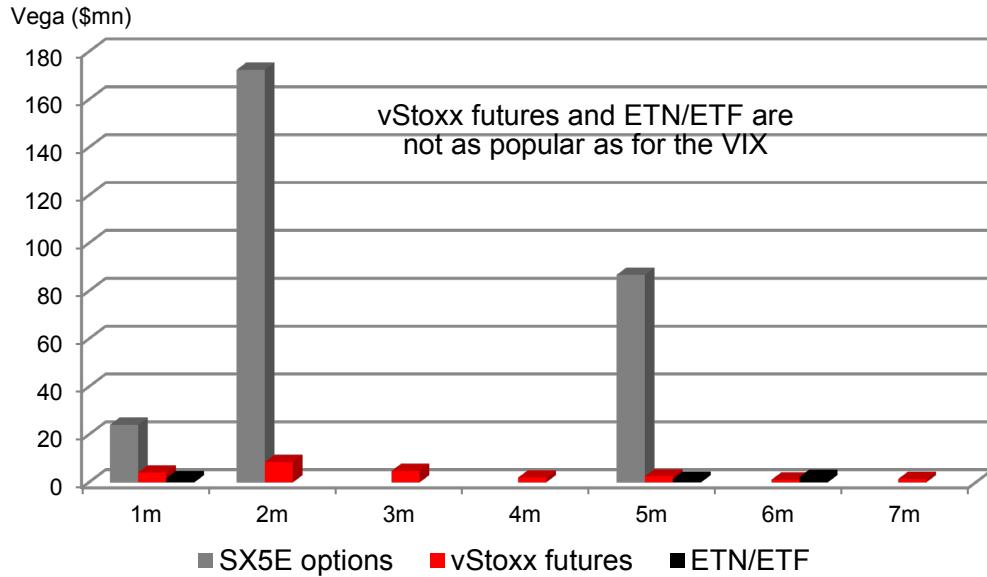
While futures on a volatility index have the advantage of being a listed instrument, they have the disadvantage of having an expiry and, therefore, a longer-term position needs to be rolled. In response to investor demand, many investment banks sold products based on having a fixed maturity exposure on an underlying volatility index. As time passes, these banks hedge their exposure by selling a near-dated expiry and buying a far-dated expiry. The weighted average maturity is therefore kept constant, but the flow puts upward pressure on the term structure. For products of sufficient size, the impact of structured products on the market ensures the market moves against them. Products on short-dated VIX futures which have an average 1-month maturity (by selling front month and buying the second expiry) are now sufficiently large to be moving the volatility market for the S&P500.

Figure 74. Size of VIX ETN/ETF**Size of vStoxx ETN/ETF**

Volatility products suffer ‘roll-down cost’ due to positive term structure

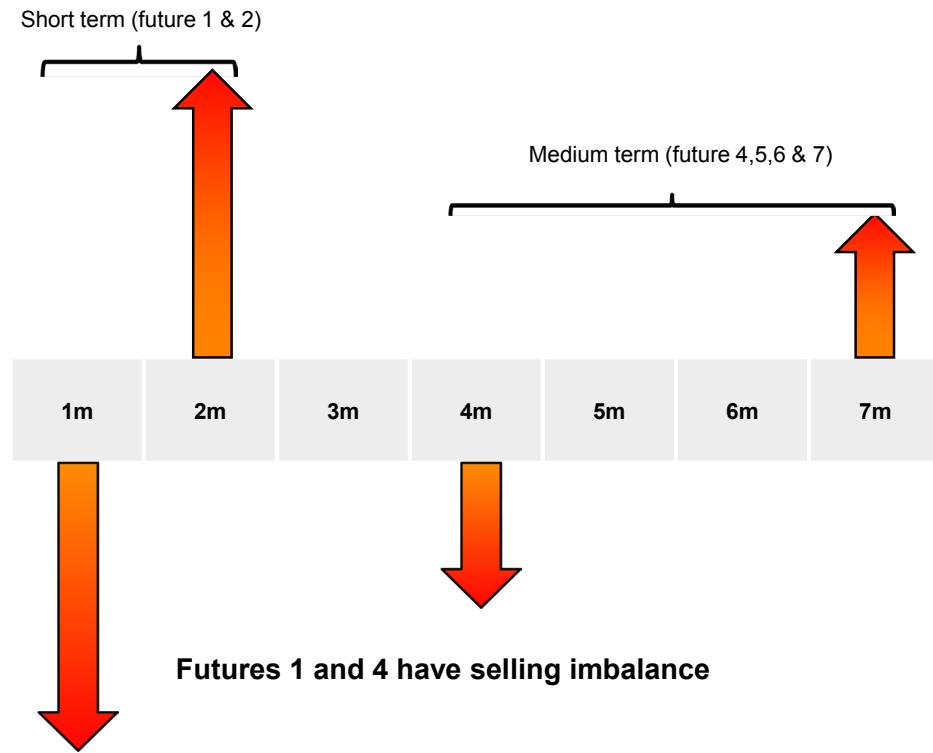
Historically, due to risk aversion and supply demand imbalances, the average term structure of implied volatility has been positive. Index implied volatility term structure will also be lifted by positive implied correlation term structure. The launch of volatility futures and their related ETN/ETF products has increased the supply demand imbalance and supported term structure for the S&P500. VStoxx products are not yet large enough to have an impact on SX5E term structure. As long volatility ETN/ETF products are always selling near-dated implied and buying far-dated implied, there is a roll-down cost if the term structure is positive (as a low near-dated volatility future is sold and a high far-dated volatility future is bought). The higher the positive term structure, the greater the roll-down cost. Conversely, volatility ETN/ETF products will benefit from negative term structure. Investors tend not to benefit from the periods of time there is positive roll-down cost, as these products are often used as a hedge (or view on volatility increasing) and the position is typically closed if equity markets decline and volatility spikes.

Figure 75. SX5E Vega by Maturity for Options, Vol Futures and ETN/ETFs



Longer dated maturity ETN/ETF products have less roll-down cost, but lower beta

Given the fact term structure flattens out as maturity increases, the gradient of volatility surfaces is steeper for near-term expiries than far-term expiries. This means structured products based on longer-dated volatility suffer less roll-down cost. This led to the creation of medium-dated ETN/ETFs, which have an average maturity of five months as they invest in futures of maturity 4, 5, 6 and 7 month (hence sells 4th future to buy 7th). However, as the beta of volatility futures decreases with maturity, longer-dated volatility products benefit less from volatility spikes. While the ratio of beta to roll-down cost is similar across different maturity volatility products, near-dated products do have a worse ratio. There are some products that try to benefit from the excess demand for near-term 1-month ETN/ETFs (medium-dated 5-month ETN/ETFs are less popular) by going short a 1-month volatility product, and at the same time going long approximately twice that size of a 5-month volatility product (as $\sqrt{5} \approx 2$ and volatility often moves in a square root of time manner).

Figure 76. SX5E Vega by Maturity for Options, Vol Futures and ETN/ETFs

BUYING FRONT MONTH OR 4TH MONTH VOL FUTURES IS BEST

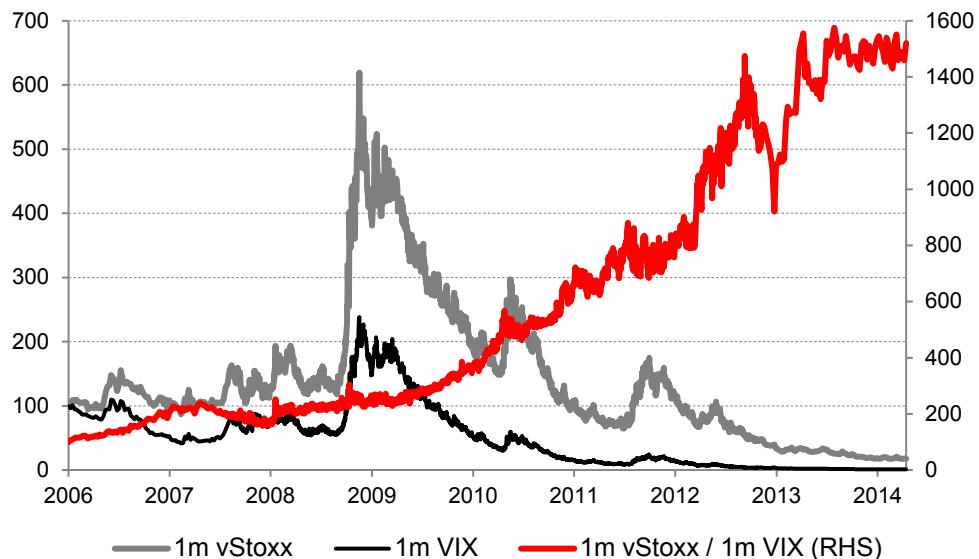
If an investor wishes to initiate a long volatility future position, the best future is the front-month future as this has the most selling pressure from ETN/ETFs (and hence is likely to be relatively cheap). For investors who wish to have a longer-dated exposure, we would recommend the fourth volatility future as our second favourite. This future benefits from the selling of medium-dated ETN/ETFs. A long fourth future position should be closed when it becomes the second month future (as this price is supported by the short dated ETN/ETF). While the eighth future is also a viable investment, the liquidity at this maturity is lower than the others.

INVESTORS COULD SHORT VIX BASED ETN/ETFs

Given the imbalance in the VIX futures market resulting from the size of ETN/ETF products for near-dated VIX products, investors could short these products. The XXV ETF (inverse of the VXX, whose ticker is also the letters of the VXX backwards) based on short 1-month VIX futures is also a viable method of profiting from this imbalance. As the size of the XXV is only c20% of the size of the VXX, a significant imbalance still remains in our

view. As vStoxx-based products are not sufficiently large to be causing an imbalance, a short VIX product long vStoxx product is an attractive way to profit from the VIX imbalance while hedging the overall level of volatility (we note that this trade does not hedge any US or Europe specific volatility). As can be seen in Figure 77 below, the profile of 1m vStoxx/1m VIX (proxy for long 1m vStoxx and short 1m VIX rebalanced every day) offers an attractive performance.

Figure 77. VIX and vStoxx 1-Month Rolling Volatility Future Performance



ETN HAS COUNTERPARTY RISK UNLIKE ETF

There are both ETNs and ETFs based on volatility futures, the primary difference being counterparty risk. Despite the fact the underlying of the product is listed (and hence has no counterparty risk), an investor in an ETN suffers the counterparty risk of the provider. ETFs do not suffer this problem.

EXCESS RETURN PRODUCTS ARE SUPERIOR TO TOTAL RETURN

For investors who are able to trade them, swaps based on excess return indices (eg, VST1ME for vStoxx 1-month futures) are better than ETN/ETF based on total return indices (eg, VST1MT for vStoxx 1-month futures). This is because the returns received from a total return product (EONIA) are likely to be less than the funding levels of a client.

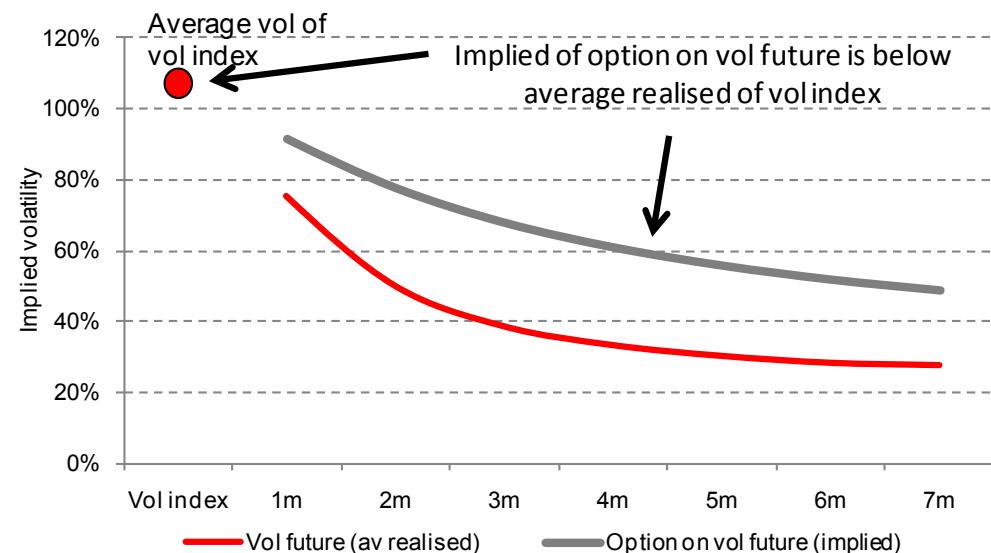
4.5: OPTIONS ON VOLATILITY FUTURES

The arrival of options on volatility futures has encouraged trading on the underlying futures. It is important to note that an option on a volatility future is an option on future implied volatility, whereas an option on a variance swap is an option on realised volatility. As implieds always trade at a lower level to peak realised (as you never know when peak realised will occur) the volatility of implied is lower than the volatility of realised, hence options on volatility futures should trade at a lower implied than options on var. Both have significantly downward sloping term structure and positive skew. We note that the implied for options on volatility futures should not be compared to the realised of volatility indices.

OPTIONS ARE ON THE FUTURE, NOT THE VOL INDEX ITSELF

As volatility markets have become more liquid, investors became increasingly interested in purchasing options on volatility. As it is impossible to buy a volatility index itself, options on volatility have to be structured as an option on a volatility future. For equities there is not much difference between the volatility of spot and the volatility of a future (as futures are near dated, the effect of interest rate and implied dividend volatility is small). However, there is a very significant difference between the volatility of a vol index, and the volatility of a vol future.

Figure 78. Realised of Vol Futures and Implied of Option on Vol Future



VOLATILITY TERM STRUCTURE IS VERY NEGATIVE

The term structure of implied volatility of vanilla equity options is on average relatively flat¹⁷. In contrast, the term structure of implied volatility of option on vol futures is sharply negative. The volatility of a vol future is significantly less than the volatility of the vol index, but does converge as it approaches expiry (when it becomes as volatile as the vol index itself).

COMPARING IMPLIED VOL AND REALISED VOL IS DIFFICULT

The realised volatility of a vol future increases as it approaches expiry (near-dated volatility is more volatile than far-dated volatility). The implied volatility of an option on a vol future should trade roughly in line with the average realised volatility of the vol future over the life of the option. Hence the average realised volatility of a vol future will be a blend of the initial low realised volatility, and the higher realised volatility close to expiry. The implied volatility term structure of an option on vol future will therefore be less negative than the realised volatility term structure of the vol future.

Realised of vol futures < implied of option on vol futures < realised of vol index

The implied volatility level of options on vol futures is also higher than the realised volatility of the vol future for that expiry (eg, implied of 6-month option on vol futures is above current realised of 6-month vol future). The implied volatility of options on vol futures will, however, be less than the realised volatility of the vol index, which makes options on vol futures look cheap if an investor mistakenly compares its implied to the realised of the vol index.

OPTION ON VOL FUTURE CHEAPER THAN OPTIONS ON VAR

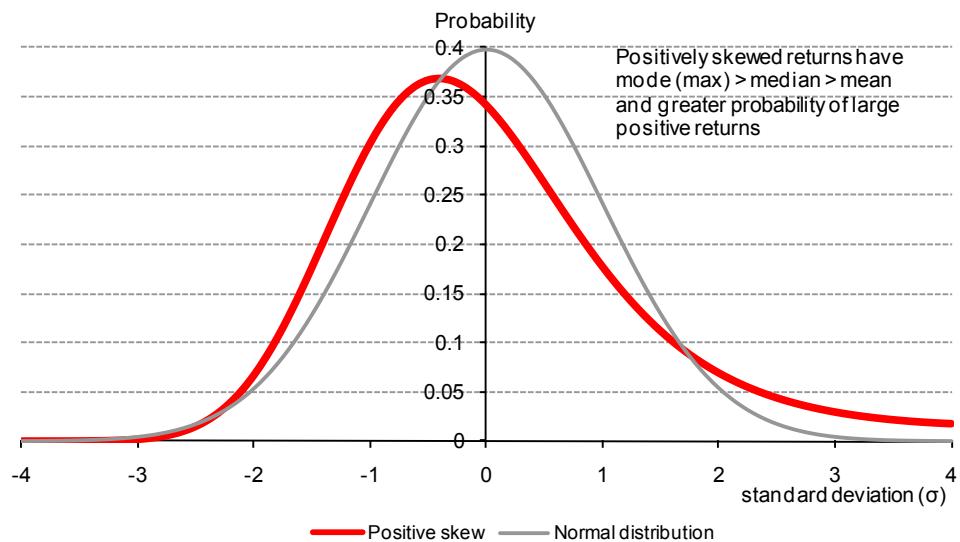
Implied volatility is less volatile than realised volatility, as implied volatility will never trade at the min or max level of realised (as it is an estimate of future volatility, and there is never a time that the market can be 100% certain realised will reach its min or max). As implied volatility is less volatile than realised volatility, an option on a vol future should be at a lower implied than an option on realised variance (particularly for near-dated expiries). They will, however, have a similar negative term structure.

¹⁷ On average slightly upward sloping, but at a far shallower gradient to the negative term structure of options on vol futures.

Options on vol future have positive skew, just like options on var

When there is an equity market panic, there tends to be large negative returns for equities and a volatility spike. As the probability distribution of equity prices has a greater probability of large negative returns, it has a negative skew. Volatility, on the other hand, tends to have a larger probability of large positive returns and hence has positive skew (just like options on realised variance).

Figure 79. Probability Distribution of Options on Vol Futures



Need to model volatility with high vol of vol and mean reversion

As vol futures have a high near-term volatility, and low far-dated volatility, they have to be modelled with a high vol of vol and high mean reversion.

OPTIONS ON VOLATILITY FUTURE PRODUCTS ALSO EXIST

At present there are only options on VIX and vStoxx futures. There are, however, also options on structured products based on VIX volatility futures. The list of underlyings for options is given below.

Figure 80. Volatility Securities with Listed Options

Ticker	Underlying Type
VIX	Vol index
V2X	Vol index
VXX US	ETN
VXZ US	ETN
VIIX US	ETN
SVXY US	ETF
UVXY US	ETF
VIXM US	ETF
VIXY US	ETF

CHAPTER 5

LIGHT EXOTICS

Advanced investors can make use of more exotic equity derivatives. Some of the most popular are light exotics, such as barriers, worst-of/best-of options, outperformance options, look-back options, contingent premium options, composite options and quanto options.

5.1: BARRIER OPTIONS

Barrier options are the most popular type of light exotic product, as they are used within structured products or to provide cheap protection. The payout of a barrier option knocks in or out depending on whether a barrier is hit. There are eight types of barrier option, but only four are commonly traded, as the remaining four have a similar price to vanilla options. Barrier puts are more popular than calls (due to structured product and protection flow), and investors like to sell visually expensive knock-in options and buy visually cheap knock-out options. Barrier options (like all light exotics) are always European (if they were American, the price would be virtually the same as a vanilla option, as the options could be exercised just before the barrier was hit).

BARRIER OPTIONS CAN HAVE DELTA OF MORE THAN $\pm 100\%$

The hedging of a barrier option is more involved than for vanilla options, as the delta near the barrier can be significantly more than $\pm 100\%$ near expiry. The extra hedging risk of barriers widens the bid-offer spread in comparison with vanilla options. Barrier options are always European and are traded OTC.

THERE ARE THREE KEY VARIABLES FOR BARRIER OPTIONS

There are three key variables to a barrier option, each of which has two possibilities. These combinations give eight types of barrier option ($8=2\times 2\times 2$).

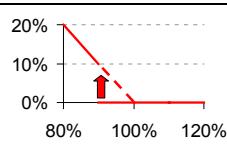
- **Down/up.** The direction of the barrier in relation to spot. Almost all put barriers are down barriers and, similarly, almost all call barriers are up barriers.
- **Knock in/out.** Knock-out (KO) options have a low premium and give the impression of being cheap; hence, they are usually bought by investors. Conversely, knock-in (KI) options are visually expensive (as knock-in options are a similar price to a vanilla) and are usually sold by investors (through structured products). For puts, a knock-in is the most popular barrier (structured product selling of down and knock-in puts). However, for calls this is reversed and knock-outs are the most popular. Recent volatility has made knock-out products less popular than they once were, as many hit their barrier and became worthless.
- **Put/call.** The type of payout of the option. Put barriers are three to four times more popular than call barriers, due to the combination of selling from structured products (down and knock-in puts) and cheap protection buying (down and knock-out puts).

ONLY FOUR OF THE EIGHT TYPES OF BARRIER ARE TRADED

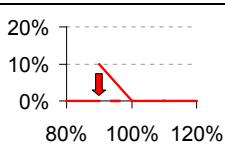
The difference in price between a vanilla option and barrier option is only significant if the barrier occurs when the option has intrinsic value. If the only value of the option when the barrier knocks in/out is time value, then the pricing for the barrier option will be roughly equal to the vanilla option. Because of this, the naming convention for barrier options can be shortened to knock in (or out) followed by call/put (as puts normally have a down barrier, and calls an up barrier). The four main types of barrier option and their uses are shown below.

- **Knock-in put (down and knock-in put).** Knock-in puts are the most popular type of barrier option, as autocallables are normally hedged by selling a down and knock-in put to fund the high coupon. They have a barrier which is below both spot and strike and give an identical payoff to a put only once spot has gone below the down barrier. Until spot reaches the down barrier there is no payout. However, as this area has the least intrinsic value, the theoretical price is similar to a vanilla and therefore visually expensive.
- **Knock-out put (down and knock-out put).** Knock-out puts are the second most popular barrier option after knock-in puts (although knock-in puts are three times as popular as knock-out puts due to structured product flow). Knock-out puts give an identical payout to a put, until spot declines through the down barrier (which is below both spot and strike), in which case the knock-out option becomes worthless. As the maximum payout for a put lies below the knock-out barrier, knock-out puts are relatively cheap and are often thought of as a cheap method of gaining protection.
- **Knock-in call (up and knock-in call).** Knock-in calls give an identical payout to a call, but only when spot trades above the up barrier, which lies above spot and the strike. They are the least popular barrier option, as their high price is similar to the price of a call and structured product flow is typically less keen on selling upside than downside.
- **Knock-out call (up and knock-out call).** Knock-out calls are the most popular barrier option for calls, but their popularity still lags behind both knock-in and knock-out puts. As they give the same upside participation as a vanilla call until the up barrier (which is above spot and strike) is reached, they can be thought of as a useful way of gaining cheap upside.

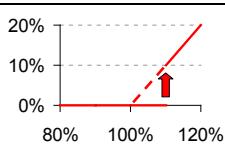
Figure 81. KI put



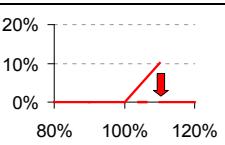
KO put



KI call



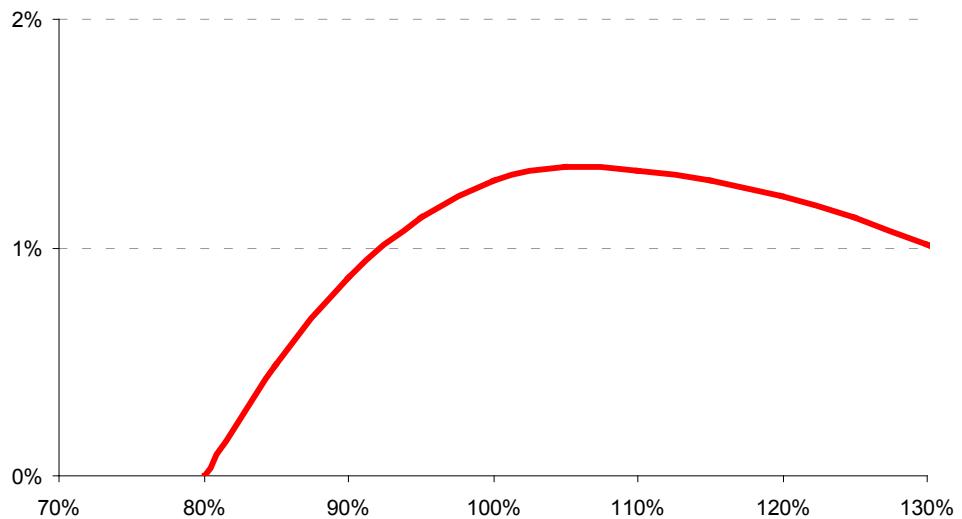
KO call



KNOCK-OUT DECREASE IN VALUE AS STRIKE APPROACHES

While vanilla options (and knock-in options) will increase in value as spot moves further in the money, this is not the case for knock-out options, where spot is near the strike. This effect is caused by the payout equalling zero at the barrier, which can cause delta to be of opposite sign to the vanilla option. This effect is shown below for a one-year ATM put with 80% knock-out. The peak value of the option is at c105%; hence, for values lower than that value the delta is positive not negative. This is a significant downside to using knock-out puts for protection, as their mark to market can increase (not decrease) equity sensitivity to the downside.

Figure 82. Price of One-Year ATM Put with 80% Knock-Out

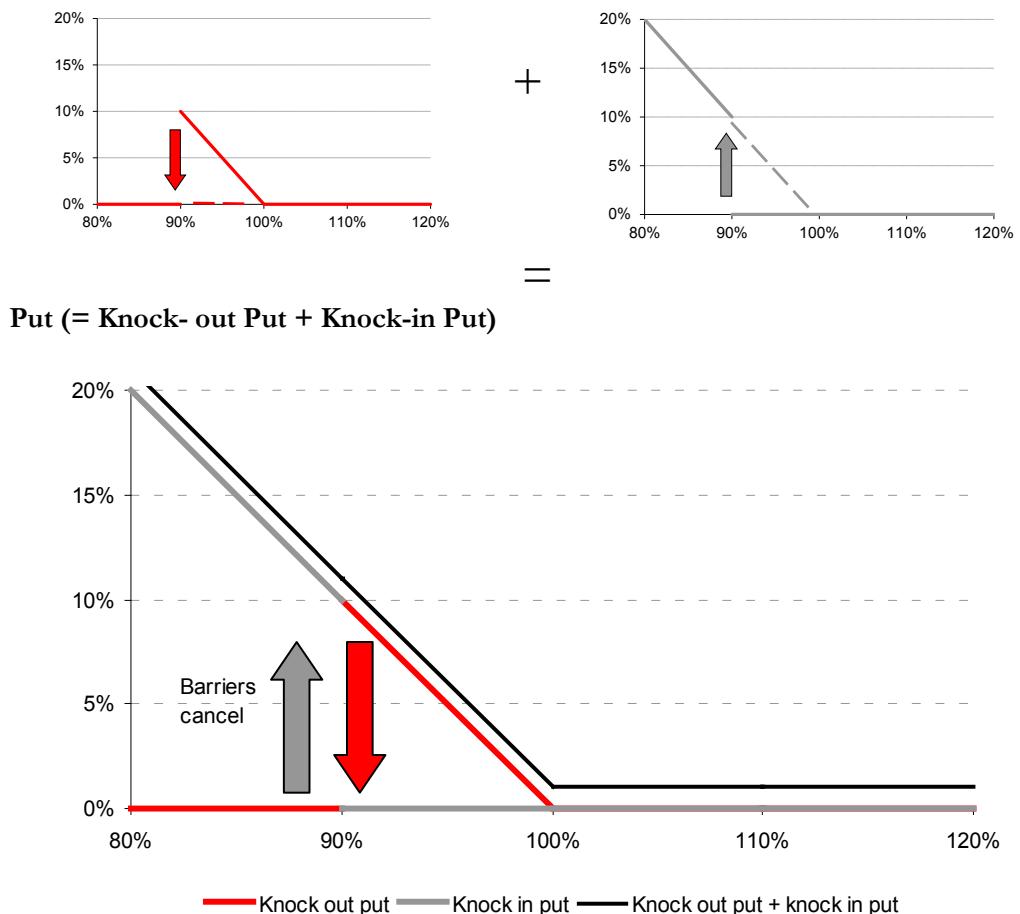


KNOCK-OUT PUT + KNOCK-IN PUT = (VANILLA) PUT

If a knock-out put and a knock-in put have the same strike and barrier, then together the combined position is equal to a long vanilla put ($P_{KO} + P_{KI} = P$). This is shown in the charts below. The same argument can apply to calls ($P_{KO} + P_{KI} = P$). This relationship allows us to see mathematically that if knock-out options are seen as visually cheap, then knock-in options must be visually expensive (as a knock-in option must be equal to the price of a vanilla less the value of a visually cheap knock-out option).

Figure 83. Knock-out put

Knock-in put



KNOCK-OUTS COST C15%-25% OF PUT SPREAD COST

The payout of a knock-out put is equal to a ‘shark fin’ (see top left chart in Figure 83 above) until the barrier is reached. A ‘shark fin’ is equal to a short digital position (at the barrier) plus a put spread (long put at strike of knock-out put, short put at barrier of knock-out put). The price of a knock-out put can therefore be considered to be the cost of a put spread, less a digital and less the value of the knock-out. As pricing digitals and barriers is not trivial, comparing the price of a knock-out put as a percentage of the appropriate put spread can be a quick way to evaluate value (the knock-out will have a lower value as it offers less payout to the downside). For reasonable barriers between 10% and 30% below the strike, the price of the knock-out option should be between c15% and c25% of the cost of the put spread.

CONTINUOUS BARRIERS ARE CHEAPER THAN DISCRETE

There are two types of barriers, continuous and discrete. A continuous barrier is triggered if the price hits the barrier intraday, whereas a discrete barrier is only triggered if the closing price passes through the barrier. Discrete knock-out barriers are more expensive than continuous barriers, while the reverse holds for knock-in barriers (especially during periods of high volatility). There are also additional hedging costs to discrete barriers, as it is possible for spot to move through the barrier intraday without the discrete barrier being triggered (ie, if the close is the correct side of the discrete barrier). As these costs are passed on to the investor, discrete barriers are far less popular than continuous barriers for single stocks (c10%-20% of the market), although they do make up almost half the market for indices.

Jumps in stock prices between close and open is a problem for all barriers

While the hedge for a continuous barrier should, in theory, be able to be executed at a level close to the barrier, this is not the case should the underlying jump between close and open. In this case, the hedging of a continuous barrier suffers a similar problem to the hedging of a discrete barrier (delta hedge executed at a significantly different level to the barrier).

DOUBLE BARRIERS ARE POSSIBLE, BUT RARE

Double barrier options have both an up barrier and a down barrier. As only one of the barriers is significant for pricing, they are not common (as their pricing is similar to an ordinary single-barrier option). They make up less than 5% of the light exotic market.

REBATES CAN COMPENSATE FOR TRIGGER OF BARRIER

The main disadvantage of knock-out barrier options is that the investor receives nothing for purchasing the option if they are correct about the direction of the underlying (option is ITM) but incorrect about the magnitude (underlying passes through barrier). In order to provide compensation, some barrier options give the long investors a rebate if the barrier is triggered: for example, an ATM call with 120% knock-out that gives a 5% rebate if the barrier is touched. Rebates comprise approximately 20% of the index barrier market but are very rare for single-stock barrier options.

5.2: WORST-OF/BEST-OF OPTIONS

Worst-of (or best-of) options give payouts based on the worst (or best) performing asset. They are the second most popular light exotic due to structured product flow. Correlation is a key factor in pricing these options, and investor flow typically buys correlation (making uncorrelated assets with low correlation the most popular underlyings). The underlyings can be chosen from different asset classes (due to low correlation), and the number of underlyings is typically between three and 20. They are always European, and normally ATM options.

NORMALLY 1 YEAR MATURITY AND CAN BE CALLS OR PUTS

Worst-of/best-of options can be any maturity. Although the most popular is one-year maturity, up to three years can trade. As an option can be a call or a put, and either ‘worst-of’ or ‘best-of’; there are four types of option to choose from. However, the most commonly traded are worst-of options (call or put). The payouts of the four types are given below:

Worst-of call = $\text{Max} (\text{Min} (r_1, r_2, \dots, r_N), 0)$ where r_i is the return of N assets

Worst-of put = $\text{Max} (-\text{Min} (r_1, r_2, \dots, r_N), 0)$ where r_i is the return of N assets

Best-of call = $\text{Max} (\text{Max} (r_1, r_2, \dots, r_N), 0)$ where r_i is the return of N assets

Best-of put = $\text{Max} (-\text{Max} (r_1, r_2, \dots, r_N), 0)$ where r_i is the return of N assets

WORST-OF CALLS ARE POPULAR TO BUY (AS CHEAP)

The payout of a worst-of call option will be equal to the lowest payout of individual call options on each of the underlyings. As it is therefore very cheap, they are popular to buy. If all the assets are 100% correlated, then the value of the worst-of call is equal to the value of calls on all the underlyings (hence, in the normal case of correlation less than 100%, a worst-of call will be cheaper than any call on the underlying). If we lower the correlation, the price of the worst-of call also decreases (eg, the price of a worst-of call on two assets with -100% correlation is zero, as one asset moves in the opposite direction to the other). A worst-of call option is therefore long correlation. As worst-of calls are cheap, investors like to buy them and, therefore, provide buying pressure to implied correlation.

Rumour of QE2 lifted demand for worst-of calls on cross assets

Before QE2 (second round of quantitative easing) was announced, there was significant buying flow for worst-of calls on cross assets. The assets chosen were all assets that were likely to be correlated should QE2 occur but that would normally not necessarily be correlated (giving attractive pricing). QE2 was expected to cause USD weakening (in favour of other G10 currencies like the JPY, CHF and EUR), in addition to lifting ‘risk-on’ assets, like equities and commodities. The buying of worst-of calls on these three assets would therefore be a cheap way to gain exposure to the expected movements of markets if quantitative easing was extended (which it was).

WORST-OF PUTS ARE EXPENSIVE AND USUALLY SOLD

A worst-of put will have a greater value than any of the puts on the underlying assets and is therefore very expensive to own. However, as correlation increases towards 100%, the value of the worst-of put will decrease towards the value of the most valuable put on either of the underlyings. A worst-of put is therefore short correlation. As selling (expensive) worst-of puts is popular, this flow puts buying pressure on implied correlation (the same effect as the flow for worst-of calls).

BEST-OF CALLS AND BEST-OF PUTS ARE RELATIVELY RARE

While worst-of options are popular, there is relatively little demand for best-of options. There are some buyers of best-of puts (which again supports correlation); however, best-of calls are very rare. Figure 84 below summarises the popularity and direction of investor flows (normally from structured products) and the effect on implied correlation. A useful rule of thumb for worst-of/best-of options is that they are short correlation if the price of the option is expensive (worst-of put and best-of call) and the reverse if the price of the option is cheap. This is why the buying of cheap and selling of expensive worst-of/best-of options results in buying flow to correlation.

Figure 84. Best-of/Worst-of Options

Option	Correlation	Flow	Cost	Notes
Worst-of put	Short	Sellers	Expensive	Popular structure to sell as cost is greater than that of most expensive put
Worst-of call	Long	Buyers	Cheap	Popular way to buy upside as low cost is less than cheapest call on any of the assets
Best-of put	Long	Some buyers	Cheap	Some buyers as cost is lower than cheapest put
Best-of call	Short	Rare	Expensive	Benefits from correlation falling as markets rise

LIGHT EXOTIC OPTIONS FLOW LIFTS IMPLIED CORRELATION

As the flow from worst-of/best-of products tends to support the levels of implied correlation, implied correlation typically trades above fair value. While other light exotic flow might not support correlation (eg, outperformance options, which are described below), worst-of/best-of options are the most popular light exotic, whose pricing depends on correlation and are therefore the primary driver for this market. We would point out that the most popular light exotics – barrier options – have no impact on correlation markets. In addition, worst-of/best-of flow is concentrated in uncorrelated assets, whereas outperformance options are usually on correlated assets.

5.3: OUTPERFORMANCE OPTIONS

Outperformance options are an option on the difference between returns on two different underlyings. They are a popular method of implementing relative value trades, as their cost is usually cheaper than an option on either underlying. The key unknown parameter for pricing outperformance options is implied correlation, as outperformance options are short correlation. The primary investor base for outperformance options is hedge funds, which are usually buyers of outperformance options on two correlated assets (to cheapen the price). Outperformance options are European and can always be priced as a call. Unless they are struck with a hurdle, they are an ATM option.

OUTPERFORMANCE OPTIONS ARE SHORT-DATED CALLS

Outperformance options give a payout based on the difference between the *returns* of two underlyings. While any maturity can be used, they tend to be for maturities up to a year (maturities less than three months are rare). The payout formula for an outperformance option is below – by convention always quoted as a call of ' r_A over r_B ' (as a put of ' r_A over r_B ' can be structured as a call on ' r_B over r_A '). Outperformance options are always European (like all light exotics) and are traded OTC.

$$\text{Payout} = \text{Max} (r_A - r_B, 0) \text{ where } r_A \text{ and } r_B \text{ are the returns of assets A and B, respectively}$$

OPTIONS CAN HAVE HURDLE AND ALLOWABLE LOSS

While outperformance options are normally structured ATM, they can be cheapened by making it OTM through a hurdle or by allowing an allowable loss at maturity (which simply defers the initial premium to maturity). While outperformance options can be structured ITM by having a negative hurdle, as this makes the option more expensive, this is rare. The formula for outperformance option payout with these features is:

$$\text{Payout} = \text{Max} (r_A - r_B - \text{hurdle}, -\text{allowable loss})$$

OUTPERFORMANCE OPTIONS ARE SHORT CORRELATION

The pricing of outperformance options depends on both the volatility of the two underlyings and the correlation between them. As there tends to be a more liquid and visible market for implied volatility than correlation, it is the implied correlation that is the key factor in determining pricing. Outperformance options are short correlation, which can be intuitively seen as: the price of an outperformance option must decline to zero if one assumes correlation rises towards 100% (two identical returns give a zero payout for the outperformance option).

As flow is to the buy side, some hedge funds outperformance call overwrite

Outperformance options are ideal for implementing relative value trades, as they benefit from the upside, but the downside is floored to the initial premium paid. The primary investor base for outperformance options are hedge funds. While flow is normally to the buy side, the overpricing of outperformance options due to this imbalance has led some hedge funds to call overwrite their relative value position with an outperformance option.

MARGRABE'S FORMULA CAN BE USED FOR PRICING

An outperformance option volatility σ_{A-B} can be priced using Margrabe's formula given the inputs of the volatilities σ_A and σ_B of assets A and B, respectively, and their correlation ρ . This formula is shown below.

$$\sigma_{A-B} = \sqrt{\sigma_A^2 + \sigma_B^2 - 2\rho\sigma_A\sigma_B}$$

TEND TO BE USED FOR CORRELATED ASSETS

The formula above confirms mathematically that outperformance options are short correlation (due to the negative sign of the final term with correlation ρ). From an investor perspective, it therefore makes sense to sell correlation at high levels; hence, outperformance options tend to be used for correlated assets (so cross-asset outperformance options are very rare). This is why outperformance options tend to be traded on indices with a 60%-90% correlation and on single stocks that are 30%-80% correlated. The pricing of an outperformance option offer tends to have an implied correlation 5% below realised for correlations of c80%, and 10% below realised for correlations of c50% (outperformance option offer is a bid for implied correlation).

Best pricing is with assets of similar volatility

The price of an outperformance is minimised if volatilities σ_A and σ_B of assets A and B are equal (assuming the average of the two volatilities is kept constant). Having two assets of equal volatility increases the value of the final term $2\rho\sigma_A\sigma_B$ (reducing the outperformance volatility σ_{A-B}).

LOWER FORWARD FLATTERS OUTPERFORMANCE PRICING

Assuming that the two assets have a similar interest rate and dividends, the forwards of the two assets approximately cancel each other out, and an ATM outperformance option is also ATMf (ATM forward or At The Money Forward). When comparing relative costs of outperformance options with call options on the individual underlyings, ATMf strikes must be used. If ATM strikes are used for the individual underlyings, the strikes will usually be lower than ATMf strikes and the call option will appear to be relatively more expensive compared to the ATMf (= ATM) outperformance option.

Pricing of ATM outperformance options is usually less than ATMf on either underlying

If two assets have the same volatility ($\sigma_A = \sigma_B$) and are 50% correlated ($\rho = 50\%$), then the input for outperformance option pricing σ_{A-B} is equal to the volatilities of the two underlyings ($\sigma_{A-B} = \sigma_A = \sigma_B$). Hence, ATMf (ATM forward) options on either underlying will be the same as an ATMf (\approx ATM) outperformance option. As outperformance options tend to be used on assets with higher than 50% correlation and whose volatilities are similar, outperformance options are usually cheaper than similar options on either underlying.

5.4: LOOK-BACK OPTIONS

There are two types of look-back options, strike look-back and payout look-back, and both are usually multi-year options. Strike reset (or look-back) options have their strike set to the highest, or lowest, value within an initial look-back period (of up to three months). These options are normally structured so the strike moves against the investor in order to cheapen the cost. Payout look-back options conversely tend to be more attractive and expensive than vanilla options, as the value for the underlying used is the best historical value. As with all light exotics, these options are European and OTC.

STRIKE OF RESET OPTIONS MOVES AGAINST INVESTOR

There are two main strike reset options, and both have an initial look-back period of typically one to three months, where the strike is set to be the highest (for a call) or lowest (for a put) traded value. While the look-back optionality moves against the investor, as the expiry of these options is multi-year (typically three), there is sufficient time for spot to move back in the investor's favour, and the strike reset cheapens the option premium. While having a strike reset that moves the strike to be the most optimal for the investor is possible, the high price means they are unpopular and rarely trade. While the cheaper form of strike reset options does attract some flow due to structured products, they are not particularly popular.

Strike resets are cheaper alternative to ATM option at end of reset period

There are three possible outcomes to purchasing a strike reset option. Strike reset options can be considered a cheaper alternative to buying an ATM option at the end of the strike reset period, as the strike is roughly identical for two of the three possible outcomes (but at a lower price).

- **Spot moves in direction of option payout.** If spot moves in a direction that would make the option ITM, the strike is reset to be equal to spot as it moves in a favourable direction, and the investor is left with a roughly ATM option.
- **Range-trading markets.** Should markets range trade, the investor will similarly receive a virtually ATM option at the end of the strike reset period.
- **Spot moves in opposite direction to option payout.** If spot initially moves in the opposite direction to the option payout (down for calls, up for puts), then the option strike is identical to an option that was initially ATM (as the key value of the underlying for the strike reset is the initial value) and, hence, OTM at the end of the strike rest period. The downside of this outcome is why strike reset options can be purchased for a lower cost than an ATM option.

Strike resets perform best when there is an initial period of range trading

Strike reset options are therefore most suitable for investors who believe there will be an initial period of range trading, before the underlying moves in a favourable direction.

PAYOUT LOOK-BACK OPTIONS

Having a look-back option that selects the best value of the underlying (highest for calls, lowest for puts) increases the payout of an option – and cost. These options typically have a five-year maturity and typically use end-of-month or end-of-year values for the selection of the optimal payout.

5.5: CONTINGENT PREMIUM OPTIONS

Contingent premium options are initially zero-premium and only require a premium to be paid if the option becomes ATM on the close. The contingent premium to be paid is, however, larger than the initial premium would be, compensating for the fact that it might never have to be paid. Puts are the most popular, giving protection with zero initial premium. These typically one-year put options are OTM (or the contingent premium would almost certainly have to be paid immediately) and European.

CONTINGENT PREMIUM OPTION HAS ZERO UPFRONT COST

While contingent premium calls are possible, the most popular form is for a contingent premium put to allow protection to be bought with no initial cost. The cost of the premium to be paid is roughly equal to the initial premium of the vanilla option, divided by the probability of spot trading through the strike at some point during the life of the option (eg, an 80% put whose contingent premium has to be paid if the underlying goes below 80%). Using contingent premium options for protection has the benefit that no cost is suffered if the protection is not needed, but if spot dips below the strike/barrier, then the large premium has to be paid (which is likely to be more than the put payout unless there was a large decline). These can be thought of as a form of ‘crash put’.

Conditional premium on a level other than strike is possible, but rare

The usual structure for contingent premium options is to have the level at which the premium is paid equal to the strike. The logic is that although investors have to pay a large premium, they do have the benefit of holding an option that is slightly ITM. Having the conditional premium at a level other than strike is possible, but rare (eg, an 80% put whose contingent premium has to be paid if the underlying reaches 110%).

5.6: COMPOSITE AND QUANTO OPTIONS

There are two types of option involving different currencies. The simplest is a composite option, where the strike (or payoff) currency is in a different currency to the underlying. A slightly more complicated option is a quanto option, which is similar to a composite option, but the exchange rate of the conversion is fixed.

COMPOSITE OPTIONS USE DIFFERENT VOLATILITY INPUT

A composite option is a cash or physical option on a security whose currency is different from the strike or payoff currency (eg, Euro strike option on Apple). If an underlying is in a foreign currency, then its price in the payout (or strike) currency will usually be more volatile (and hence more expensive) due to the additional volatility associated with currency fluctuations. Only for significantly negative correlations will a composite option be less expensive than the vanilla option (if there is zero correlation the effect of FX still lifts valuations). The value of a composite option can be calculated using Black-Scholes as usual, by substituting the volatility of the asset with the volatility of the asset in payout currency terms. The payout (or strike) currency risk-free rate should be used rather than the (foreign) security currency risk-free rate. The dividend yield assumption is unchanged (as it has no currency) between a composite option and a vanilla option.

$$\sigma_{Payout} = \sqrt{\sigma_{Security}^2 + \sigma_{FX}^2 + 2\rho\sigma_{Security}\sigma_{FX}}$$

where

σ_{Payout} = volatility of asset in payout (strike) currency

$\sigma_{Security}$ = volatility of asset in (foreign) security currency

σ_{FX} = volatility of FX rate (between payout currency and security currency)

ρ = correlation of FX rate (security currency in payoff currency terms) and security price

Composite options are long correlation (if FX is foreign currency in domestic terms)

The formula to calculate the volatility of the underlying is given above. As the payoff increases with a positive correlation between FX and the underlying, a composite option is long correlation (the positive payout will be higher due to FX, while FX moving against the investor is irrelevant when the payout is zero). Note that care has to be taken when considering the definition of the FX rate; it should be the (foreign) security currency given in (domestic) payoff currency terms.

For example, if we are pricing a euro option on a dollar-based security and assume an extreme case of $\rho = 100\%$, the volatility of the USD underlying in EUR will be the sum of the volatility of the underlying and the volatility of USD.

QUANTO OPTIONS USE DIFFERENT DIVIDEND INPUT

Quanto options are similar to a composite option, except the payout is always cash settled and a fixed FX rate is used to determine the payout. Quanto options can be modelled using Black-Scholes. As the FX rate for the payout is fixed, quanto options are modelled using the normal volatility of the underlying (as FX volatility has no effect). The payout is simply the fixed FX rate multiplied by the price of a vanilla option with the same volatility, but a different carry. The carry (risk-free rate - dividend) to be used is shown below (the risk-free rate for quanto options is assumed to be the risk-free rate of the security currency, ie, it is not the same as for composite options).

$$c_{\text{Quanto}} = rfr_{\text{Security}} - d - \rho\sigma_{\text{Security}}\sigma_{FX}$$

$$\Rightarrow d_{\text{Quanto}} = d + \rho\sigma_{\text{Security}}\sigma_{FX} \text{ as } d_{\text{quanto}} = rfr_{\text{Security}} - c_{\text{Quanto}}$$

where

c_{Quanto} = carry for quanto pricing

d_{Quanto} = dividend for quanto pricing

d = dividend yield

rfr_{Security} = risk free rate of security currency

rfr_{Payout} = risk free rate of payout currency

Quanto options are either long or short correlation depending on the sign of the delta

The correlation between the FX and the security has an effect on quanto pricing, the direction (and magnitude) of which depends on the delta of the option. This is because the dividend risk of an option is equal to its delta, and the dividend used in quanto pricing increases as correlation increases.

Quanto option calls are short correlation (if FX is foreign currency in domestic terms)

As a call option is short dividends (call is an option on the price of underlying, not the total return of the underlying), a quanto call option is short correlation. A quanto put option is therefore slightly long correlation. In both cases, we assume the FX rate is the foreign security currency measured in domestic payout terms.

Intuitively, we can see a quanto call option is short correlation by assuming the dividend yield and both currency risk-free rates are all zero and comparing its value to a vanilla call option priced in the (foreign) security currency. If correlation is high, the vanilla call option is worth more than the quanto call option (as FX moves in favour of the investor if the price of the security rises). The reverse is also true (negative correlation causes a vanilla call option to be worth less than a quanto call option). As the price of a vanilla (single currency) call does not change due to the correlation of the underlying with the FX rate, this shows a quanto call option is short correlation.

CHAPTER 6

RELATIVE VALUE AND CORRELATION TRADING

Advanced investors often use equity derivatives to gain different exposures; for example, relative value or the jumps on earnings dates. We demonstrate how this can be done and also reveal how profits from equity derivatives are both path dependent and dependent on the frequency of delta hedging.

6.1: RELATIVE VALUE TRADING

Relative value is the name given to a variety of trades that attempt to profit from the mean reversion of two related assets that have diverged. The relationship between the two securities chosen can be fundamental (different share types of same company or significant cross-holding) or statistical (two stocks in same sector). Relative value can be carried out via cash (or delta-1), options or outperformance options.

TRADES ARE USUALLY CHOSEN ON CORRELATED ASSETS

The payout of a relative value trade on two uncorrelated securities is completely random, and the investor on average gains no benefit. However, if two securities have a strong fundamental or statistical reason to be correlated, they can be thought of as trading in a similar direction with a random noise component. Assuming the correlation between the securities is sufficiently strong, the noise component should mean revert. Relative value trades attempt to profit from this mean reversion. There are five main types of relative value trades.

Random
Noise
Big

- **Dual listing.** If a share trades on different exchanges (eg, an ADR), the two prices should be equal. This is not always the case due to execution risk (different trading times) and perhaps due to indexation flow. Non-fungible shares or those with shorting restrictions are most likely to show the largest divergence in price. Of all relative value trades, dual-listing ones are likely to show the strongest correlation.
- **Share class.** If there is more than one type of share, perhaps with voting or ownership restrictions, then the price of these shares can diverge from one another. For example, preference shares typically have a higher dividend to compensate for lack of voting rights, but suffer from less liquidity and (normally) exclusion from equity indices. During special situations, for example, during the Porsche/VW saga, the difference in price between the two shares can diverge dramatically.
- **Cross-holding.** If one company (potentially a holding company) owns a significant amount of another company, the prices of the two companies will be linked. Sometimes putting on a cross-holding trade is difficult in practice due to the high borrow cost of the smaller company. This trade is also known as a stub trade when the investor wants pure exposure to the larger company, and hedges out the unwanted exposure to the equity holdings of the larger company. Potentially, these trades can occur when a larger company spins off a subsidiary but keeps a substantial stake post spin-off.

- **Event-driven.** In the event of a takeover that is estimated to have a significant chance of succeeding, the share prices of the acquiring and target company should be correlated. The target will usually trade at a discount to the bid price, to account for the probability the deal falls through (although if the offer is expected to be improved, or beaten by another bidder, the target could trade above the offer price).
- **Long-short.** If a long and short position is initiated in two securities that do not have one of the above four reasons to be correlated, it is a long-short trade. The correlation between the two securities of a long-short trade is likely to be lower than for other relative values trades. Because of this, often two stocks within a sector are chosen, as they should have a very high correlation and the noise component is likely to be bounded (assuming market share and profitability is unlikely to change substantially over the period of the relative value trade).

Long-short can focus returns on stock picking ability (which is c10% of equity return)

General market performance is typically responsible for c70% of equity returns, while c10% is due to sector selection and the remaining c20% due to stock picking. If an investor wishes to focus returns on the proportion due to sector or stock picking, they can enter into a long position in that security and a short position in the appropriate market index (or vice versa). This will focus returns on the c30% due to sector and stock selection. Typically, relatively large stocks are selected, as their systematic risk (which should cancel out in a relative value trade) is usually large compared to specific risk. Alternatively, if a single stock in the same sector (or sector index) is used instead of the market index, then returns should be focused on the c20% due to stock picking within a sector.

Long/short allows a lot of market Risk to be hedged and focus only on company performance

SIZE OF POSITIONS SHOULD BE WEIGHTED BY BETA

If the size of the long-short legs are chosen to have equal notional (share price \times number of shares \times FX), then the trade will break even if both stock prices go to zero. However, the legs of the trade are normally weighted by beta to ensure the position is market neutral for more modest moves in the equity market. The volatility (historical or implied) of the stock divided by the average volatility of the market can be used as an estimate of the beta.

DELTA-1, OPTIONS AND OUTPERFORMANCE OPTIONS

Relative value trades can be implemented via cash/delta-1, vanilla options or outperformance options. They have very different trade-offs between liquidity and risk. No one method is superior to others, and the choice of which instrument to use depends on the types of liquidity and risk the investor is comfortable with.

Figure 85. Different Methods of Relative Value Trading

Asset Class	Position	Benefits	Disadvantages
Cash/delta-1	Long A, short B using stock/CFD, future, forwards, total return swap or ETF	High liquidity (volatility products might not be available)	Unlimited risk
Options	Long call on A, short call on B (or put/call spread/put spread)	Limited downside on long leg and convex payoff	Unlimited risk on short side (unless call spreads/put spreads)
Outperformance option	Long outperformance option on A vs B	Limited downside and convex payoff	Poor liquidity/wide bid-offer spreads

(1) CASH/DELTA-1: BEST LIQUIDITY, BUT UNLIMITED RISK

The deepest and most liquid market for relative value trades is the cash (or delta-1) market. While there are limited restrictions in the size or stocks available, the trade can suffer potentially unlimited downside. While there are many similarities between cash or delta-1 instruments, there are also important differences.

Benefits of more beneficial taxation can be shared

For many delta-1 products, the presence of investors with more beneficial taxation can result in more competitive pricing. Products that have to be based in one location, such as ETFs, suffer from being unable to benefit from the different taxation of other investors.

(2) OPTIONS: CONVEX PAYOFF AND CAN LIMIT DOWNSIDE ON LONG LEG

Options can be used in place of stock or delta-1 for either the long or short leg, or potentially both. Options offer convexity, allowing a position to profit from the expected move while protecting against the potentially unlimited downside. Often a relative value trade will be put on in the cash/delta-1 market, and the long leg rotated into a call once the long leg is profitable (in order to protect profits). While volatility is a factor in determining the attractiveness of using options, the need for safety or convexity is normally the primary driver for using options (as relative value traders do not delta hedge, the change in implied volatility is less of a factor in profitability than the delta/change in equity market). Investors who are concerned about the cost of options can cheapen the trade by using call spreads or put spreads in place of vanilla calls or puts.

Weighting options by volatility is similar to weighting by beta and roughly zero cost

The most appropriate weighting for two relative value legs is beta weighting the size of the delta hedge of the option (ie, same beta \times number of options \times delta \times FX), rather than having identical notional (share price \times number of options \times FX). Beta weighting ensures the position is market neutral. Volatility weighting can be used as a substitute for beta weighting, as volatility divided by average volatility of the market is a reasonable estimate for beta. Volatility weighting ATM (or ATMf) options is roughly zero cost, as the premium of ATM options is approximately linear in volatility.

Choosing strike and maturity of option is not trivial

One disadvantage of using options in place of equity is the need to choose a maturity, although some investors see this as an advantage as it forces a view to be taken on the duration or exit point of the trade at inception. If the position has to be closed or rolled before expiry, there are potentially mark-to-market risks. Similarly, the strike of the option needs to be chosen, which can be ATM (at the money), ATMf (ATM forward), same percentage of spot/forward or same delta. Choosing the same delta of an OTM option means trading a strike further away from spot/forward for the more volatile asset (as delta increases as volatility increases). We note that trading the same delta option is not the same as volatility weighting the options traded as delta is not linear in volatility. Delta also does not take into account the beta of the underlyings.

(3) OUTPERFORMANCE OPTIONS: LIMITED DOWNSIDE BUT LOW LIQUIDITY

Outperformance options are ideally suited to relative value trades, as the maximum loss is the premium paid and the upside is potentially unlimited. However, outperformance options suffer from being relatively illiquid. While pricing is normally cheaper than vanilla options (for normal levels of correlation), it might not be particularly appealing depending on the correlation between the two assets. As there are usually more buyers than sellers of outperformance options, some hedge funds use outperformance options to overwrite their relative value trades.

6.2: RELATIVE VALUE VOLATILITY TRADING

Volatility investors can trade volatility pairs in the same way as trading equity pairs. For indices, this can be done via options, variance swaps or futures on a volatility index (such as the VIX or vStoxx). For indices that are popular volatility trading pairs, if they have significantly different skews this can impact the volatility market. Single-stock relative value volatility trading is possible, but less attractive due to the wider bid-offer spreads.

TWO WAYS TO PROFIT FROM VOLATILITY PAIR TRADING

When a pair trade between two equities is attempted, the main driver of profits is from a mean reversion of the equity prices. With volatility relative value trading, there are two ways of profiting:

- **Mean reversion.** In the same way an equity pair trade profits from a mean reversion of stock prices, a volatility pair trade can profit from a mean reversion of implied volatility. For short-term trades, mean reversion is the primary driver for profits (or losses). For relative value trades using forward starting products (eg, futures on volatility indices), this is the only driver of returns as forward starting products have no carry. The method for finding suitable volatility pair trades that rely on a short-term mean reversion are similar to that for a vanilla pair trade on equities.
- **Carry.** For an equity pair trade, the carry of the position is not as significant as, typically, the dividend yields of equities do not differ much from one another and are relatively small compared to the movement in spot. However, the carry of a volatility trade (difference between realised volatility and implied volatility) can be significant. As the duration of a trade increases, the carry increases in importance. Hence, for longer term volatility pair trades it is important to look at the difference between realised and implied volatility.

IMPLIED VOL SPREAD BETWEEN PAIRS IS KEPT STABLE

While the skew of different indices is dependent on correlation, traders tend to keep the absolute difference in implied volatility stable due to mean reversion. This is why if equity markets move down, the implied volatility of the S&P500 or FTSE (as they are large diversified indices that hence have high skew) tends to come under pressure, while the implied volatility of country indices with fewer members, such as the DAX, are likely to be supported. The SX5E tends to lie somewhere in between, as it has fewer members than the S&P500 or FTSE but is more diverse than other European country indices. Should markets rise, the reverse tends to occur (high skew indices implieds are lifted, low skew implieds are weighed on).

Difference between implieds is key, not the absolute level of each implied

We note that for returns due to mean reversion, it is not the absolute level of volatility that is key but the difference. For example, let us assume stock A implieds trade between 20% and 25% while stock B implieds trade between 30% and 35%. If stock A is at 25% implied (top of range) while stock B implied is at 30% implied (bottom of range), a short A volatility long B volatility position should be initiated. This is despite the 25% implied of A being less than the 30% implied of B.

VOLATILITY PAIR TRADING GREEKS SIMILAR TO DISPERSION

In dispersion trading, a (normally short) index volatility position is traded against a basket of (normally long) single-stock volatility positions. This position of index volatility vs basket could be considered to be a pair trade where one leg is the index and the other leg is the basket. A pair trade can be carried out via straddles / strangles or variance swaps, just like dispersion. We shall assume that the pair trade is being carried out by delta hedging options, for trading via variance swaps simply replaces notional in the table below with the vega of the variance swap. The weighting of the legs in order to be vega / theta or gamma flat is similar to dispersion trading, as can be seen below.

Figure 86. Greeks of Option Pair Trades with Different Weightings (shorting low vol, long high vol)

Greeks	Theta-Weighted	Vega-Weighted	Dollar Gamma-Weighted
Theta	0	Pay (or negative/short)	Pay a lot (very negative/short)
Vega	Short	0	Long
Gamma	Very short	Short	0
Ratio high vol to low vol notional	$\sigma_{\text{low vol}} / \sigma_{\text{high vol}}$	1	$\sigma_{\text{high vol}} / \sigma_{\text{low vol}}$
Notional of high vol stock	Less than low vol	Equal to low vol	More than low vol

Sign of theta, vega and gamma depends on which way round the pair trade is initiated

The sign of theta, vega and gamma are based on a trade of shorting the lower volatility security and going long the higher volatility security (on an absolute basis) in order for easy comparison to dispersion trading (where, typically, the lower absolute volatility of the index is shorted against a long of the higher absolute volatility of the single stocks). For the reverse trade (short the higher absolute volatility security and long the lower absolute volatility security), the signs of the Greeks need to be reversed.

PAIR TRADES CAN BE THETA OR VEGA WEIGHTED

Theta and vega weighted are the most common methods of weighting pair trades. Dollar gamma weighted is rarely used and is included for completeness purposes only. Theta-weighted trades assume proportional volatility changes (eg, if stock A has 20% implied and stock B has 25% implied, if stock A rises from 20% to 30% implied that is a 50% increase and stock B rises 50% to 37.5% implied). Vega-weighted trades assume absolute volatility changes (eg, if stock A has 20% implied and stock B has 25% implied, if stock A rises from 20% to 30% that is a 10 volatility point increase and stock B rises 10 volatility points to 35% implied).

Pair trade between two securities of same type should be theta weighted

If a pair trade between two securities of the same type (ie, two indices, or two single stocks) is attempted, theta weighting is the most appropriate. This is because the difference between a low volatility security and a high volatility security (of the same type) usually increases as volatility increases (ie, a proportional move). If a pair trade between an index and a single stock is attempted, vega weighting is the best as the implied volatility of an index is dependent not only on single-stock implied volatility but also on implied correlation. As volatility and correlation tend to move in parallel, this means the payout of a vega-weighted pair trade is less dependent on the overall level of volatility (hence the volatility mispricing becomes a more significant driver of the P&L of the trade)¹⁸.

¹⁸ There is evidence to suggest that vega-weighted index vs single-stock pair trades on average associate 2%-5% too much weight to the single-stock leg compared to the index leg. However, as this is so small compared to stock specific factors, it should be ignored.

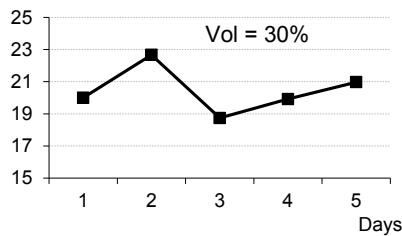
6.3: CORRELATION TRADING

The volatility of an index is capped at the weighted average volatility of its constituents. Due to diversification (or less than 100% correlation), the volatility of indices tends to trade significantly less than its constituents. The flow from both institutions and structured products tends to put upward pressure on implied correlation, making index implied volatility expensive. Hedge funds and proprietary trading desks try to profit from this anomaly by either selling correlation swaps, or through dispersion trading (going short index implied volatility and long single stock implied volatility). Selling correlation became an unpopular strategy following losses during the credit crunch, but demand is now recovering.

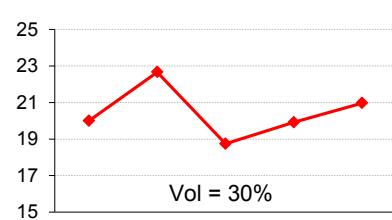
INDEX IMPLIED LESS THAN SINGLE STOCKS

The volatility of an index is capped by the weighted average volatility of its members. In order to show this we shall construct a simple index of two equal weighted members who have the same volatility. If the two members are 100% correlated with each other, then the volatility of the index is equal to the volatility of the members (as they have the same volatility and weight, this is the same as the weighted average volatility of the constituents).

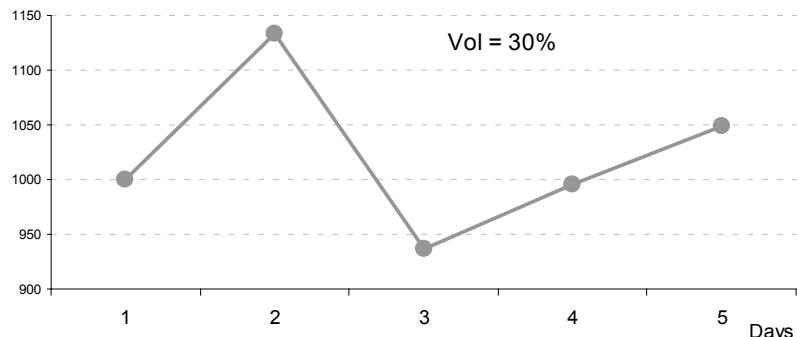
Figure 87. Stock 1



Stock 2 (100% correlation)



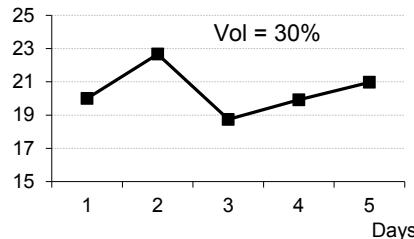
Index (of stock 1 and stock 2)



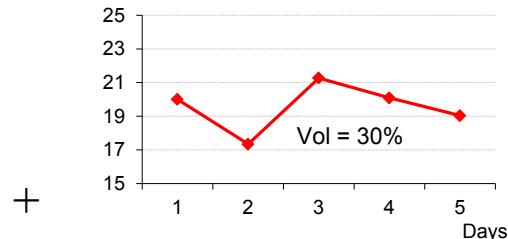
Volatility of index has floor at zero when there is very low correlation

If we take a second example of two equal weighted index members with the same volatility, but with a negative 100% correlation (ie, as low as possible), then the index is a straight line with zero volatility.

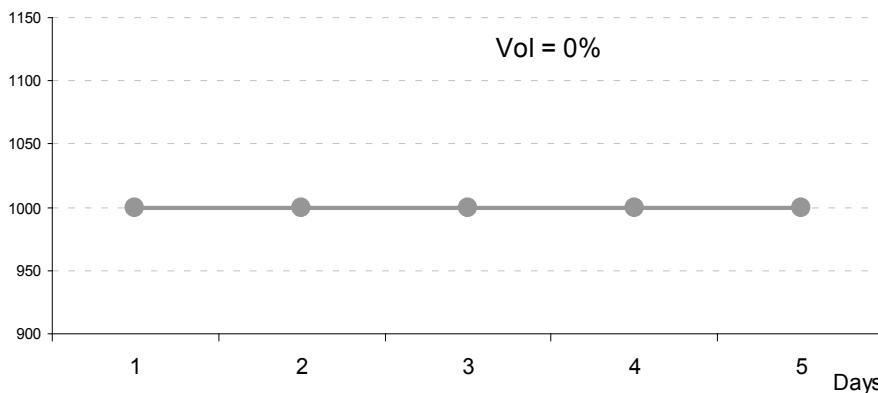
Figure 88. Stock 1



Stock 2 (-100% correlation)



Index (of stock 1 and stock 2)



Index vol is bounded by zero and weighted average single stock vol

While the simple examples above have an index with only two members, results for a bigger index are identical. Therefore, the equation below is true. While we are currently examining historical volatility, the same analysis can be applied to implied volatility. In this way, we can get an implied correlation surface from the implied volatility surfaces of an index and its single-stock members. However, it is usually easiest to look at variance swap levels rather than implied volatility to remove any strike dependency.

$$0 \leq \sigma_I^2 \leq \sum_{i=1}^n w_i^2 \sigma_i^2$$

where

σ_I = index volatility

σ_i = single stock volatility (of i^{th} member of index)

w_i = single stock weight in index (of i^{th} member of index)

n = number of members of index

INDEX CORRELATION CAN BE ESTIMATED FROM VARIANCE

If the correlation of all the different members of an index is assumed to be identical (a heroic assumption, but a necessary one if we want to have a single measure of correlation), the correlation implied by index and single-stock implied volatility can be estimated as the variance of the index divided by the weighted average single-stock variance. This measure is a point or two higher than the actual implied correlation but is still a reasonable approximation.

$$\rho_{\text{imp}} = \frac{\sigma_I^2}{\sum_{i=1}^n w_i \sigma_i^2}$$

where

ρ_{imp} = implied correlation (assumed to be identical between all index members)

Proof implied correlation can be estimated by index variance divided by single stock variance

The formula for calculating the index volatility from the members of the index is given below.

$$\sigma_I^2 = \sum_{i=1}^n w_i^2 \sigma_i^2 + \sum_{i=1, j \neq i}^n w_i w_j \sigma_i \sigma_j \rho_{ij}$$

where

ρ_{ij} = correlation between single stock i and single stock j

If we assume the correlations between each stock are identical, then this correlation can be implied from the index and single stock volatilities.

$$\rho_{imp} = \frac{\sigma_I^2 - \sum_{i=1}^n w_i^2 \sigma_i^2}{\sum_{i=1, j \neq i}^n w_i w_j \sigma_i \sigma_j}$$

Assuming reasonable conditions (correlation above 15%, c20 members or more, reasonable weights and implied volatilities), this can be rewritten as the formula below.

$$\rho_{imp} = \frac{\sigma_I^2}{(\sum_{i=1}^n w_i \sigma_i)^2}$$

This can be approximated by the index variance divided by the weighted average single-stock variance.

$$\rho_{imp} \approx \frac{\sigma_I^2}{\sum_{i=1}^n w_i \sigma_i^2} \text{ eg, if index variance}=20\% \text{ and members average variance}=25\%, \\ Q \approx 64\%.$$

This approximation is slightly too high (c2pts) due to Jensen's inequality (shown below).

$$(\sum_{i=1}^n w_i \sigma_i)^2 \leq \sum_{i=1}^n w_i \sigma_i^2$$

STRUCTURED PRODUCTS LIFT IMPLIED CORRELATION

Using correlation to visually cheapen payouts through worst-of/best-of options is common practice for structured products. Similarly, the sale of structured products, such as Altiplano (which receives a coupon provided none of the assets in the basket has fallen), Everest (payoff on the worst performing) and Himalayas (performance of best share of index), leave their vendors short implied correlation. This buying pressure tends to lift implied correlation above fair value. We estimate that the correlation exposure of investment banks totals c€200mn per percentage point of correlation. The above formulae can show that two correlation points is equivalent to 0.3 to 0.5 (single-stock) volatility points. Similarly, the fact that institutional investors tend to call overwrite on single stocks but buy protection on an index also leads to buying pressure on implied correlation. The different methods of trading correlation are shown below.

- **Dispersion trading.** Going short index implied volatility and going long single-stock implied volatility is known as a dispersion trade. As a dispersion trade is short Volga, or vol of vol, the implied correlation sold should be c10pts higher value than for a

correlation swap. A dispersion trade was historically put on using variance swaps, but the large losses from being short single stock variance led to the single stock market becoming extinct. Now dispersion is either put on using straddles, or volatility swaps. Straddles benefit from the tighter bid-offer spreads of ATM options (variance swaps need to trade a strip of options of every strike). Using straddles does imply greater maintenance of positions, but some firms offer delta hedging for 5-10bp. A disadvantage of using straddles is that returns are path dependent. For example, if half the stocks move up and half move down, then the long single stocks are away from their strike and the short index straddle is ATM.

- **Correlation swaps.** A correlation swap is simply a swap between the (normally equal weighted) average pairwise correlation of all members of an index and a fixed amount determined at inception. Market value-weighted correlation swaps are c5 correlation points above equal weighted correlation, as larger companies are typically more correlated than smaller companies. While using correlation swaps to trade dispersion is very simple, the relative lack of liquidity of the product is a disadvantage. We note the levels of correlation sold are typically c5pts above realised correlation.
- **Covariance swaps.** While correlation swaps are relatively intuitive and are very similar to trading correlation via dispersion, the risk is not identical to the covariance risk of structured product sellers (from selling options on a basket). Covariance swaps were invented to better hedge the risk on structure books, and they pay out the correlation multiplied by the volatility of the two assets.
- **Basket options.** Basket options (or options on a basket) are similar to an option on an index, except the membership and weighting of the members does not change over time. The most popular basket option is a basket of two equal weighted members, usually indices.
- **Worst-of/best-of option.** The pricing of worst-of and best-of options has a correlation component. These products are discussed in the section [5.2 Worst-of/Best-of Options](#) in the Forward Starting Products and Light Exotics chapter.
- **Outperformance options.** Outperformance options pricing has as an input the correlation between the two assets. These products are also discussed in the section [5.3 Outperformance Options](#) in the Forward Starting Products and Light Exotics chapter.

Implied correlation of dispersion and level of correlation swap are not the same measure

We note that the profit from theta-weighted (explained later in section) variance dispersion is roughly the difference between implied and realised correlation multiplied by the average single-stock volatility. As correlation is correlated to volatility, this means the payout when correlation is high is increased (as volatility is high) and the payout when correlation is low is decreased (as volatility is low). A short correlation position from going long dispersion (short index variance, long single-stock variance) will suffer from this as profits are less than expected and losses are greater. Dispersion is therefore short vol of vol; hence, implied correlation tends to trade c10 correlation points more than correlation swaps (which is c5 points above realised correlation). We note this does not necessarily mean a long dispersion trade should be profitable (as dispersion is short vol of vol, the fair price of implied correlation is above average realised correlation).

Implied vs realised correlation increases for low levels of correlation

For example, in normal market conditions the SX5E and S&P500 will have an implied correlation of 50-70 and a realised of 30-60. If realised correlation is 30, implied will tend to be at least 50 (as investors price in the fact correlation is unlikely to be that low for very long; hence, the trade has more downside than upside). The NKY tends to have correlation levels ten points below the SX5E and SPX.

Figure 89. Different Types of Correlation

Traded level (c5pts above payout)	Payout	Formula
Implied correlation (market value weighted)	Realised correlation (using realised vol rather than implied vol in formula)	$\frac{\sigma_I^2 - \sum_{i=1}^n w_i^2 \sigma_i^2}{\sum_{i=1, j \neq i}^n w_i w_j \sigma_i \sigma_j} \approx \frac{\sigma_I^2}{\sum_{i=1}^n w_i \sigma_i}$
c10pts above implied correlation (as not short volga)		
Correlation swap (market value weighted)	Pairwise realised correlation (market value weighted)	$\frac{\sum_{i=1, j > i}^n w_i w_j \rho_{ij}}{\sum_{i=1}^n w_i w_j}$
c5pts above market value weighted correlation		
Correlation swap (equal weighted)	Pairwise realised correlation (equal weighted)	$\frac{2}{n(n-1)} \rho_{ij}$

DISPERSION MARKET SHRANK POST CREDIT CRUNCH

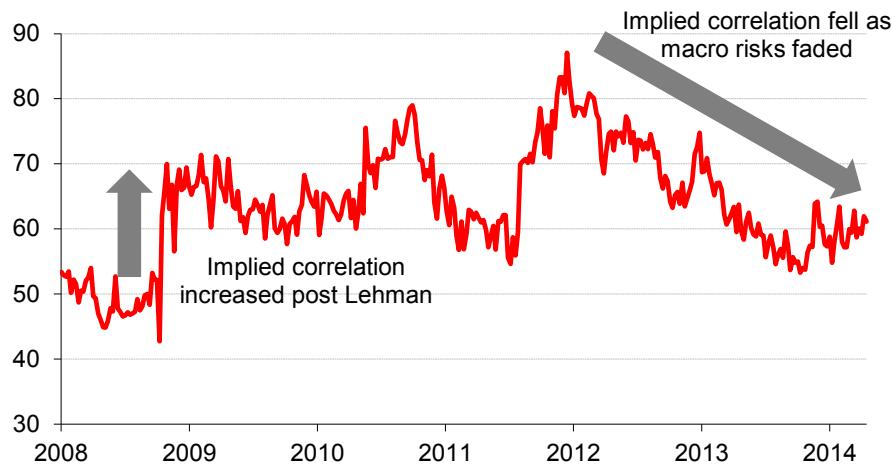
Selling correlation led to severe losses when the market collapsed in 2008, as implied correlation spiked to c90%, which led many investors to cut back exposures or leave the market. Similar events occurred in the market during the May 2010 correction. The amount of crossed vega has been reduced from up to €100mn at some firms to €5-20mn now (crossed vega is the amount of offsetting single-stock and index vega, ie, €10mn crossed vega is €10mn on single stock and €10mn on index). Similarly, the size of trades has declined from a peak of €2.5mn to €0.5mn vega now.

Figure 90. CBOE Implied Correlation Tickers and Expiries

Expiry S&P500	Expiry Top 50 Stocks	Ticker	Start Date	End Date
Dec-09	Jan-10	ICJ	Nov-07	Nov-09
Dec-10	Jan-11	JCJ	Nov-08	Nov-10
Dec-11	Jan-12	KCJ	Nov-09	Nov-11
Dec-12	Jan-13	ICJ	Nov-10	Nov-12
Dec-13	Jan-14	JCJ	Nov-11	Nov-13
Dec-14	Jan-15	KCJ	Nov-12	Nov-14
Dec-15	Jan-16	ICJ	Nov-13	Nov-15
Dec-16	Jan-17	JCJ	Nov-14	Nov-16
Dec-17	Jan-18	KCJ	Nov-15	Nov-17

CBOE INDICES ALLOW CORRELATION TO BE PLOTTED

There are now correlation indices for calculating the implied correlation of dispersion trades calculated by the CBOE. As there are 500 members of the S&P500, the CBOE calculation only takes the top 50 stocks (to ensure liquidity). There are three correlation indices tickers (ICJ, JCJ and KCJ), but only two correlation indices are calculated at any one time. On any date one correlation index has a maturity up to one year, and another has a maturity between one and two years. The calculation uses December expiry for S&P500 options, and the following January expiry for the top 50 members as this is the only listed expiry (US single stocks tend to be listed for the month after index triple witching expiries). The index is calculated until the previous November expiry, as the calculation tends to be very noisy for maturities only one month to December index expiry. On the November expiry, the one month maturity (to S&P500 expiry) index ceases calculation, and the previously dormant index starts calculation as a two-year (and one-month) maturity index. For the chart below, we use the longest dated available index.

Figure 91. CBOE Implied Correlation (rolling maturity between 1Y and 2Y)

CORRELATION SWAPS HAVE PURE CORRELATION EXPOSURE

Correlation swaps (which, like variance swaps, are called swaps but are actually forwards) simply have a clean payout of the (normally equal-weighted) correlation between every pair in the basket less the correlation strike at inception. Correlation swaps usually trade on a basket, not an index, to remove the names where a structured product has a particularly high correlation risk. Half of the underlyings are typically European, a third US and the final sixth Asian stocks. The product started trading in 2002 as a means for investment banks to reduce their short correlation exposure from their structured products books. While a weighted pairwise correlation would make most sense for a correlation swap on an index, the calculation is typically equal-weighted as it is normally on a basket.

Equal-weighted correlation is c5 points below market value-weighted

Market value-weighted correlation swaps tends to trade c5 correlation points above realised correlation (a more sophisticated methodology is below). This level is c10 correlation points below the implied correlation of dispersion (as dispersion payout suffers from being short volga). In addition, the correlation levels for equal-weighted correlations tends to be c5 correlation points lower than for market value-weighted, due to the greater weight allocated to smaller – and hence less correlated – stocks. The formula for the payout of a correlation swap is below.

Correlation swap payoff

$$(\rho_K - \rho) \times \text{Notional}$$

where

Notional = notional paid (or received) per correlation point

ρ_K = strike of correlation swap (agreed at inception of trade)

$$\rho = \frac{2}{n(n-1)} \rho_{ij} \quad (\text{equal weight correlation swap})$$

$$\rho = \frac{\sum_{i=1, j>i}^n w_i w_j \rho_{ij}}{\sum_{i=1, j>i}^n w_i w_j} \quad (\text{market value weight correlation swap})$$

n = number of stocks in basket

Correlation swaps tend to trade c5 correlation points above realised

A useful rule of thumb for the level of a correlation swap is that it trades c5 correlation points above realised correlation (either equal-weighted or market value-weighted, depending on the type of correlation swap). However, for very high or very low values of correlation, this formula makes less sense. Empirically, smaller correlations are typically more volatile than higher correlations. Therefore, it makes sense to bump the current realised correlation by a larger amount for small correlations than for higher correlations (correlation swaps should trade above realised due to demand from structured products). The bump should also tend to zero as correlation tends to zero, as having a correlation swap above 100% would result in arbitrage (can sell correlation swap above 100% as max correlation is 100%). Hence, a more accurate rule of thumb (for very high and low correlations) is given by the formula below.

Correlation swap level = $\rho + \alpha(1 - \rho)$

where

ρ = realised correlation

α = bump factor (typical $\alpha = 0.1$)

Maturity of correlation swap is typically between one and three years

Structured products typically have a maturity of 5+ years; however, many investors close their positions before expiry. The fact that a product can also delete a member within the lifetime of the product has led dealers to concentrate on the three-year maturity rather than

five-year. As the time horizon of hedge funds is short dated, correlation swaps typically trade between one and three years. The size is usually between €250k and €1,000k.

Correlation swaps suffer from lack of liquidity

The market for correlation swaps has always been smaller than for dispersion. Whereas the variance swap or option market has other market participants who ensure liquidity and market visibility, the investor base for correlation swaps is far smaller. This can be an issue should a position wish to be closed before expiry. It can also cause mark-to-market problems. The correlation swap market grew from 2002 onwards until the credit crunch, when investor appetite for exotic products disappeared. At its peak, it is estimated that some structured derivative houses shed up to c10% of their short correlation risk to hedge funds using correlation swaps.

DISPERSION IS THE MOST POPULAR METHOD OF TRADING CORRELATION

As the levels of implied correlation are usually overpriced (a side effect of the short correlation position of structured product sellers), index implied volatility is expensive when compared with the implied volatility of single stocks. A long dispersion trade attempts to profit from this by selling index implied and going long single-stock implied¹⁹. Such a long dispersion trade is short implied correlation. While dispersion is the most common method of trading implied correlation, the payoff is also dependent on the level of volatility. The payout of (theta-weighted) dispersion is shown below. Because of this, and because correlation is correlated to volatility, dispersion trading is short vol of vol (volga).

$$P & L_{\text{dispersion}} = \sum_{i=1}^n w_i \sigma_i^2 (\rho_{imp} - \rho)$$

There are four instruments that can be used to trade dispersion:

- **Straddle (or call) dispersion.** Using ATM straddles to trade dispersion is the most liquid and transparent way of trading. Because it uses options, the simplest and most liquid volatility instrument, the pricing is usually the most competitive. Trading 90% strike rather than ATM allows higher levels of implied correlation to be sold. Using options is very labour intensive, however, as the position has to be delta-hedged (some firms offer delta hedging for 5-10bp). In addition, the changing vega of the positions needs to be monitored, as the risks are high given the large number of options that have to be traded. In a worst-case scenario, an investor could be right about the correlation position but suffer a loss from lack of vega monitoring. We believe that using OTM strangles rather than straddles is a better method of using vanilla options to trade dispersion as OTM strangles have a flatter vega profile. This means that spot moving

¹⁹ Less liquid members of an index are often excluded, eg, CRH for the Euro STOXX 50.

away from strike is less of an issue, but we acknowledge that this is a less practical way of trading.

- **Variance swap dispersion.** Because of the overhead of developing risk management and trading infrastructure for straddle dispersion, many hedge funds preferred to use variance swaps to trade dispersion. With variance dispersion it is easier to see the profits (or losses) from trading correlation than it is for straddles. Variance dispersion suffers from the disadvantage that not all the members of an index will have a liquid variance swap market. Since 2008, the single-stock variance market has disappeared due to the large losses suffered from single-stock variance sellers (as dispersion traders want to go long single-stock variance, trading desks were predominantly short single-stock variance). It is now rare to be able to trade dispersion through variance swaps.
- **Volatility swap dispersion.** Since liquidity disappeared from the single-stock variance market, investment banks have started to offer volatility swap dispersion as an alternative. Excluding dispersion trades, volatility swaps rarely trade.
- **Gamma swap dispersion.** Trading dispersion via gamma swaps is the only ‘fire and forget’ way of trading dispersion. As a member of an index declines, the impact on the index volatility declines. As a gamma swap weights the variance payout on each day by the closing price on that day, the payout of a gamma swap similarly declines with spot. For all other dispersion trades, the volatility exposure has to be reduced for stocks that decline and increased for stocks that rise. Despite the efforts of some investment banks, gamma swaps never gained significant popularity.

Need to decide on weighting scheme for dispersion trades

While a dispersion trade always involves a short index volatility position and a long single-stock volatility position, there are different strategies for calculating the ratio of the two trade legs. If we assume index implied is initially 20%, if it increases to 30% the market could be considered to have risen by ten volatility points or risen by 50%. If the market is considered to rise by ten volatility points and average single-stock implied is 25%, it would be expected to rise to 50% (vega-weighted). If the market is considered to rise by 50% and average single-stock implied is 30%, it would be expected to rise to 45% (theta- or correlation-weighted). The third weighting, gamma-weighted, is not often used in practice.

- **Vega-weighted.** In a vega-weighted dispersion, the index vega is equal to the sum of the single-stock vega. If both index and single-stock vega rise one volatility point, the two legs cancel and the trade neither suffers a loss or reveals a profit.
- **Theta- (or correlation-) weighted.** Theta weighting means the vega multiplied by $\sqrt{\text{variance}}$ (or volatility for volatility swaps) is equal on both legs. This means there is a smaller single-stock vega leg than for vega weighting (as single-stock volatility is larger than index volatility, so it must have a smaller vega for vega \times volatility to be equal). Under theta-weighted dispersion, if all securities have zero volatility, the theta of both

the long and short legs cancels (and total theta is therefore zero). Theta weighting can be thought of as correlation-weighted (as correlation \approx index var / average single stock var = ratio of single-stock vega to index vega). If volatility rises 1% (relative move) the two legs cancel and the dispersion breaks even.

- **Gamma-weighted.** Gamma weighting is the least common of the three types of dispersion. As gamma is proportional to vega/vol, then the vega/vol of both legs must be equal. As single-stock vol is larger than index vol, there is a larger single-stock vega leg than for vega-weighted.

Greeks of dispersion trading depend on weighting used

The Greeks of a dispersion trade²⁰ are very much dependent on the vega weighting of the two legs. The easiest weighting to understand is a vega-weighted dispersion, which by definition has zero vega (as the vega of the short index and long single-stock legs are identical). A vega-weighted dispersion is, however, short gamma and short theta (ie, have to pay theta).

Theta-weighted dispersion needs a smaller long single-stock leg than the index leg (as reducing the long position reduces theta paid on the long single-stock leg to that of the theta earned on the short index leg). As the long single-stock leg is smaller, a theta-weighted dispersion is very short gamma (as it has less gamma than vega-weighted, and vega-weighted is short gamma).

Gamma-weighted dispersion needs a larger long single-stock leg than the index leg (as increasing the long position increases the gamma to that of the short index gamma). As the long single-stock leg is larger, the theta paid is higher than that for vega-weighted.

Figure 92. Greeks of Dispersion Trades with Different Weightings

Greeks	Theta-Weighted	Vega-Weighted	Dollar Gamma-Weighted
Theta	0	Short/pay	Very short/pay a lot
Vega	Short	0	Long
Gamma	Very short	Short	0
Ratio single stock vega to index vega	$\sigma_{\text{Index}} / \sigma_{\text{Single Stock}}$	1	$\sigma_{\text{Single Stock}} / \sigma_{\text{Index}}$
Total single-stock vega	Less than index	Equal to index	More than index

²⁰ The mathematical proof of the Greeks is outside of the scope of this report.

Figure 93. Breakevens for Theta, Vega and Gamma-Weighted Dispersion

	Theta-Weighted	Vega-Weighted	Dollar Gamma-Weighted
Start of trade			
Index vol (vol pts)	20.0	20.0	20.0
Average single-stock vol (vol pts)	25.0	25.0	25.0
Implied correlation (correlation pts)	64.0	64.0	64.0
Trade size			
Index vega (k)	100	100	100
Single-stock vega (k)	80	100	125
End of trade if index rises 10 vol pts and trade breaks even			
Index vol (vol pts)	30.0	30.0	30.0
Avg single-stock vol (vol pts)	37.5	35.0	33.0
Implied correlation (correlation pts)	64.0	73.5	82.6
Change			
Change in index vol (vol. pts)	10.0	10.0	10.0
Change in single-stock vol (vol pts)	12.5	10.0	8.0
Change in implied correlation (correlation pts)	0.0	9.5	18.6
Change (%)			
Change in index vol (%)	50%	50%	50%
Change in single-stock vol (%)	50%	40%	32%
Change in implied correlation (%)	0.0%	14.8%	29.1%

DISPERSION CAN BE THETA, VEGA AND GAMMA-WEIGHTED

The Figure 93 above shows the different weightings for theta, vega and dollar-gamma-weighted dispersion. The change in volatility for the different trades to break even is shown. As can be seen, only theta-weighted dispersion gives correlation exposure (ie, if realised correlation is equal to implied correlation, theta-weighted dispersion breaks even).

Theta-weighted dispersion is best weighting for almost pure correlation exposure

The sole factor that determines if theta-weighted dispersion makes a profit or loss is the difference between realised and implied correlation. For timing entry points for theta-weighted dispersion, we believe investors should look at the implied correlation of an index (as theta-weighted dispersion returns are driven by correlation). Note that theta-weighted dispersion breaks even if single stock and index implied moves by the same percentage amount (eg, index vol of 20%, single-stock vol of 25% and both rise 50% to 30% and 37.5%, respectively).

Vega-weighted dispersion gives hedged exposure to mispricing of correlation

When a dispersion trade is vega-weighted, it can be thought of as being the sum of a theta-weighted dispersion (which gives correlation exposure), plus a long single-stock volatility position. This volatility exposure can be thought of as a hedge against the short correlation position (as volatility and correlation are correlated); hence, a vega-weighted dispersion gives greater exposure to the mispricing of correlation. When looking at the optimal entry point for vega-weighted dispersion, it is better to look at the difference between average single-stock volatility and index volatility (as this applies an equal weight to both legs, like in a vega-weighted dispersion). Note that vega-weighted dispersion breaks even if single stock and index implied moves by the same absolute amount (eg, index vol of 20%, single-stock vol of 25% and both rise ten volatility points to 30% and 35%, respectively). Empirically, the difference between single-stock and index volatility (ie, vega-weighted dispersion) is not correlated to volatility²¹, which supports our view of vega-weighted dispersion being the best.

Gamma-weighted dispersion is rare, and not recommended

While gamma weighting might appear mathematically to be a suitable weighting for dispersion, in practice it is rarely used. It seems difficult to justify a weighting scheme where more single-stock vega is bought than index (as single stocks have a higher implied than index and, hence, should move more). We include the details of this weighting scheme for completeness, but do not recommend it.

²¹ Single-stock leg is arguably 2%-5% too large; however, slightly over-hedging the implicit short volatility position of dispersion could be seen as an advantage.

DISPERSION TRADES ARE SHORT VOL OF VOL (VOLGA)

The P&L of a theta-weighted dispersion trade is proportional to the spread between implied and realized market value-weighted correlation (ρ), multiplied by a factor that corresponds to a weighted average variance of the components of the index²².

$$P \& L_{\text{theta weighted dispersion}} = \sum_{i=1}^n w_i \sigma_i^2 (\rho_{imp} - \rho)$$

where:

ρ = market value weighted correlation

The payout of a theta-weighted dispersion is therefore equal to the difference in implied and realised correlation (market value-weighted pairwise realised correlation) multiplied by the weighted average variance. If vol of vol was zero and volatility did not change, then the payout would be identical to a correlation swap and both should have the same correlation price. If volatility is assumed to be correlated to correlation (as it is, as both volatility and correlation increase in a downturn) and the correlation component is profitable, the profits are reduced (as it is multiplied by a lower volatility). Similarly, if the correlation suffers a loss, the losses are magnified (as it is multiplied by a higher volatility). Dispersion is therefore short volga (vol of vol) as the greater the change in volatility, the worse the payout. To compensate for this short volga position, the implied correlation level of dispersion is c10 correlation points above the level of correlation swaps.

BASKET OPTIONS ARE MOST LIQUID CORRELATION PRODUCT

The most common product for trading correlation is a basket option (otherwise known as an option on a basket). If the members of a basket are identical to the members of an index and have identical weights, then the basket option is virtually identical to an option on the index. The two are not completely identical, as the membership and weight of a basket option does not change²³, but it can for an index (due to membership changes, rights issues, etc). The formula for basket options is below.

$$\text{Basket} = \sum_{i=1}^n w_i S_i \text{ where } S_i \text{ is the } i^{\text{th}} \text{ security in the basket}$$

$$\text{Basket call payoff at expiry} = \text{Max}(0, \sum_{i=1}^n w_i^2 S_i - K) \text{ where } K \text{ is the strike}$$

²² Proof of this result is outside the scope of this publication.

²³ Weighting for Rainbow options is specified at maturity based on the relative performance of the basket members, but discussion of these options is outside of the scope of this publication.

MOST POPULAR BASKET OPTIONS ARE ON TWO INDICES

While the above formula can be used for all types of basket, the most popular is a basket on two equal weighted indices. In this case the correlation traded is not between multiple members of a basket (or index) but the correlation between only two indices. As the options usually want the two indices to have identical value, it is easier to define the basket as the equal weighted sum of the two security returns (see the below formula setting $n = 2$). The previous formula could be used, but the weight w would not be 0.5 (would be $0.5 / S_i \text{ at inception}$).

$$\text{Basket} = \sum_{i=1}^n w \frac{S_i \text{ at expiry}}{S_i \text{ at inception}} \text{ where } S_i \text{ is the } i^{\text{th}} \text{ security (and } w \text{ normally } = 1/n)$$

PAYOUT IS BASED ON COVARIANCE, NOT CORRELATION

The payout of basket options is based on the correlation multiplied by the volatility of the two securities, which is known as covariance. The formula for covariance is shown below. As basket options are typically the payout of structured products, it is better to hedge the exposure using products whose payout is also based on covariance. It is therefore better to use covariance swaps rather than correlation swaps or dispersion to offset structured product risk.

$$\text{Covariance}(A,B) = \rho \sigma_A \sigma_B \text{ where } \rho \text{ is the correlation between A and B}$$

COVARIANCE SWAPS BETTER REPRESENT STRUCTURED PRODUCT RISK

The payout of structured products is often based on a basket option. The pricing of an option on a basket involves covariance, not correlation. If an investment bank sells an option on a basket to a customer and hedges through buying correlation (via correlation swaps or dispersion) there is a mismatch²⁴. Because of this, attempts were made to create a covariance swap market, but liquidity never took off.

Correlation swap payoff

$$[\text{Covariance}(A,B) - K_{\text{covariance}}] \times \text{Notional}$$

where

Notional = notional paid (or received) per covariance point

²⁴ Results in being short cross-gamma. Cross-gamma is the effect a change in the value of one underlying has on the delta of another.

ρ = correlation between A and B

σ_i = volatility of i

Covariance(A,B) = $\rho\sigma_A\sigma_B$ (note if A = B then covariance = variance as $\rho = 1$)

$K_{\text{covariance}}$ = strike of covariance swap (agreed at inception of trade)

6.4: TRADING EARNINGS ANNOUNCEMENTS/JUMPS

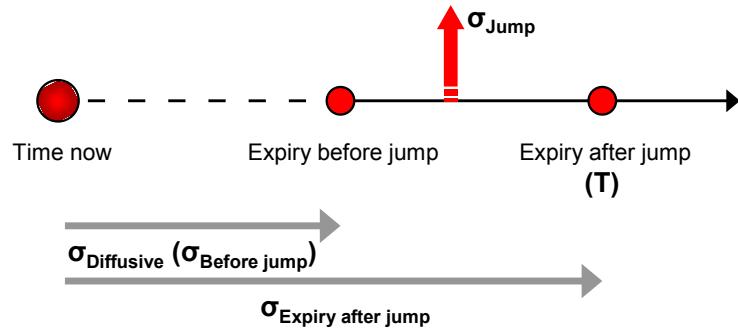
From the implied volatilities of near dated options, it is possible to calculate the implied jump on key dates. Trading these options in order to take a view on the likelihood of unanticipated (low or high) volatility on reporting dates is a very common strategy. We examine the different methods of calculating the implied jump, and show how the jump calculation should normalise for index term structure.

TOTAL VOL = DIFFUSIVE VOL + JUMP VOL

While stock prices under Black-Scholes are modelled as having a GBM (Geometric Brownian Motion) with constant volatility, in reality there are certain dates where there is likely to be more volatility than average. These key dates are usually reporting dates, but could also coincide with conference dates or investor days (in fact, any day where material non-public information is released to the public). The implied volatility of an option whose expiry is after a key date can be considered to be the sum of the normal diffusive volatility (normal volatility for the stock in the absence of any significantly material information being released) and the volatility due to the anticipated jump on the key date. While options of any expiry after the key date could be used, we shall assume the expiry chosen is the expiry just after the key date (to ensure the greatest percentage of the options' time value is associated with the jump).

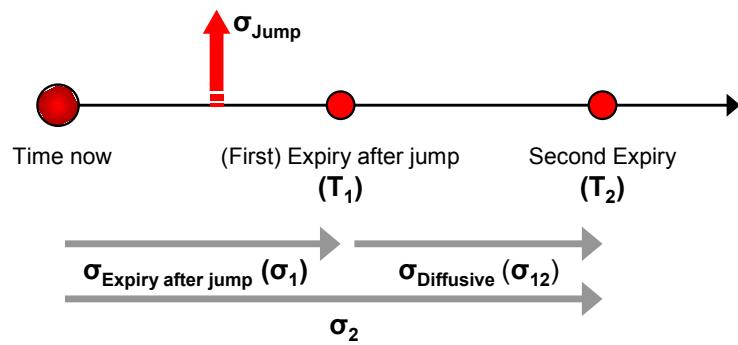
Implied jumps can be traded as part of a relative value trade

Trading an option whose expiry is just after the key date gives exposure not only to the implied jump, but also the normal diffusive volatility. If the expiry is far away then there should be another expiry just before the key date. In this case the long position in the expiry just after the key date can be hedged by shorting the expiry just before the key date. This relative value trade assumes the investor wishes to be long the implied volatility jump. If an investor wishes to short an expensive implied volatility jump then the reverse position (short expiry just after key date, long expiry just before key date) should be put on.

Figure 94. Implied Vol As Diffusive Vol Assumption

ESTIMATING DIFFUSIVE VOLATILITY IS NOT TRIVIAL

In order to calculate the implied jump due to a key date, the diffusive (normal) volatility of the stock needs to be estimated. While the diffusive volatility could be estimated by looking at historical volatility, it is usual to look at implied volatility (as there are several measures of historical volatility, but only one implied volatility). If there is an option that expires just before the key date, then the implied volatility of this option can be used. If not, the forward volatility after the key date is used as the estimate for the normal volatility of the security.

Figure 95. Forward Vol As Diffusive Vol Assumption

Implied jumps normally calculated for near-dated events

Implied jumps are normally only calculated for near-dated events, as the effect of the jump tends to be too diluted for far dated expiries (and hence would be untradeable taking bid-offer spreads into account). Forward starting options could be used to trade jumps after the first expiry, but the wider bid-offer spread could be greater than potential profits.

Forward volatility can be calculated with implied of two options

The calculation for forward volatility is derived from the fact variance (time weighted) is additive. The formula is given below (σ_x is the implied volatility for options of maturity T_x).

$$\sigma_{12} = \sqrt{\frac{\sigma_2^2 T_2 - \sigma_1^2 T_1}{T_2 - T_1}} = \text{forward volatility } T_1 \text{ to } T_2$$

JUMP VOLATILITY CAN BE CALCULATED FROM DIFFUSIVE VOLATILITY

As variance is additive, the volatility due to the jump can be calculated from the total volatility and the diffusive volatility. We note this assumes the normal diffusive volatility is constant, whereas volatility just after a reporting date is, in fact, typically $\frac{3}{4}$ of the volatility just before a reporting date (as the previously uncertain earnings are now known).

$$\sigma_{\text{Expiry after jump}}^2 T = \sigma_{\text{Jump}}^2 + \sigma_{\text{Diffusive}}^2 (T - 1)$$

$$\Rightarrow \sigma_{\text{Jump}} = \sqrt{(\sigma_{\text{Expiry after jump}}^2 T - \sigma_{\text{Diffusive}}^2 (T - 1))}$$

where

$\sigma_{\text{Expiry after jump}}$ = implied volatility of option whose expiry is after the jump

T = time to the expiry after jump ($= T_1$)

$\sigma_{\text{Diffusive}}$ = diffusive volatility ($\sigma_{\text{Before jump}}$ if there is an expiry before the jump, if not it is σ_{12})

σ_{Jump} = implied volatility due to the jump

IMPLIED JUMP CALCULATED FROM JUMP VOLATILITY

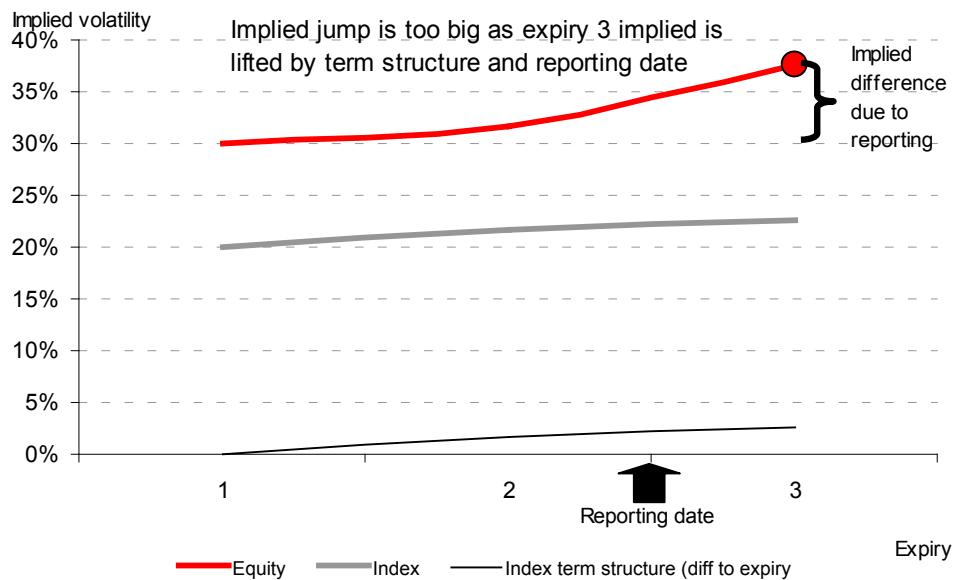
From the above implied volatility due to jump (σ_{Jump}) it is possible to calculate the implied daily return on the day of the jump (which is a combination of the normal daily move and the effect of the jump).

$$\text{Expected daily return} = e^{\frac{\sigma_{\text{Jump}}^2}{2}} \left[(2 \times N(\sigma_{\text{Jump}}) - 1) \right]$$

EQUITY JUMP ASSUMES FLAT TERM STRUCTURE

The methodology for extracting jumps from the difference between the front-month implieds is simply a case of mathematics, assuming the volatility of a stock is equal to a ‘normal’ volatility on every day plus an ‘abnormal’ jump on reporting. In order to calculate the ‘abnormal’ jump, we need to estimate the ‘normal’ volatility, and this estimate usually requires a flat term structure to be assumed.

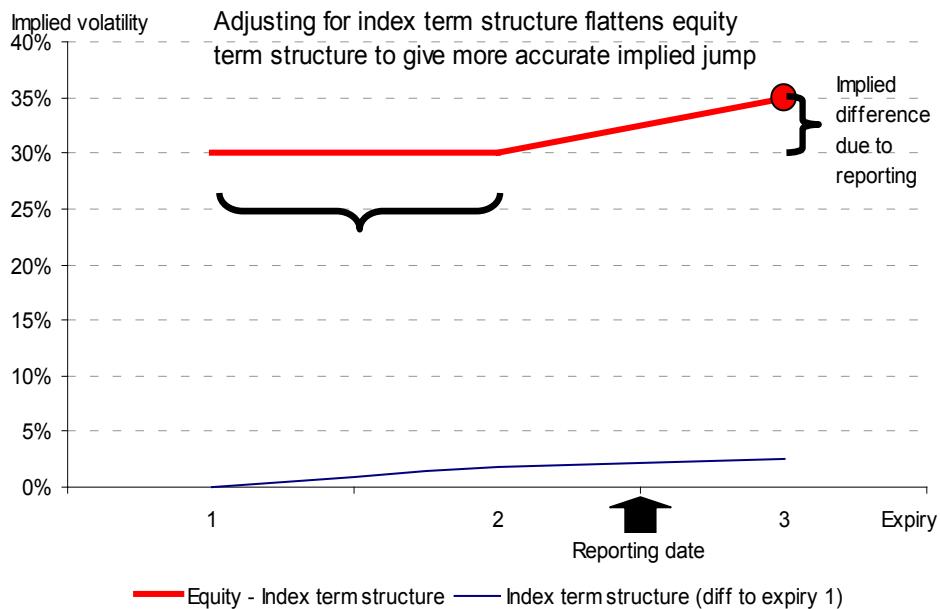
Figure 96. Equity and Index Term Structure



Adjusting for index term structure removes macro effects

Single stock term structure is the sum of index (or macro) term structure and the impact of the jump on reporting. If the index term structure is used to adjust the single-stock term structure, then a more accurate implied jump can be calculated²⁵ (assuming the single-stock term structure would be identical to index term structure without the effect of a reporting date). For simplicity, the diagrams below assume reporting is between expiry 2 and 3, but the effect will be similar if earnings is between expiry 1 and 2.

Figure 97. Equity Term Structure Adjusted by Index Term Structure



²⁵ This assumes a flat implied correlation term structure, which is a reasonable assumption for the very near-dated expiries.

CHAPTER 7

SKEW AND TERM STRUCTURE TRADING

We examine how skew and term structure are linked and the effect on volatility surfaces of the square root of time rule. The correct way to measure skew and smile is examined, and we show how skew trades only breakeven when there is a static local volatility surface.

7.1: SKEW AND TERM STRUCTURE ARE LINKED

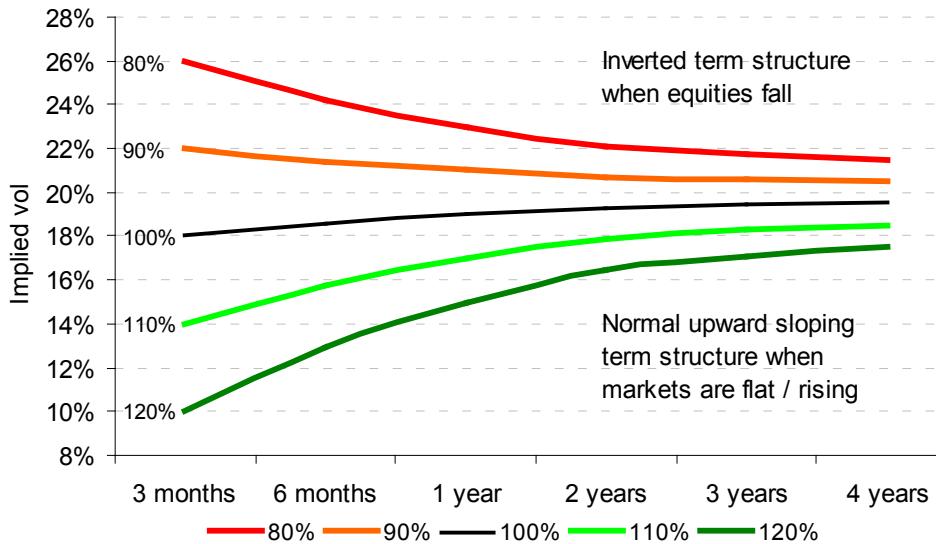
When there is an equity market decline, there is normally a larger increase in ATM implied volatility at the near end of volatility surfaces than the far end. Assuming sticky strike, this causes near-dated skew to be larger than far-dated skew. The greater the term structure change for a given change in spot, the higher skew is. Skew is also positively correlated to term structure (this relationship can break down in panicked markets). For an index, skew (and potentially term structure) is also lifted by the implied correlation surface. Diverse indices tend to have higher skew for this reason, as the ATM correlation is lower (and low strike correlation tends to 100% for all indices).

SKEW & TERM STR CUT SURFACE IN DIFFERENT DIMENSIONS

A volatility surface has three dimensions (strike, expiry and implied volatility), which is difficult to show on a two dimensional page. For simplicity, a volatility surface is often plotted as two separate two dimensional graphs. The first plots implied volatility vs expiry (similar to the way in which a yield curve plots credit spread against expiry) in order to show term structure (the difference in implied volatility for options with different maturities and the same strike). The second plots implied volatility vs strike to show skew (the difference in implied volatility for options with different strikes and the same maturity). We examine a volatility surface in both these ways (ie, term structure and skew) and show how they are related.

TERM STRUCTURE IS NORMALLY UPWARD SLOPING

When there is a spike in realised volatility, near-dated implied volatility tends to spike in a similar way (unless the spike is due to a specific event such as earnings). This is because the high realised volatility is expected to continue in the short term. Realised volatility can be expected to mean revert over a c8-month period, on average. Hence far-dated implied volatilities tend to rise by a smaller amount than near-dated implied volatilities (as the increased volatility of the underlying will only last a fraction of the life of a far-dated option). Near-dated implieds are therefore more volatile than far-dated implieds. The theoretical term structure for different strikes is shown in Figure 98 below, which demonstrates that near-dated implieds are more volatile. We have shown ATM (100%) term structure as upward sloping as this is how it trades on average (for the same reasons credit spread term structure is normally upward sloping, ie, risk aversion and supply-demand imbalances for long maturities).

Figure 98. Term Structure for Options of Different Strikes

If equity markets decline, term structure becomes inverted

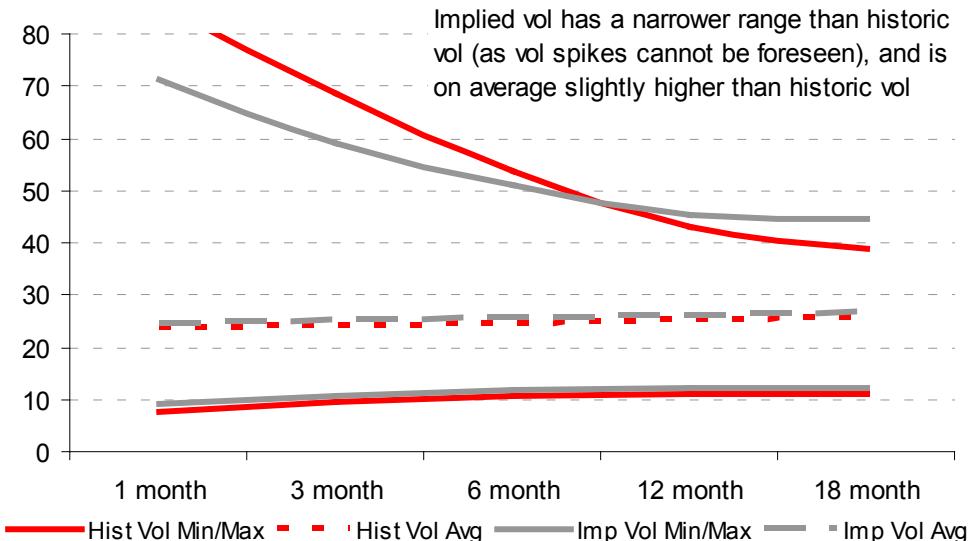
Typically, an increase in volatility tends to be accompanied by a decline in equity markets, while a decline in volatility tends to occur in periods of calm or rising markets. If volatility surfaces are assumed not to move as spot moves (ie, sticky strike), then this explains why the term structure of low strike implied volatility is normally downward sloping (as the 80% strike term structure will be the ATM term structure when equities fall 20%). Similarly, this explains why the term structure of high strike implieds is normally upward sloping (as the 120% strike term structure will become the ATM term structure when equities rise 20%).

Slope of rising term structure is shallower than slope of inverted term structure

While Figure 98 above shows the term structure of a theoretical volatility surface, in practice the slope of rising term structure is shallower than the slope of inverted term structure. This can be seen by looking at a volatility cone (Figure 99).

Despite the fact that the inverted term structure is steeper, the more frequent case of upward sloping term structure means the average term structure is slightly upward sloping²⁶.

²⁶ Positive implied correlation term structure will also lift index term structure relative to single stock.

Figure 99. Implied and Historic Volatility Cone

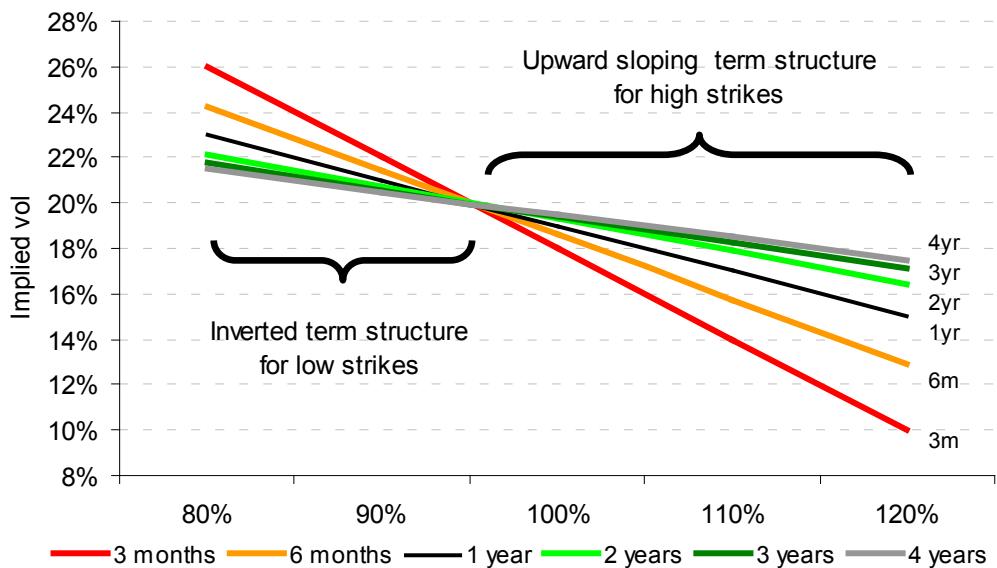
Implied volatility is usually greater than realised volatility and less volatile

While historic and realised volatility are linked, there are important differences which can be seen when looking at empirical volatility cones. Average implied volatility lies slightly above average realised volatility, as implieds are on average slightly expensive. Implied volatility is also less volatile (it has a smaller min-max range) than realised volatility for near-dated maturities. This is because implieds are forward looking (ie, similar to an average of possible outcomes) and there is never 100% probability of the maximum or minimum possible realised. This effect fades for longer maturities, potentially due to the additional volatility caused by supply-demand imbalances (eg, varying demand for put protection). This causes inverted implied volatility term structure to be less steep than realised volatility term structure.

SKEW IS INVERTED & LARGEST FOR NEAR-DATED EXPIRIES

Assuming volatility surfaces stay constant (ie, sticky strike), the effect of near-dated ATM implieds moving further than far-dated implieds for a given change in spot is priced into volatility surfaces by having a larger near-dated skew. The example data given in Figure 98 is plotted in Figure 100 below with a change of axes to show skew for options of different maturity. This graph shows that near-dated implieds have higher skew than far-dated implieds. The more term structure changes for a given change in spot, the steeper skew is. As near-dated ATM volatility is more volatile than far-dated ATM volatility, near-dated implied volatility has higher skew.

Figure 100. Skew for Options of Different Maturity



Skew for equities is normally inverted

Unless there is a high likelihood of a significant jump upwards (eg, if there were a potential takeover event), equities normally have negative skew (low strike implied greater than high strike implied). There are many possible explanations for this, some of which are listed below.

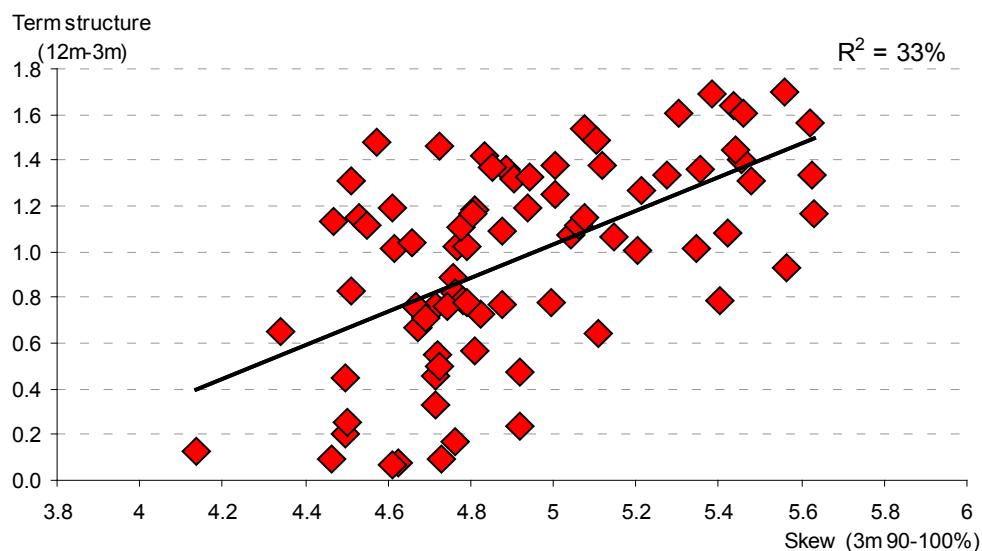
- **Big jumps in spot tend to be down, rather than up.** If there is a jump in the stock price, this is normally downwards as it is more common for an unexpectedly bad event to occur (bankruptcy, tsunami, terrorist attack, accident, loss or death of key personnel, etc) than an unexpectedly good event to occur (positive drivers are normally planned for).

- **Volatility is a measure of risk and leverage (hence risk) increases as equities decline.** If we assume no change in the number of shares in issue or amount of debt, then as a company's stock price declines its leverage (debt/equity) increases. Both leverage and volatility are a measure of risk and, hence, they are correlated, with volatility rising as equities fall.
- **Demand for protection and call overwriting.** Typically, investors are interested in buying puts for protection, rather than selling them. This lifts low strike implieds. Additionally, some investors like to call overwrite their positions, which weighs on higher strike implieds.

REASONS WHY SKEW & TERM STR ARE CORRELATED

The correlation between skew and term structure is shown below. The diagram only shows data for positive term structure, as the relationship tends to break down during a crisis.

Figure 101. SX5E Skew and Positive Term Structure



There are three reasons why skew and term structure are correlated

- Credit events, such as bankruptcy, lift both skew and term structure
- Implied volatility is ‘sticky’ for low strikes and long maturities
- Implied correlation is ‘sticky’ for low strikes and long maturities (only applies to index)

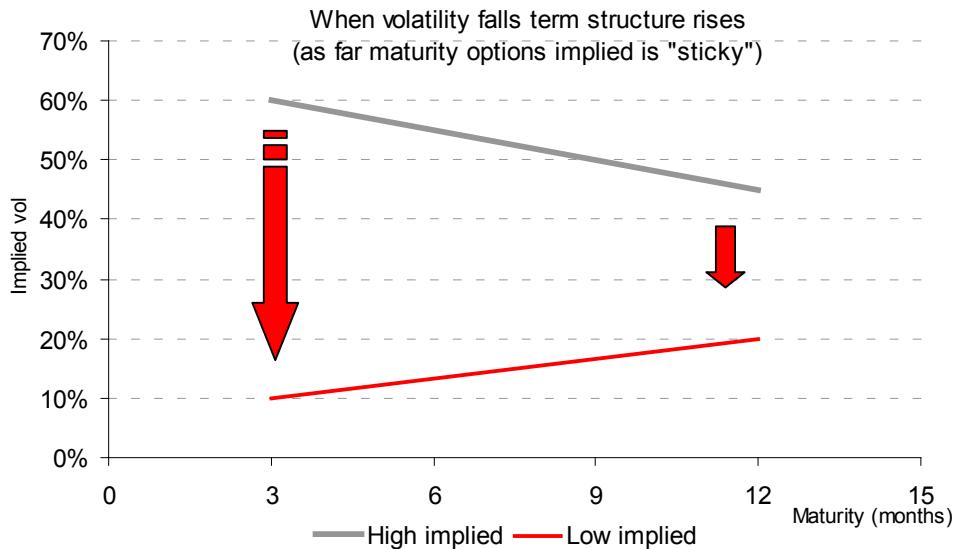
(1) BANKRUPTCY LIFTS BOTH SKEW AND TERM STRUCTURE

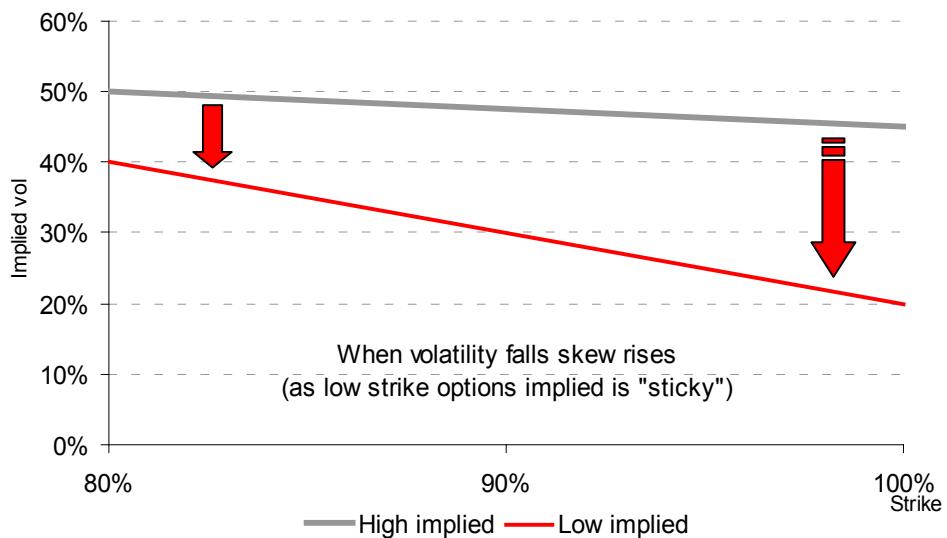
There are various models that show the effect of bankruptcy (or credit risk) lifting both skew and term structure. As implieds with lower strikes have a greater sensitivity to credit risk (as most of the value of low strike puts is due to the risk of bankruptcy), their implieds rise more, which causes higher skew. Similarly, options with longer maturity are more sensitive to credit risk (causing higher term structure, as far-dated implieds rise more). Longer-dated options have a higher sensitivity to credit risk as the probability of entering default increases with time (hence a greater proportion of an option’s value will be associated with credit events as maturity increases). More detail on the link between volatility and credit can be seen in section [A12 capital Structure Arbitrage](#) in the *Appendix*.

(2) IMPLIED VOL IS ‘STICKY’ FOR LOW STRIKES AND LONG MATURITIES

Term structure should rise if near-dated ATM implieds fall, as far-dated ATM implieds are relatively constant (as they tend to include complete economic cycles). This is shown in Figure 102 below.

Figure 102. Term Structure Rising with Falling Volatility

**Figure 103. Skew Rising with Falling Volatility**



Skew tends to rise as ATM implieds fall as low strike implieds are ‘sticky’

If there is a sudden decline in equity markets, it is reasonable to assume realised volatility will jump to a level in line with the peak of realised volatility. Therefore, low-strike, near-dated implieds should be relatively constant (as they should trade near the all-time highs of realised volatility). If a low-strike implied is constant, the difference between a low-strike implied and ATM implied increases as ATM implieds falls. This means near-dated skew should rise if near-dated ATM implieds decline (see Figure 103 above).

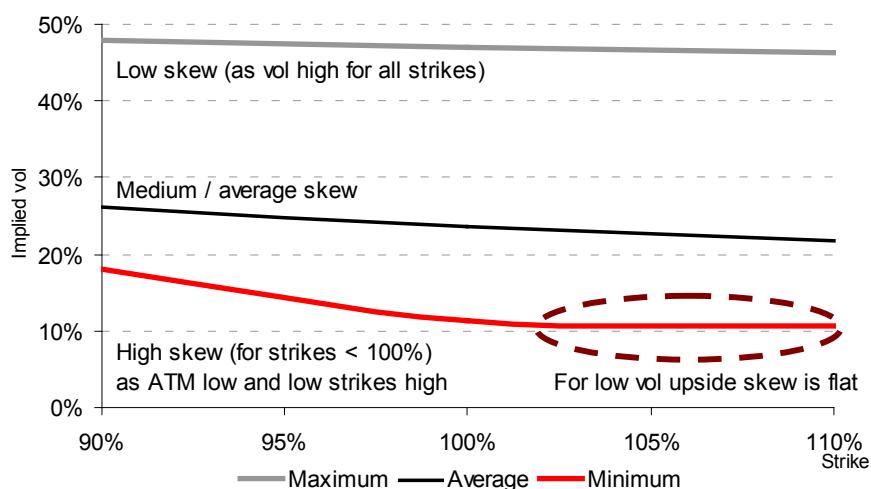
Term structure & skew are correlated as both rise as implied volatility falls

A decline in near dated ATM implied volatility lifts term structure and skew as low strike and long maturity implied volatility is ‘sticky’. Hence skew and term structure should be correlated as a fall in near-dated ATM implied lifts both of them.

Skew is often mistakenly used as a risk indicator

We do not view skew as a reliable risk indicator, as it can be inversely correlated to ATM volatility²⁷. The effect of falling implieds causing an increase in 90%-100% skew is shown with empirical data in Figure 104 below (many investors prefer to use 90%-100% skew rather than 90%-110%, as upside 100%-110% skew flattens as implieds reach a bottom).

Figure 104. SX5E 1 Year Max, Min and Average Implied Vol



(3) CORRELATION SURFACE CAUSES INDEX SKEW AND TERM STRUCTURE TO BE CORRELATED

²⁷ ATM volatility is a risk measure; hence, a measure often inversely correlated to ATM volatility, such as skew, is not a reliable risk measure.

In the same way implied volatility is ‘sticky’ for low strikes and long maturities, so is implied correlation (in a crisis correlation tends to 100% hence low strike correlation is ‘sticky’, and in the long run macro trends are the primary driver for all stock prices hence long dated correlation is also ‘sticky’). This can be an additional reason why index skew and index term structure are correlated.

CORRELATION LIFTS INDEX SKEW ABOVE SINGLE-STOCK SKEW

An approximation for implied correlation is the index volatility squared divided by the average single-stock volatility squared [$Q = \sigma_{\text{Index}}^2 \div \text{average}(\sigma_{\text{Single stock}})^2$]. Implied correlation is assumed to tend towards 100% for low strikes, as all stocks can be expected to decline in a crisis. This causes index skew to be greater than single stock skew. Index skew can be thought of as being caused by both the skew of the single stock implied volatility surface, and the skew of the implied correlation surface.

Example of how index skew can be positive with flat single-stock skew

We shall assume all single stocks in an index have the same (flat) implied volatility and single-stock skew is flat. Low strike index volatility will be roughly equal to the constant single-stock volatility (as implied correlation is close to 100%), but ATM index volatility will be less than this value due to diversification (as implied correlation Q for ATM strikes is less than 100% and $\sigma_{\text{Index}}^2 = Q \times \text{average}(\sigma_{\text{Single stock}})^2$). Despite single stocks having no skew, the index has a skew (as low strike index implieds > ATM index implieds) due to the change in correlation. For this reason, index skew is always greater than the average single-stock skew.

Implied correlation is likely to be sticky for low strikes and long maturities

A correlation surface can be constructed for options of all strikes and expiries, and this surface is likely to be close to 100% for very low strikes. The surface is likely to be relatively constant for far maturities; hence, implied correlation term structure and skew will be correlated (as both rise when near-dated ATM implied correlation falls, similar to volatility surfaces). This also causes skew and term structure to be correlated for indices.

DIVERSE INDICES HAVE HIGHER SKEW

As index skew is caused by both single-stock skew and implied correlation skew, a more diverse index should have a higher skew than a less diverse index (assuming there is no significant difference in the skew of the single-stock members). This is due to the fact that diverse indices have a lower ATM implied, but low strike implieds are in line with (higher) average single-stock implieds for both diverse and non-diverse indices.

7.2: SQUARE ROOT OF TIME RULE CAN COMPARE DIFFERENT TERM STRUCTURES AND SKEWS

When implied volatility changes, the change in ATM volatility multiplied by the square root of time is generally constant. This means that different $(T_2 - T_1)$ term structures can be compared when multiplied by $\sqrt{(T_2 - T_1)} / (\sqrt{T_2} - \sqrt{T_1})$, as this normalises against 1Y-3M term structure. Skew weighted by the square root of time should also be constant. Looking at the different term structures and skews, when normalised by the appropriate weighting, can allow us to identify calendar and skew trades in addition to highlighting which strike and expiry is the most attractive to buy (or sell).

REALISED VOLATILITY MEAN REVERTS AFTER 8 MONTHS

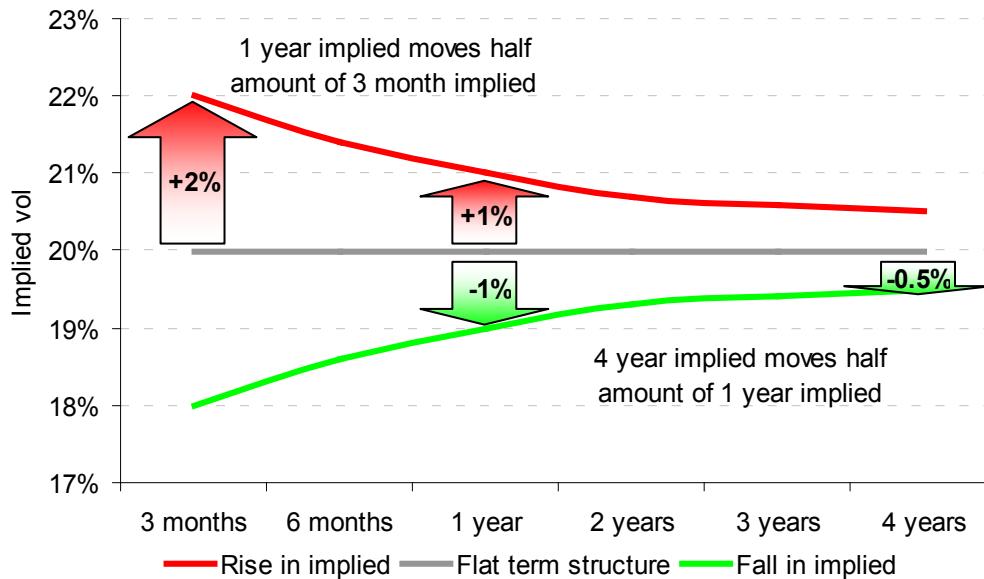
When there is a spike in realised volatility, it takes on average eight months for three-month realised volatility to settle back down to levels seen before the spike. The time taken for volatility to normalise is generally longer if the volatility is caused by a negative return, than if it is caused by a positive return (as a negative return is more likely to be associated with an event that increases uncertainty than a positive return). This mean reversion is often modelled via the square root of time rule.

VOL MOVE MULTIPLIED BY $\sqrt{\text{TIME}}$ IS USUALLY CONSTANT

The near-dated end of volatility surfaces is highly correlated to realised volatility, as hedge funds and prop desks typically initiate long/short gamma positions should there be a significant divergence. As volatility mean reverts, the far-dated end of volatility surfaces is more stable (as investors know that any spike in volatility will be short-lived and not last for the full length of a far-dated option). A common way to model the movement of volatility surfaces, is to define the movement of one-year implied and then adjust the rest of the curve by that move divided by time (in years) to the power of p. Only two parameters (the one-year move and p) are needed to adjust the whole surface. Fixing the power (or p) at 0.5 is the most common and is known as the square root of time rule (which only has one parameter, the one-year change).

$$\text{Implied vol move for maturity } T \text{ years} = \frac{\text{One year implied volatility move}}{T^p}$$

Figure 105. ATM Implied Volatility Moving in a Square Root of Time Manner



Square root of time rule has power 0.5, parallel moves are power 0

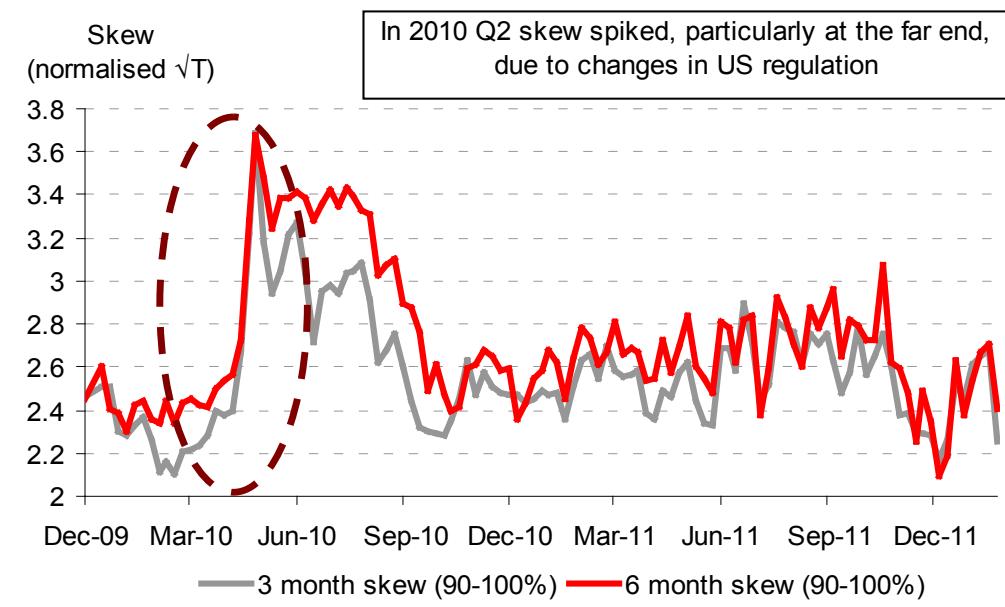
While the above method is usually used with power 0.5 (square root of time rule), any power can be used. If there is a parallel movement in volatility surfaces (all maturities move the same amount), then a power of 0 should be used. In practice, implied volatility tends to move with power 0.44, suggesting that surfaces move primarily in a square root of time manner but at times also in parallel. If implieds rise (or decline) in a square root of time manner when equities decline (or rise), then this causes skew to decay by the square root of time as well (assuming sticky strike). This means that the skews of different maturities can be compared with each other by simply multiplying the skew by the square root of the maturity (see Figure 104 below).

Figure 106. Same skew when multiplied by square root of time

POSITIVE PUT/CALL SPREADS IMPLY $\sqrt{\text{TIME}}$ RULE FOR SKEW

Structures such as put spreads or call spreads, which can only have a positive payout, must have a cost associated with them, or investors would simply purchase an infinite amount of them for zero cost (or small profit) and enter a position which could never suffer a loss. This means that when the strike of a put is increased, its premium must rise too (intuitively correct, as the strike is the amount of money you receive when you ‘put’ the stock, so the higher the strike the better). Conversely, as the strike of a call increases, its premium must decrease. It can be shown that enforcing positive put/call spreads puts a cap/floor on skew, which decays by the square root of time. This provides mathematical support for the empirical evidence, suggesting implied volatility should normally move in a power weighted (by square root of time) manner²⁸. For more details, see the section [A6 Modelling Volatility Surfaces](#) in the *Appendix*.

Figure 107. SX5E Skew Multiplied by the Square Root of Time ($R^2=83\%$)



²⁸ Looking at ratio put spreads, it can be shown that for long maturities (five years) skew should decay by time, ie, $1/T$ or power=1 (rather than \sqrt{T} or power 0.5).

√TIME RULE CAN COMPARE DIFFERENT TERM STRUCTURES

ATM term structure can be modelled as flat volatility, with a square root of time adjustment on top. With this model, flat volatility is equal to the volatility for an option of infinite maturity. There are, therefore, two parameters to this model, the volatility at infinity (V_∞) and the scale of the square root of time adjustment ,which we define to be z (for one-year implied).

$$\text{Volatility} = V_\infty - z / \sqrt{T}$$

where z = scale of the square root of time adjustment (which we define as normalised term structure)

We have a negative sign in front of z , so that a positive z implies an upward sloping term structure and a negative z is a downward sloping term structure.

Different term structures are normalised by multiplying by $\sqrt{(T_2 T_1)} / (\sqrt{T_2} - \sqrt{T_1})$

Using the above definition, we can calculate the normalised term structure z from two volatility points V_1 and V_2 (whose maturity is T_1 and T_2).

$$V_1 = V_\infty - z / \sqrt{T_1}$$

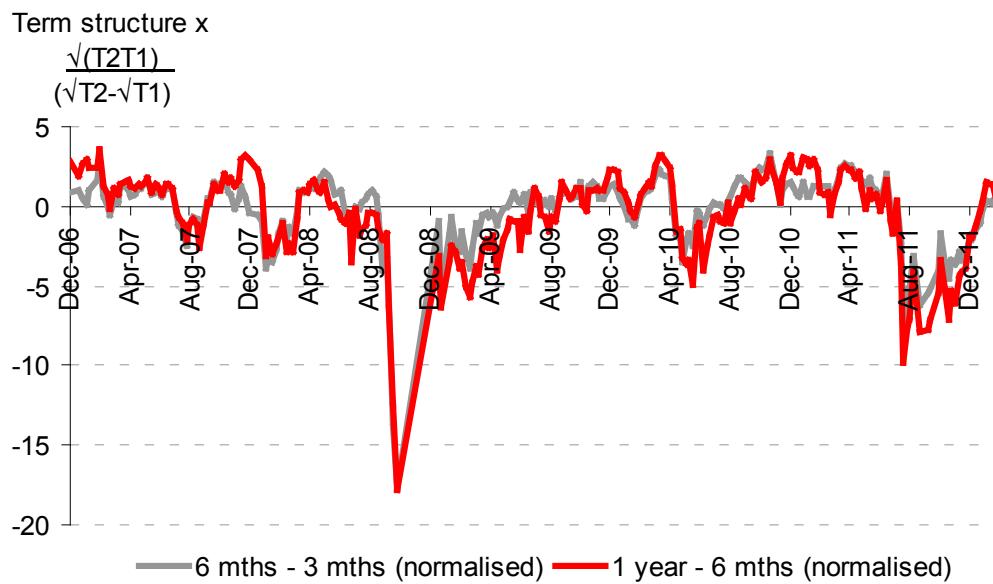
$$V_2 = V_\infty - z / \sqrt{T_2}$$

$$V_1 + z / \sqrt{T_1} = V_2 + z / \sqrt{T_2} = V_\infty$$

$$z (1/\sqrt{T_1} - 1/\sqrt{T_2}) = V_2 - V_1$$

$$z = (V_2 - V_1) \times \frac{\sqrt{T_2 T_1}}{\sqrt{T_2} - \sqrt{T_1}}$$

$V_2 - V_1$ is the normal definition for term structure. Hence, term structure can be normalised by multiplying by $\sqrt{(T_2 T_1)} / (\sqrt{T_2} - \sqrt{T_1})$. We note that the normalisation factor for 1Y-3M term structure is 1. Therefore, normalising allows all term structure to be compared to 1Y-3M term structure.

Figure 108. SX5E Normalised Term Structure (R2=80%)

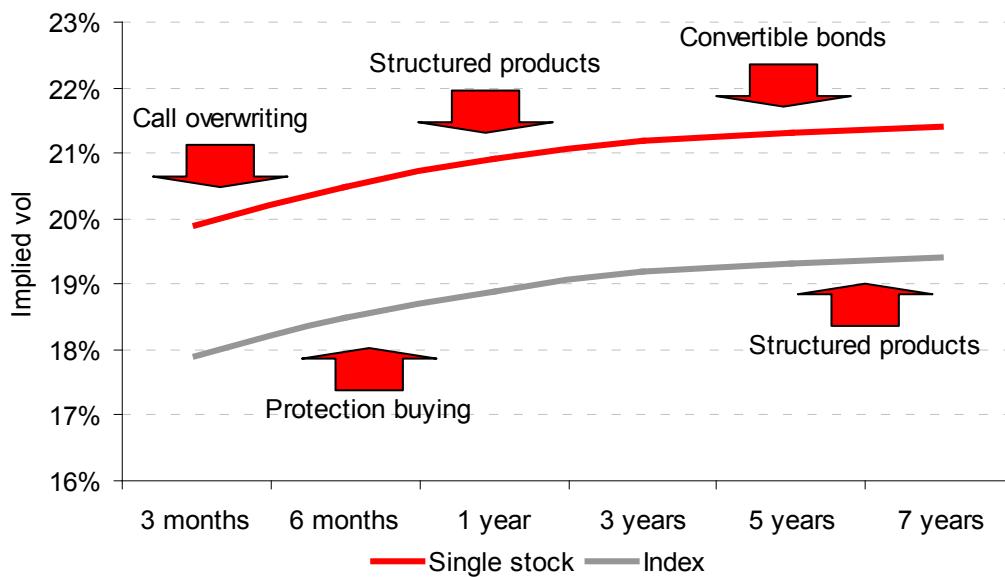
7.3: TERM STRUCTURE TRADING

The supply-demand imbalances of different products on implied volatility surfaces can create opportunity for other investors. The degree of the imbalance depends on the popularity of the product at the time. Investors who are willing to take the other side of the trade should be able to profit from the imbalance, and the risk taken can be hedged with other maturities or related securities.

TERM STRUCTURE TRADES CAN PROFIT FROM IMBALANCES

The graph below shows the imbalances of the different products by maturity for both single stock and index. The demand for long dated hedges typically lifts term structure, an imbalance that an investor can profit from by selling term structure.

Figure 109. Implied Volatility Imbalances by Maturity



CALENDARS CONSTANT IF SURFACES FOLLOW $\sqrt{\text{TIME}}$ RULE

Given that the square root of time appears in the Black-Scholes formula for premium, the price of a 1x1 calendar (long one far-dated option, short one near-dated option) remains approximately constant if implied volatility surfaces move in a square root of time manner. Calendars can therefore be used to trade term structure imbalances as the trade is indifferent to the level of volatility as long as volatility moves in a power weighted manner.

IDENTIFYING WHEN TO GO LONG, OR SHORT, CALENDARS

When examining term structure trades, the power of the movement in volatility surfaces can be compared to the expected 0.5 power of the square root of time rule. If the movement has a power significantly different from 0.5, then a long (or short) calendar position could be initiated to profit from the anticipated correction. This method assumes calendars were previously fairly priced (otherwise the move could simply be a mean reversion to the norm).

If vol rises with power less than 0.5, investors should short calendars

If surfaces rise with a power less than 0.5 (ie, a more parallel move) then near-dated implieds have not risen as much as expected and a short calendar (long near-dated, short far-dated) position should be initiated. This position will profit from the anticipated correction. Should surfaces fall with a power less than 0.5, a long calendar (short near-dated, long far-dated) would profit from the anticipated further decline of near-dated implieds.

If vol rises with power more than 0.5, investors should go long calendars

Conversely, if surfaces rise with a power greater than 0.5, near-dated implieds have risen too far and a long calendar position should be initiated. On the other hand, if surfaces fall with a power greater than 0.5, a short calendar position should be initiated (as near-dated implieds have fallen too far).

POWER VEGA IS VEGA DIVIDED BY THE $\sqrt{\text{TIME}}$

As volatility surfaces tend to move in a square root of time manner, many systems report power vega (vega divided by square root of time). Power vega takes into account the fact that the implied volatility of near-dated options is more volatile than far-dated options.

VARIANCE TERM STRUCTURE CAN IDENTIFY TRADES

To determine if a term structure trade is needed, we could look at variance term structure rather than implied volatility term structure. Using variance term structure eliminates the need to choose a strike (an ATM term structure will not be ATM as soon as the spot moves, so it is effectively strike dependent, but simply delegates the choice of strike to the equity market). Variance term structure is similar to ATM term structure, despite variance being long skew and skew being greater for near-dated implieds. This is because the time value of an OTM option increases with maturity. Hence, the increased weight associated with OTM options cancels the effect of smaller skew for longer maturities.

Forward starting var swaps (or options) can be used to trade term structure

Trading term structure via a long and short variance swap is identical to a position in a forward starting variance swap (assuming the weights of the long and short variance swap are correct; if not, there will be a residual variance swap position left over). The correct weighting for long and short variance swaps to be identical to a forward starting variance swap is detailed in the section [*4.1 Forward Starting Products*](#). If an investor wants to trade term structure, but does not want to have exposure to current volatility (ie, wants to have zero theta and gamma), then forward starting products (variance swaps or options) can be used. Note that while forward starting products have no exposure to current realised volatility, they do have exposure to future expectations of volatility (ie, implied volatility hence has positive vega).

7.4: HOW TO MEASURE SKEW AND SMILE

The implied volatilities for options of the same maturity, but of different strike, are different from each other for two reasons. Firstly, there is skew, which causes low strike implieds to be greater than high strike implieds due to the increased leverage and risk of bankruptcy. Secondly, there is smile (or convexity/kurtosis), when OTM options have a higher implied than ATM options. Together, skew and smile create the ‘smirk’ of volatility surfaces. We look at how skew and smile change by maturity in order to explain the shape of volatility surfaces both intuitively and mathematically. We also examine which measures of skew are best, and why.

MOMENTS DESCRIBE THE PROBABILITY DISTRIBUTION

In order to explain skew and smile, we shall break down the probability distribution of log returns into moments. Moments can describe the probability distribution²⁹. From the formula below we can see that the zero-th moment is 1 (as the sum of a probability distribution is 1, as the probability of all outcomes is 100%). The first moment is the expected value (ie, mean or forward) of the variable. The second, third and fourth moments are variance, skew and kurtosis, respectively (see table on the left below). For moments of two or greater it is usual to look at central moments, or moments about the mean (we cannot for the first moment as the first central moment would be 0). We shall normalise the central moment by dividing it by σ^n in order to get a dimensionless measure. The higher the moment, the greater the number of data points that are needed in order to get a reasonable estimate.

$$\text{Raw moment} = E(X^k) = \int_{-\infty}^{\infty} x^k f(x)$$

$$\text{Normalised central moment} = E((X - \mu)^k) / \sigma^k = \int_{-\infty}^{\infty} (x - \mu)^k f(x) / \sigma^k$$

where

$f(x)$ is the probability distribution function

²⁹ The combination of all moments can perfectly explain any distribution as long as the distribution has a positive radius of convergence or is bounded (eg, a sine wave is not bounded; hence, it cannot be explained by moments alone).

Figure 110. Moments

Moment	Name	Decay/Movement over Time	Maturity Where Dominates Surface
1†	Forward (expected price)	Random walk	NA
2	Variance (volatility ²)	Mean reverts	Far-dated/all maturities
3	Skew	Decay square root	Medium-dated
4	Kurtosis	Decay by time	Near-dated

(†) First raw moment (other moments are normalised central moments).

Figure 111. Related Option Position of Moments

Moment	Related (Long) Position	Key Greek	P&L Driver
1†	Stock/futures	Delta	Price
2	ATM options	Vega	Implied volatility
3	Risk reversal (long put)	Vanna	Skew
4	Butterfly (long wings)	Volga (gamma of volatility)	Vol of vol

(†) First raw moment (other moments are normalised central moments).

VEGA, VANNA AND VOLGA MEASURE 2ND, 3RD & 4TH MOMENTS

While there is not an exact mapping, the Greeks that best measure the 2nd, 3rd and 4th moments are vega, vanna and volga (VOL GAMMA) respectively. The driver for these Greeks are implied volatility, skew and vol of vol.

Options with high volga benefit from volatility of volatility

Just as an option with high gamma benefits from high stock price volatility, an option with high volga benefits from volatility of volatility. The level of volatility of volatility can be calculated in a similar way to how volatility is calculated from stock prices (taking log returns is recommended for volatility as well). The more OTM an option is, the greater the volatility of volatility exposure. This is because the more implied volatility can change, the greater the chance of it rising and allowing an OTM option to become ITM. This gives the appearance of a ‘smile’, as the OTM option’s implied volatility is lifted while the ATM implied volatility remains constant.

Stock returns have positive excess kurtosis and are leptokurtic

Kurtosis is always positive³⁰. Hence, excess kurtosis (kurtosis -3) is usually used. The kurtosis (or normalised fourth moment) of the normal distribution is three; hence, normal distributions have zero excess kurtosis (and are known as mesokurtic). High kurtosis distributions (eg, stock price log returns) are known as leptokurtic, whereas low kurtosis distributions (pegged currencies that change infrequently by medium-sized adjustments) are known as platykurtic.

³⁰ Kurtosis is only zero for a point distribution.

VEGA MEASURES SIZE OF VOLATILITY POSITION

Vega measures the change in price of an option for a given (normally 1%) move in implied volatility. Implied volatility for far-dated options is relatively flat compared to near-dated, as both skew and kurtosis decay with maturity. Vega is highest for ATM options, as can be seen in the right hand chart in Figure 113 below.

Figure 112. Moment 2 = Variance



Distribution for Constant Volatility

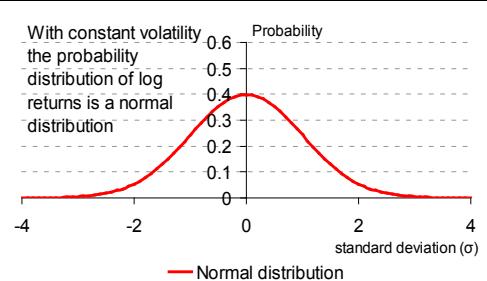
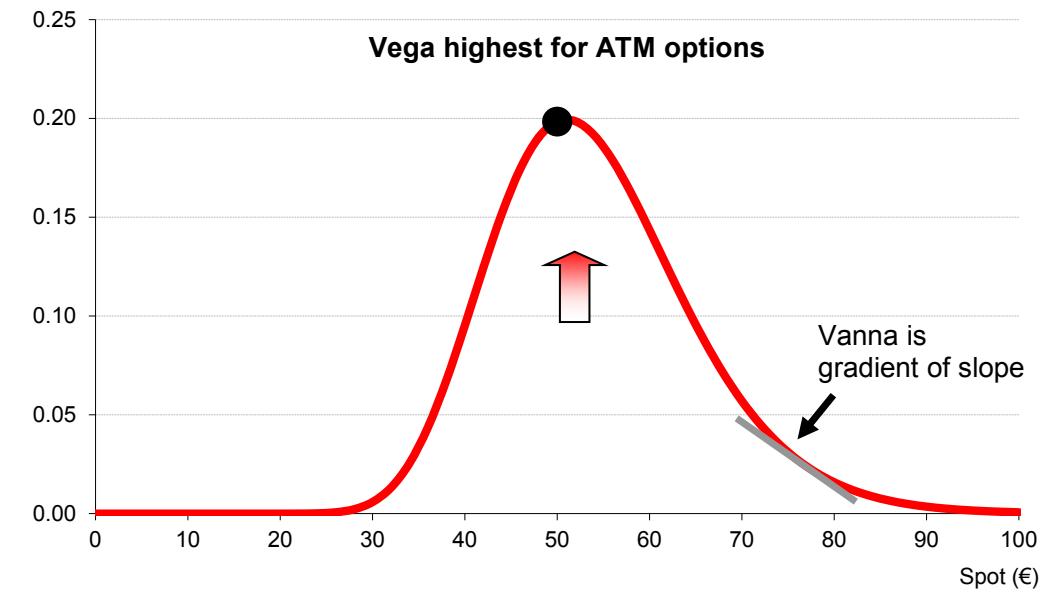


Figure 113. Vega is Size of Volatility Position



VANNA MEASURES SIZE OF SKEW POSITION

Vanna ($d\text{Vega}/d\text{Spot}$ which is equal to $d\Delta/d\text{Vol}$) measures the size of a skew position³¹, and is shown on the right side of Figure 110 below. Vanna is the slope of vega plotted against spot (see graph on right in Figure 113 above).

Figure 114. Moment 3 = Skew³²

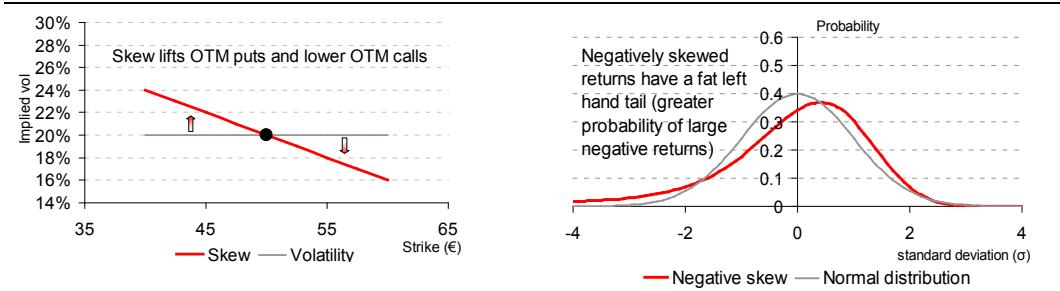
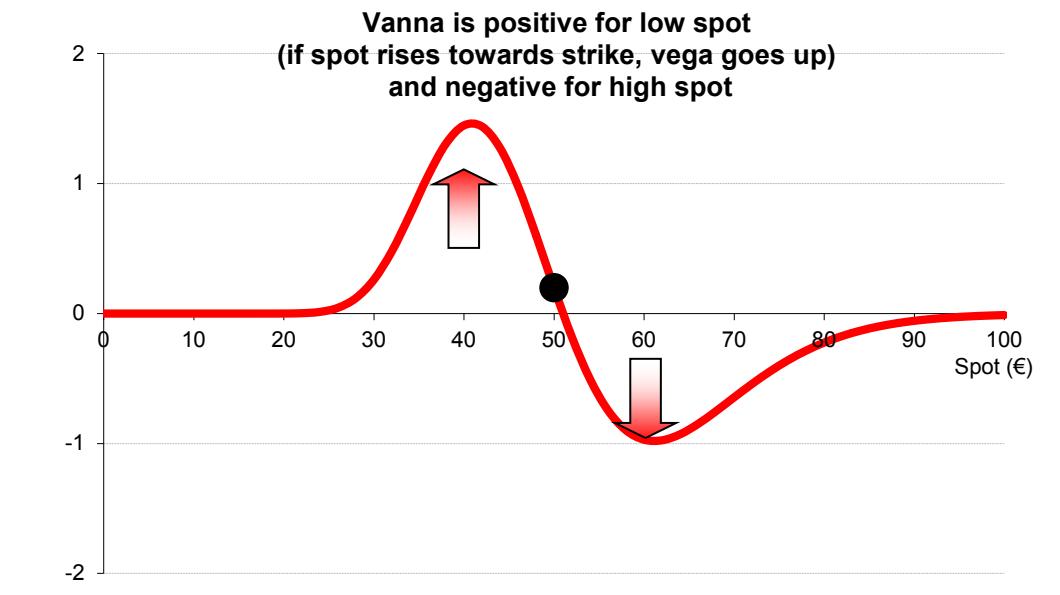


Figure 115. Vanna is Size of Skew Position



³¹ For details, see the next section 7.5 *Skew Trading*.

³² This is an approximation as the effect of moments on slope and convexity are intertwined.

VOLGA MEASURES SIZE OF GAMMA OF VOLATILITY

The gamma of volatility is measured by Volga ($d\text{Vega}/d\text{Volatility}$), which is also known as volatility gamma or vega convexity. Volga is always positive (similar to option gamma always being positive) and peaks for c10-15 delta options (like Vanna).

Figure 116. Moment 4 = Kurtosis³³

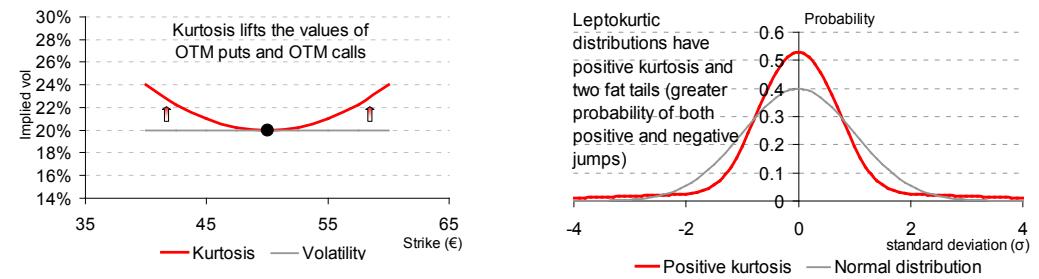
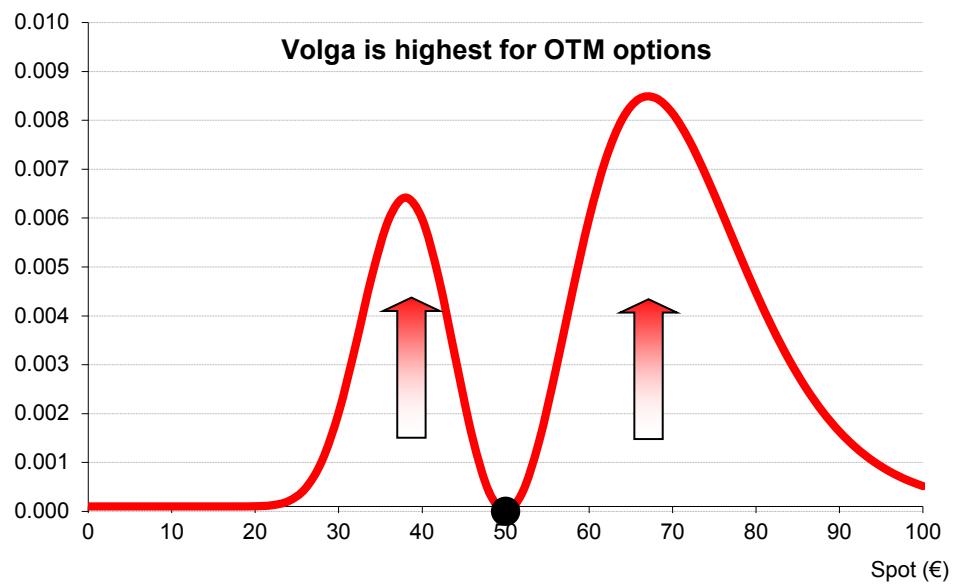


Figure 117. Volga = VOLatility GAMma

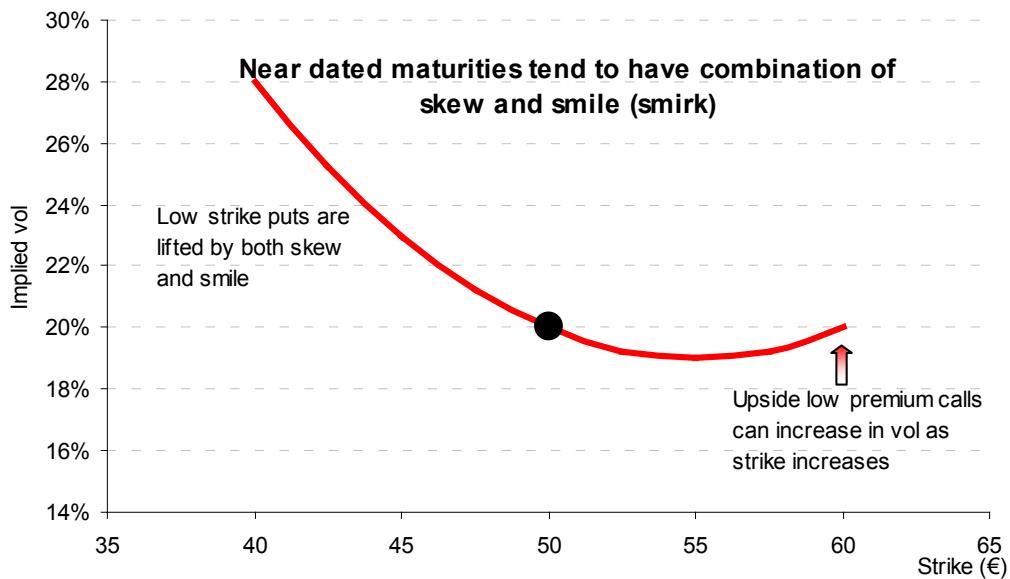


³³ This is an approximation as the effect of moments on slope and convexity are intertwined.

IMPLIED VOL SMIRK IS A COMBINATION OF SKEW AND SMILE

The final ‘smirk’ for options of the same maturity is the combination of skew (3rd moment) and smile (4th moment). The exact smirk depends on maturity. Kurtosis (or smile) can be assumed to decay with maturity by dividing by time³⁴ and, hence, is most important for near-dated expiries. For medium- (and long-) dated expiries, the skew effect will dominate kurtosis, as skew usually decays by the reciprocal of the square root of time (for more details, see the section *A6 Modelling Volatility Surfaces* in the *Appendix*). Skew for equities is normally negative and therefore have mean < median < mode (max) and a greater probability of large negative returns (the reverse is true for positively skewed distributions). For far-dated maturities, the effect of both skew and kurtosis fades; hence, implieds converge to a flat line for all strikes. Skew can be thought of as the effect of changing volatility as spot moves, while smile can be thought of as the effect of jumps (up or down).

Figure 118. Near-Dated Implied Volatilities with Smirk (Skew and Smile)



³⁴ Assuming stock price is led by Lévy processes (eg, accumulation of independent identical shocks).

THERE ARE 3 WAYS TO MEASURE SKEW

There are 3 main ways skew can be measured. While the first is the most mathematical, in practice the other 2 are more popular with market participants.

- Third moment
- Strike skew (eg, 90%-110%)
- Delta skew (eg, [25 delta put – 25 delta call] / 50 delta)

(1) THIRD MOMENT IS DEFINITION OF CBOE SKEW INDEX

CBOE have created a skew index on the S&P500. This index is based on the normalised third central moment; hence, it is strike independent. The formula for the index is given below. For normal negative skew, if the size of skew increases, so does the index (as negative skew is multiplied by -10).

$$\text{SKEW} = 100 - 10 \times 3^{\text{rd}} \text{ moment}$$

(2) STRIKE SKEW SHOULD NOT BE DIVIDED BY VOLATILITY

The most common method of measuring skew is to look at the difference in implied volatility between two strikes, for example 90%-110% skew or 90%-ATM skew. It is a common mistake to believe that strike skew should be divided by ATM volatility in order to take into account the fact that a 5pt difference is more significant for a stock with 20% volatility than 40% volatility. This ignores the fact that the strikes chosen (say 90%-110% for 20% volatility stocks) should also be wider for high volatility stocks (say 80%-120%, or two times wider, for 40% volatility stocks as the volatility is $2 \times 20\%$). The difference in implied volatility should be taken between two strikes whose width between the strikes is proportional to the volatility (similar to taking the implied volatility of a fixed delta, eg, 25% delta). An approximation to this is to take the fixed strike skew, and multiply by volatility, as shown below. As the two effects cancel each other out, we can simply take a fixed strike skew without dividing by volatility.

$$\text{Difference in vol between 2 strikes} = 90-110\%$$

$$\text{Difference in vol between 2 strikes whose width increases with vol} = 90-110\% \times \text{ATM}$$

$$\text{Skew} = \frac{\text{Difference in vol between 2 strikes whose width increases with vol}}{\text{ATM}}$$

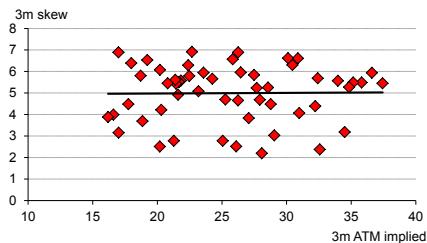
$$\Rightarrow \text{Skew} = \frac{90-110\% \times \text{ATM}}{\text{ATM}}$$

$$\Rightarrow \text{Skew} = 90-110\%$$

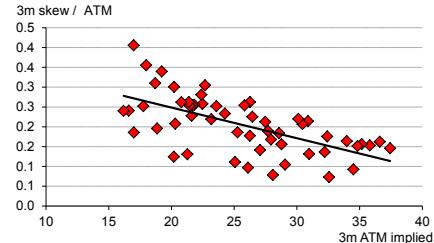
90%-100% (or 90%-110%) skew is correct measure for fixed strike skew

The best measure of skew is one that is independent of the level of volatility. If this were not the case, then the measure would be partly based on volatility and partly on skew, which would make it more difficult to determine if skew was cheap or expensive. We have shown mathematically that an absolute difference (eg, 90%-110% or 90%-100%) is the correct measure of skew, but we can also show it empirically. The left-hand chart in Figure 119 below shows that there is no correlation between volatility and skew (90%-110%) for any European stocks that have liquid equity derivatives. If skew is divided by volatility, there is unsurprisingly a negative correlation between this measure and volatility (see right-hand chart below).

Figure 119. Strike Skew (90%-110%)

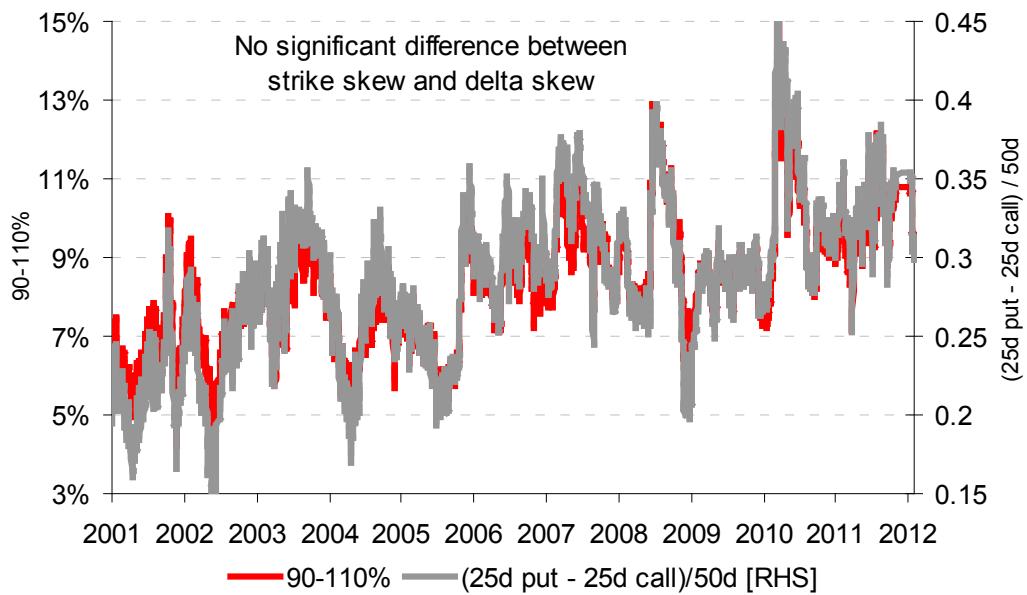


Strike Skew Divided by Volatility



(3) DELTA SKEW IS VIRTUALLY IDENTICAL TO STRIKE SKEW

Arguably the best measure of skew is delta skew, where the difference between constant delta puts and calls is divided by 50 delta implied. An example of skew measured by delta is $[25 \text{ delta put} - 25 \text{ delta call}] / 50 \text{ delta}$. As this measure widens the strikes examined as vol rises, in addition to normalising (ie, dividing) by the level of volatility, it is a ‘pure’ measure of skew (ie, not correlated to the level of volatility). While delta skew is theoretically the best measure, in practice it is virtually identical to strike skew. As there is a R^2 of 93% between delta skew and strike skew, we believe both are viable measures of skew (although strike skew is arguably more practical as it represents a more intuitive measure).

Figure 120. Strike Skew vs Delta Skew

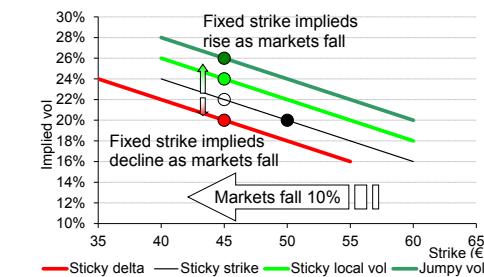
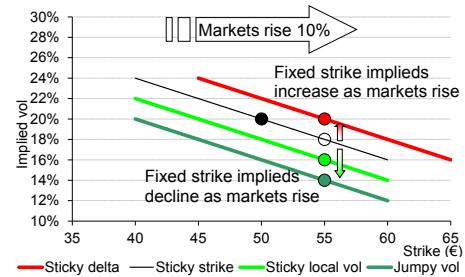
7.5: SKEW TRADING

The profitability of skew trades is determined by the dynamics of a volatility surface. We examine sticky delta (or ‘moneyness’), sticky strike, sticky local volatility and jumpy volatility regimes. Long skew suffers a loss in both a sticky delta and sticky strike regimes due to the carry cost of skew. Long skew is only profitable with jumpy volatility. We also show how the best strikes for skew trading can be chosen.

4 IDEALISED REGIMES DESCRIBE MOVEMENT OF VOL SURFACE

There are four idealised regimes for a volatility surface. While sticky delta, sticky strike and (sticky) local volatility are well known and widely accepted names, we have added ‘jumpy volatility’ to define volatility with a high negative correlation with spot. These regimes are summarised below, and more details are given on pages 216-222 of this section.

- (1) **Sticky delta (or sticky moneyness).** Sticky delta assumes a constant volatility for options of the same strike as a percentage of spot. For example, ATM or 100% strike volatility has constant volatility. As this model implies there is a positive correlation between volatility and spot, the opposite of what is usually seen in the market, it is not a particularly realistic model (except over a very long time horizon).
- (2) **Sticky strike.** A sticky strike volatility surface has a constant volatility for options with the same fixed currency strike. Sticky strike is usually thought of as a stable (or unmoving) volatility surface as real-life options (with a fixed currency strike) do not change their implied volatility.
- (3) **Sticky local volatility.** Local volatility is the instantaneous volatility of stock at a certain stock price. When local volatility is static, implied volatility rises when markets fall (ie, there is a negative correlation between stock prices and volatility). Of all the four volatility regimes, it is arguably the most realistic and fairly prices skew.
- (4) **Jumpy volatility.** We define a jumpy volatility regime as one in which there is an excessive jump in implied volatility for a given movement in spot. There is a very high negative correlation between spot and volatility. This regime usually occurs over a very short time horizon in panicked markets (or a crash).

Figure 121. Vol When Equities Fall 10%**Vol When Equities Rise 10%**

LONG SKEW TRADES HAVE A COST (SKEW THETA)

If an investor initiates a long skew position by buying an OTM put and selling an OTM call, the implied volatility of the put purchased has a higher implied volatility than the implied volatility sold through the call. The long skew position therefore has a cost associated with it, which we shall define as ‘skew theta’. Skew theta is the difference between the cost of gamma (theta per unit of dollar gamma) of an OTM option compared to an ATM option. If skew is flat, then all strikes have an identical cost of gamma, but as OTM puts have a higher implied volatility than ATM ones they pay more per unit of gamma. Skew theta is explained in greater depth at the end of this section. If the long skew position does not give the investor enough additional profit to compensate for the skew theta paid, then skew can be sold at a profit.

Skew trades profit from negative spot volatility correlation

If there is a negative correlation between the movement of a volatility surface and spot (as is usually seen in practice), then this movement will give a long skew position a profit when the volatility surface is re-marked. For example, let us assume an investor is long skew via a risk reversal (long an OTM put and short an OTM call). If equity markets decline, the put becomes ATM and is the primary driver of value for the position (as the OTM call becomes further OTM it is far less significant). The rise in the volatility surface (due to negative correlation between spot and volatility) boosts the value of the (now ATM) put and, hence, the value of the risk reversal.

SKEW TRADES BREAK EVEN IF LOCAL VOL IS CONSTANT

If the local volatility surface stays constant, the amount volatility surfaces move for a change in spot is equal to the skew (ie, ATM volatility moves by twice the skew, once for moving up the skew and another by the movement of the volatility surface itself). This movement is exactly the correct amount for the profit (or loss) on a volatility surface re-mark to compensate for the cost (or benefit) of skew theta³⁵. The profit (or loss) caused by skew trades given the four volatility regimes are shown below.

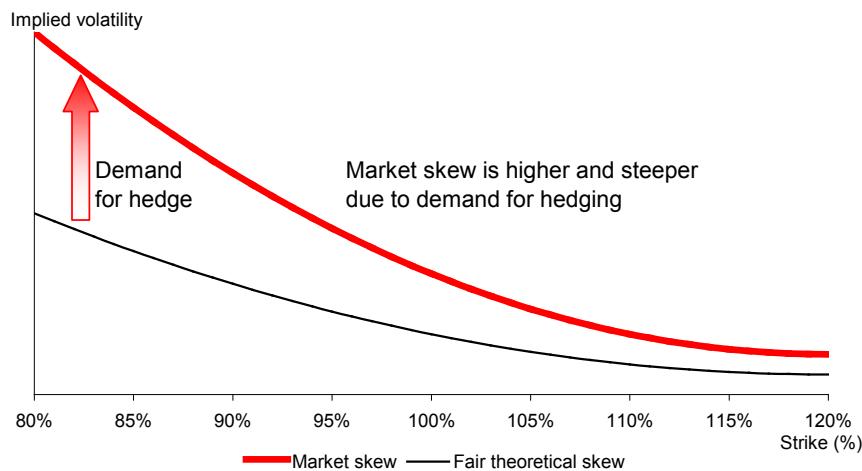
Figure 122. Different Vol Regimes and Breakdown of P&L for Skew Trades

Volatility regime	Fixed strike implied volatility change		P&L breakdown for long skew (e.g. long put, short call)		
	Equity decline	Equity rise	Remark	Skew theta	Total
Sticky delta	Falls	Rises			=
Sticky strike	-	-			=
Sticky local volatility	Rises	Falls			=
Jumpy volatility	Rises significantly	Falls significantly			=

SKEW IS USUALLY OVERPRICED DUE TO HEDGING

As volatility markets tend to trade between a static strike and static local volatility regime, long skew trades are usually unprofitable (usually there is negative spot volatility correlation, but not enough to compensate for the skew theta). As long skew trades break even during static local volatility regimes, they are only profitable in periods of jumpy volatility. This overpricing of skew can be considered to be a result of excessive demand for downside put options, potentially caused by hedging. Another reason for the overpricing of skew could be the popularity of short volatility long (downside) skew trades (traders often hedge a short volatility position with a long skew (OTM put) position, in order to protect themselves should markets suddenly decline). The profits from shorting expensive volatility are likely to more than compensate for paying an excessive amount for the long skew position (OTM put).

³⁵ More details on local volatility can be found in the section *A1 Local Volatility* in the *Appendix*.

Figure 123. Market and Theoretical Skew

VOL REGIME DETERMINED BY TIME AND SENTIMENT

Implied volatility can be thought of as the market's estimate of future volatility³⁶. It is therefore investor sentiment that determines which implied volatility regime the market trades in, and this choice is largely determined by how much profit (or loss) a long skew position is expected to reveal over a certain time period. The choice of regime is also determined by the time horizon chosen.

Sticky delta regimes occur over long time horizon or trending markets

A sticky delta regime is typically one in which markets are trending in a stable manner (either up or down, with ATM volatility staying approximately constant) or over a very long time horizon of months or years (as over the long term the implied volatility mean reverts as it cannot go below zero or rise to infinity).

Jumpy vol regimes occur over short time horizons and panicked markets

It is rare to find a jumpy volatility regime that occurs over a long time horizon, as they tend to last for periods of only a few days or weeks. Markets tend to react in a jumpy volatility manner after a sudden and unexpected drop in equity markets (large increase in implied volatility given a decline in spot) or after a correction from such a decline (a bounce in the markets causing implied volatility to collapse). Figure 124 below summarises the different time horizons for the different volatility regimes.

³⁶ In the absence of any supply-demand imbalance in the market.

Markets tend to trade between a sticky strike and sticky local vol regime

Sticky delta and jumpy volatility are the two extremes of volatility regimes. Sticky strike and sticky local volatility are far more common volatility regimes. Sticky strike is normally associated with calmer markets than sticky local volatility (as it is closer to a sticky delta model than jumpy volatility).

Figure 124. Characteristics of Different Volatility Regimes

Characteristic	Sticky Delta	Sticky Strike	Sticky Local Vol	Jumpy Vol
Sentiment	Calm/trending		Normal	Panicked
Time horizon	Long term		Medium term	Short term
Spot vol correlation	Positive	Zero	Negative	Very Neg
Call delta	$\delta_{call} >$	$\delta_{Black-Scholes} >$	$\delta_{call} >$	δ_{call}
Put delta	$\delta_{put} >$	$\delta_{Black-Scholes} >$	$\delta_{put} >$	δ_{put}
Abs(Put delta)	$Abs(\delta_{put}) <$	$Abs(\delta_{BS}) <$	$Abs(\delta_{put}) <$	$Abs(\delta_{put})$

DELTA OF OPTION DEPENDS ON VOLATILITY REGIME

How a volatility surface reacts to a change in spot changes the value of the delta of the option. For sticky strike, as implied volatilities do not change, the delta is equal to the Black-Scholes delta.

However, if we assume a sticky delta volatility regime if an investor is long a call option, then the implied volatility of that option will decline if there is a fall in the market. The value of the call is therefore lower than expected for falls in the market. The reverse is also true as implied volatility increases if equities rise. As the value of the call is lower for declines and higher for rises (as volatility is positively correlated to spot), the delta is higher than that calculated by Black-Scholes (which is equal to the sticky strike delta).

A similar argument can be made for sticky local volatility (as volatility is negatively correlated to spot, the delta is less than the Black-Scholes delta).

EXAMINE VOL IN RELATIVE OR ABSOLUTE DIMENSIONS

To evaluate the profit – or loss – from a skew trade, assumptions have to be made regarding the movement of volatility surfaces over time. As we assume a skew trader always delta hedges, we are not concerned with the change in premium only the change in volatility.

Typically, traders use two main ways to examine implied volatility surfaces. Absolute dimensions tend to be used when examining individual options, a snapshot of volatilities, or plotting implied volatilities over a relatively short period of time. Relative dimensions tend to be used when examining implied volatilities over relatively long periods of time³⁷.

- **Absolute dimensions.** In absolute dimensions, implied volatility surfaces are examined in terms of fixed maturity (eg, Dec14 expiry) and fixed strike (eg, €4,000). This surface is a useful way of examining how the implied volatility of actual traded options changes.
- **Relative dimensions.** An implied volatility surface is examined in terms of relative dimensions when it is given in terms of relative maturity (eg, three months or one year) and relative strike (eg, ATM, 90% or 110%). Volatility surfaces tend to move in relative dimensions over a very long period of time, whereas absolute dimensions are more suitable for shorter periods of time.

Care must be taken when examining implieds in relative dimensions

As the options (and variance swaps) investors buy or sell are in fixed dimensions with fixed expiries and strikes, the change in implied volatility in absolute dimensions is the key driver of volatility profits (or losses). However, investors often use ATM volatility to determine when to enter (or exit) volatility positions, which can be misleading. For example, if there is a skew (downside implieds higher than ATM) and equity markets decline, ATM implieds will rise even though volatility surfaces remain stable. A plot of ATM implieds will imply buying volatility was profitable over the decline in equity markets; however, in practice this is not the case.

Absolute implied volatility is the key driver for equity derivative profits

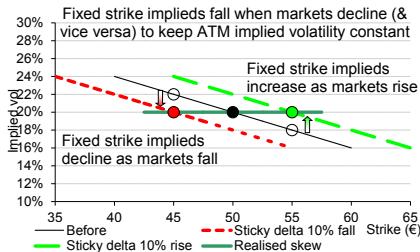
As options that are traded have a fixed strike and expiry, it is absolute implied volatility that is the driver for equity derivative profits and skew trades. However, we accept that relative implied volatility is useful when looking at long-term trends. For the volatility regimes (1) sticky delta and (2) sticky strike, we shall plot implieds using both absolute and relative dimensions in order to explain the difference. For the remaining two volatility regimes (sticky local volatility and jumpy volatility), we shall only plot implied volatility using absolute dimensions (as that is the driver of profits for traded options and variance swaps).

³⁷ This is usually for liquidity reasons, as options tend to be less liquid for maturities greater than two years (making implied volatility plots of more than two years problematic in absolute dimensions).

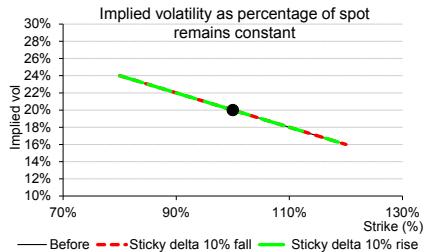
(1) STICKY DELTA ASSUMES ATM VOL IS CONSTANT

A sticky delta model assumes a constant implied volatility for strikes as a percentage of spot (eg, ATM stays constant). How a volatility surface moves with a change in spot is shown below for both absolute/fixed strike and relative strike in Figure 125.

Figure 125. Sticky Delta Fixed Strike



Sticky Delta Relative (%) Strike



RANGE-BOUNDED VOL SUPPORTS A STICKY DELTA MODEL

As implied volatility cannot be negative, it is therefore usually floored close to the lowest levels of realised volatility. Although an infinite volatility is theoretically possible, in practice implied volatility is typically capped close to the all-time highs of realised volatility. Over a long period of time, ATM implied volatility can be thought of as being range bound and likely to trend towards an average value (although this average value will change over time as the macro environment varies). As the trend towards this average value is independent of spot, the implied volatility surface in absolute dimensions (fixed currency strike) has to move to keep implied volatility surface in relative dimensions (strike as percentage of spot) constant. Thinking of implied volatility in this way is a sticky delta (or sticky moneyness) implied volatility surface model.

Sticky delta most appropriate over long term (many months or years)

While over the long term implied volatility tends to return to an average value, in the short term volatility can trade away from this value for a significant period of time. Typically, when there is a spike in volatility it takes a few months for volatility to revert back to more normal levels. This suggests a sticky delta model is most appropriate for examining implied volatilities for periods of time of a year or more. As a sticky delta model implies a positive correlation between (fixed strike) implied volatility and spot, the opposite of what is normally seen, it is not usually a realistic model for short periods of time. Trending markets (calmly rising or declining) are usually the only situation when a sticky delta model is appropriate for short periods of time. In this case, the volatility surface tends to reset to keep ATM volatility constant, as this implied volatility level is in line with the realised volatility of the market.

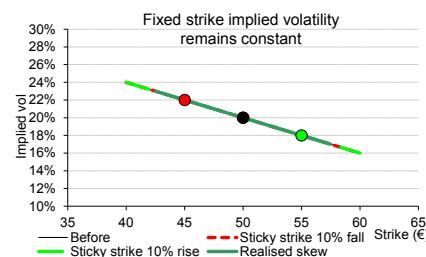
LONG SKEW UNPROFITABLE IN STICKY DELTA VOL REGIME

In a sticky delta volatility regime the fixed strike implied volatility (of traded options) has to be re-marked when spot moves. The direction of this re-mark for long skew positions causes a loss, as skew should be flat if ATM volatility is going to remain unchanged as markets move (we assume skew is negative). Additionally, the long skew position carries the additional cost of skew theta, the combination of which causes long skew positions to be very unprofitable.

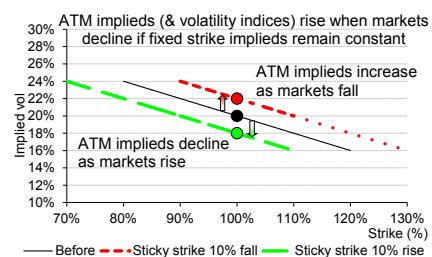
(2) STICKY STRIKE HAS NO SPOT VOL CORRELATION

A sticky strike model assumes that options of a fixed currency strike are fixed (absolute dimensions). The diagrams below show how a volatility surface moves in both absolute/fixed strike and relative strike due to a change in spot.

Figure 126. Sticky Strike Fixed Strike



Sticky Strike Relative (%) Strike



TRADER'S SYSTEMS CAN GIVE ILLUSION OF STICKY STRIKE

While Figure 124 above on page 214 describes which volatility regime normally applies in any given environment, there are many exceptions. A particular exception is that for very small time horizons volatility surfaces can seem to trade in a sticky strike regime. We believe this is due to many trading systems assuming a static strike volatility surface, which then has to be re-marked by traders (especially for less liquid instruments, as risk managers are likely to insist on volatilities being marked to their last known traded implied volatility)³⁸. As the effect of these trading systems on pricing is either an illusion (as traders will re-mark their surface when asked to provide a firm quote) or well within the bid-offer arbitrage channel, we believe this effect should be ignored.

³⁸ Anchor delta measures the effect of re-marking a volatility surface and is described in the section *A9 Advanced (Practical or Shadow) Greeks* in the *Appendix*.

LONG SKEW IS UNPROFITABLE WITH STICKY STRIKE

While there is no profit or loss from re-marking a surface in a sticky strike model, a long skew position still has to pay skew theta. Overall, a long skew position is still unprofitable with sticky strike, but it is less unprofitable than with sticky delta.

(3) STICKY LOCAL VOLATILITY PRICES SKEW FAIRLY

As a sticky local volatility causes a negative correlation between spot and Black-Scholes volatility (shown below), this re-mark is profitable for long skew positions. As the value of this re-mark is exactly equal to the cost of skew theta, skew trades break even in a sticky local volatility regime. If volatility surfaces move as predicted by sticky local volatility, then skew is priced fairly (as skew trades do not make a loss or profit).

BLACK-SCHOLES VOL IS AVERAGE OF LOCAL VOL

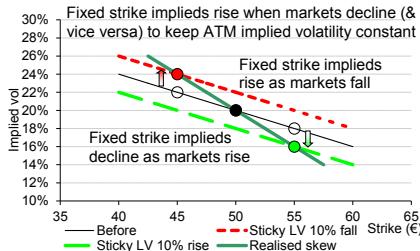
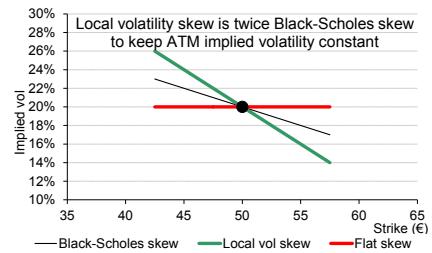
Local volatility is the name given for the instantaneous volatility of an underlying (ie, the exact volatility it has at a certain point). The Black-Scholes volatility of an option with strike K is equal to the average local (or instantaneous) volatility of all possible paths of the underlying from spot to strike K. This can be approximated by the average of the local volatility at spot and the local volatility at strike K. This approximation gives two results³⁹:

- The ATM Black-Scholes volatility is equal to the ATM local volatility.
- Black-Scholes skew is half the local volatility skew (due to averaging).

Example of local volatility skew = 2x Black-Scholes skew

The second point can be seen if we assume the local volatility for the 90% strike is 22% and the ATM local volatility is 20%. The 90%-100% local volatility skew is therefore 2%. As the Black-Scholes 90% strike option will have an implied volatility of 21% (the average of 22% and 20%), it has a 90%-100% skew of 1% (as the ATM Black-Scholes volatility is equal to the 20% ATM local volatility).

³⁹ As this is an approximation, there is a slight difference which we shall ignore.

Figure 127. Local Vol Fixed Strike**Black-Scholes and Local Vol Skew**

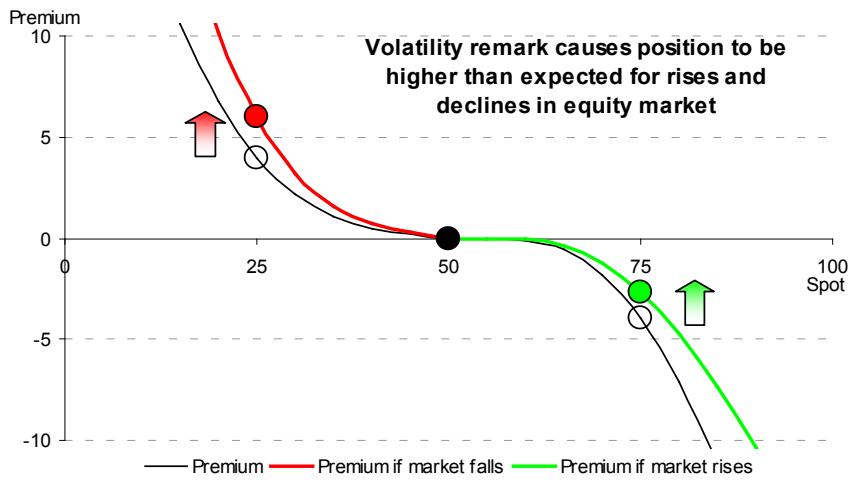
STICKY LOCAL VOL => NEGATIVE SPOT VOL CORRELATION

As local volatility skew is twice the Black-Scholes skew, and ATM volatilities are the same, a sticky local volatility surface implies a negative correlation between spot and implied volatility. This can be seen by the ATM Black-Scholes volatility resetting higher if spot declines and is shown in the diagrams above.

Example of negative correlation between spot and Black-Scholes volatility

We shall use the values from the previous example, with the local volatility for the 90% strike = 22%, Black-Scholes of the 90% strike = 21% and the ATM volatility for both (local and Black-Scholes) = 20%. If markets decline 10%, then the 90% strike option Black-Scholes volatility will rise 1% from 21% to 22% (as ATM for both local and Black-Scholes volatility must be equal). This 1% move will occur in parallel over the entire surface (as the Black-Scholes skew has not changed). Similarly, should markets rise 10%, the Black-Scholes volatility surface will fall 1% (assuming constant skew).

Figure 128. Premium of Long Put, Short Call (long skew) Risk Reversal



LONG SKEW PROFITS FROM VOLATILITY SURFACE RE-MARK

In order to demonstrate how the negative correlation between spot and (Black-Scholes) implied volatility causes long skew positions to profit from volatility surfaces re-mark, we shall assume an investor is long a risk reversal (long OTM put, short OTM call). This position is shown in Figure 128 above.

When markets fall, the primary driver of the risk reversal's value is the put (which is now more ATM than the call), and the put value will increase due to the rise in implied volatility (due to negative correlation with spot). Similarly, the theoretical value of the risk reversal will rise (as the call is now more ATM – and therefore the primary driver of value – and, as implied volatilities decline as markets rise, the value of the short call will rise as well). The long skew position therefore profits from both a movement up or down in equity markets, as can be seen in the diagram below as both the long call and short put position increase in value.

VOL RE-MARK WITH STICKY LOCAL VOL = SKEW THETA

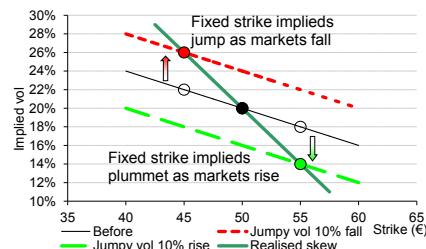
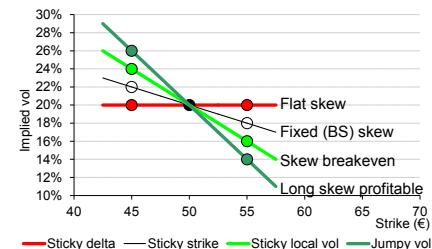
While a sticky local volatility regime causes long skew positions to profit from (Black-Scholes) implied volatility changes, the position still suffers from skew theta. The combination of these two cancel exactly, causing a long (or short) skew trade to break even. As skew trades break even under a static local volatility model, and as there is a negative spot vol correlation, it is arguably the most realistic volatility model.

Figure 129. Breakdown of P&L for Skew Trades

P&L breakdown for long skew (e.g. long put, short call)			
Volatility regime	Remark	Skew theta	Total
Sticky delta			
Sticky strike			
Sticky local volatility			
Jumpy volatility			

(4) JUMPY VOLATILITY IS ONLY REGIME WHERE LONG SKEW IS PROFITABLE

During very panicked markets, or immediately after a crash, there is typically a very high correlation between spot and volatility. During this volatility regime (which we define as jumpy volatility) volatility surfaces move in excess of that implied by sticky local volatility. As the implied volatility surface re-mark for a long skew position is in excess of skew theta, long skew positions are profitable. A jumpy volatility regime tends to last for a relatively short period of time.

Figure 130. Jumpy Vol Fixed Strike**Realised Skew of 4 Vol Regimes**

Example of volatility regimes and skew trading

If one-year 90%-100% skew is 25bp per 1% (ie, 2.5% for 90-100%) and markets fall 1%, volatility surfaces have to rise by 25bp for the profit from realised skew to compensate for the cost of skew theta. If surfaces move by more than 25bp, surfaces are moving in a jumpy volatility way and skew trades are profitable. If surfaces move by less than 25bp then skew trades suffer a loss.

SKEW ONLY FAIRLY PRICED IF ATM MOVES BY TWICE THE SKEW

For a given movement in spot from S_0 to S_1 , we shall define the movement of the (Black-Scholes) implied volatility surface divided by the skew (implied volatility of strike S_1 – implied volatility of strike S_0) to be the realised skew. The realised skew can be thought of as the profit due to re-marking the volatility surface. Defining realised skew to be the movement in the volatility surface is similar to the definition of realised volatility, which is the movement in spot.

realised skew = movement of surface/skew

where:

movement of surface = movement of surface when spot moves from S_0 to S_1

skew = difference in implied volatility between S_1 and S_0

The ATM volatility can then be determined by the below equation:

$$\text{ATM}_{\text{time } 1} = \text{ATM}_{\text{time } 0} + \text{skew} + \text{movement of surface}$$

$$\text{ATM}_{\text{time } 1} = \text{ATM}_{\text{time } 0} + \text{skew} + (\text{skew} \times \text{realised skew})$$

$$\text{ATM}_{\text{time } 1} = \text{ATM}_{\text{time } 0} + \text{skew} \times (1 + \text{realised skew})$$

The realised skew for sticky delta is therefore -1 in order to keep ATM constant (and hence skew flat) for all movements in spot. A sticky strike regime has a realised skew of 0, as there is no movement of the volatility surface and skew is fixed. A local volatility model has a realised skew of 1, which causes ATM to move by twice the value implied by a fixed skew. As local volatility prices skew fairly, skew is only fairly priced if ATM moves by twice the skew. We shall assume the volatility surface for jumpy volatility moves more than it does for sticky local volatility, hence has a realised skew of more than 1.

Skew profit is proportional to realised skew – 1 (due to skew theta)

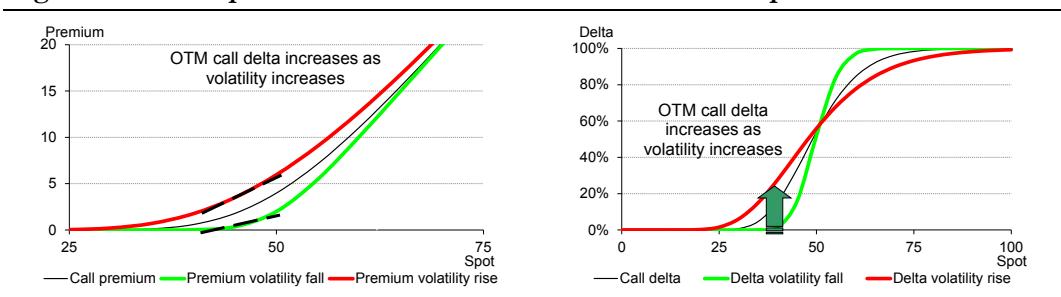
In order to calculate the relative profit (or loss) of trading skew, the value of skew theta needs to be taken away, and this value can be thought of as -1. Skew profit is then given by the formula below:

$$\text{Skew profit} \propto \text{realised skew} - 1$$

SKEW TRADING IS EQUIVALENT TO TRADING 2ND ORDER GAMMA

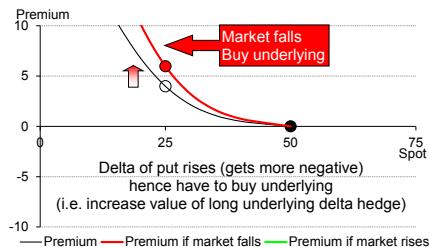
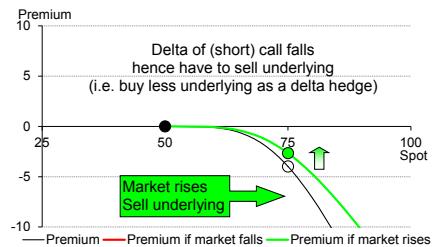
Determining the current volatility regime helps a trader decide if skew trades are likely to be profitable. In order to determine the strikes used to initiate long or short skew positions, a trader needs to evaluate the richness or cheapness of skew across different strikes. It is possible to show intuitively, and mathematically, that skew trading is very similar to delta hedging gamma. Given this relationship, comparing vanna ($d\text{Vega}/d\text{Spot}$), weighted by the square root of time, to skew theta can be a useful rule of thumb to identify potential trading opportunities.

Figure 131. Call Option with 50 Strike



MOVEMENT OF IMPLIED VOL SURFACE CHANGES DELTA

We shall assume we are in a sticky local vol (or jumpy vol) market, ie, volatility rises if markets fall, and a trader is trading skew using a long OTM put and short OTM call (ie, a risk reversal). As the delta of OTM options increases in value if implied volatility increases, and vice versa, the delta hedging of the long skew position is impacted by the movement in volatility surfaces.

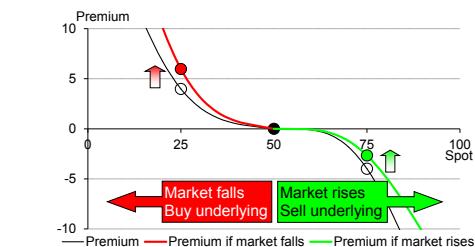
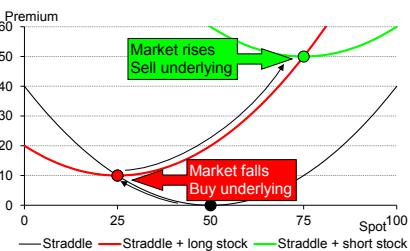
Figure 132. Delta of Put if Market Falls**Delta of Short Call if Market Rises**

Trader needs to buy stock (or futures) if market declines

If there is a decline in spot, the volatility of the long put (which is now more ATM and the primary driver of value) increases. This causes the delta of the position to decrease (absolute delta of put increases and, as delta of put is negative, the delta decreases). A trader has to buy more stock (or futures) than expected in order to compensate for this change, as shown on the left of Figure 132 above.

Conversely, trader needs to sell stock (or futures) if the market rises

The opposite trade occurs if markets rise as, for an increase in spot, the volatility of the short call (which is now more ATM and the primary driver of value) decreases. This causes the delta of the position to increase (delta of call decreases as delta of short call increases). Traders have to sell more stock (or futures) than they expect to compensate for this change (as shown on the right of Figure 132 above.), which is the reverse trade of that which occurs for a decline in the market.

Figure 133. Delta Hedging Skew**Delta Hedging Gamma**

DELTA HEDGING SKEW SIMILAR TO DELTA HEDGING GAMMA

Let us assume a negative correlation between spot and volatility (ie, for sticky local volatility or jumpy volatility) and that a trader is initially delta hedged⁴⁰ and intends to remain so. The movement of the volatility surface means the trader has to buy more stock (or futures) than he expects if markets fall and sell more stock (or futures) if markets rise. Buying low and selling high locks in the profit from the long skew position. This trade is identical to delta hedging a long gamma position, which can be seen in Figure 133 above.

If there is a positive relationship between spot and volatility (ie, a sticky delta volatility regime), then the reverse trade occurs with stock (or futures) being sold if markets decline and bought if markets rise. For sticky delta regimes, a long skew position is similar to being short gamma (and hence very unprofitable, given skew theta has to be paid as well).

MATHEMATICALLY, SKEW TRADING α GAMMA TRADING

It is possible to show mathematically the relationship between skew trading and gamma trading if one assumes a correlation between spot and volatility. Vanna, the rate of change in vega for a change in spot ($d\text{Vega}/d\text{Spot}$) measures the size of a skew position. This can be seen intuitively from the arguments above; as markets decline, the OTM put becomes more ATM and hence the primary driver of value. It is this change in vega (long put dominating the short call) for a change in spot, that causes volatility surface re-marks to be profitable for skew trading. Vanna is not only equal to $d\text{Vega}/d\text{Spot}$, but is also equal to $d\text{Delta}/d\text{Vol}$ ⁴¹. The equations below show that this relationship, when combined with spot being correlated to volatility, links skew and gamma trading.

$$\text{Vanna} = d\text{Delta}/d\text{Vol} \text{ (and } = d\text{Vega}/d\text{Spot})$$

As $\text{Vol} \propto \text{Spot}$

⁴⁰ A long put and short call risk reversal would be delta hedged with a long stock (or futures) position.

⁴¹ The proof of this relationship is outside of the scope of this publication

Vanna \propto dDelta/dSpot

As Gamma = dDelta/dSpot

Vanna \propto Gamma

Therefore, gamma can be considered to be second order gamma due to the negative correlation between volatility and spot.

SKEW THETA PAYS FOR SKEW, GAMMA THETA PAYS FOR GAMMA

In order to break down an option's profit into volatility and skew, the total theta paid needs to be separated into gamma theta and skew theta. Gamma theta pays for gamma (or volatility) while skew theta pays for skew. We note that skew across the term structure can be compared with each other if weighted by the square root of time. As skew is measured by vanna, skew theta should therefore be compared to power vanna (vanna weighted by the square root of time) to identify skew trading opportunities. This is equivalent to comparing gamma to gamma theta. The method for calculating skew theta is given below.

Total theta = gamma theta + skew theta (all measured in theta per units of cash gamma)

Cash (or dollar) gamma = $\gamma \times S^2 / 100$ = notional cash value bought (or sold) per 1% spot move

GAMMA THETA IDENTICAL FOR ALL OPTIONS IF IMPLIEDS IDENTICAL

Gamma theta is the cost (or income) from a long (or short) gamma position. To calculate the cost of gamma, we shall assume an index has a volatility of 20% for all strikes and maturities. We shall ignore interest rates, dividends and borrowing costs and assume spot is currently at 3000pts.

Figure 134. Theta (per Year)

Strike	3 Months	1 Year	4 Years
80%	-0.070	-0.227	-0.178
90%	-0.517	-0.390	-0.213
100%	-0.949	-0.473	-0.233
110%	-0.632	-0.442	-0.237
120%	-0.197	-0.342	-0.230

Cash Gamma per 1% Move

Strike	3 Months	1 Year	4 Years
80%	9	29	22
90%	65	49	27
100%	120	60	29
110%	80	56	30
120%	25	43	29

Both gamma and theta are high for short-dated ATM options

As can be seen in Figure 134 above, both cash gamma and theta are highest for near-dated and ATM options. When the cost per unit of cash gamma is calculated, it is identical for all strikes and expiries as the implied volatility is 20% for them all (see Figure 135 below). This is, essentially, the values on the left in Figure 134 divided by the values on the right in Figure 134. We shall define the ATM theta cost per unit of cash gamma to be gamma theta (in units of 1 million cash gamma to have a reasonably sized number).

Figure 135. Theta per 1 million Cash Gamma with Const Vol

Strike	3 Months	1 Year	4 Years
80%	-7,937	-7,937	-7,937
90%	-7,937	-7,937	-7,937
100% = Gamma (γ) Theta	-7,937	-7,937	-7,937
110%	-7,937	-7,937	-7,937
120%	-7,937	-7,937	-7,937

TERM STRUCTURE CHANGES GAMMA THETA BY MATURITY

In order to have a more realistic volatility surface we shall introduce positive sloping term structure (TS), while keeping the implied volatility of one-year maturity options identical. As there is no skew in the surface, all the theta is solely due to the cost of gamma or gamma theta. The gamma theta is now lower for near-dated maturities, which is intuitively correct as near-dated implieds are now lower than the far-dated implieds.

Figure 136. Vol with Term Structure

Strike	3 Months	1 Year	4 Years
80%	19%	20%	21%
90%	19%	20%	21%
100%	19%	20%	21%
110%	19%	20%	21%
120%	19%	20%	21%

Theta per 1mn Cash Gamma

Strike	3 Months	1 Year	4 Years
80%	-7,163	-7,937	-8,338
90%	-7,163	-7,937	-8,338
γ Theta	-7,163	-7,937	-8,338
110%	-7,163	-7,937	-8,338
120%	-7,163	-7,937	-8,338

SKEW MAKES IT MORE EXPENSIVE TO OWN PUTS THAN CALLS

If we introduce skew to the volatility surface we increase the cost of gamma for puts and decrease it for calls. This can be seen on the right of Figure 137; the ATM options have the same cost of gamma as before but the wings now have a different value.

Figure 137. Vol with Skew and TS

Strike	3 Months	1 Year	4 Years
80%	27%	24%	23%
90%	23%	22%	22%
100%	19%	20%	21%
110%	15%	18%	20%
120%	11%	16%	19%

Theta per 1mn Cash Gamma

Strike	3 Months	1 Year	4 Years
80%	-14,464	-11,429	-10,045
90%	-10,496	-9,603	-9,172
γ Theta	-7,163	-7,937	-8,338
110%	-4,464	-6,429	-7,545
120%	-2,401	-5,079	-6,791

SKEW THETA IS THE COST OF GOING LONG SKEW

As we have defined the theta paid for ATM option gamma (or gamma theta) as the fair price for gamma, the difference between this value and other options' cost of gamma is the cost of skew (or skew theta). Skew theta is therefore calculated by subtracting the cost of ATM gamma from all other options (and hence skew theta is zero for ATM options by definition).

Figure 138. Theta per 1mn Cash Gamma

Strike	3 Months	1 Year	4 Years
80%	-14,464	-11,429	-10,045
90%	-10,496	-9,603	-9,172
γ Theta	-7,163	-7,937	-8,338
110%	-4,464	-6,429	-7,545
120%	-2,401	-5,079	-6,791

Skew Theta

Strike	3 Months	1 Year	4 Years
80%	-7,302	-3,492	-1,706
90%	-3,333	-1,667	-833
100%	0	0	0
110%	2,698	1,508	794
120%	4,762	2,857	1,548

Example of skew theta calculation

The annual cost for a million units of cash gamma for three-month 90% strike options is €10,496, whereas ATM options only have to pay €7,163. The additional cost of being long 90% options (rather than ATM) is therefore €10,496 - €7,163 = €3,333. This additional €3,333 cost is the cost of being long skew, or skew theta.

Strikes lower than ATM suffer from skew theta

For low strike options there is a cost (negative sign) to owning the option and hence being long skew. High strike options benefit from an income of skew theta (which causes the lower cost of gamma) to compensate for being short skew (hence they have a positive sign).

VOL SLIDE THETA HAS A MINOR EFFECT ON SKEW TRADING

If one assumes volatility surfaces have relative time (one-year skew stays the same) rather than absolute time (ie, Dec14 skew stays the same) then one needs to take into account volatility slide theta (to factor in the increase in skew as the maturity of the option decreases). Volatility slide theta partly compensates for the cost of skew theta. For more details, see the section [A9 Advanced \(Practical or Shadow\) Greeks](#) in the Appendix.

APPENDIX

This includes technical detail and areas related to volatility trading that do not fit into earlier sections.

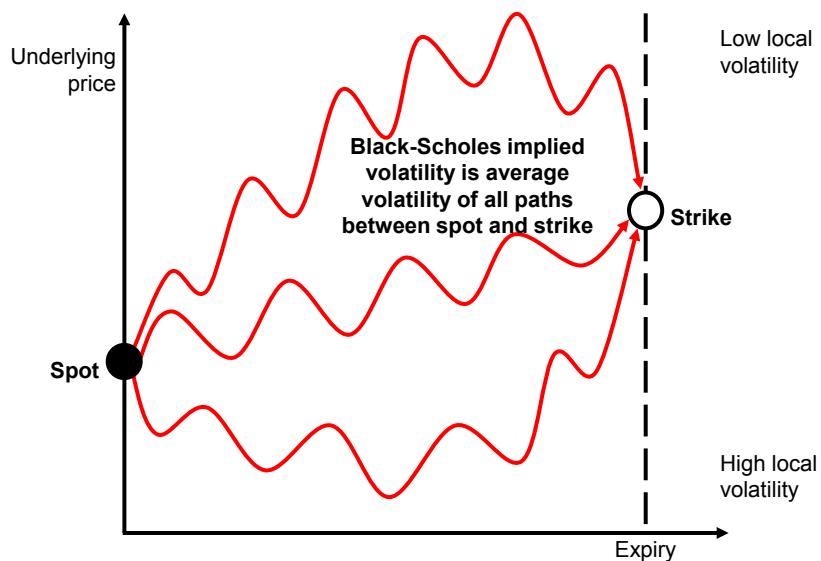
A.1: LOCAL VOLATILITY

While Black-Scholes is the most popular method for pricing vanilla equity derivatives, exotic equity derivatives (and ITM American options) usually require a more sophisticated model. The most popular model after Black-Scholes is a local volatility model as it is the only completely consistent volatility model⁴². A local volatility model describes the instantaneous volatility of a stock, whereas Black-Scholes is the average of the instantaneous volatilities between spot and strike.

LOCAL VOL IS INSTANTANEOUS VOL OF UNDERLYING

Instantaneous volatility is the volatility of an underlying at any given local point, which we shall call the local volatility. We shall assume the local volatility is fixed and has a normal negative skew (higher volatility for lower spot prices). There are many paths from spot to strike and, depending on which path is taken, they will determine how volatile the underlying is during the life of the option (see Figure 139 below).

Figure 139. Different Paths between Spot and Strike



⁴² Strictly speaking, this is true only for deterministic models. However, as the expected volatility of non-deterministic models has to give identical results to a local volatility model to be completely consistent, they can be considered to be a ‘noisy’ version of a local volatility model.

BLACK-SCHOLES VOL IS AVERAGE OF LOCAL VOLATILITIES

It is possible to calculate the local (or instantaneous) volatility surface from the Black-Scholes implied volatility surface. This is possible as the Black-Scholes implied volatility of an option is the average of all the paths between spot (ie, zero maturity ATM strike) and the maturity and strike of the option. A reasonable approximation is the average of all local volatilities on a direct straight-line path between spot and strike. For a normal relatively flat skew, this is simply the average of two values, the ATM local volatility and the strike local volatility.

Black-Scholes skew is half local volatility skew as it is the average

If the local volatility surface has a 22% implied at the 90% strike, and 20% implied at the ATM strike, then the Black-Scholes implied volatility for the 90% strike is 21% (average of 22% and 20%). As ATM implieds are identical for both local and Black-Scholes implied volatility, this means that 90%-100% skew is 2% for local volatility but 1% for Black-Scholes. Local volatility skew is therefore twice the Black-Scholes skew.

ATM volatility is the same for both Black-Scholes and local volatility

For ATM implieds, the local volatility at the strike is equal to ATM, hence the average of the two identical numbers is simply equal to the ATM implied. For this reason, Black-Scholes ATM implied is equal to local volatility ATM implied.

LOCAL VOL IS THE ONLY COMPLETE CONSISTENT VOL MODEL

A local volatility model is complete (it allows hedging based only on the underlying asset) and consistent (does not contain a contradiction). It is often used to calculate exotic option implied volatilities to ensure the prices for these exotics are consistent with the values of observed vanilla options and hence prevent arbitrage. A local volatility model is the only complete consistent volatility model; a constant Black-Scholes volatility model (constant implied volatility for all strikes and expiries) can be considered to be a special case of a static local volatility model (where the local volatilities are fixed and constant for all strikes and expiries).

A.2: MEASURING HISTORICAL VOLATILITY

The implied volatility for a certain strike and expiry has a fixed value. There is, however, no single calculation for historical volatility. The number of historical days for the historical volatility calculation changes the calculation, in addition to the estimate of the drift (or average amount stocks are assumed to rise). There should, however, be no difference between the average daily or weekly historical volatility. We also examine different methods of historical volatility calculation, including close-to-close volatility and exponentially weighted volatility, in addition to advanced volatility measures such as Parkinson, Garman-Klass (including Yang-Zhang extension), Rogers and Satchell and Yang-Zhang. We also show that it is best to assume a zero drift assumption for close-to-close volatility, and that under this condition variance is additive.

DEFINITION OF VOLATILITY

Assuming that the probability distribution of the log returns of a particular security is normally distributed (or follows a normal ‘bell-shape distribution’), volatility σ of that security can be defined as the standard deviation of the normal distribution of the log returns. As the mean absolute deviation is $\sqrt{2/\pi}$ (≈ 0.8) \times volatility, the volatility can be thought of as c1.25 \times the expected percentage change (positive or negative) of the security.

$$\sigma = \text{standard deviation of log returns} \times \sqrt{1/\Delta t}$$

CLOSE-TO-CLOSE HISTORICAL VOLATILITY IS THE MOST COMMON

Volatility is defined as the annualised standard deviation of log returns. For historical volatility the usual measure is close-to-close volatility, which is shown below.

$$\text{Log return } = x_i = \ln\left(\frac{c_i + d_i}{c_{i-1}}\right) \text{ where } d_i = \text{ordinary dividend and } c_i \text{ is close price}$$

$$\text{Volatility}^{43} \text{ (not annualised)} = \sigma_x = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2}$$

where \bar{x} = drift = Average (x_i)

Historical volatility calculation is an estimate from a sample

⁴³ We take the definition of volatility of John Hull in *Options, Futures and Other Derivatives* in which n day volatility uses n returns and n+1 prices. We note Bloomberg uses n prices and n-1 returns.

Historical volatility is calculated as the standard deviation of the log returns of a particular securities' time series. If the log returns calculation is based on daily data, we have to multiply this number by the square root of 252 (the number of trading days in a calendar year) in order to annualise the volatility calculation (as $\Delta t = 1/252$ hence $\sqrt{1/\Delta t} = \sqrt{252}$). As a general rule, to annualise the volatility calculation, regardless of the periodicity of the data, the standard deviation has to be multiplied by the square root of the number of days/weeks/months within a year (ie, $\sqrt{252}$, $\sqrt{52}$, $\sqrt{12}$).

$$\sigma_{\text{Annualised}} = \sigma_x \times \sqrt{\text{values in year}}$$

VARIANCE IS ADDITIVE IF ZERO MEAN IS ASSUMED

Frequency of returns in a year = F (eg, 252 for daily returns)

$$\sigma_{\text{Annualised}} = \sqrt{F \times \sigma_x^2} = \sqrt{F} \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \bar{x})^2}$$

As $\bar{x} \approx 0$ if we assume zero average returns

$$\sigma_{\text{Annualised}} = \sqrt{F} \sqrt{\frac{1}{N} \sum_{i=1}^N x_i^2}$$

$$\sigma_{\text{Annualised}}^2 = \frac{F}{N} \sum_{i=1}^N x_i^2$$

Now if we assume that the total sample N can be divided up into period 1 and period 2 where period 1 is the first M returns then:

$$\sigma_{\text{Total}}^2 = \frac{F}{N_{\text{Total}}} \times \sum_{i=1}^N x_i^2$$

$$\sigma_{\text{Period 1}}^2 = \frac{F}{N_{\text{Period 1}}} \times \sum_{i=1}^M x_i^2 \quad (\text{where } N_{\text{Period 1}} = M)$$

$$\sigma_{\text{Period 2}}^2 = \frac{F}{N_{\text{Period 2}}} \times \sum_{i=M+1}^N x_i^2 \quad (\text{where } N_{\text{Period 2}} = N - M)$$

then

$$\sigma_{\text{Total}}^2 = \frac{F}{N_{\text{Total}}} \times \sum_{i=1}^N x_i^2 = \frac{F}{N_{\text{Total}}} \left(\sum_{i=1}^M x_i^2 + \sum_{i=M+1}^N x_i^2 \right)$$

$$\sigma_{\text{Total}}^2 = \frac{F}{N_{\text{Total}}} \times \sum_{i=1}^M x_i^2 + \frac{F}{N_{\text{Total}}} \times \sum_{i=M+1}^N x_i^2$$

$$\sigma_{\text{Total}}^2 = \frac{N_{\text{Period 1}}}{N_{\text{Total}}} \left(\frac{F}{N_{\text{Period 1}}} \times \sum_{i=1}^M x_i^2 \right) + \frac{N_{\text{Period 2}}}{N_{\text{Total}}} \left(\frac{F}{N_{\text{Period 2}}} \times \sum_{i=M+1}^N x_i^2 \right)$$

$$\sigma_{\text{Total}}^2 = \frac{N_{\text{Period 1}}}{N_{\text{Total}}} \sigma_{\text{Period 1}}^2 + \frac{N_{\text{Period 2}}}{N_{\text{Total}}} \sigma_{\text{Period 2}}^2$$

Hence variance is additive (when weighted by the time in each period / total time)

BEST TO ASSUME ZERO DRIFT FOR VOL CALCULATION

The calculation for standard deviation calculates the deviation from the average log return (or drift). This average log return has to be estimated from the sample, which can cause problems if the return over the period sampled is very high or negative. As over the long term very high or negative returns are not realistic, the calculation of volatility can be corrupted by using the sample log return as the expected future return. For example, if an underlying rises 10% a day for ten days, the volatility of the stock is zero (as there is zero deviation from the 10% average return). This is why volatility calculations are normally more reliable if a zero return is assumed. In theory, the expected average value of an underlying at a future date should be the value of the forward at that date. As for all normal interest rates (and dividends, borrow cost) the forward return should be close to 100% (for any reasonable sampling frequency, ie, daily/weekly/monthly). Hence, for simplicity reasons it is easier to assume a zero log return as $\ln(100\%) = 0$.

WHICH HISTORICAL VOLATILITY SHOULD I USE?

When examining how attractive the implied volatility of an option is, investors will often compare it to historical volatility. However, historical volatility needs two parameters.

- Length of time (eg, number of days/weeks/months)
- Frequency of measurement (eg, daily/weekly)

LENGTH OF TIME FOR HISTORICAL VOLATILITY

Choosing the historical volatility number of days is not a trivial choice. Some investors believe the best number of days of historical volatility to look at is the same as the implied volatility of interest. For example, one-month implied should be compared to 21 trading day historical volatility (and three-month implied should be compared to 63-day historical volatility, etc). While an identical duration historical volatility is useful to arrive at a realistic minimum and maximum value over a long period of time, it is not always the best period of time to determine the fair level of long-dated implieds. This is because volatility mean reverts over a period of c8 months. Using historical volatility for periods longer than c8 months is not likely to be the best estimate of future volatility (as it could include volatility caused by earlier events, whose effect on the market has passed). Arguably a multiple of three months should be used to ensure that there is always the same number of quarterly reporting dates in the historical volatility measure. Additionally, if there has been a recent jump in the share price that is not expected to reoccur, the period of time chosen should exclude that jump.

The best historical volatility period does not have to be the most recent

If there has been a rare event which caused a volatility spike, the best estimate of future volatility is not necessarily the current historical volatility. A better estimate could be the past historical volatility when an event that caused a similar volatility spike occurred. For example, the volatility post credit crunch could be compared to the volatility spike after the Great Depression or during the bursting of the tech bubble.

FREQUENCY OF HISTORICAL VOLATILITY

While historical volatility can be measured monthly, quarterly or yearly, it is usually measured daily or weekly. Normally, daily volatility is preferable to weekly volatility as five times as many data points are available. However, if volatility over a long period of time is being examined between two different markets, weekly volatility could be the best measure to reduce the influence of different public holidays (and trading hours⁴⁴). If stock price returns are independent, then the daily and weekly historical volatility should on average be the same. If stock price returns are not independent, there could be a difference. Autocorrelation is the correlation between two different returns so independent returns have an autocorrelation of 0%.

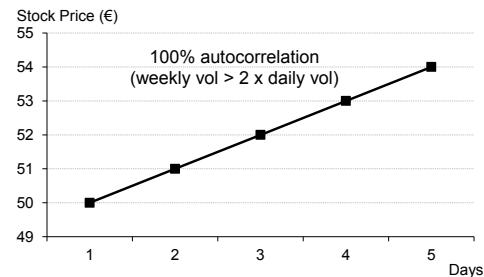
Trending markets imply weekly volatility is greater than daily volatility

With 100% autocorrelation, returns are perfectly correlated (ie, trending markets). Should autocorrelation be -100% correlated, then a positive return is followed by a negative return (mean reverting or range trading markets). If we assume markets are 100% daily correlated with a 1% daily return, this means the weekly return is 5%. The daily volatility is therefore

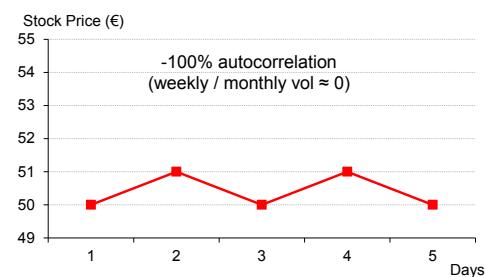
⁴⁴ Advanced volatility measures could be used to remove part of the effect of different trading hours.

c16% ($1\% \times \sqrt{252}$), while the weekly volatility of c35% ($5\% \times \sqrt{52}$) is more than twice as large.

Figure 140. 100% Daily Autocorrelation



Negative 100% Daily Autocorrelation



High market share of HFT should prevent autocorrelation

Historically (decades ago), there could have been positive autocorrelation due to momentum buying, but once this became understood this effect is likely to have faded. Given the current high market share of HFT or high frequency trading (accounting for up to three-quarters of US equity trading volume), it appears unlikely that a simple trading strategy such as ‘buy if security goes up, sell if it goes down’ will provide above-average returns over a significant period of time⁴⁵.

Panicked markets could cause temporary negative autocorrelation

While positive autocorrelation is likely to be arbitraged out of the market, there is evidence that markets can overreact at times of stress as market panic (rare statistical events can occur under the weak form of efficient market hypotheses). During these events human traders and some automated trading systems are likely to stop trading (as the event is rare, the correct response is unknown), or potentially exaggerate the trend (as positions get ‘stopped out’ or to follow the momentum of the move). A strategy that is long daily variance and short weekly variance will therefore usually give relatively flat returns, but occasionally give a positive return.

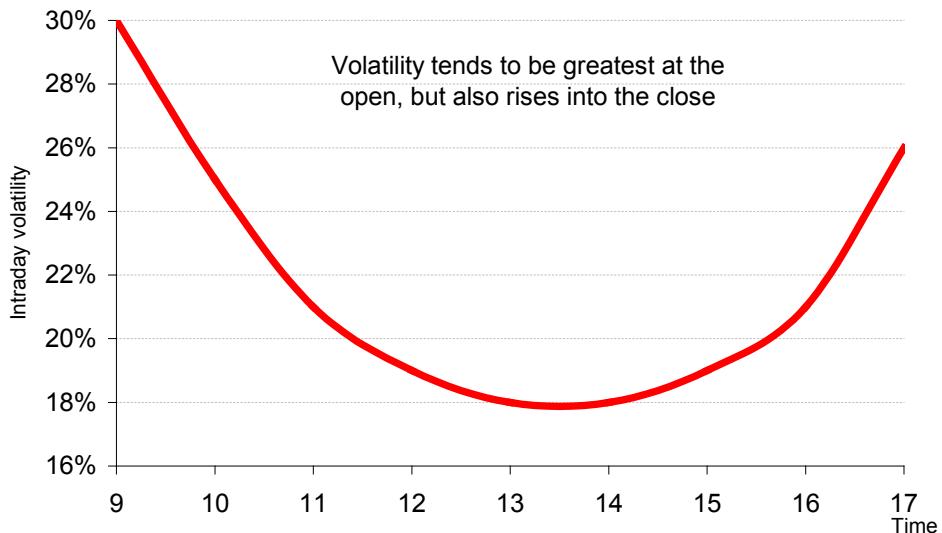
INTRADAY VOLATILITY IS NOT CONSTANT

For most markets, intraday volatility is greatest just after the open (as results are often announced around the open) and just before the close (performance is often based upon closing prices). Intraday volatility tends to sag in the middle of the day due to the combination of a lack of announcements and reduced volumes/liquidity owing to lunch breaks. For this reason, using an estimate of volatility more frequent than daily tends to be

⁴⁵ Assuming there are no short selling restrictions.

very noisy. Traders who wish to take into account intraday prices should instead use an advanced volatility measure.

Figure 141. Intraday Volatility



EXPONENTIALLY WEIGHTED VOL IS RARELY USED

An alternate measure could be to use an exponentially weighted moving average model, which is shown below. The parameter λ is between zero (effectively one-day volatility) and one (ignore current vol and keep vol constant). Normally, values of c0.9 are used. Exponentially weighted volatilities are rarely used, partly due to the fact they do not handle regular volatility-driving events such as earnings very well. Previous earnings jumps will have least weight just before an earnings date (when future volatility is most likely to be high) and most weight just after earnings (when future volatility is most likely to be low). It could, however, be of some use for indices.

$$\sigma_i^2 = \lambda \sigma_{i-1}^2 + (1 - \lambda)x_i^2$$

Exponentially weighted volatility avoids volatility collapse of historic volatility

Exponential volatility has the advantage over standard historical volatility in that the effect of a spike in volatility gradually fades (as opposed to suddenly disappearing causing a collapse in historic volatility). For example, if we are looking at the historical volatility over the past month and a spike in realised volatility suddenly occurs the historical volatility will be high for a month, then collapse. Exponentially weighted volatility will rise at the same time as

historical volatility and then gradually decline to lower levels (arguably in a similar way to how implied volatility spikes, then mean reverts).

ADVANCED VOLATILITY MEASURES

Close-to-close volatility is usually used as it has the benefit of using the closing auction prices only. Should other prices be used, then they could be vulnerable to manipulation or a ‘fat fingered’ trade. However, a large number of samples need to be used to get a good estimate of historical volatility, and using a large number of closing values can obscure short-term changes in volatility. There are, however, different methods of calculating volatility using some or all of the open (O), high (H), low (L) and close (C). The methods are listed in order of their maximum efficiency (close-to-close var divided by alternative measure var).

- **Close to close (C).** The most common type of calculation that benefits from only using reliable prices from closing auctions. By definition its efficiency is one at all times.
- **Parkinson (HL).** As this estimate only uses the high and low price for an underlying, it is less sensitive to differences in trading hours. For example, as the time of the EU and US closes are approximately half a trading day apart, they can give very different returns. Using the high and low means the trading over the whole day is examined, and the days overlap. As it does not handle jumps, on average it underestimates the volatility, as it does not take into account highs and lows when trading does not occur (weekends, between close and open). Although it does not handle drift, this is usually small. While other measures are more efficient based on simulated data, some studies have shown it to be the best measure for actual empirical data.
- **Garman-Klass (OHLC).** This estimate is the most powerful for stocks with Brownian motion, zero drift and no opening jumps (ie, opening price is equal to closing price of previous period). Like Parkinson, it also underestimates the volatility (as it assumes no jumps).
- **Rogers-Satchell (OHLC).** The efficiency of the Rogers-Satchell estimate is similar to that for Garman-Klass; however, it benefits from being able to handle non-zero drift. Opening jumps are not handled well though, which means it underestimates the volatility.
- **Garman-Klass Yang-Zhang extension (OHLC).** Yang-Zhang extended the Garman-Klass method that allows for opening jumps hence it is a fair estimate, but does assume zero drift. It has an efficiency of eight times the close-to-close estimate.
- **Yang-Zhang (OHLC).** The most powerful volatility estimator which has minimum estimation error. It is a weighted average of Rogers-Satchell, the close-open volatility and

the open-close volatility. It is up to a maximum of 14 times as efficient (for two days of data) as the close-to-close estimate.

Figure 142. Summary of Advanced Volatility Estimates

Estimate	Prices Taken	Handle Drift?	Handle Overnight Jumps?	Efficiency (max)
Close to close	C	No	No	1
Parkinson	HL	No	No	5.2
Garman-Klass	OHLC	No	No	7.4
Rogers-Satchell	OHLC	Yes	No	8
Garman-Klass Yang-Zhang ext.	OHLC	No	Yes	8
Yang-Zhang	OHLC	Yes	Yes	14

EFFICIENCY AND BIAS DETERMINE BEST VOL MEASURE

There are two measures that can be used to determine the quality of a volatility measure: efficiency and bias. Generally, for small sample sizes the Yang-Zhang measure is best overall, and for large sample sizes the standard close to close measure is best.

- **Efficiency.** Efficiency(σ_x^2) = $\frac{\sigma_{cc}^2}{\sigma_x^2}$ where σ_x is the volatility of the estimate and σ_{cc} is the volatility of the standard close to close estimate.
- **Bias.** Difference between the estimated variance and the average (ie, integrated) volatility.

Efficiency measures the volatility of the estimate

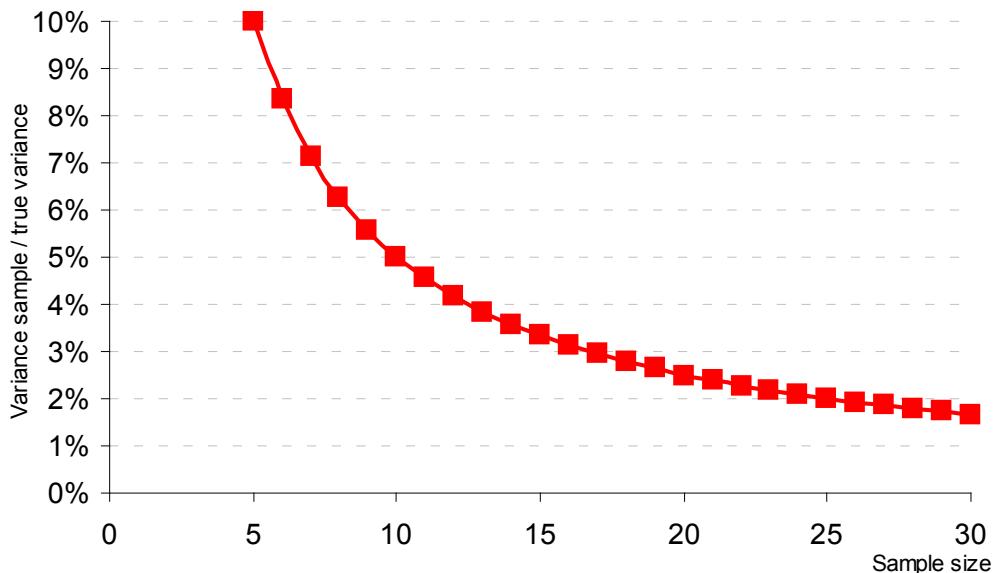
The efficiency describes the variance, or volatility of the estimate. The efficiency is dependent on the number of samples, with efficiency decreasing the more samples there are (as close-to-close will converge and become less volatile with more samples). The efficiency is the theoretical maximum performance against an idealised distribution, and with real empirical data a far smaller benefit is usually seen (especially for long time series). For example, while the Yang-Zhang based estimators deal with overnight jumps if the jumps are large compared to the daily volatility the estimate will converge with the close-to-close volatility and have an efficiency close to one.

Close-to-close vol should use at least five samples (and ideally 20 or more)

The variance of the close-to-close volatility can be estimated as a percentage of the actual variance by the formula $1/(2N)$ where N is the number of samples. This is shown in Figure 143 below and demonstrates that at least five samples are needed (or the estimate has an

error of over 10%) and that only marginal extra accuracy is gained for each additional sample above 20.

Figure 143. Variance of Close-To-Close Volatility/Actual Variance



Variance, volatility and gamma swaps should look at close to close

As the payout of variance, volatility and gamma swaps are based on close-to-close prices, the standard close-to-close volatility (or variance) should be used for comparing their price against realised. Additionally, if a trader only hedges at the close (potentially for liquidity reasons) then again the standard close-to-close volatility measure should be used.

Bias depends on the type of distribution of the underlying

While efficiency (how volatile the measure is) is important, so too is bias (whether the measure is, on average, too high or low). Bias depends on the sample size, and the type of distribution the underlying security has. Generally, the close-to-close volatility estimator is too big⁴⁶ (as it does not model overnight jumps), while alternative estimators are too small (as they assume continuous trading, and discrete trading will have a smaller difference between the maximum and minimum). The key variables that determine the bias are:

- **Sample size.** As the standard close-to-close volatility measure suffers with small sample sizes, this is where alternative measures perform best (the highest efficiency is reached for only two days of data).
- **Volatility of volatility.** While the close-to-close volatility estimate is relatively insensitive to a changing volatility (vol of vol), the alternative estimates are far more sensitive. This bias increases the more vol of vol increases (ie, more vol of vol means a greater underestimate of volatility).
- **Overnight jumps between close and open.** Approximately one-sixth of equity volatility occurs outside the trading day (and approximately twice that amount for ADRs). Overnight jumps cause the standard close-to-close estimate to overestimate the volatility, as jumps are not modelled. Alternative estimates that do not model jumps (Parkinson, Garman Klass and Rogers-Satchell) underestimate the volatility. Yang-Zhang estimates (both Yang-Zhang extension of Garman Klass and the Yang-Zhang measure itself) will converge with standard close-to-close volatility if the jumps are large compared to the overnight volatility.
- **Drift of underlying.** If the drift of the underlying is ignored as it is for Parkinson and Garman Klass (and the Yang Zhang extension of Garman Glass), then the measure will overestimate the volatility. This effect is small for any reasonable drifts (ie, if we are looking at daily, weekly or monthly data).
- **Correlation daily volatility and overnight volatility.** While Yang-Zhang measures deal with overnight volatility, there is the assumption that overnight volatility and daily volatility are uncorrelated. Yang-Zhang measures will underestimate volatility when there is a correlation between daily return and overnight return (and vice versa), but this effect is small.

⁴⁶ Compared to integrated volatility.

CLOSE-TO-CLOSE

The simplest volatility measure is the standard close-to-close volatility. We note that the volatility should be the standard deviation multiplied by $\sqrt{N}/(N-1)$ to take into account the fact we are sampling the population (or take standard deviation of the sample)⁴⁷. We ignored this in the earlier definition as for reasonably large n it $\sqrt{N}/(N-1)$ is roughly equal to one.

$$\text{Standard dev of } x = s_x = \sqrt{\frac{F}{N} \sqrt{\sum_{i=1}^N (x_i - \bar{x})^2}}$$

As $\sigma = \sqrt{\sigma^2} = \sqrt{E(s^2)} < E(\sqrt{s^2}) = E(s)$ by Jensens's inequality

$$\text{Volatility} = \sigma_x = s_x \times \sqrt{\frac{N}{N-1}}$$

$$\text{Volatility}_{\text{close to close}} = \sigma_{cc} = \sqrt{\frac{F}{N-1} \sqrt{\sum_{i=1}^N (x_i - \bar{x})^2}} = \sqrt{\frac{F}{N-1} \sqrt{\sum_{i=1}^N \left(\ln\left(\frac{c_i}{c_{i-1}}\right) \right)^2}}$$

with zero
drift

PARKINSON

The first advanced volatility estimator was created by Parkinson in 1980, and instead of using closing prices it uses the high and low price. One drawback of this estimator is that it assumes continuous trading, hence it underestimates the volatility as potential movements when the market is shut are ignored.

$$\text{Volatility}_{\text{Parkinson}} = \sigma_p = \sqrt{\frac{F}{N} \sqrt{\frac{1}{4 \ln(2)} \sum_{i=1}^N \left(\ln\left(\frac{h_i}{l_i}\right) \right)^2}}$$

⁴⁷ As the formula for standard deviation has N-1 degrees of freedom (as we subtract the sample average from each value of x)

GARMAN-KLASS

Later in 1980 the Garman-Klass volatility estimator was created. It is an extension of Parkinson which includes opening and closing prices (if opening prices are not available the close from the previous day can be used instead). As overnight jumps are ignored the measure underestimates the volatility.

$$\text{Volatility}_{\text{Garman-Klass}} = \sigma_{\text{GK}} = \sqrt{\frac{F}{N} \sqrt{\sum_{i=1}^N \frac{1}{2} \left(\ln\left(\frac{h_i}{l_i}\right)^2 - (2\ln(2)-1) \left(\ln\left(\frac{c_i}{o_i}\right)^2 \right)}^2}}$$

ROGERS-SATCHELL

All of the previous advanced volatility measures assume the average return (or drift) is zero. Securities that have a drift, or non-zero mean, require a more sophisticated measure of volatility. The Rogers-Satchell volatility created in the early 1990s is able to properly measure the volatility for securities with non-zero mean. It does not, however, handle jumps; hence, it underestimates the volatility.

$$\text{Volatility}_{\text{Rogers-Satchell}} = \sigma_{\text{RS}} = \sqrt{\frac{F}{N} \sqrt{\sum_{i=1}^N \ln\left(\frac{h_i}{c_i}\right) \ln\left(\frac{h_i}{o_i}\right) + \ln\left(\frac{l_i}{c_i}\right) \ln\left(\frac{l_i}{o_i}\right)}}$$

GARMAN-KLASS YANG-ZHANG EXTENSION

Yang-Zhang modified the Garman-Klass volatility measure in order to let it handle jumps. The measure does assume a zero drift; hence, it will overestimate the volatility if a security has a non-zero mean return. As the effect of drift is small, the fact continuous prices are not available usually means it underestimates the volatility (but by a smaller amount than the previous alternative measures).

$$\text{Volatility}_{\text{GKYZ}} =$$

$$\sigma_{\text{GKYZ}} = \sqrt{\frac{F}{N} \sqrt{\sum_{i=1}^N \left(\ln\left(\frac{o_i}{c_{i-1}}\right)^2 + \frac{1}{2} \left(\ln\left(\frac{h_i}{l_i}\right)^2 - (2\ln(2)-1) \left(\ln\left(\frac{c_i}{o_i}\right)^2 \right) \right) \right)}}$$

YANG-ZHANG

In 2000 Yang-Zhang created a volatility measure that handles both opening jumps and drift. It is the sum of the overnight volatility (close-to-open volatility) and a weighted average of the Rogers-Satchell volatility and the open-to-close volatility. The assumption of continuous prices does mean the measure tends to slightly underestimate the volatility.

Volatility_{Yang-Zhang} =

$$\sigma_{YZ} = \sqrt{\sigma_{\text{overnight volatility}}^2 + k \sigma_{\text{open to close volatility}}^2 + (1-k)\sigma_{RS}^2}$$

$$\text{where } k = \frac{0.34}{1.34 + \frac{N+1}{N-1}}$$

$$\sigma_{\text{overnight volatility}}^2 = \frac{F}{N-1} \sum_{i=1}^N \left[\ln\left(\frac{o_i}{c_{i-1}}\right) - \overline{\ln\left(\frac{o_i}{c_{i-1}}\right)} \right]^2$$

$$\sigma_{\text{open to close volatility}}^2 = \frac{F}{N-1} \sum_{i=1}^N \left[\ln\left(\frac{c_i}{o_i}\right) - \overline{\ln\left(\frac{c_i}{o_i}\right)} \right]^2$$

A.3: PROOF IMPLIED JUMP FORMULA

In the section 6.4 Trading earnings announcements/jumps we showed that the implied jump from an earnings announcement (or any one off event) can be backed out from the implied volatility of two options. We prove the formula for the expected daily return from the implied volatility jump.

ADDITIVE VAR ALLOWS JUMP VOL TO BE CALCULATED

The formula for calculating the jump volatility (which relies on variance being additive) is shown below.

$$\sigma_{\text{Jump}} = \sqrt{(\sigma_{\text{Expiry after jump}}^2 T - \sigma_{\text{Diffusive}}^2 (T - 1))}$$

where

$\sigma_{\text{Expiry after jump}}$ = implied volatility of option whose expiry is after the jump

T = time to the expiry after jump (= T_1)

$\sigma_{\text{Diffusive}}$ = diffusive volatility ($\sigma_{\text{Before jump}}$ if there is an expiry before the jump, if not it is σ_{12})

σ_{Jump} = implied volatility due to the jump

PROOF OF IMPLIED JUMP (EXPECTED DAILY RETURN)

The derivation of how to calculate the implied daily return on the day of the jump (which is a combination of the normal daily move and the effect of the jump) from the implied volatility due to jump (σ_{Jump}) is below.

$$\Delta S = S_1 - S_0 = S_0 \left(\frac{S_1}{S_0} - 1 \right) = S_0 (e^r - 1) \text{ where } r = \ln \left(\frac{S_1}{S_0} \right)$$

$$\Rightarrow E \left(\frac{|\Delta S|}{S_0} \right) = E(|e^r - 1|)$$

$$\Rightarrow \text{Expected daily return} = E(|e^r - 1|)$$

as r is normally distributed

$$\Rightarrow \text{Expected daily return} = \frac{1}{\sqrt{2\pi\sigma^2}} \int_{-\infty}^{\infty} (e^r - 1) e^{\frac{-r^2}{2\sigma^2}} dr$$

$$\Rightarrow \text{Expected daily return} = \frac{1}{\sqrt{2\pi\sigma^2}} \int_0^{\infty} (e^r - 1) e^{\frac{-r^2}{2\sigma^2}} dr + \frac{1}{\sqrt{2\pi\sigma^2}} \int_{-\infty}^0 (e^r - 1) e^{\frac{-r^2}{2\sigma^2}} dr$$

$$\Rightarrow \text{Expected daily return} = \frac{1}{\sqrt{2\pi\sigma^2}} \int_0^{\infty} (e^r - e^{-r}) e^{\frac{-r^2}{2\sigma^2}} dr$$

if we define x such that $r = x\sigma$

$$\Rightarrow \text{Expected daily return} = \frac{1}{\sqrt{2\pi}} \int_0^{\infty} (e^{x\sigma} - e^{-x\sigma}) e^{\frac{-x^2}{2}} dx$$

$$\Rightarrow \text{Expected daily return} = \frac{1}{\sqrt{2\pi}} \int_0^{\infty} e^{\frac{x\sigma - \frac{x^2}{2}}{2}} - e^{\frac{-x\sigma - \frac{x^2}{2}}{2}} dx$$

$$\Rightarrow \text{Expected daily return} = \frac{1}{\sqrt{2\pi}} \int_0^{\infty} e^{\frac{x\sigma - \frac{x^2}{2}}{2}} dx - \frac{1}{\sqrt{2\pi}} \int_0^{\infty} e^{\frac{-x\sigma - \frac{x^2}{2}}{2}} dx$$

$$\Rightarrow \text{Expected daily return} = \frac{1}{\sqrt{2\pi}} \int_0^{\infty} e^{\frac{-(x-\sigma)^2 + \frac{\sigma^2}{2}}{2}} dx - \frac{1}{\sqrt{2\pi}} \int_0^{\infty} e^{\frac{-(x+\sigma)^2 + \frac{\sigma^2}{2}}{2}} dx$$

$$\Rightarrow \text{Expected daily return} = e^{\frac{\sigma^2}{2}} \frac{1}{\sqrt{2\pi}} \int_0^{\infty} e^{\frac{-(x-\sigma)^2}{2}} dx - e^{\frac{\sigma^2}{2}} \frac{1}{\sqrt{2\pi}} \int_0^{\infty} e^{\frac{-(x+\sigma)^2}{2}} dx$$

$$\Rightarrow \text{Expected daily return} = e^{\frac{\sigma^2}{2}} \frac{1}{\sqrt{2\pi}} \int_{-\infty}^0 e^{\frac{-(x-\sigma)^2}{2}} dx - e^{\frac{\sigma^2}{2}} \frac{1}{\sqrt{2\pi}} \int_0^{\infty} e^{\frac{-(x+\sigma)^2}{2}} dx$$

$$\begin{aligned}\Rightarrow \text{Expected daily return} &= e^{\frac{\sigma^2}{2}} \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\sigma} e^{\frac{-x^2}{2}} dx - e^{\frac{\sigma^2}{2}} \frac{1}{\sqrt{2\pi}} \int_{\sigma}^{\infty} e^{\frac{-x^2}{2}} dx \\ \Rightarrow \text{Expected daily return} &= e^{\frac{\sigma^2}{2}} N(\sigma) - e^{\frac{\sigma^2}{2}} [1 - N(\sigma)] \\ \Rightarrow \text{Expected daily return} &= e^{\frac{\sigma^2}{2}} [2 \times N(\sigma) - 1]\end{aligned}$$

A.4: PROOF VARIANCE SWAPS CAN BE HEDGED BY LOG CONTRACT ($=1/K^2$)

A log contract is a portfolio of options of all strikes (K) weighted by $1/K^2$. When this portfolio of options is delta hedged on the close, the payoff is identical to the payoff of a variance swap. We prove this relationship and hence show that the volatility of a variance swap can be hedged with a static position in a log contract.

PORFOLIO OF OPTIONS WITH CONST VEGA WEIGHTED $1/K^2$

In order to prove that a portfolio of options with flat vega has to be weighted $1/K^2$, we will define the variable x to be K/S (strike K divided by spot S). With this definition and assuming zero interest rates, the standard Black-Scholes formula for vega of an option simplifies to:

$$\text{Vega of option} = \tau \times S \times f(x, v)$$

where

$$x = K / S \text{ (strike a ratio of spot)}$$

$$\tau = \text{time to maturity}$$

$$v = \sigma^2 \tau \text{ (total variance)}$$

$$f(x, v) = \frac{1}{\sqrt{2\pi}} \times e^{\frac{-d_1^2}{2}}$$

$$d_1 = \frac{\ln(\frac{1}{x}) + \frac{v}{2}}{\sqrt{v}}$$

If we have a portfolio of options where the weight of each option is $w(K)$, then the vega of the portfolio of options $V(S)$ is:

$$V(S) = \tau \int_{K=0}^{\infty} w(K) \times S \times f(x, v) dK$$

As $K = xS$ this means $dK / dx = S$, hence $dK = S \times dx$ and we can change variable K for x .

$$V(S) = \tau \int_{x=0}^{\infty} w(xS) \times S^2 \times f(x, v) dx$$

In order for the portfolio of options to have a constant vega – no matter what the level of spot – $dV(S)/dS$ has to be equal to zero.

$$\frac{dV}{dS} = \tau \int_{x=0}^{\infty} \frac{d}{dS} [S^2 w(xS)] \times f(x, v) dx = 0$$

And by the chain rule:

$$\begin{aligned} &\Rightarrow \tau \int_{x=0}^{\infty} \left[2Sw(xS) + S^2 \frac{d}{dS} w(xS) \right] \times f(x, v) dx = 0 \\ &\Rightarrow \tau \int_{x=0}^{\infty} S \left[2w(xS) + S \frac{d}{dS} w(xS) \right] \times f(x, v) dx = 0 \end{aligned}$$

As $d/dS = (d/dK) \times (dK/dS)$, and $dK/dS = x$

$$\Rightarrow \tau \int_{x=0}^{\infty} S \left[2w(xS) + xS \frac{d}{dK} w(xS) \right] \times f(x, v) dx = 0$$

As $K = xS$

$$\Rightarrow \tau \int_{x=0}^{\infty} S \left[2w(xS) + K \frac{d}{dK} w(K) \right] \times f(x, v) dx = 0$$

$$\Rightarrow 2w + K \frac{d}{dK} w(K) = 0 \text{ for all values of } S$$

$$\Rightarrow w(K) = \frac{\text{constant}}{K^2}$$

A.5: PROOF VARIANCE SWAP NOTIONAL = VEGA/2 σ

For small differences between the future volatility and current (implied) swap volatility, the payout of a volatility swap can be approximated by a variance swap. We show how the difference in their notionals should be weighted by 2σ .

Proof that Variance Swap Notional = vega/ 2σ

We intend to calculate the relative size (Z) of the variance swap notional compared to volatility swap notional (volatility swap notional = vega by definition) so they have a similar payout (for small differences between realised and implied volatility).

$$\text{Notional}_{\text{variance swap}} \approx Z \times \text{Notional}_{\text{volatility swap}}$$

$$\Rightarrow (\sigma_F - \sigma_S) \approx Z (\sigma_F^2 - \sigma_S^2)$$

where:

σ_F = future volatility (that occurs over the life of contract)

σ_S = swap rate volatility (fixed at the start of contract)

As there is a small difference between future (realised) volatility and swap rate (implied) volatility, then we can define $\sigma_F = \sigma_S + x$ where x is small.

$$\begin{aligned} \Rightarrow ((\sigma_S + x) - \sigma_S) &\approx Z ((\sigma_S + x)^2 - \sigma_S^2) \text{ for simplification we shall replace } \sigma_S \text{ with } \sigma \\ \Rightarrow x &\approx Z ((\sigma^2 + 2\sigma x + x^2) - \sigma^2) \\ \Rightarrow x &\approx Z (2\sigma x + x^2) \\ \Rightarrow 1 &\approx Z (2\sigma + x) \text{ and as } x \text{ is small} \\ \Rightarrow 1/2\sigma &\approx Z \end{aligned}$$

Hence $\text{Notional}_{\text{variance swap}} = \text{vega} / 2\sigma$ (as $\text{vega} = \text{Notional}_{\text{volatility swap}}$)

A.6: MODELLING VOLATILITY SURFACES

There are a variety of constraints on the edges of a volatility surface, and this section details some of the most important constraints from both a practical and theoretical point of view. We examine the considerations for very short-dated options (a few days or weeks), options at the wings of a volatility surface and very long-dated options.

IMPLIED VOL IS LESS USEFUL FOR NEAR-DATED OPTIONS

Options that only have a few days or a few weeks to expiry have a very small premium. For these low-value options, a relatively small change in price will equate to a relatively large change in implied volatility. This means the implied volatility bid-offer arbitrage channel is wider, and hence less useful. The bid-offer spread is more stable in cash terms for options of different maturity, so shorter-dated options should be priced more by premium rather than implied volatility.

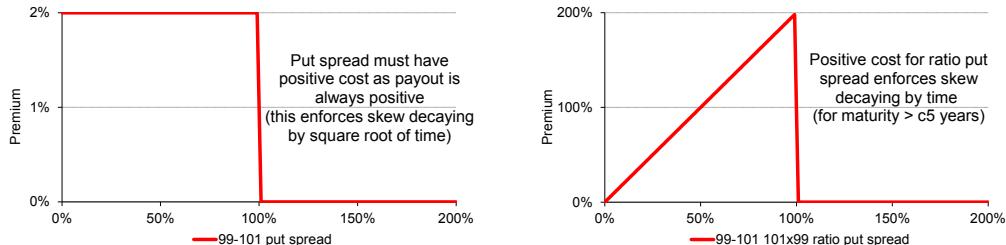
Need to price short-dated options with a premium after a large collapse in the market

If there has been a recent dip in the market, there is a higher than average probability that the markets could bounce back to their earlier levels. The offer of short-dated ATM options should not be priced at a lower level than the size of the decline. For example, if markets have dropped 5%, then a one-week ATM call option should not be offered for less than c5% due to the risk of a bounce-back.

SKEW SHOULD DECAY BY SQUARE ROOT OF TIME

The payout of a put spread (and call spread) is always positive; hence, it should always have a positive cost. If it was possible to enter into a long put (or call) spread position for no cost (or potentially earning a small premium), any rational investor would go long as large a position as possible and earn risk-free profits (as the position cannot suffer a loss). A put spread will have a negative cost if the premium earned by selling the lower strike put is more than the premium of the higher strike put bought. This condition puts a cap on how negative skew can be: for high (negative) skew, the implied of the low strike put could be so large the premium is too high (ie, more than the premium of higher strike puts). The same logic applies for call spreads, except this puts a cap on positive skew (ie, floor on negative skew). As skew is normally negative, the condition on put spreads (see

Figure 144 below on the left) is usually the most important. As time increases, it can be shown that the cap and floor for skew (defined as the gradient of first derivative of volatility with respect to strike, which is proportional to 90%-100% skew) decays by roughly the square root of time. This gives a mathematical basis for the ‘square root of time rule’ used by traders.

Figure 144. Put Spread**Ratio put spread**

Far-dated skew should decay by time for long maturities (c5 years)

It is possible to arrive at a stronger limit to the decay of skew by considering leveraged ratio put spreads (see chart above on the right). For any two strikes A and B (assume A < B), then the payout of going long A \times puts with strike B, and going short B \times puts with strike A creates a ratio put spread whose value cannot be less than zero. This is because the maximum payouts of both the long and short legs (puts have maximum payout with spot at zero) is A \times B. This can be seen in the

Figure 144 above on the right (showing a 99-101 101x99 ratio put spread). Looking at such leveraged ratio put spreads enforces skew decaying by time, not by the square root of time. However, for reasonable values of skew this condition only applies for long maturities (c5 years).

PROOF SKEW IS CAPPED AND FLOORED BY $\sqrt{\text{TIME}}$

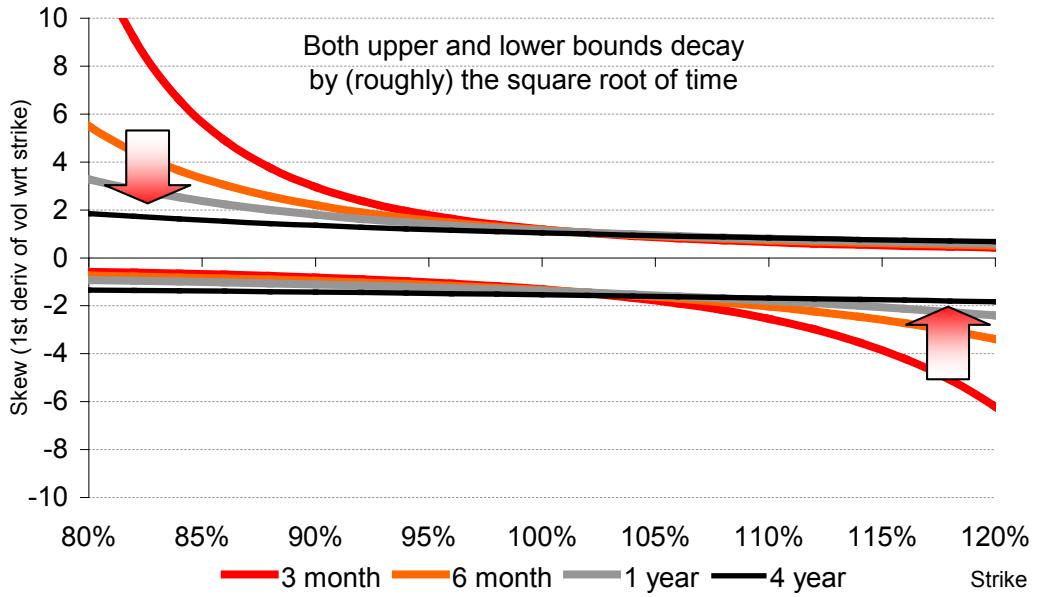
Enforcing positive values for put and call spreads is the same as the below two conditions:

- **Change in price of a call when strike increases has to be negative** (intuitively makes sense, as you have to pay more to exercise the higher strike call).
- **Change in price of a put when strike increases has to be positive** (intuitively makes sense, as you receive more value if the put is exercised against you).

These conditions are the same as saying the gradient of x (=Strike/Forward) is bound by:

$$\text{Lower bound} = -\sqrt{2\pi} \cdot e^{\frac{d_1^2}{2}} [1 - N(d_2)] \leq x \leq \sqrt{2\pi} \cdot e^{\frac{d_1^2}{2}} \cdot N(d_2) = \text{upper bound}$$

It can be shown that these bounds decay by (roughly) the square root of time. This is plotted below.

Figure 145. Upper and Lower Bound for Skew (given 25% volatility)

Proof of theoretical cap for skew works in practice

In the above example, for a volatility of 25% the mathematical lower bound for one-year skew (gradient of volatility with respect to strike) is -1.39. This is the same as saying that the maximum difference between 99% and 100% strike implied is 1.39% (ie, 90%-100% or 95%-105% skew is capped at 13.9%). This theoretical result can be checked by pricing one-year put options with Black-Scholes.

- Price 100% put with 25% implied = 9.95%
- Price 99% put with 26.39% implied = 9.95% (difference of implied of 1.39%)

In practice, skew is likely to be bounded well before mathematical limits

While a 90%-100% one-year skew of 13.9% is very high for skew, we note buying cheap put spreads will appear to be attractive long before the price is negative. Hence, in practice, traders are likely to sell skew long before it hits the mathematical bounds for arbitrage (as a put spread's price tends to zero as skew approaches the mathematical bound). However, as the mathematical bound decays by the square root of time, so too should the 'market bound'.

OTM IMPLIEDS AT WINGS HAVE TO BE FLAT IN LOG SPACE

While it is popular to plot implieds vs delta, it can be shown for many models⁴⁸ that implied volatility must be linear in log strike (ie, $\ln[K/F]$) as log strike goes to infinity. Hence a parameterisation of a volatility surface should, in theory, be parameterised in terms of log strike, not delta. In practice, however, as the time value of options for a very high strike is very small, modelling implieds against delta can be used as the bid-offer should eliminate any potential arbitrage.

⁴⁸ Eg, stochastic volatility plus jump models.

A.7: BLACK-SCHOLES FORMULA

The most popular method of valuing options is the Black-Scholes-Merton model. We show the key formulas involved in this calculation. The assumptions behind the model are also discussed.

BLACK-SCHOLES MAKES A NUMBER OF ASSUMPTIONS

It is often joked that Black-Scholes is the wrong model with the wrong assumptions that gets the right price. The simplicity of the model has ensured that it is still used despite the competition from other, more complicated models. The assumptions are below:

- Constant (known) volatility
- Constant interest rates
- No dividends (a constant dividend yield can, however, be incorporated into the interest rate)
- Zero borrow cost, zero trading cost and zero taxes
- Constant trading
- Stock price return is log normally distributed
- Can trade infinitely divisible amounts of securities
- No arbitrage

BLACK-SCHOLES PRICE OF EUROPEAN OPTIONS

$$\text{Call option price} = S \times N(d_1) - Ke^{-rT} N(d_2)$$

$$\text{Put option price} = -S \times N(-d_1) + Ke^{-rT} N(d_2)$$

where

$$d_1 = \frac{\ln\left(\frac{S}{K}\right) + (r + \frac{\sigma^2}{2})T}{\sigma\sqrt{T}}$$

$$d_2 = d_1 - \sigma\sqrt{T} = \frac{\ln\left(\frac{S}{K}\right) + (r - \frac{\sigma^2}{2})T}{\sigma\sqrt{T}}$$

S = Spot

K = Strike

r = risk free rate (– dividend yield)

σ = volatility

T = time (years)

A.8: GREEKS AND THEIR MEANING

Greeks is the name given to the (usually) Greek letters used to measure risk. We give the Black-Scholes formula for the key Greeks and describe which risk they measure.

VEGA IS NOT A GREEK LETTER

Although Vega is a Greek, it is not a Greek letter. It is instead the brightest star in the constellation Lyra. The main Greeks and their definition are in the table below.

Figure 146. Greeks and their definition

Greek	Symbol	Measures	Definition
Delta	δ or Δ	Equity exposure	Change in option price due to spot
Gamma	γ or Γ	Convexity of payout	Change in delta due to spot
Theta	θ or Θ	Time decay	Change in option price due to time passing
Vega	ν	Volatility exposure	Change in option price due to volatility
Rho	ω or Ω	Interest rate exposure	Change in option price due to interest rates
Volga	λ or Λ	Vol of vol exposure	Change in vega due to volatility
Vanna	ψ or Ψ	Skew	Change in vega due to spot OR change in delta due to volatility

The variables for the below formulae are identical to the earlier definitions in the previous section [A7 Black-Scholes Formula](#). In addition:

$$N'(z) \text{ is the normal density function, } \frac{e^{-\frac{z^2}{2}}}{\sqrt{2\pi}}$$

$N(z)$ is the cumulative normal distribution, ie, $N(0) = 0.5$.

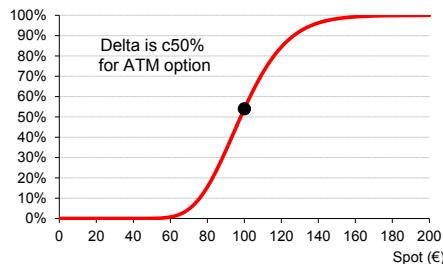
DELTA MEASURES EQUITY EXPOSURE

The most commonly examined Greek is delta, as it gives the equity sensitivity of the option (change of option price due to change in underlying price). Delta is normally quoted in percent. For calls it lies between 0% (no equity sensitivity) and 100% (trades like a stock). The delta of puts lies between -100% (trades like short stock) and 0%. If a call option has a delta of 50% and the underlying rises €1, the call option increases in value €0.50 (= €1 * 50%). Note the values of the call and put delta in the formula below give the equity sensitivity of a forward of the same maturity as the option expiry. The equity sensitivity to spot is slightly different. Please note that there is a (small) difference between the probability that an option expires ITM and delta.

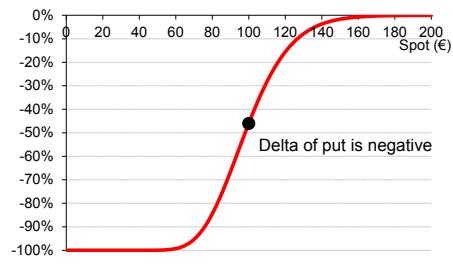
Call delta: $N(d_1)$

Put delta: $-N(-d_1) = N(d_1) - 1$

Figure 147. Delta (for Call)



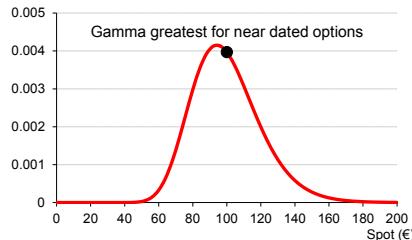
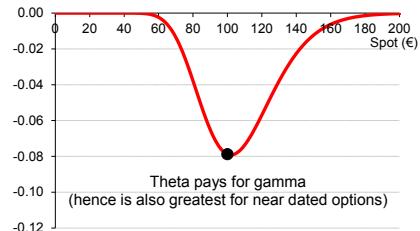
Delta (for Put)



GAMMA MEASURES CONVEXITY (AMOUNT EARN HEDGING)

Gamma measures the change in delta due to the change in underlying price. The higher the gamma, the more convex is the theoretical payout. Gamma should not be considered a measure of value (low or high gamma does not mean the option is expensive or cheap); implied volatility is the measure of an option's value. Options are most convex, and hence have the highest gamma, when they are ATM and also about to expire. This can be seen intuitively as the delta of an option on the day of expiry will change from c0% if spot is just below expiry to c100% if spot is just above expiry (a small change in spot causes a large change in delta; hence, the gamma is very high).

$$\text{Gamma: } -\frac{N'(d_1)}{S\sqrt{T}\sigma}$$

Figure 148. Gamma**Theta**

THETA MEASURES TIME DECAY (COST OF LONG GAMMA)

Theta is the change in the price of an option for a change in time to maturity; hence, it measures time decay. In order to find the daily impact of the passage of time, its value is normally divided by 252 (trading days in the year). If the second term in the formula below is ignored, the theta for calls and puts are identical and proportional to gamma. Theta can therefore be considered the cost of being long gamma.

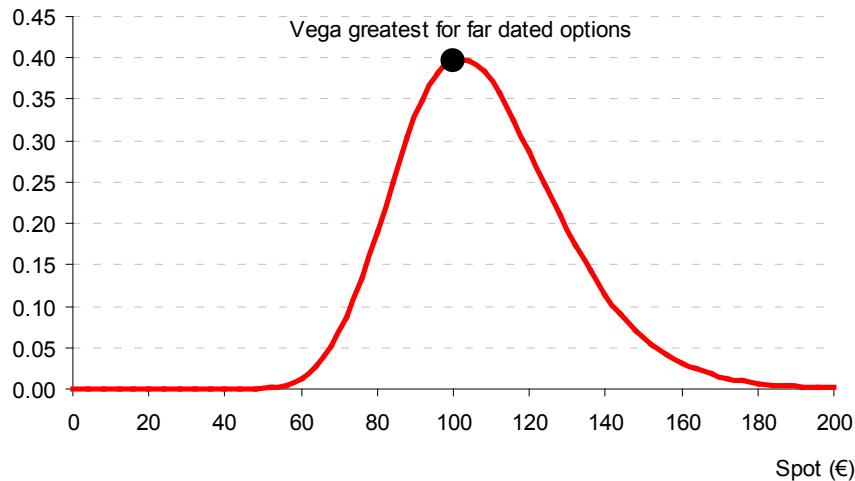
$$\text{Call theta: } -\frac{S\sigma \times N'(d_1)}{2\sqrt{T}} - rKe^{-rT}N(d_2)$$

$$\text{Put theta: } -\frac{S\sigma \times N'(d_1)}{2\sqrt{T}} + rKe^{-rT}N(-d_2)$$

VEGA MEASURES VOL EXPOSURE (AVERAGE OF GAMMAS)

Vega gives the sensitivity to volatility of the option price. Vega is normally divided by 100 to give the price change for a 1 volatility point (ie, 1%) move in implied volatility. Vega can be considered to be the average gamma (or non-linearity) over the life of the option. As vega has a \sqrt{T} in the formula power vega (vega divided by square root of time) is often used as a risk measure (to compensate for the fact that near dated implieds move more than far-dated implieds).

$$\text{Vega: } S\sqrt{T} \times N'(d_1)$$

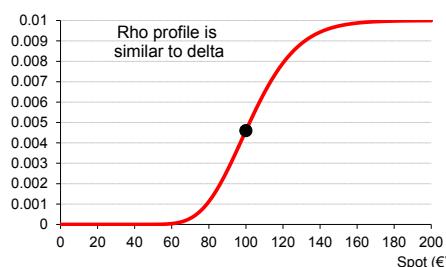
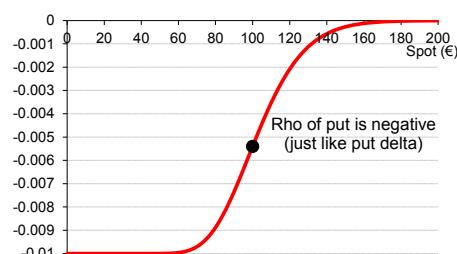
Figure 149. Vega

RHO MEASURES INTEREST RATE RISK (RELATIVELY SMALL)

Rho measures the change in the value of the option due to a move in the risk-free rate. The profile of rho vs spot is similar to delta, as the risk-free rate is more important for more equity-sensitive options (as these are the options where there is the most benefit in selling stock and replacing it with an option and putting the difference in value on deposit). Rho is normally divided by 10,000 to give the change in price for a 1bp move.

Call rho: $KTe^{-rT} N(d_2)$

Put rho: $- KTe^{-rT} N(-d_2)$

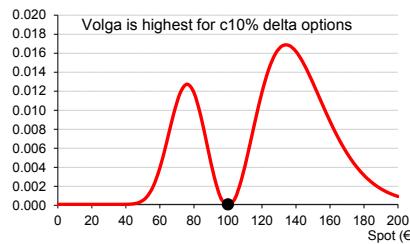
Figure 150. Rho (for call)**Rho (for put)**

VOLGA MEASURES VOLATILITY OF VOLATILITY EXPOSURE

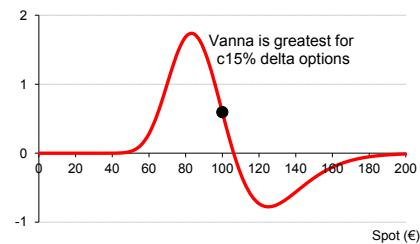
Volga is short for VOLatility GAMMA, and is the rate of change of vega due to a change in volatility. Volga (or Vomma/vega convexity) is highest for OTM options (approximately 10% delta), as these are the options where the probability of moving from OTM to ITM has the greatest effect on its value. For more detail on Volga, see the section [7.4 How to Measure Skew and Smile](#).

$$\text{Volga: } \frac{S\sqrt{T}d_1d_2N'(d_1)}{\sigma}$$

Figure 151. Volga



Vanna



VANNA MEASURES SKEW EXPOSURE

Vanna has two definitions as it measures the change in vega given a change in spot and the change in delta due to a change in volatility. The change in vega for a change in spot can be considered to measure the skew position, as this will lead to profits on a long skew trade if there is an increase in volatility as spot declines. The extreme values for vanna occur for c15 delta options, similar to volga's c10 delta peaks. For more detail on vanna, see the section [7.4 How to Measure Skew and Smile](#).

$$\text{Vanna: } \frac{-d_2N'(d_1)}{\sigma}$$

A.9: ADVANCED (PRACTICAL OR SHADOW) GREEKS

How a volatility surface changes over time can impact the profitability of a position. While the most important aspects have already been covered (and are relatively well understood by the market) there are ‘second order’ Greeks that are less well understood. Two of the most important are the impact of the passage of time on skew (volatility slide theta), and the impact of a movement in spot on OTM options (anchor delta). These Greeks are not mathematical Greeks, but are practical or ‘shadow’ Greeks.

SKEW INCREASE AS TIME PASSES CAUSES ‘VOL SLIDE THETA’

As an option approaches expiry, its maturity decreases. As near-dated skew is larger than far-dated skew, the skew of a fixed maturity option will increase as time passes. This can be seen by assuming that skew by maturity (eg, three-month or one-year) is constant (ie, relative time, the maturity equivalent of sticky moneyness or sticky delta). We also assume that three-month skew is larger than the value of one year skew. If we buy a low strike one year option (ie, we are long skew) then, assuming spot and ATM volatility stay constant, when the option becomes a three-month option its implied will have risen (as three-month skew is larger than one-year skew and ATM volatility has not changed). We define ‘volatility slide theta’ as the change in price of an option due to skew increasing with the passage of time⁴⁹.

VOL SLIDE THETA IS MOST IMPORTANT FOR NEAR EXPIRIES

Given that skew increases as maturity decreases, this change in skew will increase the value of long skew positions (as in the example) and decrease the value of short skew positions. The effect of ‘volatility slide theta’ is negligible for medium- to far-dated maturities, but increases in importance as options approach expiry. If a volatility surface model does not take into account ‘volatility slide theta’, then its impact will be seen when a trader re-marks the volatility surface.

VOL SLIDE THETA MEASURES IMPACT OF CONST SMILE RULE

The constant smile rule (CSR) details how forward starting options should be priced. The impact of this rule on valuations is given by the ‘volatility slide theta’ as they both assume a fixed maturity smile is constant. The impact of this assumption is more important for forward starting options than for vanilla options.

⁴⁹ While we concentrate on Black-Scholes implied volatilities, volatility slide theta also affects local volatility surfaces.

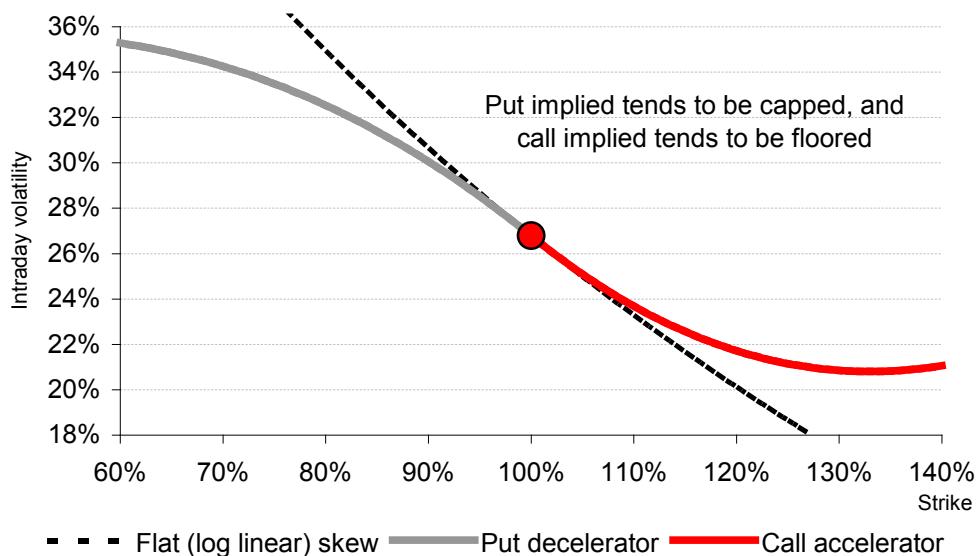
WHEN TRADERS CHANGE THEIR ‘ANCHOR’ THIS INTRODUCES A SECOND ORDER DELTA (‘ANCHOR DELTA’)

Volatility surfaces are normally modelled via a parameterisation. One of the more popular parameterisations is to set the ATMf volatility from a certain level of spot, or ‘anchor’, and then define the skew (slope). While this builds a reasonable volatility surface for near ATM options, the wings will normally need to be slightly adjusted. Normally a fixed skew for both downside puts and upside calls will cause upside calls to be too cheap (as volatility will be floored) and downside puts to be too expensive (as volatility should be capped at some level, even for very low strikes). As the ‘anchor’ is raised, the implied volatility of OTM options declines (assuming the wing parameters for the volatility surface stay the same). We call this effect ‘anchor delta’.

Implied volatility has to be floored, and capped, for values to be realistic

There are many different ways a volatility surface parameterisation can let traders correct the wings, but the effect is usually similar. We shall simply assume that the very OTM call implied volatility is lifted by a call accelerator, and very OTM put implied volatility is lowered by a put decelerator. This is necessary to prevent call implieds going too low (ie, below minimum realised volatility), or put implieds going too high (ie, above maximum realised volatility). The effect of these wing parameters is shown in Figure 152 below.

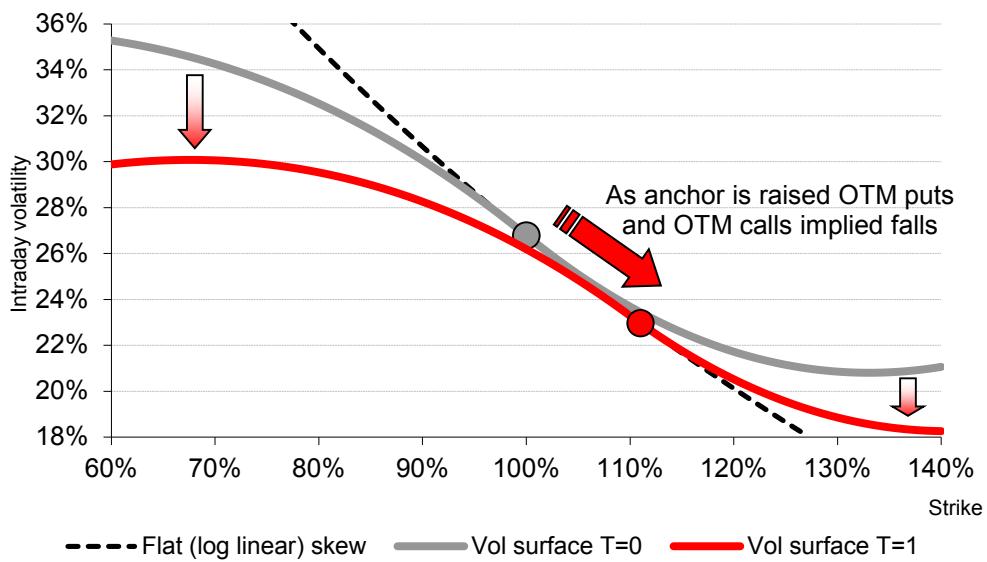
Figure 152. Skew with Put Decelerators and Call Accelerators



Traders tend to refresh a surface by only changing the key parameters

For liquid underlyings such as indices, a volatility surface is likely to be updated several times a day (especially if markets are moving significantly). Usually only the key parameters will be changed, and the less key parameters such as the wing parameters are changed less frequently. We shall assume that there will be many occasions where there is a movement in spot along the skew (ie, static strike for near ATM strikes). In these cases a trader is likely to change the anchor (and volatility at the anchor, which has moved along the skew), but leave the remaining skew and wing parameters (which are defined relative to the anchor) unchanged. In order to have the same implied volatility for OTM options after changing the anchor, the call accelerator should be increased and the put decelerator decreased. In practice this does not always happen, as wing parameters are typically changed less frequently. The effect of an increase in anchor along the (static strike) skew while leaving the wing parameters unchanged is shown below.

Figure 153. Moving Anchor 10% Higher Along the Skew



OTM options have a second order ‘anchor delta’

To simplify the example we shall assume the call wing parameter increases the implied volatility for strikes 110% and more, and the put wing parameter decreases the implied volatility for strikes 90% or lower. If spot rises 10%, the 120% call implied volatility will suffer when the anchor is re-marked 10% higher, because the call implied volatility is initially lifted by the call wing parameter (which no longer has an effect). OTM calls therefore have a negative ‘anchor delta’ as they lose value as anchor rises. Similarly, as anchor rises the effect of the put wing will increase, lowering the implied volatility of puts of strike 90% or less as anchor rises. So, under this scenario all options that are OTM have a negative ‘anchor delta’ that needs to be hedged.

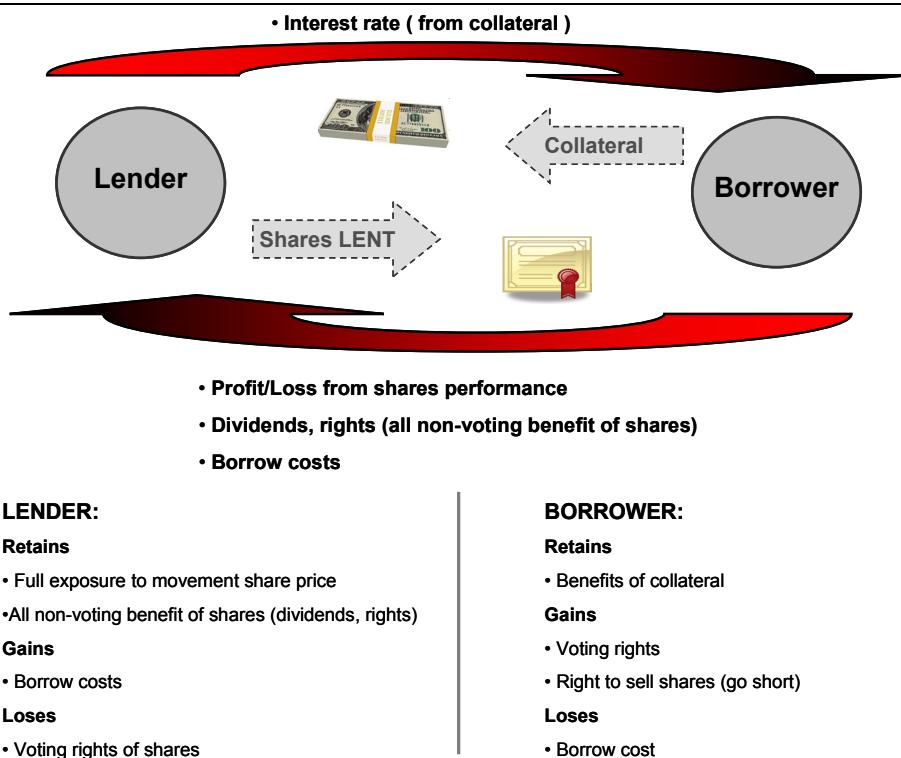
A.10: SHORTING STOCK BY BORROWING SHARES

The hedging of equity derivatives assumes you can short shares by borrowing them. We show the processes involved in this operation. The disadvantages, and benefits, to an investor who lends out shares are also explained.

NO COUNTERPARTY RISK WHEN YOU LEND SHARES

To short shares initially, the shares must first be borrowed. In order to remove counterparty risk, when an investor lends out shares he/she receives collateral (cash, stock, bonds, etc) for the same value. Both sides retain the beneficial ownership of both the lent security and the collateral, so any dividends, coupons, rights issues are passed between the two parties. If cash is used as collateral, the interest on the cash is returned. Should a decision have to be made, ie, to receive a dividend in cash or stock, the decision is made by the original owner of the security. The only exception is that the voting rights are lost, which is why lent securities are often called back before votes.

Figure 154. Borrowing Shares



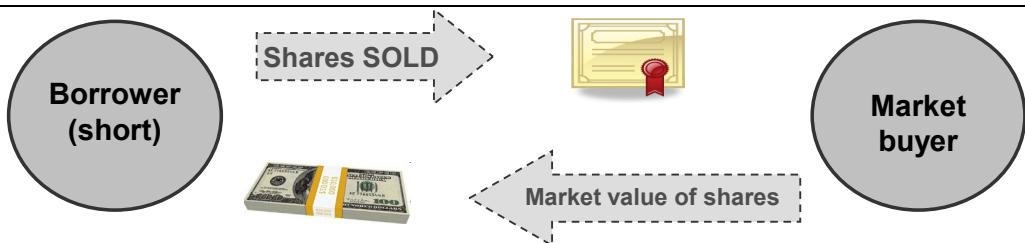
Mark to market P&L has to be settled in cash

To ensure there is no counterparty risk during the time the security is lent out, the collateral and lent security is marked to market and the difference settled for cash (while a wide range of securities can be used as initial collateral, only cash can be used for the change in value of the lent security).

SELLING STOCK YOU HAVE BORROWED GIVES SHORT POSITION

Once an investor has borrowed shares, these shares can be sold in the market. The proceeds from this sale can be used as the collateral given to the lender. Selling borrowed shares gives a short position, as profits are earned if the stock falls (as it can be bought back at a lower price than it was sold for, and then returned to the original owner).

Figure 155. Borrowing Shares



BORROWER:

Retains

- Benefits of collateral
- Responsibility to return shares to lender

Gains

- Profits from share price declines
- Cash value of shares (can be used as collateral)

Loses

- Borrow cost

EARN SHORT REBATE WHEN YOU SHORT

The investor who has shorted the shares receives interest on the collateral, but has to pass dividends and borrow cost to the original owner. The net of these cash flows is called the short rebate, as it is the profit (or loss for high dividend paying stocks) that occurs if there is no change in the price of the shorted security. Shorting shares is therefore still profitable if shares rise by less than the short rebate.

Short rebate = interest rate (normally central bank risk free rate) – dividends – borrow cost

Figure 156. Initial and Final Position of Lender, Borrower and Market Following Shorting of Shares

	Before lending			After lending			
	Lender	Borrow	Market	Lender	Borrow	Market	
Equity							
Holds shares	+					+	Share passes from lender to market (as receives collateral there is no credit risk)
P&L from share movement	+				+	-	Lender still has equity exposure, borrower is short to market
Economic benefit of shares (dividends, rights)	+				+	-	Borrower is short dividends and rights to lender
Voting rights	+					+	Lender loses voting rights (only difference to owning equity)
Borrow cost					+	-	Lender earns borrow cost (compensates loss of voting rights)
Cash (or collateral)							
Cash (or collateral) value of shares			+		+		Borrower passes value of share sale to lender as collateral (compensates for value of shares borrowed)
Interest on cash value of shares (or benefits of collateral, e.g. coupons)			+			+	Borrower gets interest on collateral (less dividends and borrow) as short rebate

A.11: SORTINO RATIO

If an underlying is distributed normally, standard deviation is the perfect measure of risk. For returns with a skewed distribution, such as with option trading or call overwriting, there is no one perfect measure of risk; hence, several measures of risk should be used. The Sortino is one of the most popular measures of skewed risk, as it only takes into account downside risk.

SORTINO RATIO IS MODIFICATION OF SHARPE RATIO

The Sharpe ratio measures the excess return, or amount of return (R) that is greater than the target rate of return (T). Often zero or risk-free rate is used as the target return. To take volatility of returns into account, this excess return is divided by the standard deviation. However, this takes into account both upside and downside risks. As investors are typically more focused on downside risks, the Sortino ratio modifies the Sharpe ratio calculation to only divide by the downside risk (DR). The downside risk is the square root of the target semivariance, which can be thought of as the amount of standard deviation due to returns less than the target return. The Sortino ratio therefore only penalises large downside moves, and is often thought of as a better measure of risk for skewed returns.

$$\text{Sortino} = \frac{R - T}{DR}$$

where

$$DR = \left[\int_{-\infty}^T (T - x) f(x) dx \right]^{0.5}$$

R = Return

T = Target return

A.12: CAPITAL STRUCTURE ARBITRAGE

When Credit Default Swaps were created in the late 1990's, they traded independently of the equity derivative market. The high levels of volatility and credit spreads during the bursting of the TMT bubble demonstrated the link between credit spreads, equity, and implied volatility. We examine four models that demonstrate this link (Merton model, jump diffusion, put vs CDS, and implied no-default volatility).

NORMALLY TRADE CREDIT VS EQUITY, NOT VOLATILITY

Capital structure arbitrage models can link the price of equity, credit and implied volatility. However, the relatively wide bid-offer spreads of equity derivatives mean trades are normally carried out between credit and equity (or between different subordinations of credit and preferred shares vs ordinary shares). The typical trade is for an investor to go long the security that is highest in the capital structure, for example, a corporate bond (or potentially a convertible bond), and short a security that is lower in the capital structure, for example, equity. Reverse trades are possible, for example, owning a subordinated higher yielding bond and shorting a senior lower yielding bond (and earning the positive carry as long as bankruptcy does not occur). Only for very wide credit spreads and high implied volatility is there a sufficiently attractive opportunity to carry out an arbitrage between credit and implied volatility. We shall concentrate on trading credit vs equity, as this is the most common type of trade.

Credit spread is only partly due to default risk

The OAS (Option Adjusted Spread) of a bond over the risk-free rate can be divided into three categories. There is the expected loss from default; however, there is also a portion due to general market risk premium and additionally a liquidity cost. Tax effects can also have an effect on the corporate bond market. Unless a capital structure arbitrage model takes into account the fact that not all of a bond's credit spread is due to the risk of default, the model is likely to fail. The fact that credit spreads are higher than they should be if bankruptcy risk was the sole risk of a bond was often a reason why long credit short equity trades have historically been more popular than the reverse (in addition to the preference to being long the security that is highest in the capital structure in order to reduce losses in bankruptcy).

CDS usually better than bonds for credit leg, as they are unfunded and easier to short

Using CDS rather than corporate bonds can reduce many of the discrepancies in spread that a corporate bond suffers and narrow the difference between the estimated credit spread and the actual credit spread. We note that CDS are an unfunded trade (ie, leveraged), whereas corporate bonds are a funded trade (have to fund the purchase of the bond) that has many advantages when there is a funding squeeze (as occurred during the credit crunch). CDS also allow a short position to be easily taken, as borrow for corporate bonds is not always available, is usually expensive and can be recalled at any time. While borrow for bonds was c50bp before the credit crunch it soared to c5% following the crisis.

Credit derivatives do not have established rules for equity events

While credit derivatives have significant language against credit events, they have no language for equity events, such as special dividends or rights issues. Even for events such as takeovers and mergers, where there might be relevant documentation, credit derivatives are likely to behave differently than equity (and equity derivatives).

CREDIT CAN LEAD EQUITY MARKET AND VICE VERSA

We note that there are occasions when corporate bond prices lag a movement in equity prices, simply as traders have not always updated levels (but this price would be updated should an investor request a firm price). CDS prices suffer less from this effect, and we note for many large companies the corporate bond market is driven by the CDS market and not vice versa (the tail wags the dog). Although intuitively the equity market should be more likely to lead the CDS market than the reverse (due to high frequency traders and the greater number of market participants), when the CDS market is compared to the equity market on average neither consistently leads the other. Even if the CDS and equity on average react equally as quickly to new news, there are still occasions when credit leads equity and vice versa. Capital structure arbitrage could therefore be used on those occasions when one market has a delayed reaction to new news compared to the other.

GREATEST OPPORTUNITY ON BBB OR BB RATED BONDS

In order for capital structure arbitrage to work, there needs to be a strong correlation between credit and equity. This is normally found in companies that are rated BBB or BB. The credit spread for companies with ratings of A or above is normally more correlated to the general credit supply and interest rates than the equity price. For very speculative companies (rated B or below), the performance of their debt and equity is normally very name-specific, and often determined by the probability of takeover or default.

Capital structure arbitrage works best when companies don't default

Capital structure arbitrage is a bet on the convergence of equity and credit markets. It has the best result when a company in financial distress recovers, and the different securities it has issued converge. Should the company enter bankruptcy, the returns are less impressive. The risk to the trade is that the company becomes more distressed, and as the likelihood of bankruptcy increases the equity and credit markets cease to function properly. This could result in a further divergence or perhaps closure of one of the markets, potentially forcing a liquidation of the convergence strategy.

FUNDAMENTALS CAN DWARF STATISTICAL RELATIONSHIPS

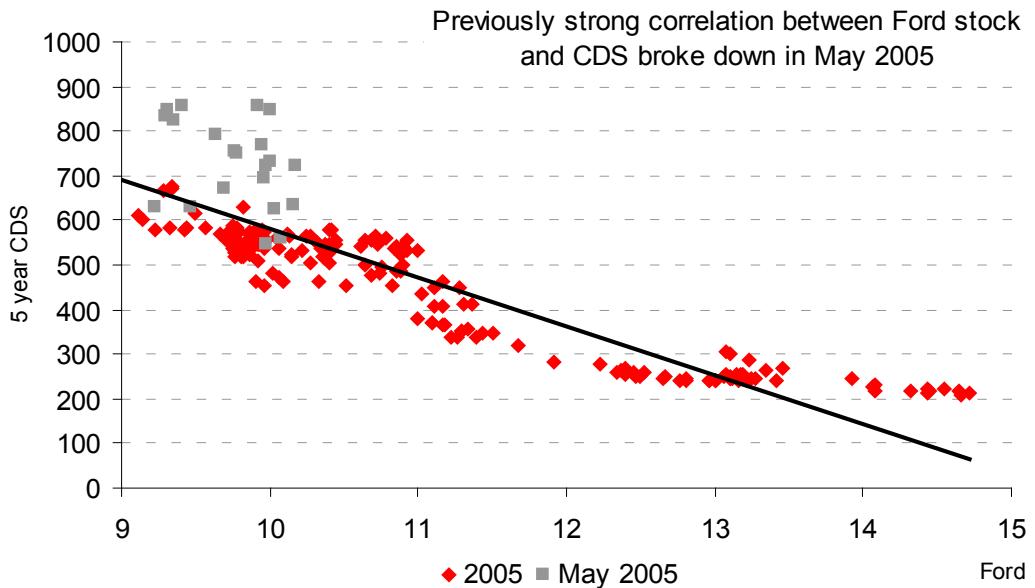
Capital structure arbitrage assumes equity and credit markets move in parallel. However, there are many events that are bullish for one class of investors and bearish for another. This normally happens when the leverage of a company changes suddenly. Takeovers and rights issues are the two main events that can quickly change leverage. Special dividends, share buybacks and a general reduction of leverage normally have a smaller, more gradual effect.

Rights issue. A rights issue will always reduce leverage, and is effectively a transfer of value from equity holders to debt holders (as the company is less risky, and earnings are now divided amongst a larger number of shares).

Takeover bid (which increases leverage). When a company is taken over, unless the acquisition is solely for equity, a portion of the acquisition will have to be financed with cash or debt (particularly during a leveraged buyout). In this case, the leverage of the acquiring company will increase, causing an increase in credit spreads and a reduction in the value of debt. Conversely, the equity price of the acquiring company is more stable. For the acquired company, the equity price should jump close to the level of the bid and, depending on the structure of the offer, the debt could fall (we note that if the acquired company is already in distress the value of debt can rise; for example, when Household was acquired by HSBC).

GM EQUITY SOARED A DAY BEFORE CREDIT SANK

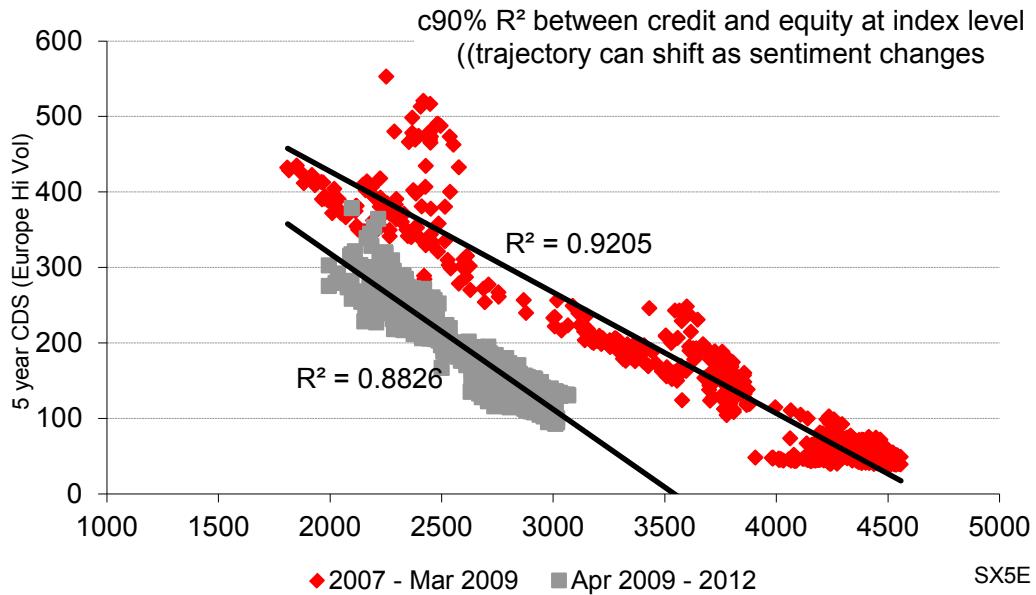
On May 4, 2005, Kirk Kerkorian announced the intention to increase his (previously unknown) stake in GM, causing the troubled company's share price to soar 18% intraday (7.3% close to close). The following day, S&P downgraded GM and Ford to 'junk', causing a collapse in the credit market and a 122bp CDS rise in two days. As many capital structure arbitrage investors had a long credit short equity position, both legs were loss making and large losses were suffered.

Figure 157. Equity vs Credit Spread (5-Year CDS) for Ford

CORRELATION BETWEEN CREDIT AND EQUITY IS LOW AT STOCK LEVEL

For many companies the correlation between equity and credit is not particularly strong, with a typical correlation between 5% and 15%. Hence it is necessary for a capital structure arbitrage investor to have many different trades on simultaneously. The correlation of a portfolio of bonds and equities is far higher (c90%).

Figure 158. SX5E vs 5-Year CDS (European HiVol)



MODELLING THE LINK BETWEEN CREDIT AND IMPLIED VOL

While there are many models that show the link between the equity, equity volatility and debt of a company, we shall restrict ourselves to four of the most popular.

- **Merton model.** The Merton model uses the same model as Black-Scholes, but applies it to a firm. If a firm is assumed to have only one maturity of debt, then the equity of the company can be considered to be a European call option on the value of the enterprise (value of enterprise = value of debt + value of equity) whose strike is the face value of debt. This model shows how the volatility of equity rises as leverage rises. The Merton model also shows that an increase in volatility of the enterprise increases the value of equity (as equity is effectively long a call on the value of the enterprise), and decreases the value of the debt (as debt is effectively short a put on the enterprise, as they suffer the downside should the firm enter bankruptcy but the upside is capped).
- **Jump diffusion.** A jump diffusion model assumes there are two parts to the volatility of a stock. There is the diffusive (no-default) volatility, which is the volatility of the equity without any bankruptcy risk, and a separate volatility due to the risk of a jump to bankruptcy. The total volatility is the sum of these two parts. While the diffusive volatility is constant, the effect on volatility due to the jump to bankruptcy is greater for options of low strike than high strike causing 'credit induced skew'. This means that as

the credit spread of a company rises, this increases the likelihood of a jump to bankruptcy and increases the skew. A jump diffusion model therefore shows a link between credit spread and implied volatility.

- **Put vs CDS.** As the share price of a company in default tends to trade close to zero, a put can be assumed to pay out its strike in the event of default. This payout can be compared to the values of a company's CDS, or its debt market (as the probability of a default can be estimated from both). As a put can also have a positive value even if a company does not default, the value of a CDS gives a floor to the value of puts. As $1 \times N$ put spreads can be constructed to never have a negative payout, various caps to the value of puts can be calculated by ensuring they have a cost. The combination of the CDS price floor, and put price cap, gives a channel for implieds to trade without any arbitrage between CDS and put options.
- **No-default implied volatility.** Using the above put vs CDS methodology, the value of a put price due to the payout in default can be estimated. If this value is taken away from the put price, the remaining price can be used to calculate a no-default implied volatility (or implied diffusive volatility). The skew and term structure of implied no-default implied volatilities are flatter than Black-Scholes implied volatility, which allows an easier comparison and potential for identifying opportunities.

(1) MERTON MODEL

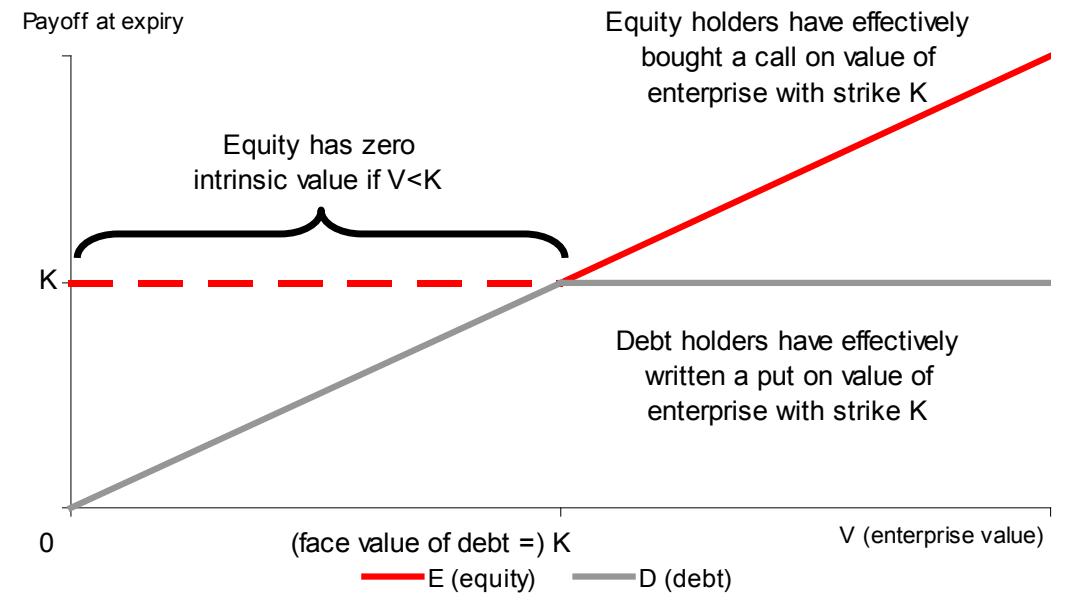
The Merton model assumes that a company has an enterprise value (V) whose debt (D) consists of only one zero coupon bond whose value at maturity is K . These assumptions are made in order to avoid the possibility of a default before maturity (which would be possible if there was more than one maturity of debt, or a coupon had to be paid). The company has one class of equity (E) that does not pay a dividend. The value of equity (E) and debt (D) at maturity is given below.

$$\text{Enterprise value} = V = E + D$$

$$\text{Equity} = \text{Max}(V - K, 0) = \text{call on } V \text{ with strike } K$$

$$\text{Debt} = \text{Min}(V, K) = K - \text{Max}(K - V, 0) = \text{Face value of debt } K - \text{put on } V \text{ with strike } K$$

Figure 159. Graph of Value of Enterprise, Equity and Debt



Enterprise value of a firm at maturity has to be at least K or it will enter bankruptcy

Before the maturity of the debt, the enterprise has obligations to both the equity and debt holders. At the maturity of the debt, if the value of the enterprise is equal to or above K , the enterprise will pay off the debt K and the remaining value of the firm is solely owned by the equity holders. If the value of the enterprise is below K then the firm enters bankruptcy. In the event of bankruptcy, the equity holders get nothing and the debt holders get the whole value of the enterprise V (which is less than K).

Equity is long a call on the value of a firm

If the value of the enterprise V is below the face value of debt K at maturity the equity holders receive nothing. However, if V is greater than K , the equity holders receive $V - K$. The equity holders therefore receive a payout equal to a call option on V of strike K .

Debt is short a put on value of firm

The maximum payout for owners of debt is the face value of debt. This maximum payout is reduced by the amount the value of the enterprise is below the face value of debt at maturity. Debt is therefore equal to the face value of debt less the value of a put on V of strike K .

DEBT HAS A DELTA THAT CAN BE ARBITRAGED VS EQUITY

As the value of the short put has a delta, debt has a delta. It is therefore possible to go long debt and short equity (at the calculated delta using the Merton model) as part of a capital structure arbitrage trade.

If enterprise value is unchanged, then if value of equity rises, value of credit falls

As enterprise value is equal to the sum of equity and debt, if enterprise value is kept constant then for equity to rise the value of debt must fall. An example would be if a company attempts to move into higher-risk activity, lifting its volatility. As equity holders are long a call on the value of the company they benefit from the additional time value. However, as debt holders are short a put they suffer should a firm move into higher-risk activities.

Merton model assumes too high a recovery rate

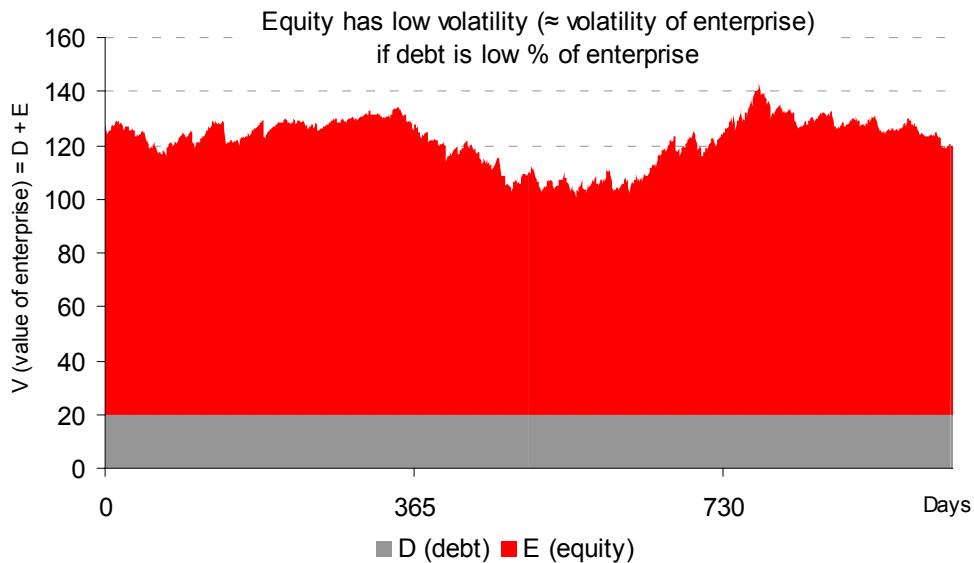
Using the vanilla Merton model gives unrealistic results with credit spreads that are too tight. This is because the recovery rate (of V/K) is too high. However, using more advanced models (eg, stochastic barrier to take into account the default point is unknown), the model can be calibrated to market data.

MERTON MODEL EXPLAINS EQUITY SKEW

The volatility of an enterprise should be based on the markets in which it operates, interest rates and other macro risks. It should, however, be independent of how it is funded. The proportion of debt to equity therefore should not change the volatility of the enterprise V ; however, it does change the volatility of the equity E . It can be shown that the volatility of equity is approximately equal to the volatility of the enterprise multiplied by the leverage (V/E). Should the value of equity fall, the leverage will rise, lifting the implied volatility. This explains skew: the fact that options of lower strike have an implied volatility greater than options of high strike.

$$\sigma_E \approx \sigma_V \times V / E (= \sigma_V \times \text{leverage})$$

Figure 160. Value of Enterprise and Equity with Low Debt

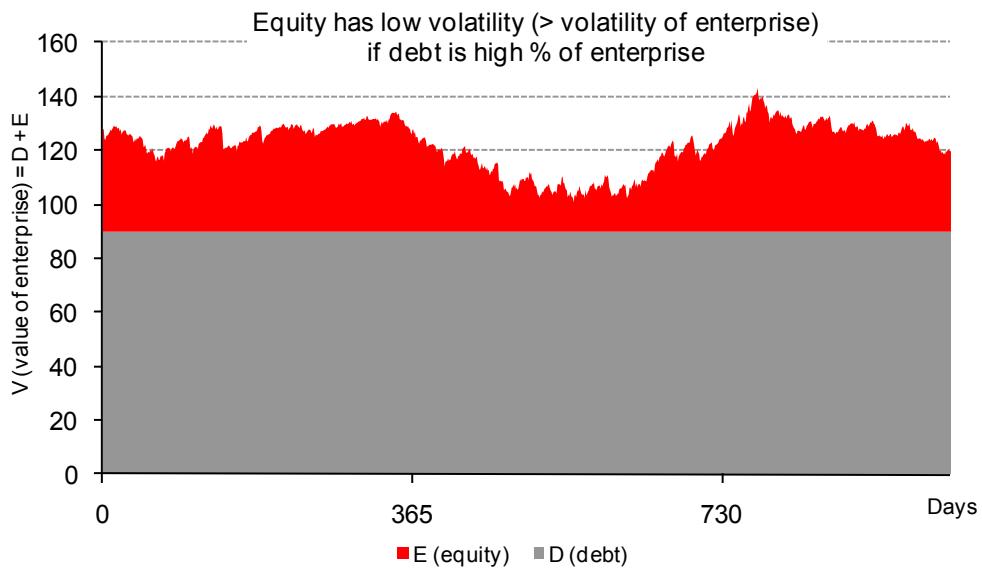


Firms with a small amount of debt have equity volatility roughly equal to firm volatility

If a firm has a very small (or zero) amount of debt, then the value of equity and the enterprise are very similar. In this case, the volatility of the equity and enterprise should be very similar (see Figure 160 above).

Firms with high value of debt to equity have very high equity volatility

For enterprises with very high levels of debt, a relatively small percentage change in the value of the enterprise V represents a relatively large percentage change in the value of equity. In these cases equity volatility will be substantially higher than the enterprise volatility (see Figure 161).

Figure 161. Value of Enterprise and Equity with High Debt

Proof equity volatility is proportional to leverage

The mathematical relationship between the volatility of the enterprise and volatility of equity is given below. The $N(d_1)$ term adjusts for the delta of the equity.

$$\sigma_E = N(d_1) \times \sigma_V \times V / E$$

If we assume the enterprise is not distressed and the equity is ITM, then $N(d_1)$ or delta of the equity should be very close to 1 (it is usually c90%). Therefore, the equation can be simplified so the volatility of equity is proportional to leverage (V / E).

$$\sigma_E \approx \sigma_V \times \text{leverage}$$

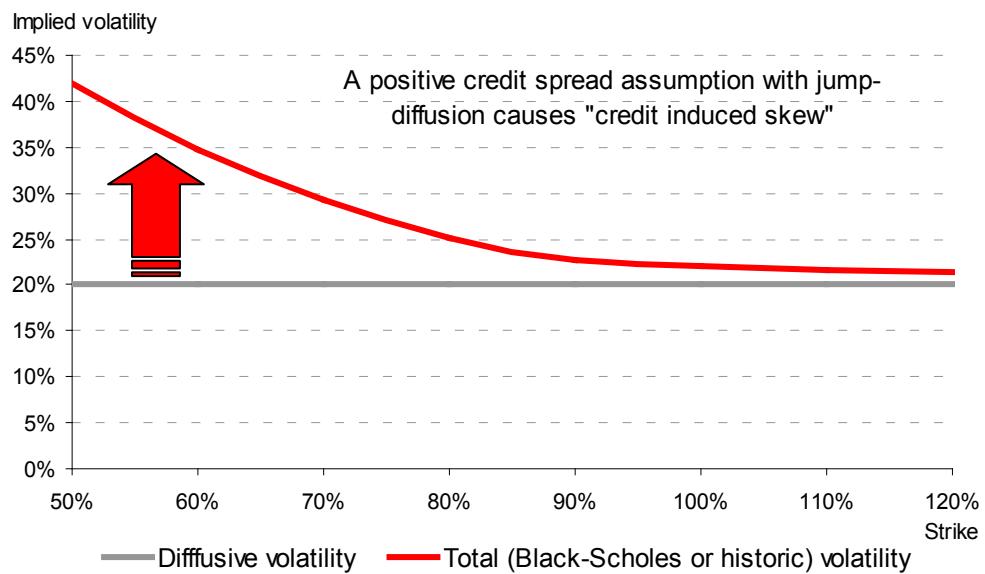
(2) JUMP DIFFUSION

A jump diffusion model separates the movement of equities into two components. There is the diffusive volatility, which is due to random log-normally distributed returns occurring continuously over time. In addition, there are discrete jumps the likelihood of which is given by a credit spread. The total of the two processes is the total volatility of the underlying. It is this total volatility that should be compared to historic volatility or Black-Scholes volatility.

Default risk explained by credit spread

For simplicity, we shall assume that in a jump diffusion model the jumps are to a zero stock price as the firm enters bankruptcy, but results are similar for other assumptions. The credit spread determines the risk of entering bankruptcy. If a zero credit spread is used, the company will never default. The probability of default increases as the credit spread increases (approximately linearly).

Figure 162. Credit-Induced Skew (with 100bp credit spread)



JUMP DIFFUSION CAUSES CREDIT-INDUCED SKEW

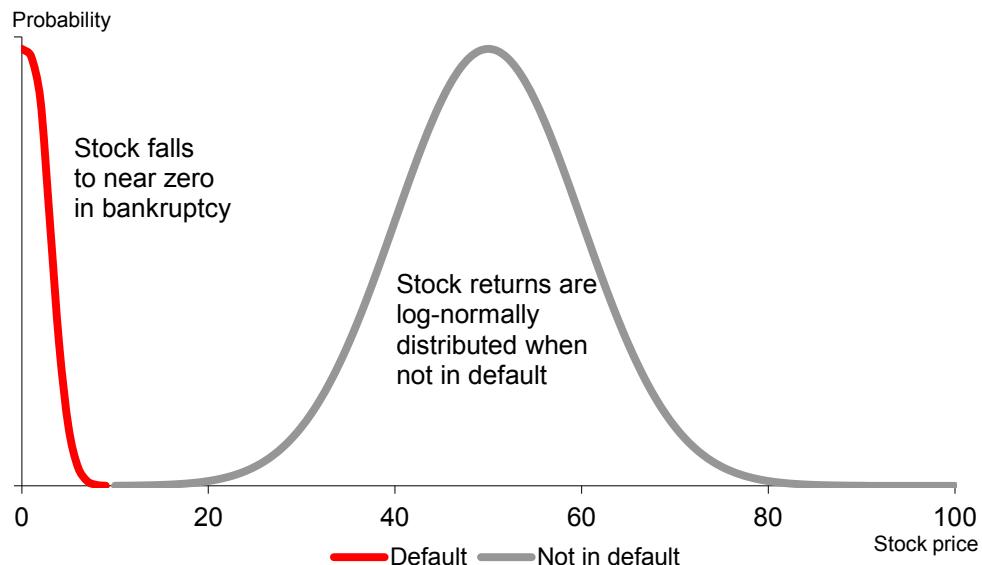
To show how credit spread (or bankruptcy) causes credit-induced skew, we shall price options of different strike with jump diffusion, keeping the diffusive volatility and credit spread constant. Using the price of the option, we shall then calculate the Black-Scholes implied volatility. The Black-Scholes implied volatility is higher for lower strikes than higher strikes, causing skew.

Credit-induced skew is caused by ‘option on bankruptcy’

The time value of an option will be divided between the time value due to diffusive volatility and the time value due to the jump to zero in bankruptcy. High strike options will be relatively unaffected by the jump to bankruptcy, and the Black-Scholes implied volatility will roughly be equal to the diffusive volatility. However, the value of a jump to a zero stock price will be relatively large for low strike put options (which, due to put call parity, is the implied for all options). The difference between the Black-Scholes implied and diffusive volatility could be considered to be the value due to the ‘option on bankruptcy’.

(3) PUT VS CDS

The probability distribution of a stock price can be decomposed into the probability of a jump close to zero due to credit events or bankruptcy, and the log-normal probability distribution that occurs when a company is not in default. While the value of a put option will be based on the whole probability distribution, the value of a CDS will be driven solely by the probability distribution due to default. The (bi-modal) probability distribution of a stock price due to default, and when not in default, is shown below.

Figure 163. Stock Price Probability Distribution

Puts can be used instead of CDS (as puts pay out strike price in event of bankruptcy)

When a stock defaults, the share price tends to fall to near zero. The recovery rate of equities can only be above zero if debt recovers 100% of face value, and most investors price in a c40% recovery rate for debt. A put can therefore be assumed to pay out the maximum level (ie, the strike) in the event of default. Puts can therefore be used as a substitute for a CDS. The number of puts needed is shown below.

$$\text{Value of puts in default} = \text{Strike} \times \text{Number of Puts}$$

$$\text{Value of CDS in default} = (100\% - \text{Recovery Rate}) \times \text{Notional}$$

In order to substitute value of puts in default has to equal value of CDS in default.

$$\text{Strike} \times \text{Number of Puts} = (100\% - \text{Recovery Rate}) \times \text{Notional}$$

$$\text{Number of Puts} = (100\% - \text{Recovery Rate}) \times \text{Notional} / \text{Strike}$$

CDS PRICES PROVIDE FLOOR FOR PUTS

As a put can have a positive value even if a stock is not in default, a CDS must be cheaper than the equivalent number of puts (equivalent number of puts chosen to have same payout in event of default, ie, using the formula above). If a put is cheaper than a CDS, an investor can initiate a long put-short CDS position and profit from the difference. This was a popular capital structure arbitrage trade in the 2000-03 bear market, as not all volatility traders were as focused on the CDS market as they are now, and arbitrage was possible.

CDS in default must have greater return than put in default (without arbitrage)

As a CDS has a lower price for an identical payout in default, a CDS must have a higher return in default than a put. Given this relationship, it is possible to find the floor for the value of a put. This assumes the price of a CDS is ‘up front’ ie, full cost paid at inception of the contract rather than quarterly.

$$\text{Puts return in default} = \text{Strike} / \text{Put Price}$$

$$\text{CDS return in default} = (100\% - \text{Recovery Rate}) / \text{CDS Price}$$

As CDS return in default must be greater than or equal to put return in default.

$$\Rightarrow (100\% - \text{Recovery Rate}) / \text{CDS Price} \geq \text{Strike} / \text{Put Price}$$

$$\Rightarrow \text{Put Price} \geq \text{Strike} \times \text{CDS Price} / (100\% - \text{Recovery Rate})$$

PUT VS CDS IS A POPULAR TRADE

As the prices of the put and CDS are known, the implied recovery rate can be backed out using the below formula. If an investor’s estimate of recovery value differs significantly from this level, a put vs CDS trade can be initiated. For a low (or zero) recovery rate, the CDS price is too high and a short CDS long put position should be initiated. Conversely, if the recovery rate is too high, a CDS price is too cheap and the reverse (long CDS, short put) trade should be initiated.

$$\text{Put Price} = \text{Strike} \times \text{CDS Price} / (100\% - \text{Implied Recovery Rate})$$

RATIO PUT SPREADS CAP VALUE OF PUTS

CDSs provide a floor to the price of a put. It is also possible to cap the price of a put by considering ratio put spreads. For example, if we have the price for the ATM put, this means we know that the value of a 50% strike put cannot be greater than half the ATM put price. If not, we could purchase an ATM-50% 1x2 put spread (whose payout is always positive) and earn a premium for free. This argument can be used for all strikes K and all 1xN put spreads, and is shown below:

$$N \times \text{put of strike } \frac{K}{N} \leq \text{put of strike } K$$

ARBITRAGE MOST LIKELY WITH LOW STRIKE AND LONG MATURITY

The combination of CDS prices providing a floor, and put prices of higher strikes providing a cap, gives a corridor for the values of puts. The width of this corridor is narrowest for low strike long maturity options, as these options have the greatest percentage of their value associated with default risk. As for all capital structure arbitrage strategies, companies with high credit spreads are more likely to have attractive opportunities and arbitrage is potentially possible for near-dated options.

(4) NO-DEFAULT IMPLIED VOLATILITY

The volatility of a stock price can be decomposed into the volatility due to credit events or bankruptcy and the volatility that occurs when a company is not in default. This is similar to the volatility due to jumps and the diffusive volatility of a jump diffusion model. As the value of a put option due to the probability of default can be calculated from the CDS or credit market, if this value was taken away from put prices this would be the ‘no-default put price’ (ie, the value the put would have if a company had no credit risk). The implied volatility calculated using this ‘no-default put price’ would be the ‘no-default implied volatility’. No-default implied volatilities are less than the vanilla implied volatility, as vanilla implied volatilities include credit risk).

No-default implied volatilities have lower skew and term structure

While we derive the no-default implied volatility from put options, due to put call parity the implied volatility of calls and puts is identical for European options. As the value of a put associated with a jump to default is highest for low-strike and/or long-dated options, no-default implied volatilities should have a lower skew and term structure than vanilla Black-Scholes implied volatilities. A no-default implied volatility surface should therefore be flatter than the standard implied volatility surface and, hence, could be used to identify potential trading opportunities.

INDEX

A

"abnormal" market moves, 7
absolute dimensions, 215
absolute time, 228
additive variance rule, 103, 104, 105, 106, 107
allowable loss, 143
alpha, 23, 24
Altiplano, 164
American options, 2, 11, 12, 13, 37, 230
anchor delta, 217n.38, 262, 263–65
annualized volatility, 42
annuities, 81
arbitrage, 169
ATM forward, 144, 145
ATM options:
 and "abnormal" market moves, 7;
 definition, 3;
 delta of, 7;
 leverage of, 7;
 liquidity of, 110;
 calculating premium, 41, 44;
 pricing with additive variance rule, 105;
 return on, 7;
 time value of, 3, 6, 110;
 vega, 203
ATM straddles, 38
ATM volatility, 222
at-the-money options. See: ATM options
autocallables, 135
autocorrelation, 235, 236
automatic exercise of options, 11, 13
AXA, pinning of, 43

B

back-testing, 76, 80
bankruptcy, 189, 201
barrier options:
 call, 134, 135;
 and correlation, 142;
 delta of, 134, 136;
 down, 134, 135;
 as European options, 134;
 hedging of, 134;
 key variables for, 134;
 naming conventions for, 135;
 payoff, 135;
 payout of, 134, 135, 136, 138;
 pricing of, 138;
 put, 134, 135, 136, 137, 138;
 rebates with, 139;
 types of, 134;
 up, 134, 135
basis, 58
basis risk, 58
basket options, 14, 68, 159, 164, 165, 167, 175–76
bear market, protection strategies for, 27, 28
bearish strategies, 3, 4, 5, 28, 31
beta, 155, 157
beta weighting, 157
bias, 239, 241
bid-offer arbitrage, 38
bid-offer spreads:
 and barrier options, 134;
 and call overwriting, 18, 22;
 and capital structure arbitrage, 270;
 and dispersion trading, 165;

- bid-offer spreads (continued):
and forward starting options, 100, 101, 102, 108;
and implied jumps, 180;
and relative value trading, 156, 158;
single stocks vs. index, 25, 158;
and volatility trading, 41
- Black, Fischer, 2
- Black-Scholes. See: Black-Scholes-Merton
- Black-Scholes-Merton formula:
assumptions behind, 2, 89, 92, 255;
and calculating ATM premium, 44;
and call option pricing, 255;
and composite options, 149;
and delta, 214;
and forward starting option pricing, 103;
and gamma, 67;
implied volatility, 220, 221, 230–31, 275, 281;
implied volatility surface, 222;
invention of, 2;
and pricing of European options, 255–56;
and put option pricing, 255;
and skew, 218, 219, 231;
and square root of time, 198;
and stock prices, 178;
and volatility, 49, 89, 218
- "black swan," 60
- bonus shares, 14
- borrow cost, 38, 154
- borrowing shares, 266–68
- bounce back, 17
- box spreads, 123
- breakevens, dispersion trading, 173, 174
- bull market, protection strategies for, 27, 28
- bullish strategies, 3, 4, 5, 26, 31
- butterflies, 34, 202
- buy-write, 15–16
- BXM index, 18, 19
- C**
- calculation agent, 58
- calendar collars, 29
- calendar put spreads, 30
- calendar trades, 34, 193, 198, 199
- call options:
American, 2, 12;
and anchor delta, 265;
ATM, 6, 12, 41, 44, 69, 139, 251;
and barrier options, 134, 135, 137;
as bearish or bullish strategy, 4, 31;
and borrow cost, 38;
breakeven, 70;
and call overwriting, 15, 16, 19, 21;
and collars, 26, 29, 30;
and contingent premium options, 148;
convexity of, 9;
cost of, 31, 32, 135, 144;
definition, 2;
delta of, 6, 7, 8, 9, 214, 223, 224, 258;
and delta hedging, 36, 37, 39, 97;
digital, 8;
and dispersion, 170;
downside of using, 9;
early exercise of, 11, 12, 13;
effect of volatility on, 3;
European, 2, 12, 13, 36, 37, 274;
and gamma scalping, 40;
high cost of, 31;
intrinsic value, 3;
ITM, 12;
and kurtosis, 205;
liquidity of, 7, 111;
long, 4, 5, 9, 12, 13, 31, 33, 40, 69, 84, 123, 156, 220;
long equity exposure of, 2, 3;
and look-back options, 146, 147;
maximum loss with, 31;
OTM, 15, 23, 24, 26, 31, 54, 63, 97, 204, 205, 211, 220, 223, 263, 264, 265;
and outperformance options, 143, 144;
payoff of, 8;

- call options (continued):
 payout, 12, 70, 135, 140, 276;
 premium of, 69, 70, 195, 220, 223;
 pricing of, 3, 44, 251, 252, 255;
 and probability of being ITM, 8, 9;
 profit conditions for, 2;
 as protection strategy, 23, 24;
 and put-call parity, 36, 37, 69, 284;
 and put underwriting, 10;
 and quanto options, 151;
 and relative value trading, 156;
 replacing stock with, 9;
 rho for, 260;
 and risk reversals, 31, 84, 85, 220,
 225n.40;
 and shadow delta, 101;
 short, 4, 5, 10, 16, 17, 19, 29, 30, 156,
 220, 223, 224, 225;
 and skew, 204, 208, 211, 223, 227;
 and sticky delta volatility, 214;
 and straddles and strangles, 33;
 strike of, 24, 31, 146, 195;
 and synthetics, 123;
 theta for, 259;
 upside, 31, 32, 54, 85, 111, 204, 263;
 on value of enterprise, 276, 277;
 vanilla, 135, 151, 156;
 and variance swaps, 54, 58, 63, 69, 70;
 vega of, 84n.9;
 and volatility trading, 36, 70;
 and VXX, 51, 109;
 when to use, 9;
 and worst-of/best-of options, 140, 141
- call overwriting:
 in bull markets, 19;
 definition, 10;
 dependence on market conditions, 19,
 20;
 to enforce disciplined investing, 16;
 enhanced, 22;
 far-dated options, 20;
 and implied volatility imbalances, 198;
 on indices, 18, 21, 22;
- as method of reducing volatility, 21;
 as method of yield enhancement, 15,
 16, 21;
 near-dated options, 20, 75;
 optimal strike for, 20–21;
 and overpricing of implied volatility,
 15, 17;
 popularity of, 19–20, 75;
 profitability of, 19;
 and put spread protection, 31;
 similarity to put underwriting, 10, 17;
 on single stocks, 18, 21, 22;
 and skew, 188;
 three-month vs. one-month, 29, 30, 76;
 vanilla, 22;
 with booster, 18
- call spreads:
 1x2, 17, 18, 31, 32;
 payout of, 195, 251;
 and relative value trading, 156;
 and square root of time rule, 195;
 vs. put, 31, 32
- capital structure arbitrage, 270–84
- carry:
 calculation of, 36, 42, 47–48;
 and call overwriting, 20;
 and forward starting products, 100,
 158;
 and quanto options, 150;
 and relative value volatility trading, 158;
 and reverse trades, 270
- cash gamma, 67, 226–27
- CBOE:
 and automatic exercise of options, 13;
 call overwriting index of, 18;
 and correlation plotting, 167, 168;
 opening of, 2;
 skew index, 207;
 and volatility indices, 110, 113, 118
- central moments, 201
- Chicago Board Options Exchange. See:
 CBOE
- cliques, 106

- close-to-close volatility, 232, 238, 239–40, 241, 242
- collars:
- calendar, 29, 30;
 - cap on performance with, 24, 26;
 - as low-cost method of protection, 24;
 - put spread, 23, 24, 25, 26, 27, 28;
 - and skew, 24, 26;
 - and volatility 24, 26, 28;
- when to use, 26.
- See also: risk reversals
- composite options, 133, 149–50
- constant smile rule, 103, 105, 106, 262
- contingent premium options, 133, 148
- continuous barriers, 138
- convertible bonds, 198
- convexity:
- and call options, 9;
 - gamma as measure of, 40, 258;
 - of options, 6;
 - and options on variance, 69, 70;
 - of payout, 40;
 - and profits from delta hedging, 41, 45;
 - and relative value trading, 156;
 - and stock replacing with calls, 9;
 - and variance, 49;
 - and variance swaps, 54, 67, 69, 109;
 - vega, 83, 86, 87, 88;
 - and volatility indices, 109
- correction, protection strategies for, 27, 28
- correlation:
- and composite options, 149, 150;
 - index, 163–64;
 - methods of trading, 164–65;
 - and outperformance options, 143, 144, 145, 157, 165;
 - plotting of, 167;
 - and quanto options, 150;
 - in relative value trading, 154, 155;
 - and skew, 158;
 - and theta-weighted dispersion trading, 174;
- types of, 166;
- and volatility, 166, 169, 170, 174, 175;
- and worst-of/best-of options, 140, 141
- correlation indices tickers, 167
- correlation surface, 192
- correlation swaps, 161, 165, 166, 168–70
- correlation trading, 161–77
- correlation-weighted dispersion trading, 171
- counterparty risk, 128, 266, 267
- covariance, 176
- covariance swaps, 165, 176
- crash, protection strategies for, 28
- crash puts, 148
- credit default swaps, 270, 271
- credit-induced skew, 274–75, 280–81
- credit spreads, 270, 273, 280, 281
- crossed vega, 167
- cross-gamma, 176n.24
- cross-holding trading, 154
- currencies, options involving, 149–51
- ## D
- daily profit and loss, calculation of, 45, 47–48, 50, 55
- debt, 275, 276, 277
- declines, types of, 28
- delta:
- anchor, 217n.38, 262, 263–65;
 - barrier options, 134, 136;
 - and borrow cost, 38;
 - call options, 8, 9, 258;
 - debt, 277;
 - definition, 6, 7, 39, 257, 258;
 - and dividend risk, 7, 37, 38;
 - intraday, 59;
 - ITM options, 6;
 - and probability of option being ITM at expiry, 8–9;
 - put options, 8, 9, 258;
 - shadow, 100–101;
 - stock, 9;

delta (continued):
 volatility futures, 119
 delta-1 products, 4, 154, 155, 156
 delta hedging:
 continuous, 40, 52, 89–93;
 daily profit and loss from, 45, 47–48,
 50;
 definition, 39;
 discrete, 89, 93–98;
 and dividend risk, 37;
 and equity exposure, 257, 258;
 every 5 percent move in spot, 94;
 every 10 percent move in spot, 94;
 example of, 39;
 gamma, 225;
 on key dates, 41;
 as method of locking in profit, 40;
 profit from, 41, 45–47, 49, 52, 90–96;
 skew, 225;
 and trending markets, 41;
 via synthetics, 123;
 when to do, 41
 delta re-hedging, 40
 delta skew, 207, 208, 209
 diffusive volatility, 178, 179, 180, 280
 digital options, 8, 138
 directional investors, 36, 39
 discrete barriers, 138
 discrete sampling, 110, 112, 113
 dispersion trading, 159, 161, 164–65,
 170–74, 175
 dividend net present value (NPV), 11, 12,
 13
 dividend risk, 7, 37
 dividends:
 discrete, 58;
 fixed, 101;
 impact on early exercise of calls, 11, 12;
 ordinary, 14;
 proportional, 58, 101;
 special, 14;
 and variance swaps, 58
 dividend yield assumption, 149

Dodd-Frank Act, the, 82
 dollar gamma, 47, 67, 226–27
 double barriers, 138
 drift, 232, 234, 238, 241, 243
 DTB (now Eurex), 118
 dual listing trading, 154

E

efficiency, 239, 241
 enhanced call overwriting, 22
 enterprise value, 275, 276, 277, 279
 EONIA, 128
 equity risk premium, 19, 20, 79
 ETN/ETF, 122–23, 125, 126, 127
 Eurex, 13, 118
 Euronext-Liffe, 13
 European options:
 automatic exercise of, 11;
 barrier options, 134;
 contingent premium options, 148;
 definition, 2, 11;
 look-back options, 146;
 maintenance of, 11;
 options on variance, 68;
 outperformance options, 143;
 payout of, 92;
 put-call parity and, 37;
 worst-of/best-of, 140
 event-driven trading, 155
 Everest, 164
 excess return indices, 128
 execution risk, 154
 expected daily return, 245–47
 expected value, 201
 expected volatility, 97
 expiry:
 choice of as enforcing investor discipline, 4;
 correlation indices, 167;
 far-dated, 4;
 key dates in choice of, 4, 5;

- expiry (continued):
linear interpolation between, 110, 113, 114;
look-back options, 146;
near-dated, 4;
non-triple witching, 123;
rolling, 110, 113;
strategies for choosing, 4–5;
triple witching, 167;
volatility futures, 118;
and volatility surfaces, 184
- exponentially weighted volatility, 232, 237–38
- F**
- fair price, 109, 110, 166
fair value, 3, 115, 142, 164
fat-fingered trade, 238
fixed annuities, 81
fixed dividends, 101
fixed leg, 57
flat market, protection strategies for, 27
floating leg, 57
forwards:
and borrow cost, 38;
correlation swaps as, 168;
definition, 6, 36;
delta of, 6;
excluding dividends, 37;
as first moment of a probability distribution, 51, 201, 202;
gamma swaps as, 58;
potential loss on, 6;
profitability of, 6;
variance swaps as, 58, 168;
volatility swaps as, 58
forward skew, 103, 105
forward starting options:
average implied volatility of, 100;
and carry, 158; definition, 99, 100, 101;
delta of, 100, 101, 102;
hedging of, 100, 102;
- and implied jump trading, 180;
maturity of, 101;
modeling of, 99, 100;
and outperformance options, 144;
pricing of, 100, 103, 262;
relative value trading of, 158;
straddles with, 102;
strike of, 101, 102;
vega hedging of, 102;
zero gamma of, 100, 102;
zero theta of, 99, 100
- forward starting variance swaps:
constructing, 107
hedging of, 102, 108;
maturity of, 101–2;
payout of, 106;
pricing of, 100, 102, 108;
spread of, 108;
and term structure trading, 200
- forward variance, 107
- forward volatility:
and diffusive volatility, 179;
methods of calculating, 103–6, 180;
methods of trading, 100, 101–2;
and pricing of forward starting options, 103;
underpricing of, 106
- FTSE, 158
- futures, on volatility indices, 2, 102, 108, 158
- G**
- gamma:
Black-Scholes, 67;
cash, 67, 226–27;
definition, 40, 257, 258–59;
and delta hedging, 52, 90, 91, 225;
dollar, 47, 67, 226–27;
and forward starting options, 100;
long position, 41;
as measure of convexity of payout, 257, 258;

- gamma (continued):
 proportionality to volatility, 47, 202;
 share, 67;
 short position, 41;
 and term structure trading, 200;
 and theta, 90;
 and volatility pair trading, 159;
 volga as measure of, 205
 gamma scalping, 40, 41
 gamma swap dispersion, 171
 gamma swaps:
 advantages of, 55–56;
 constant share gamma of, 67;
 and dispersion trading, 171, 172;
 fair price of, 52;
 as forwards, 58;
 Greeks of, 67;
 hedging of, 59, 61, 63;
 invention of, 52;
 liquidity of, 56;
 payoff, 56, 58;
 payout, 55, 57, 62, 171, 240;
 single stocks, 56;
 spot weighting of, 55–56;
 as square root of variance, 52;
 vega profile of, 62
 gamma theta, 226–27, 228
 gamma trading, 225
 gamma-weighted dispersion trading, 171, 172, 172, 174
 gamma weighting, 159, 160
 gap risk premium, 15
 Garman-Klass, 232, 238, 239, 241, 243
 gearing, 77
 geometric Brownian motion, 178
 Greeks: calculating, 96, 97;
 and convexity of payout, 257;
 definition, 257–61;
 of dispersion trading, 172;
 and equity exposure, 257;
 of gamma swaps, 67;
 and interest rate exposure, 257;
 and moments, 202;
- of option pair trades, 159;
 shadow, 67, 100–101, 262–65;
 and skew, 257;
 and time decay, 257;
 of variance swaps, 67;
 and volatility exposure, 257;
 and volatility of volatility, 257;
 of volatility swaps, 67
- ## H
- hedge funds:
 increasing popularity of, 2;
 maturity of, 7
 hedging error, 93, 94, 95
 high frequency trading (HFT), 236
 Himalayas, 164
 historical volatility:
 calculation of, 232–44,
 close-to-close, 232, 238, 239–40, 241, 242;
 exponentially weighted, 232, 237–38;
 frequency of, 234, 235–36;
 Garman-Klass, 232, 238, 239, 241, 243;
 length of time for, 234–235;
 Parkinson, 232, 238, 239, 241, 242;
 Rogers-Satchell, 232, 238, 239, 241, 243, 244;
 Yang-Zhang, 232, 238–39, 241, 243
 hurdles, 13, 143
- ## I
- implied correlation:
 and dispersion trading, 164–65, 166, 167, 170, 174;
 and outperformance options, 143, 144;
 overpricing of, 170;
 sticky, 192;
 and worst-of/best-of options, 140, 141, 142
 implied correlation surface, 162, 184
 implied diffusive volatility, 275
 implied jump, 5, 178–82, 245–47

implied no-default volatility, 270, 275, 284
implied variance, 50
implied volatility:
 absolute, 215, 216;
 Black Scholes, 220, 221, 230–31, 275,
 281;
 and calculating delta, 92, 96, 97, 98;
 and calculating implied jump, 178;
 and diffusive volatility, 179;
 discrete sampling of, 110, 112, 113;
 as estimate of future volatility, 96;
 flooring and capping of, 263;
 and forward starting options, 100;
 index, 18, 163, 170;
 overpricing of, 15, 17, 18, 29, 53, 74,
 75, 78;
 overshoot of, 83, 86;
 and put spread pricing, 26;
 and realised volatility, 186;
 relative, 215, 216;
 removal of cap on, 60, 61;
 single stock, 18, 163, 170;
 and square root of time rule, 194;
 sticky, 189–91, 192;
 term structure of, 117, 125, 186;
 undershoot of, 83;
 and variance, 49, 50;
 volatility of, 130, 186;
 and volatility surfaces, 184
implied volatility premium, 15
implied volatility surfaces, 162, 215, 216,
 222
indexation flow, 154
indices:
 call overwriting on, 18, 21, 22;
 implied volatility in, 18;
 implied volatility premium for, 15;
 liquidity of, 7;
 options on, 11
instantaneous volatility, 230–31
in-the-money options. See: ITM options
interest rates, 11, 13
intraday delta, 59

intraday volatility, 236–37
intrinsic value, 3
investors:
 directional, 36, 39;
 volatility, 36, 39
ITM options:
 automatic exercise of, 13;
 cost of, 3;
 definition 3, 6;
 delta of, 6;
 early exercise of, 11, 12, 13;
 intrinsic value, 3, 6;
 liquidity of, 7;
 and "normal" market moves, 6;
 pricing of, 12, 230;
 return on, 6;
 similarity to forward, 6

J

Jensen's inequality, 164, 242
jump diffusion, 270, 274–75, 280–81, 284
jump volatility, 180, 181, 245
jumpy volatility regimes, 210, 212, 213,
 214, 215, 221–23, 225

K

key dates:
 and choosing expiry, 4, 5, 178;
 and delta hedging, 41
knock-in options, 134, 135, 136, 137
knock-out options, 134, 135, 136, 137,
 138, 139
kurtosis, 201, 202, 203, 205, 206

L

ladders, 32, 33, 123
leptokurtic distributions, 202, 205
leverage, 7, 188, 201, 272, 279
leveraged buyout, 272
Levy processes, 206n.34

- light exotics:
- barrier options as most popular, 142;
 - definition, 133;
 - as European options, 134, 143, 146;
 - as OTC, 134, 143, 146
- liquidity:
- call options, 7;
 - correlation swaps, 170;
 - covariance swaps, 176;
 - and demand for protection, 7;
 - as factor in choosing strike, 5, 7
 - gamma swaps, 56;
 - indices, 7;
 - ITM options, 7;
 - OTM options, 7;
 - outperformance options, 157;
 - put options, 7;
 - single stock options, 7;
 - single stock variance market, 171;
 - and straddle dispersion trading, 170;
 - variance swaps, 52, 53, 55, 170
- local volatility, 212, 218, 219, 230–31
- log contract, 50, 62, 63, 115, 248
- long strategies, 2, 4, 5
- long-short trading, 155
- Long-Term Capital Management, collapse of, 52
- long volatility strategies, 4, 5, 78, 80
- look-back options, 81, 133, 146–47
- ## M
- Margrabe's formula, 144
- mark-to-market, 136, 157, 170, 267
- market disruption, 57, 58
- market risk, 93
- market skew, 213
- maturity:
- choosing, for options, 157;
 - correlation swaps, 169–70;
 - as driven by duration of likely decline, 28;
 - forward starting options, 101;
 - forward starting variance swaps, 101–2;
 - hedge funds, 7;
 - look-back options, 147;
 - outperformance options, 143;
 - and smirk, 206;
 - structured products, 7, 169;
 - worst-of/best-of options, 140
- maximum anniversary value, 81
- mean reversion, 119, 131, 154, 158, 159, 184, 193, 199
- mergers and acquisitions, effect on value of options, 14
- Merton, Robert 2
- Merton model, 270, 274, 275–79
- mesokurtic distributions, 202
- model risk, 66
- moments, 201–5
- moneyness, 103
- ## N
- naked put, selling of, 10, 17
- near-dated options:
- and call overwriting, 15, 20;
 - and capital structure arbitrage strategies, 284;
 - and implied volatility, 199, 251;
 - and multiple-expiry protection strategies, 28, 29;
 - and volatility selling strategies, 75
- NKY, 166
- noise, 77, 94, 95, 110, 113, 154, 155
- "normal" market moves, 6
- notional:
- and relative value trading, 155, 157;
 - variance, 54, 56, 62, 63, 64, 65, 66, 68, 69, 106–7;
 - variance swap, 250;
 - volatility, 53, 54, 62, 64;
 - volatility swap, 250

O

option adjusted spread (OAS), 270
option combos, 31
option on bankruptcy, 281
option structures, 31–32, 33
options:
 advantages to using, 4, 6;
 Asian, 68;
 basket, 14, 68, 159, 164, 165, 167, 175–76;
 cheap, 92, 93, 95, 96;
 choosing expiry of, 4;
 choosing maturity of, 157;
 choosing strike of, 157;
 composite, 133, 149–50;
 contingent premium, 133, 148;
 convexity of, 6;
 digital, 8, 138;
 involving currencies, 149–51;
 on indices, 11;
 knock-in, 134, 135, 136, 137;
 knock-out, 134, 135, 136, 137, 138, 139;
 look-back, 81, 133, 146–47;
 mergers and acquisitions effect on, 14;
 outperformance, 133, 142, 143–45, 154, 155, 156, 157, 165;
 potential loss on, 6;
 as preventing inertia, 4, 16;
 as protection, 15, 23–30;
 quanto, 133, 149, 150–51;
 rainbow, 175n.23;
 rolling, 11, 14, 29, 77, 78, 102, 110, 157;
 spinoffs' effect on, 14;
 time value of, 3, 14;
 worst-of/best-of, 133, 140–42, 164, 165.
See also: American options; ATM options; barrier options; call options; European options; forward starting options; ITM options; near-dated options; OTM options; put options;

single stock options; vanilla options
options on variance:
 breakeven for, 70;
 capped, 71;
 and convexity, 69;
 inverted term structure of, 71;
 payout of, 69, 70;
 poor liquidity of, 68, 71;
 price of, 68, 71;
 and put-call parity, 69;
 skew of, 70, 71;
 straddles on, 72;
 and volatility of volatility, 68
options on volatility, 70
options on volatility futures, 72
OTM options:
 and "abnormal" market moves, 7;
 and call overwriting, 20;
 definition, 3;
 delta of, 6, 7, 223;
 and implied volatility, 201, 202;
 and ladders, 32;
 leverage of, 7;
 liquidity of, 7, 51, 53, 109, 111;
 method for calculating premium, 41, 44;
 pricing with additive variance rule, 105;
 and put spread collars, 24;
 and put underwriting, 10;
 time value of, 117, 199;
 and variance swaps, 53, 54, 117;
 vega of, 86, 87;
 and volatility of volatility, 202;
 volga of, 205, 261
outperformance, 25
outperformance options, 133, 142, 143–45, 154, 155, 156, 157, 165
out-the-money options. See: OTM options

P

pair trading, weighting of, 159, 160

Parkinson, 232, 238, 239, 241, 242

payoff:

- barrier options, 135;
- call options, 8;
- correlation swaps, 168–69, 176–77;
- corridor variance swaps, 55;
- dispersion trading, 170;
- gamma swaps, 56, 58;
- variance swaps, 54, 58, 115;
- volatility swaps, 54, 58;
- volatility futures, 115

payout:

- barrier options, 134, 135, 136, 138;
- basket options, 176;
- call spreads, 195, 251;
- and close-to-close prices, 240;
- and correlation, 166;
- correlation swaps, 168–69, 175;
- delta-hedged options, 47, 91, 94;
- and dispersion trading, 168, 170, 175;
- European options, 92;
- forward starting variance swaps, 106;
- gamma swaps, 55, 57, 62, 171, 240;
- look-back options, 146, 147;
- options on variance, 69, 70;
- outperformance options, 143;
- put spreads, 195, 251, 252;
- quanto options, 150;
- relative value trading, 154;
- structured products, 176;
- variance swaps, 52, 53, 54, 55, 57, 58, 59, 64, 69, 110, 119, 240;
- vega-weighted pair trades, 160;
- volatility futures, 115;
- volatility swaps, 52, 53, 55, 57, 64, 240, 250;
- with zero volatility, 3;
- worst-of/best-of options, 140

performance caps with collars, 24, 26

perpetual structured products, 122

pinning, 42, 43, 59

platykurtic distributions, 202

power vega, 199

preference shares, 154

premium:

- of ATM options, 157;
- and contingent premium options, 148;
- equity risk, 74, 78, 79;
- importance of to directional investors, 39;
- of knock-in/knock-out options, 134;
- methods for calculating, 41–42, 44, 198;
- as method of paying for protection, 27;
- and near-dated options, 251;
- and put spreads, 195;
- as restriction on choice of strikes, 28;
- return of, 81;
- and risk reversal, 220;
- and strike reset options, 146

pricing:

- barrier options, 138;
- basket options, 176;
- calendars, 198, 199;
- forward starting options, 100, 103, 262;
- outperformance options, 143, 157, 165;
- quanto options, 151;
- straddle dispersion, 170;
- worst-of/best-of options, 140, 165

probability distribution, 201

profit:

- from delta hedging, 41, 45–47, 49, 52, 90–96;
- path independence of, 90, 91–92;
- proportionality to volatility, 41, 45–47, 49;
- from theta-weighted variance dispersion, 166

proportional dividends, 58, 101

protection:

- index vs. single stock, 15, 25;
- median maturity of, 28;
- methods of paying for, 27;

- protection (continued):
option buying for, 15, 23–30;
overpricing of, 29;
partial, 24, 25;
strategy comparison, 27;
volatility as determination of strategy
for, 28
- put-call parity, 36, 37, 69
- put options:
American, 13;
and anchor delta, 264;
ATM, 26, 28, 69, 81, 85, 136, 211, 283;
and barrier options, 134, 135, 136, 137,
138;
as bearish or bullish strategy, 3, 5;
and borrow cost, 38;
breakeven, 70;
and calendar collar, 30;
and contingent premium options, 148;
cost of, 23, 24, 25, 26, 28, 29, 31, 227;
crash, 148;
definition of, 2;
delta of, 8, 9, 39, 207, 208, 214, 224,
258, 277;
and delta hedging, 36, 39, 97;
digital, 8;
disadvantages of, 23;
downside, 23, 54, 85, 111, 119, 212,
263;
early exercise of, 11, 12, 13;
effect of volatility on, 3;
on the enterprise, 274;
European, 2, 12, 13, 36, 37;
and gamma scalping, 40;
hedging with, 80;
and implied volatility, 119, 211, 263,
284;
intrinsic value, 3;
ITM, 81, 97;
and kurtosis, 205;
liquidity of, 7, 111;
long, 2, 4, 5, 12, 13, 23, 24, 26, 29, 31,
40, 138, 202, 220, 223, 224, 225n.40,
- 251, 283;
and look-back options, 146, 147;
maximum loss with, 23;
maximum value of, 8;
methods of reducing cost of, 23;
multiple expiry strategies, 29;
and no-fault implied volatility, 284;
OTM, 12, 23, 24, 26, 31, 54, 63, 83, 88,
148, 204, 211, 212, 220, 223, 225, 263,
264;
payout of, 12, 13, 26, 70, 135, 148, 252,
282;
premium of, 41, 69, 195, 251;
price of, 3, 61, 119, 252, 255, 283;
probability expires ITM, 8;
profit conditions for, 2;
as protection strategy, 23, 24, 25, 27,
28, 29, 31, 75, 83, 136, 186, 188;
and put-call parity, 36, 37, 69;
and put underwriting, 17;
and quanto options, 151;
replacing stock with, 9, 10;
rho of, 260;
and risk reversals, 31, 84, 85, 202, 220,
225n.40;
and shadow delta, 101;
short, 2, 4, 5, 9, 17, 26, 30, 31, 69, 83,
84, 123, 138, 220, 277;
short-dated, 28, 29;
short equity exposure of, 2;
and skew, 204, 206, 211, 212, 223, 227;
and smile, 206;
and stock replacing, 9, 10, 17;
and straddles and strangles, 30, 33, 36,
38;
strike of, 24, 145;
and structured product guarantee, 88;
and synthetics, 123;
theta of, 259;
time value of, 61;
and value of jump to zero stock price,
281;
vanilla, 137, 156;

put options (continued):
 and variable annuities, 81;
 and variance swaps, 54, 58, 69, 70, 119;
 and volatility trading, 36, 70;
 vs CDS, 270, 275, 281, 282, 283;
 and VXO, 51, 109;
 and worst-of/best-of options, 140

put spreads:
 1x2, 26, 27, 28, 32, 33;
 calendar, 30;
 cost-effectiveness of, 24, 26;
 outperformance with, 25;
 as partial protection, 24, 27, 31;
 payout of, 195, 251, 252;
 price of, 26, 254;
 and relative value trading, 156;
 risk-reward profile of, 25;
 and square root of time rule, 195;
 and volatility, 28

put spread collars 23, 24, 25, 26, 27, 28

put spread vs call, 23

put vs CDs, 270, 275, 281–84

put vs strangle/straddle, 30

put underwriting, 10, 17, 31

Q

quantitative easing, 141
 quanto options, 133, 149, 150–51
 QE2, 141

R

rainbow options, 175n.23
 range trading, 16, 146, 147
 ratchet, 81
 ratio put spreads, 283
 realised dkew, 222–23
 rebates, 139, 268
 recovery rate, 277, 282
 relative dimensions, 215
 relative time, 103, 228
 relative value trading, 143, 144, 154–58,
 178

relative value volatility trading, 158–60
 re-marking of volatility, 100, 217, 218,
 220, 221, 222, 225, 262
 repo, 38
 return of premium, 81
 reverse trades, 270
 rho, 257, 260
 risk-free rate, 74, 78, 79, 80
 risk-on assets, 141
 risk reversals, 24, 31, 32, 211, 220, 223.

See also: collars

Rogers-Satchell, 232, 238, 239, 241, 243,
 244

roll-down cost, 125

rolling expiries, 110, 113

rolling options, 11, 14, 29, 77, 78, 102,
 110, 157

roll-up, 81

Russell 2000, 118

S

S&P100, 109

S&P500:

call overwriting index on, 18;
 and correlation, 166;
 implied volatility of, 158;
 OTC market of, 123–24;
 price return index (SPX), 18;
 size of options market, 122, 123;
 and skew index, 207;
 term structure, 117, 125;
 total return index (SPXI), 18;
 and VIX index, 109

Scholes, Myron, 2

shadow delta, 100–101

shadow Greeks, 67, 100–101, 262–65

share class trading, 154

share gamma, 67

shark fin, 138

Sharpe ratio, 269

short rebate, 268

- short volatility strategies:
 expiry for, 5;
 when to use, 4
- single stock options:
 as American, 11;
 call overwriting on, 18, 21, 22;
 implied volatility in, 18;
 implied volatility premium for, 15;
 liquidity of, 7;
 volatility of, 20
- skew:
 and bankruptcy, 189, 201;
 Black-Scholes, 218, 231;
 buying, 26;
 and collars, 24, 26;
 credit-induced, 274–75, 280–81;
 decay of, 106;
 definition of, 184;
 delta, 207, 208;
 delta hedging, 225;
 dependence on correlation, 158;
 with diverse indices, 192;
 fair pricing of, 218, 222–23;
 fixed strike, 207, 208;
 forward, 103, 105;
 and forward starting options, 101, 103, 105;
 index vs. single stock, 192;
 index for, 207;
 inverted, 187;
 and local volatility, 218, 219, 231;
 and log contracts, 63;
 market, 213;
 as measure of value of skew, 86;
 measuring, 183, 207–8;
 and options of different maturity, 187;
 and options on variance, 70, 71;
 and options on variance swaps, 129;
 and options on volatility futures, 129, 131;
 overpricing of, 10, 81, 212;
 and put spread pricing, 26;
 in put underwriting, 10;
- realised, 222–23;
rise in, 82;
as risk indicator, 191;
selling, 26;
and smirk, 206;
and square root of time rule, 194, 195;
sticky delta, 103;
strike, 207, 208, 209;
and term structure, 183, 184, 188–92;
theoretical, 213;
- as third moment of probability distribution, 201, 202, 204, 206;
upper and lower bound for, 252–54;
vanna as measure of, 86, 204, 257, 261;
and variable annuities, 81;
and variance, 49;
and variance swaps, 54, 59, 68, 101, 109, 117;
 and volatility indices, 109
- skew index, 207
- skew profit, 223
- skew theta, 211, 212, 217, 218, 221, 222, 223, 225, 226, 228
- skew trading, 183, 193, 210–28
- smile, 183, 201, 202, 206; measuring, 183
- smirk, 201, 206
- Sortino ratio, 21, 269
- spinoffs: effect on value of options, 14
- spreads:
 bid-offer, 25, 41, 100, 101, 102;
 box, 123;
 call, 23, 31, 195, 251;
 call 1x2, 17, 18, 32;
 call/put, 122;
 credit, 270, 273, 280, 281;
 option adjusted, 275;
 put, 23, 24, 25, 26, 27, 28, 31, 32, 195, 251, 252, 254;
 ratio put, 283
- SPX, 166
- square root of time rule, 119, 183, 193–97, 199, 251

- standard deviation, 21, 42, 58, 232, 233, 234, 242, 269
- sticky delta regimes, 210, 212, 213, 214, 215, 216–17, 218, 221, 223, 225
- sticky delta skew, 103
- sticky local volatility regimes, 210, 212, 214, 215, 218–21, 223, 225
- sticky moneyness regimes, 210, 216
- sticky strike, 184, 185, 187, 194
- sticky strike regimes, 210, 212, 214, 215, 217–18, 221, 223
- stochastic local volatility model, 116
- stock replacing, 9, 10
- straddle dispersion, 170–71
- straddles:
- and butterflies, 34;
 - delta hedging of, 59;
 - and dispersion trading, 165, 170;
 - and forward starting options, 102;
 - on options on variance, 72;
 - put vs., 30;
 - similarity to strangles, 33;
 - vanilla, 102;
 - and volatility pair trading, 159;
 - zero delta, 38, 40
- strangles, 30, 33, 34, 40, 102, 159, 170
- strike:
- liquidity as factor in choosing, 5, 7;
 - premium as factor in choosing, 28;
 - in put underwriting, 10;
 - sticky, 184, 185, 187, 194;
 - strategies for choosing, 5–7;
 - and volatility surfaces, 184;
 - zero delta straddles, 38
- strike reset, 146, 147
- strike skew, 207, 208, 209
- structured products:
- and implied correlation, 164–66;
 - capital guarantee of, 88;
 - hedging of, 7;
 - and implied volatility imbalances, 198;
 - lifting of index implied volatility, 15, 18, 83, 86;
- maturity of, 7, 169;
- perpetual, 122;
- vicious circle, 83, 87
- stub trades, 154
- supply and demand imbalances, 15, 74, 117, 125, 184, 186, 198
- Swisscom, pinning of, 43
- SX5E:
- and 1x2 put spreads, 26;
 - and ATM puts, 28;
 - and call overwriting, 20, 21;
 - implied correlation of, 166;
 - implied volatility of, 158, 191;
 - and liquidity of variance market, 55;
 - and skew, 82, 188, 195;
 - term structure, 125, 188, 197;
 - variance indices of, 118;
 - and variance swaps, 77, 79;
 - vega, 126, 127;
 - vs. 5-year CDS, 274;
 - and vStoxx, 74, 78, 125
- synthetics, 123

T

- tail risk, 53, 66
- tails, 109, 111
- takeovers, 155, 187, 271, 272
- term structure:
- ATM, 199;
 - and bankruptcy, 189;
 - definition of, 184, 196;
 - and forward starting options, 100, 103;
 - implied correlation, 192;
 - implied no-default implied volatility, 275, 284;
 - implied variance, 71;
 - implied volatility, 75, 87, 125, 130, 186, 199, 275;
 - index, 181, 182;
 - inaccurate pricing of, 81, 105;
 - inverted, 29, 68, 71, 185, 186, 187;
 - modeling of, 196;

term structure (continued):
normalising, 178, 196, 197;
and options of different strikes, 185;
options on variance swaps, 129;
risk, 34;
single stock, 182;
and skew, 183, 184, 188–92, 226;
slope of, 75, 77, 106, 113, 117, 184,
185, 187, 190, 227;
and square root of time rule, 193, 194,
196, 226;
SX5E, 125, 188, 197;
trading, 198–200;
vanilla equity options, 130;
variance, 51, 199;
volatility, 68, 110;
volatility futures, 77, 122, 124, 126, 129
theoretical skew, 213
theta:
call options, 259;
definition, 257, 259;
and dispersion trading, 172;
and forward starting options, 99, 100;
and gamma, 90, 91;
of near-dated options, 20, 29;
put options, 259;
and term structure trading, 200;
and time decay, 257, 259;
volatility slide, 106, 228;
and volatility pair trading, 159
theta-weighted dispersion trading, 171,
172, 173, 174
theta weighting, 159, 160, 166
tick value, 119
time decay, 14, 257, 259
time value, 3, 90, 117, 135, 178, 199
total return indices, 128
trending markets:
definition, 41;
and delta hedging, 41, 94;
and sticky delta regimes, 213, 216;
and volatility measurement, 235

U

upside participation strategies, 31–32

V

V1X, 109, 132
V2X, 118, 132
vanilla options:
compared to barrier options, 134, 135,
136, 137;
compared to composite options, 149;
compared to contingent premium
option, 148;
compared to forward starting options,
100, 103;
compared to look-back options, 146;
compared to quanto options, 15, 151;
and constant smile rule, 262;
and dispersion trading, 170;
implied volatility term structure for,
130;
increase in value as strike approaches,
136;
and outperformance options, 157;
pricing of, 230, 231;
and relative value trades, 155, 156, 157,
158;
and removal of implied volatility cap,
61;
similarity to options on variance, 72;
size of market for, 122–23

vanna:

definition, 85, 86, 225, 257, 261;
and determining profitability of skew
trades, 223;
as measure of third moment of
probability distribution, 202, 203, 204;
and skew, 202, 204, 257, 261

variable annuities:

definition, 81, 204;
downside to, 81;
hedging of, 81, 123;
and skew, 81;

variable annuities (continued):

- types of downside protection with, 81
- variance:
 - as additive, 58, 67, 106, 180, 232, 233–34, 245;
 - and avoiding deviation cancelling, 50;
 - and delta hedging payout, 47;
 - fair price of, 49;
 - and implied volatility, 49;
 - index, 53;
 - as measure of deviation, 41, 45, 49–51, 109;
 - options on, 68;
 - overpricing of, 74, 75;
 - as second moment of probability distribution, 201, 202, 203;
 - single stock, 53, 66;
 - term structure of, 117;
 - and volatility, 42, 47, 49, 50, 51, 53, 54, 104

variance notional, 54, 56, 62, 63, 64, 65, 66, 69

variance swap dispersion, 171

variance swaps:

- as alternative to strangles or straddles, 30;
- calculation agents for, 58;
- call on, 70;
- capped, 54, 71;
- constant cash gamma of, 67;
- and convexity, 54, 67, 69, 70, 115;
- corridor, 55;
- delta of, 67, 101;
- and dispersion trading, 165, 171;
- down, 55;
- effect of dividends on pricing of, 58;
- fair price of, 109;
- forward starting, 100, 101–2;
- as forwards, 58, 168;
- Greeks of, 67;
- hedging by log contract, 248–49;
- hedging of, 54–55, 59, 61, 62, 63, 115;
- history of, 2, 52, 53;

index, 54, 55, 58, 71;

- influence on market movement, 59, 60;
- liquidity of, 52, 53, 55, 56, 68, 170;
- maximum loss of, 66;
- notional, 250;
- one-year, 54;
- options on, 129;
- path independence of, 50;
- payoff, 54, 58, 115;
- payout of, 52, 53, 54, 55, 57, 58, 59, 64, 69, 110, 119, 240;
- price of, 65, 66, 68, 109, 110;
- put on, 70;
- and removal of cap on implied volatility, 61;
- S&P500, 55;
- single stock, 53, 54, 55, 71;
- and skew, 54, 59, 63, 67, 68, 117;
- SX5E, 55;
- and tail risk, 53;
- and term structure trading, 200;
- up, 55;
- vanilla, 54, 55, 71, 102, 107, 108;
- vega portfolio of, 62, 63, 64, 65, 67;
- and volatility of volatility, 52, 66;
- and volatility pair trading, 159

variance term structure, 51

VDAX (V1X), 51, 109

vega:

- and ATM options, 203;
- convexity, 83, 86, 87, 88;
- crossed, 167;
- decay of with time, 67;
- definition, 86, 123, 203, 257, 259–60;
- and dispersion trading, 172;
- and forward starting options, 100, 102, 106;
- of gamma swaps, 62;
- and implied volatility, 200, 203;
- as measure of second moment of probability distribution, 202, 203, 204;
- of options of different strikes, 61;
- peak, 87;

- vega (continued):
of and straddle dispersion, 170;
and variance swap notional, 250;
variance swaps, 62, 63, 64, 65;
and volatility exposure, 257, 259–60;
and volatility pair trading, 159;
of volatility swaps, 62, 64, 65
vega flat risk reversal, 84, 85
vega hedging, 102
vega-weighted dispersion trading, 171,
 172, 173, 174
vega weighting, 159, 160
VFTSE, 51, 109
VHSI, 51, 109
VIMEX, 51, 109
VIX: 77, 111;
 ETN/ETF, 125;
 forward starting, 99, 100, 101;
 futures delta, 120, 121;
 imbalance in market, 127–28;
 liquidity of, 122;
 listing of futures on, 115, 118;
 and relative value volatility trading, 158;
 settlement price of, 119;
 size of futures market, 122, 123;
 term structure, 117;
 and variance-based calculations 49, 51,
 109
VNKY, 51, 109
volatility:
 annualized, 42;
 ATM, 222;
 Black-Scholes, 90, 218;
 calculating, 202;
 of cash, 14;
 close-to-close, 232, 238, 239–40, 241,
 242;
 and correlation, 166, 169, 170, 174,
 175;
 correlation with equity, 74, 77, 78;
 and credit, 189;
 daily, 241;
 definition, 232;
 as determinator of protection strategy,
 28;
 diffusive, 178, 179, 180;
 effect on call options, 3;
 effect on put options, 3;
 expected, 97;
 exponentially weighted, 232, 237–38;
 historical, 41, 232–44;
 implied vs. realized, 130;
 index, 161–64;
 instantaneous, 230–31;
 intraday, 236–37;
 jumpy, 210;
 local, 212, 218, 230–31;
 long strategies, 4, 5;
 as measure of deviation, 41, 45, 49,
 109;
 overnight, 241;
 overpricing of, 74, 75, 77, 78, 79, 81;
 and pricing of options, 3;
 proportionality to profit, 41, 45–47, 49;
 realised, 222;
 re-marking of, 100, 217, 218, 220, 221,
 222, 225, 262;
 as risk measure, 191n.27;
 seasonality of, 117;
 as second moment of a distribution, 51;
 short strategies, 4, 5;
 sticky local, 210;
 term structure, 68;
 and variance, 42, 47, 49, 50, 51.
See also: implied volatility
volatility cone, 185, 186
volatility futures:
 constant maturity, 122;
 delta of, 119, 120;
 delta hedging of, 122;
 eighth month, 127;
 expiry of, 118;
 fair price of, 115;
 fourth month, 122, 127;
 front month, 122, 127;
 hedging of, 102;

- volatility futures (continued):
 liquidity of market for, 102;
 and mean reversion, 131;
 options on, 129–32;
 overpricing of, 102, 116;
 payoff of, 115;
 payout of, 115;
 probability distribution of options on, 131;
 and relative value volatility trading, 158;
 skew of options on, 129, 131;
 term structure of, 117, 122, 124;
 volatility of, 130;
 and volatility of volatility, 102, 115, 117, 131
- volatility indices:
 as average of 8 different options, 109;
 fair price of, 109;
 futures on, 2, 102, 108, 115–21;
 and linearly interpolating between expiries, 113, 114;
 methods of calculating, 109, 110;
 pricing of, 109, 110–14;
 providers, 51;
 and relative value volatility trading, 158;
 variance-based calculation of, 109;
 volatility of, 130
- volatility investors, 36, 39
- volatility notional, 53, 54, 62, 64
- volatility of volatility:
 and bias, 241;
 calculating level of, 202;
 and dispersion trading, 164, 166, 170, 175;
 and gamma swaps, 52, 56;
 and options on variance, 68;
 underpricing of, 116;
 and variance swaps, 52, 66;
 and volatility futures, 102, 115, 117, 131;
 and volatility swaps 52, 53, 56, 64, 65, 66, 67, 102, 115
- volatility slide theta, 106, 228, 262, 262n.49
- volatility surfaces:
 and additive variance rule, 104, 105;
 chopping tails of, 109, 111, 112;
 chopping wings of, 110, 111;
 and constant smile rule, 105;
 convexity of, 115;
 curvature of, 63n.6, 112, 117;
 discrete sampling of, 112, 113;
 and forward starting products, 100;
 idealised regimes for, 210;
 and implied correlation surfaces, 162;
 implied volatility, 60, 87, 184, 192;
 modeling of, 184, 251–54, 263;
 movement at front end of, 100;
 parallel movement of, 119, 194;
 pinning of, 59;
 and realised volatility, 59, 193;
 re-marking of, 211, 212;
 removal of cap on, 60;
 and skew, 82, 212;
 slope of, 110, 126;
 and square root of time rule, 114, 198, 199;
 and sticky delta and relative time method, 103, 104, 105;
 and sticky strike, 185, 187;
 term structure of, 103, 110;
 and theta, 67;
 three dimensions of, 184
- volatility swaps:
 and dispersion trading, 165, 171;
 fair price of, 109;
 as forwards, 58;
 Greeks of, 67;
 hedging of, 59, 61, 63, 72, 115;
 history of, 2, 52, 53;
 index, 53;
 notional, 250;
 path independence of, 53;
 payoff, 54, 58;
 payout, 52, 53, 55, 57, 64, 240, 250;
 price of, 65, 66;

- volatility swaps (continued):
single stocks, 52, 53, 55;
trade of, 53;
vega profile of, 62, 64, 65;
and volatility of volatility, 53, 56, 64,
65, 66, 67, 72, 115
- volatility trading:
carry of, 158;
forward starting options: 99, 100;
history of, 2;
payout of, 89;
put-call parity in, 36, 37
- volatility weighting, 157
- VOLAX futures, 118
- Volga:
definition, 63n.6, 87, 257, 261;
and dispersion trading, 164, 168, 170,
175;
and log contracts, 63n.6;
as measure of fourth moment of
probability distribution, 202, 205;
and OTM options, 205;
and volatility of volatility, 202, 257, 261
- Volker Rule, the, 82
- vomma, 87
- VSMI (V3X), 51, 109
- VST1ME, 128
- VST1MT, 128
- vStoxx (V2X):
correlation with SX5E, 74;
ETN/ETF, 125;
excess return indices on, 128;
liquidity of, 122;
listing of forwards on, 99, 100, 101;
and relative value volatility trading, 158;
size of market, 125, 128;
- settlement price of, 119;
and SX5E term structure, 125;
total return indices on, 128;
variance-based calculation of, 51, 109
- VXD, 118
- VXN, 118
- VXO, 49, 51, 109
- VXTH, 118
- VXX, 127, 132
- VXZ, 132
- W**
- weighting in dispersion trading, 171–74
- wings, 63, 75, 110, 111, 227, 263
- worst-of/best-of options, 133, 140–42,
164, 165
- X**
- XXV, 127
- Y**
- Yang-Zhang, 232, 238–39, 241, 243
- yield curve, 184
- Z**
- Zero cost:
1x2, 17, 18, 26, 27–31, 32
call spread vs put, 32
collar, 24, 25, 26, 27
put spread collar, 24, 25, 27
risk reversal, 32